



GREAT INDIAN **DEVELOPER** SUMMIT



™

April 25-28, 2017. Bangalore

Spark, Mesos, Akka, Cassandra and Kafka (SMACK) Stack Real Time Big Data

Rohit Bhardwaj
Principal Cloud Engineer

SMACK AGENDA

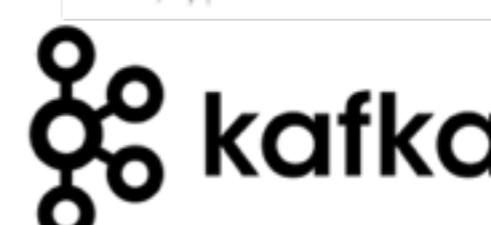
- SMACK stack overview
- Bigdata characteristics
- Cassandra NoSQL database
- Spark Streaming
- Kafka and Akka event sourcing
- Mesos Cluster management
- Mesos Scheduling and execution

Spark for Micro-Batch Processing 

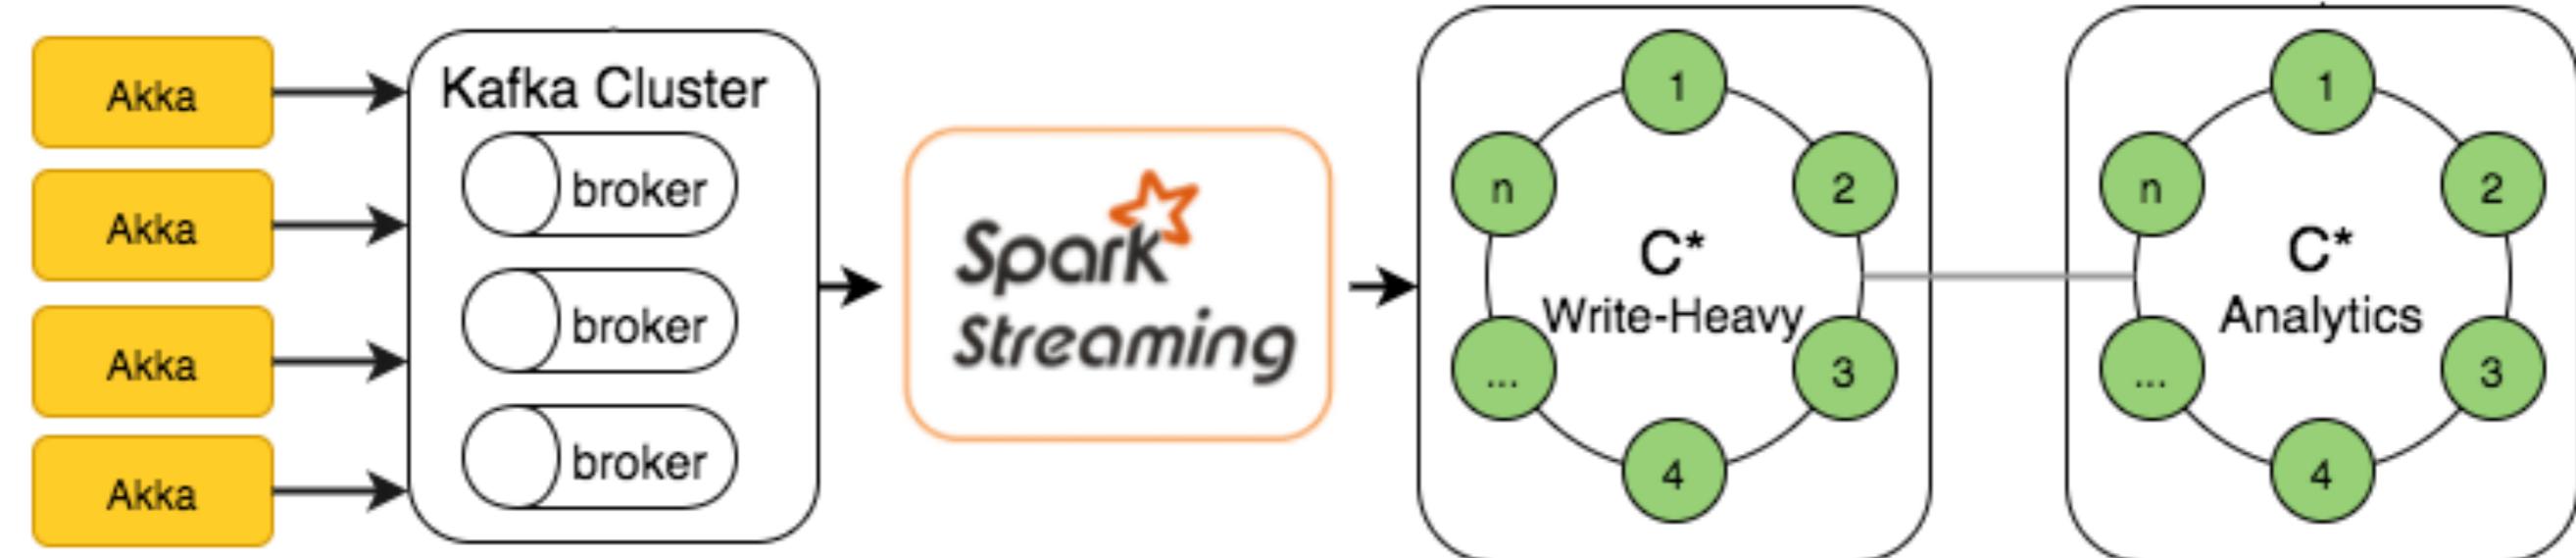
Mesos for Cluster Management 

Akka for Event Processing 

Cassandra for Persistence 

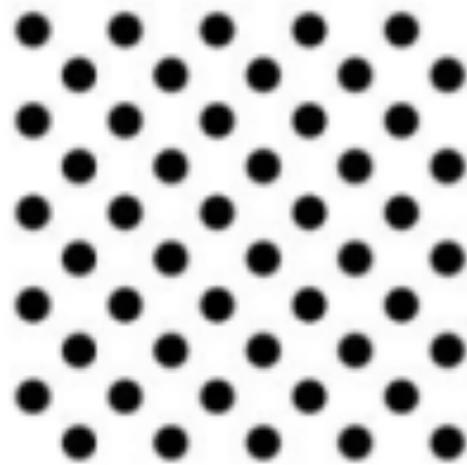
Kafka for Event Transport 

MESOS



BIG DATA CHARACTERISTICS

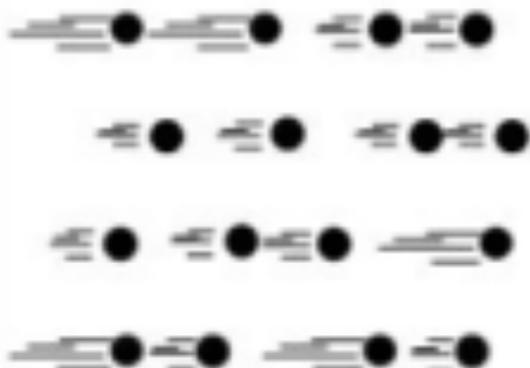
Volume



Data at Rest

Terabytes to exabytes of existing data to process

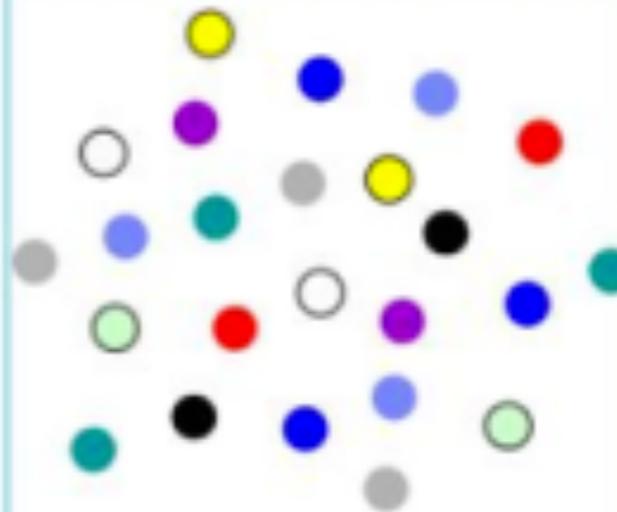
Velocity



Data in Motion

Streaming data, milliseconds to seconds to respond

Variety



Data in Many Forms

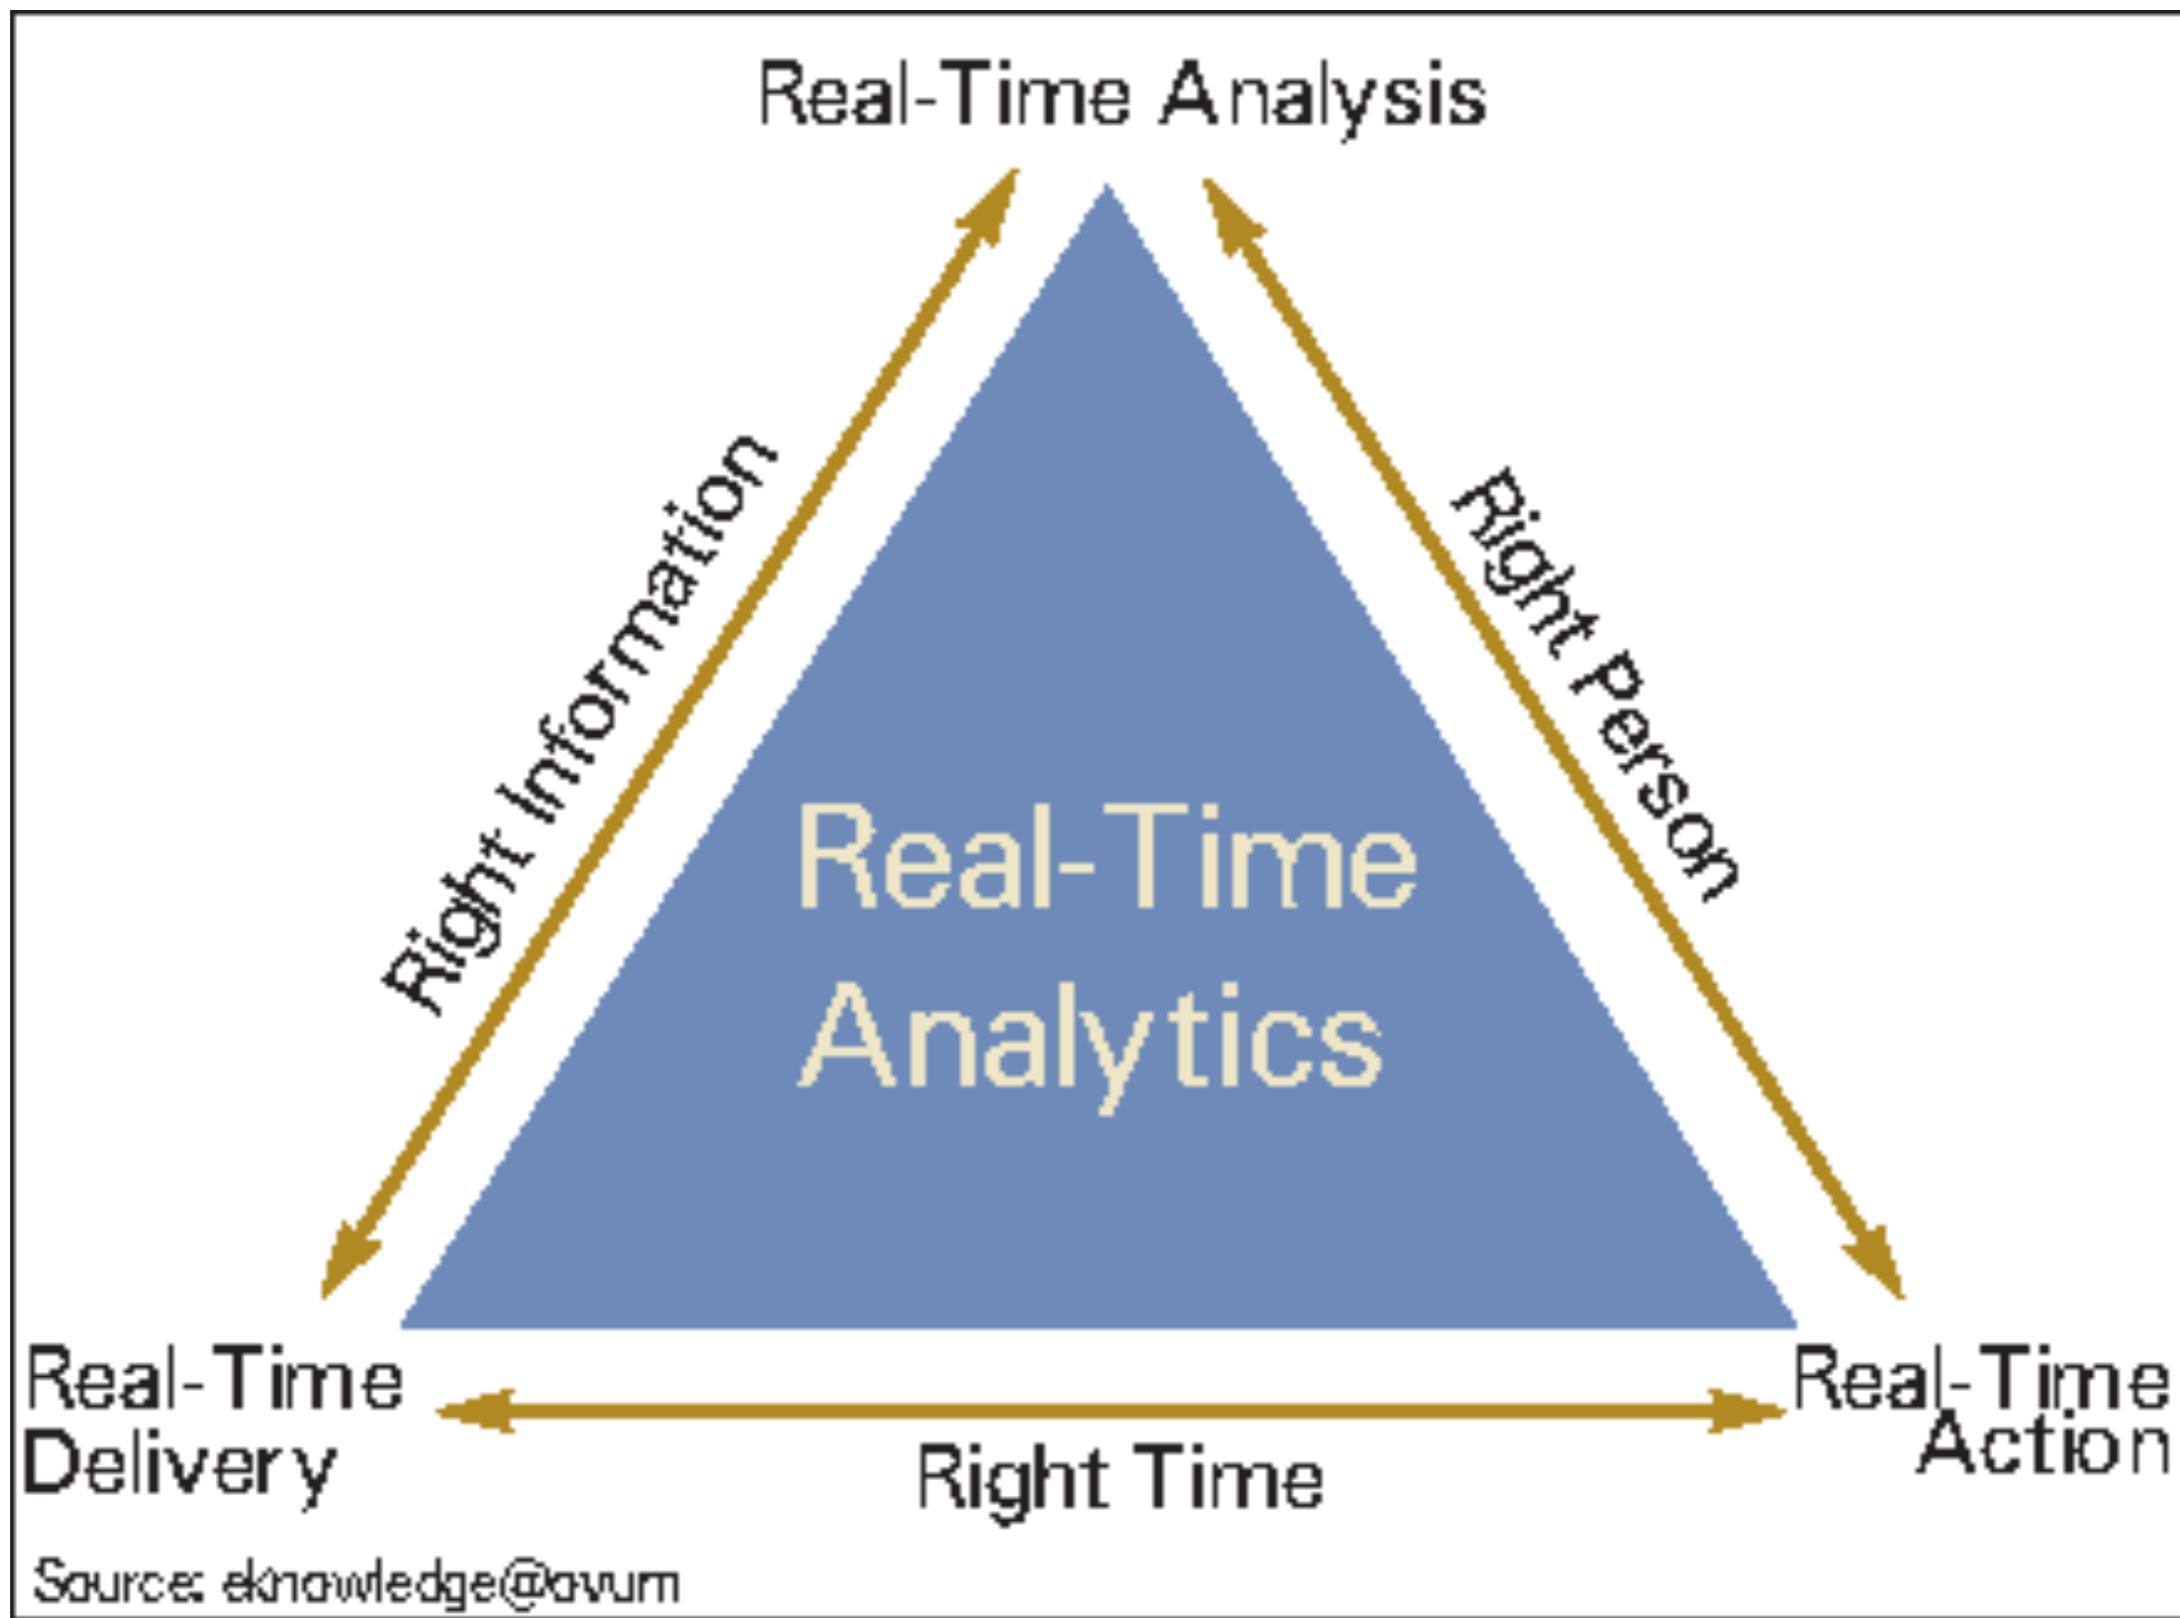
Structured, unstructured, text, multimedia

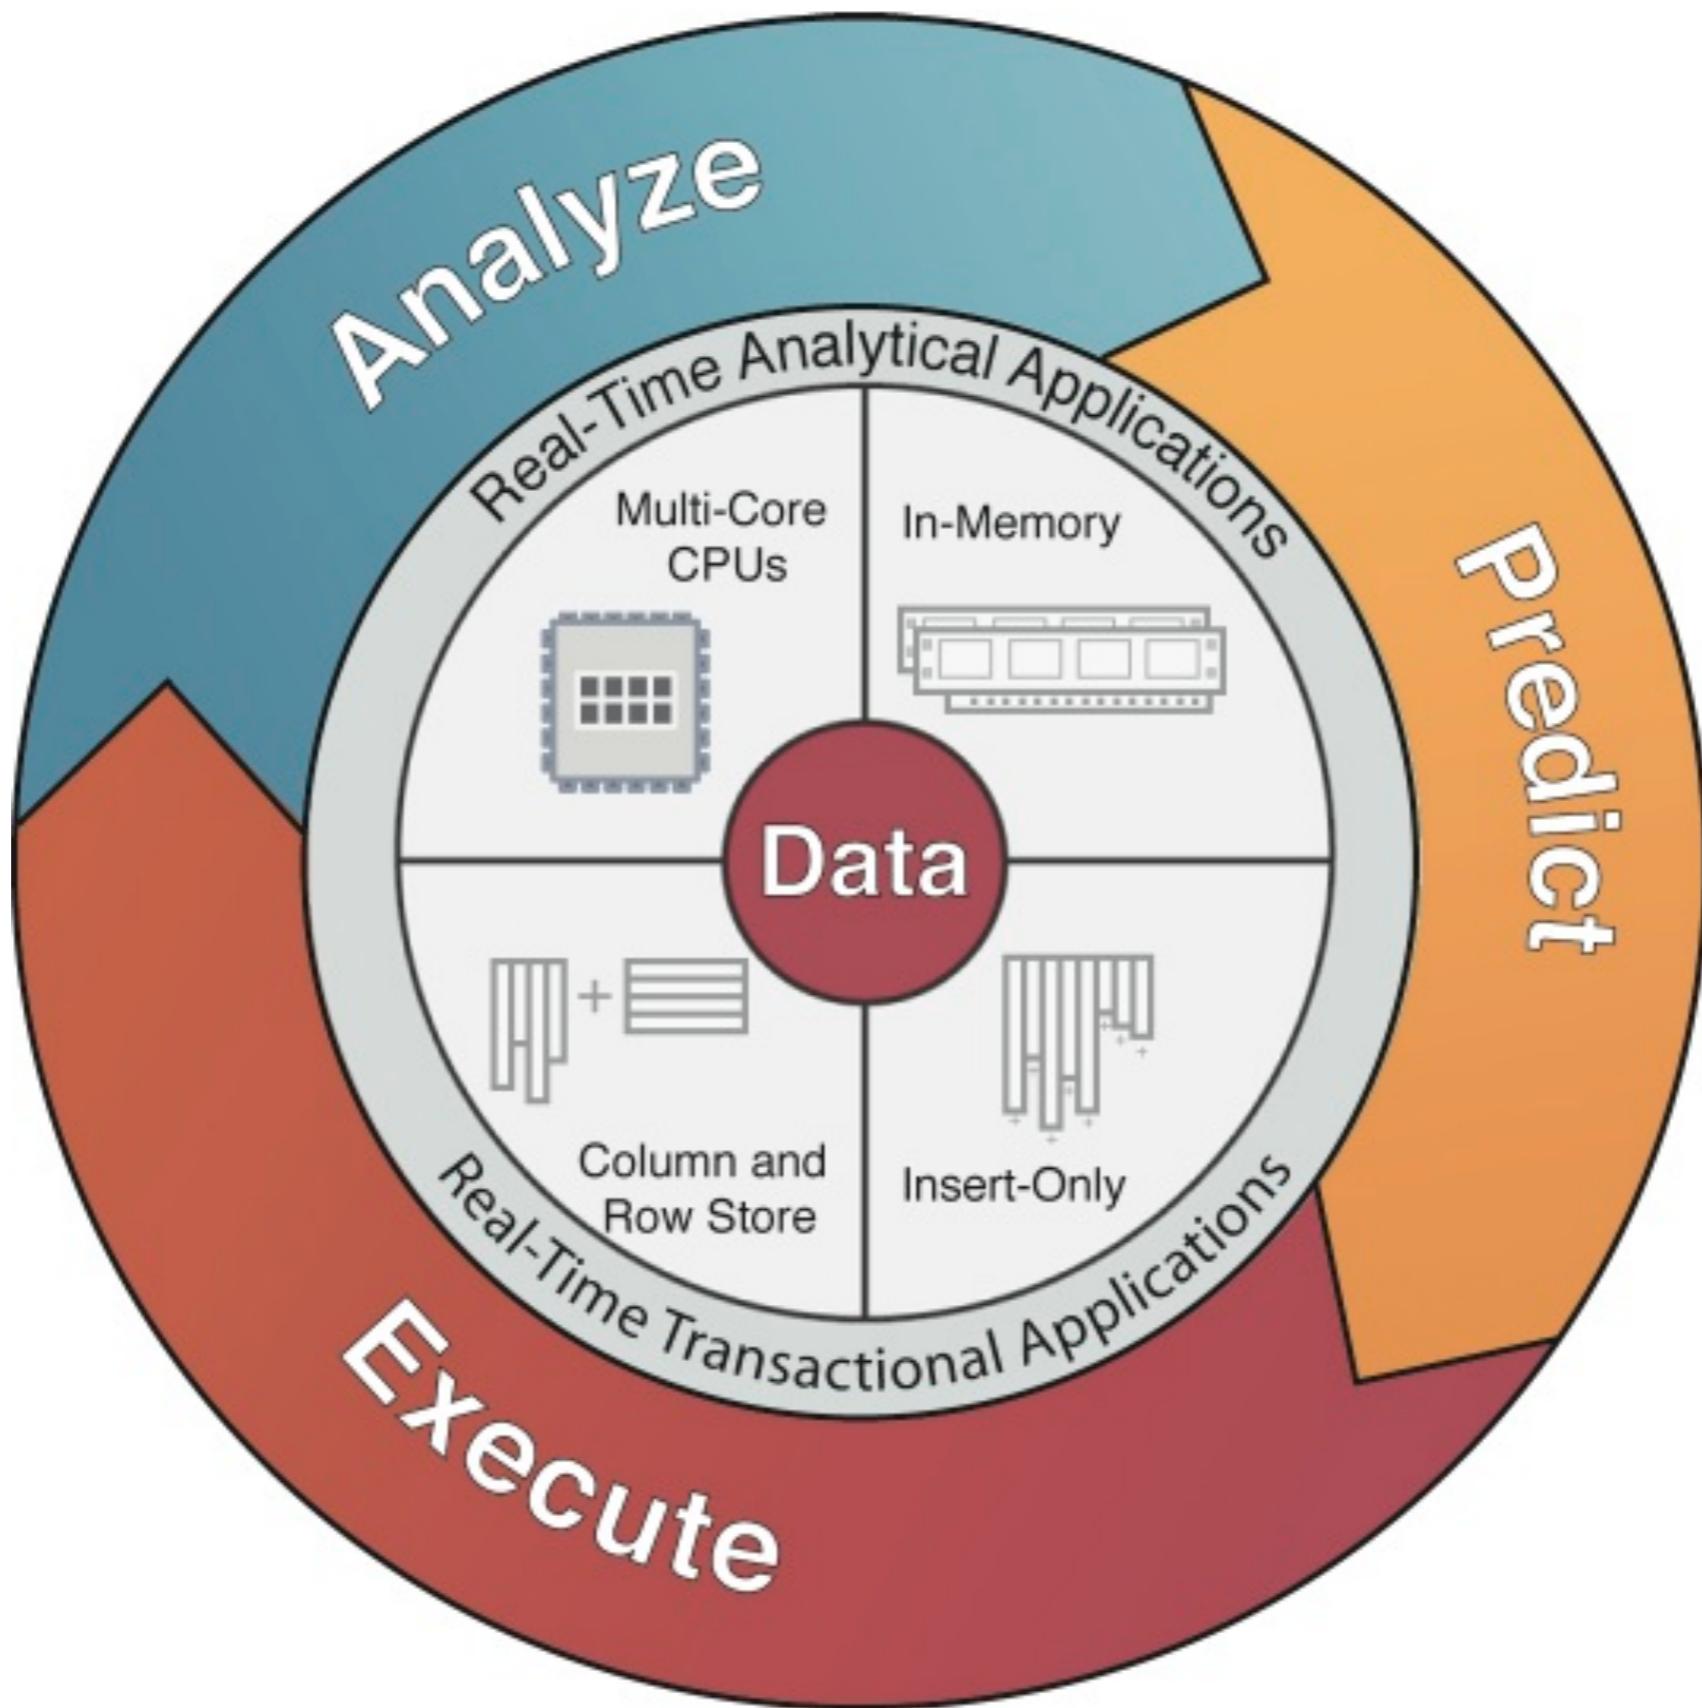
Veracity*



Data in Doubt

Uncertainty due to data inconsistency & incompleteness, ambiguities, latency, deception, model approximations





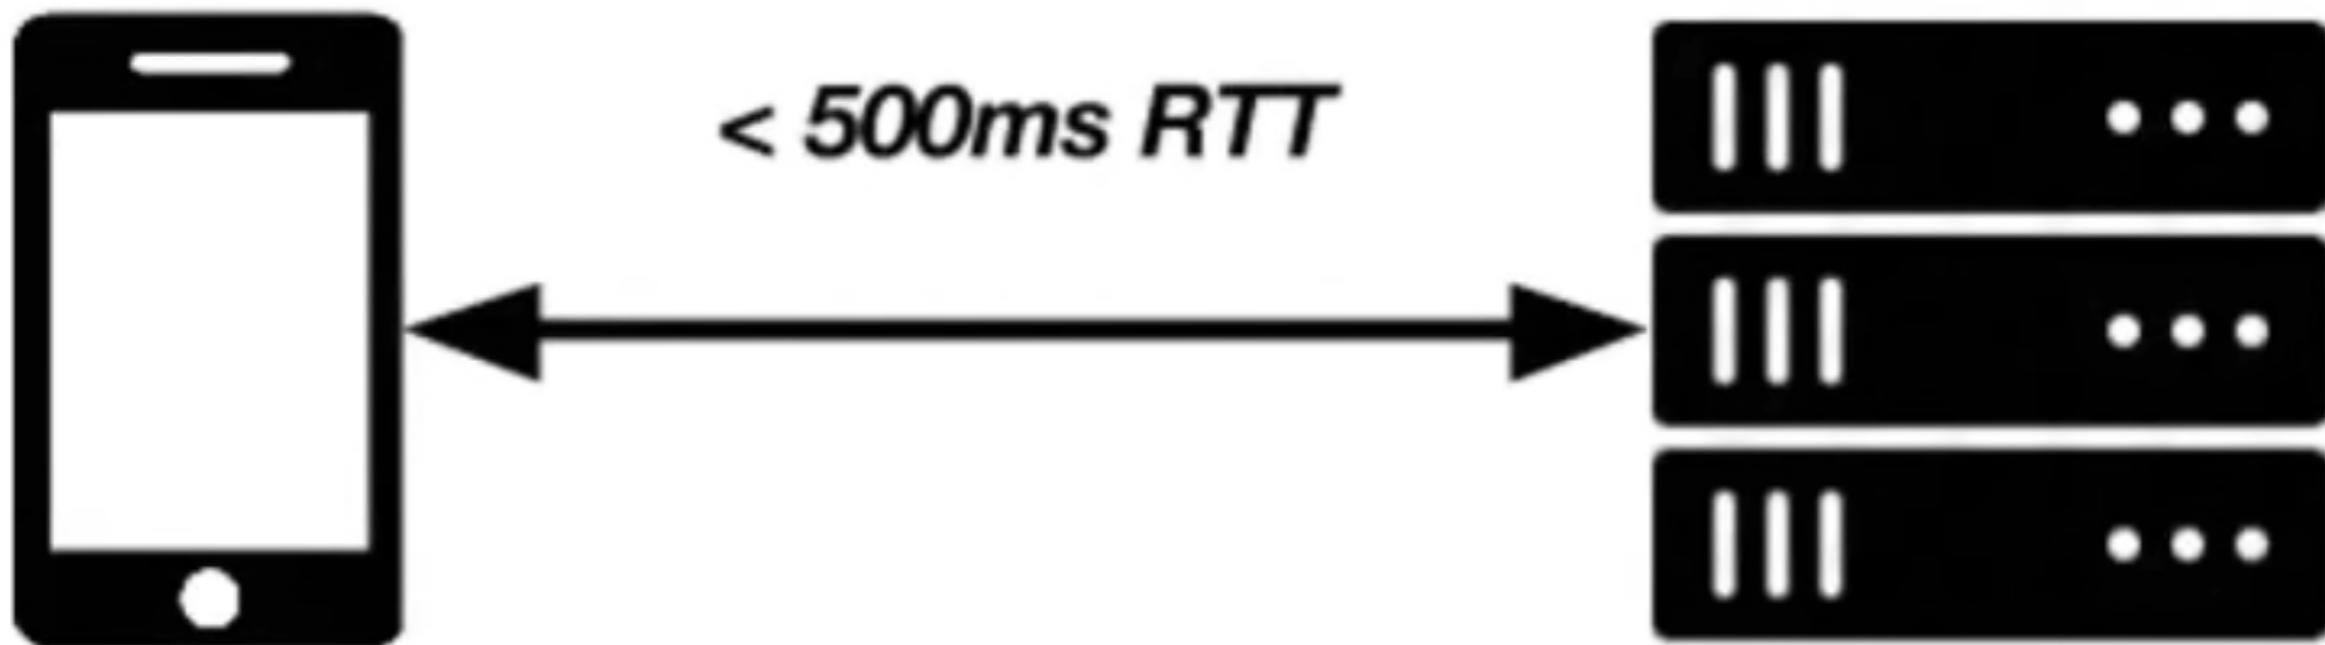
99.99% Uptime



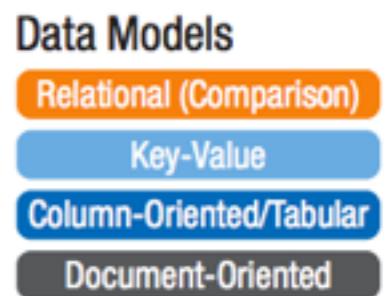
Unavailable Minutes per Month

AVAILABILITY LEVEL¹¹	ESTIMATED DOWNTIME PER YEAR
99.999%	5 minutes
99.99%	52 minutes
99.9%	8.5 hours
99%	3.5 days or about 87 hours

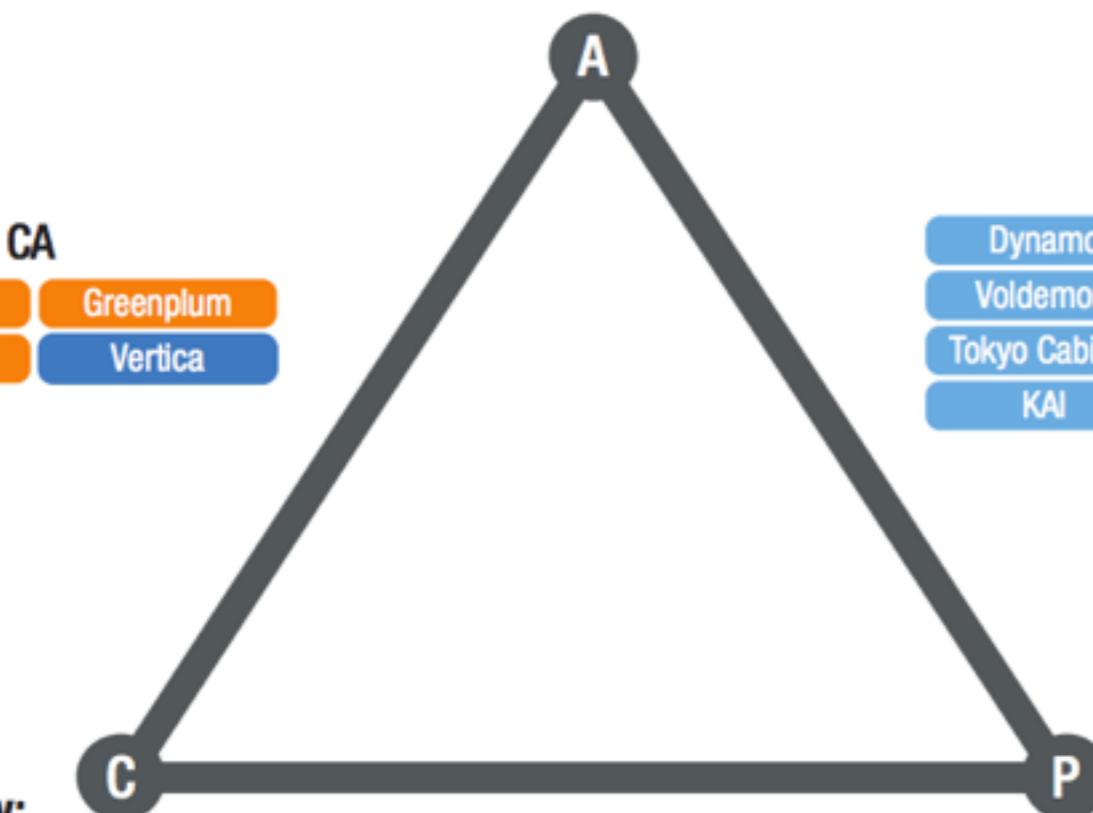
Subsecond End-User Latency



U B E R



Availability:
Each client can always read and write



Consistency:
All clients always have the same view of the data

CA

RDBMSs	Greenplum
Aster Data	Vertica

AP

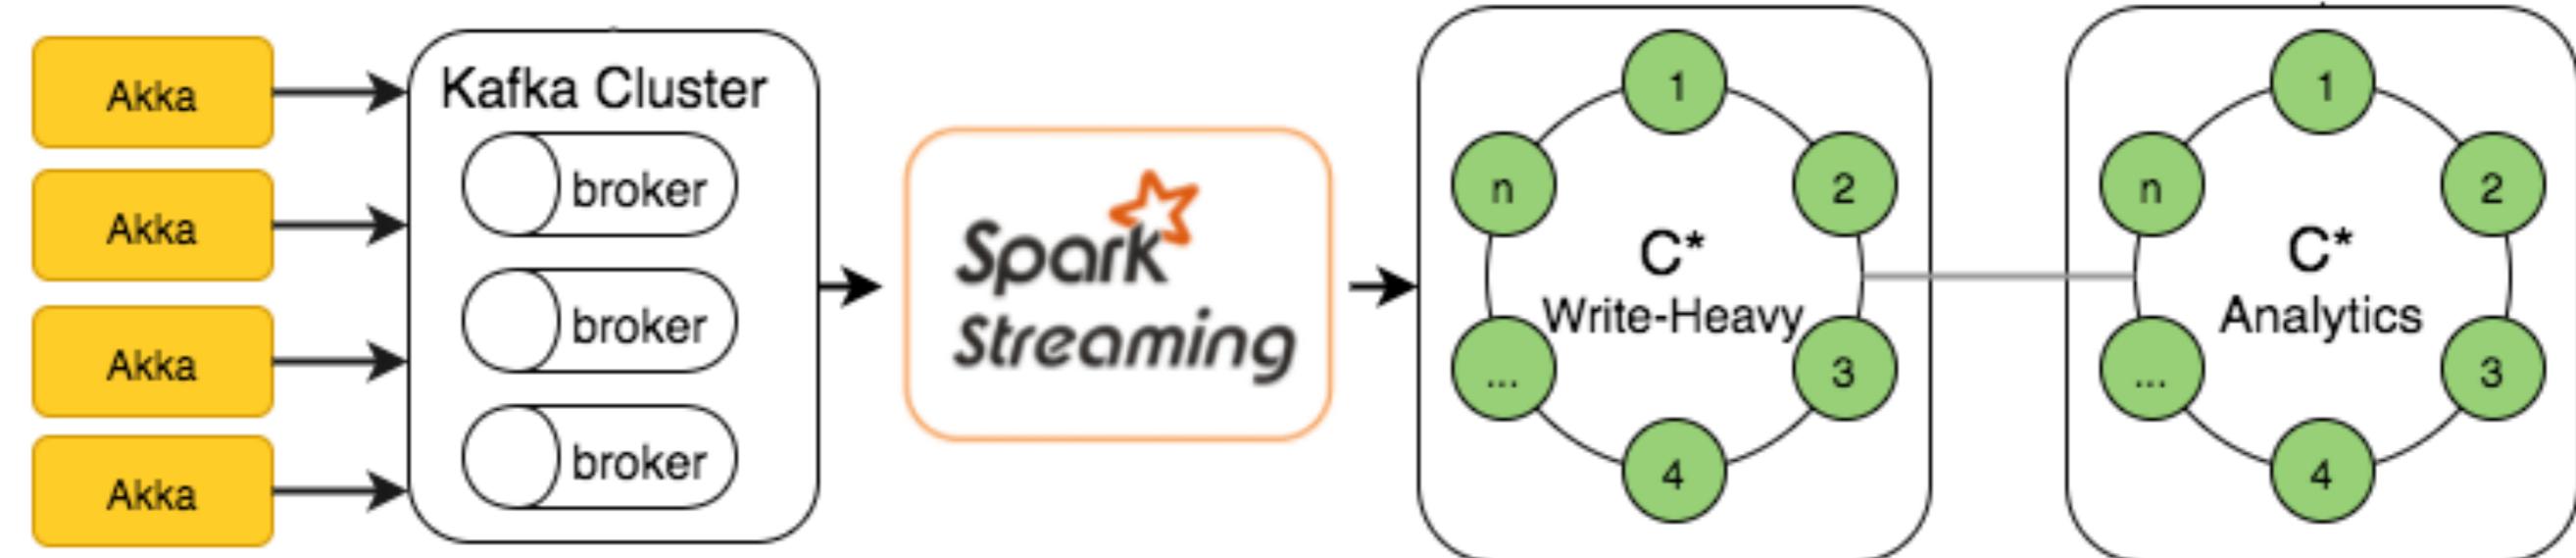
Dynamo	Cassandra
Voldemort	SimpleDB
Tokyo Cabinet	CouchDB
KAI	Riak

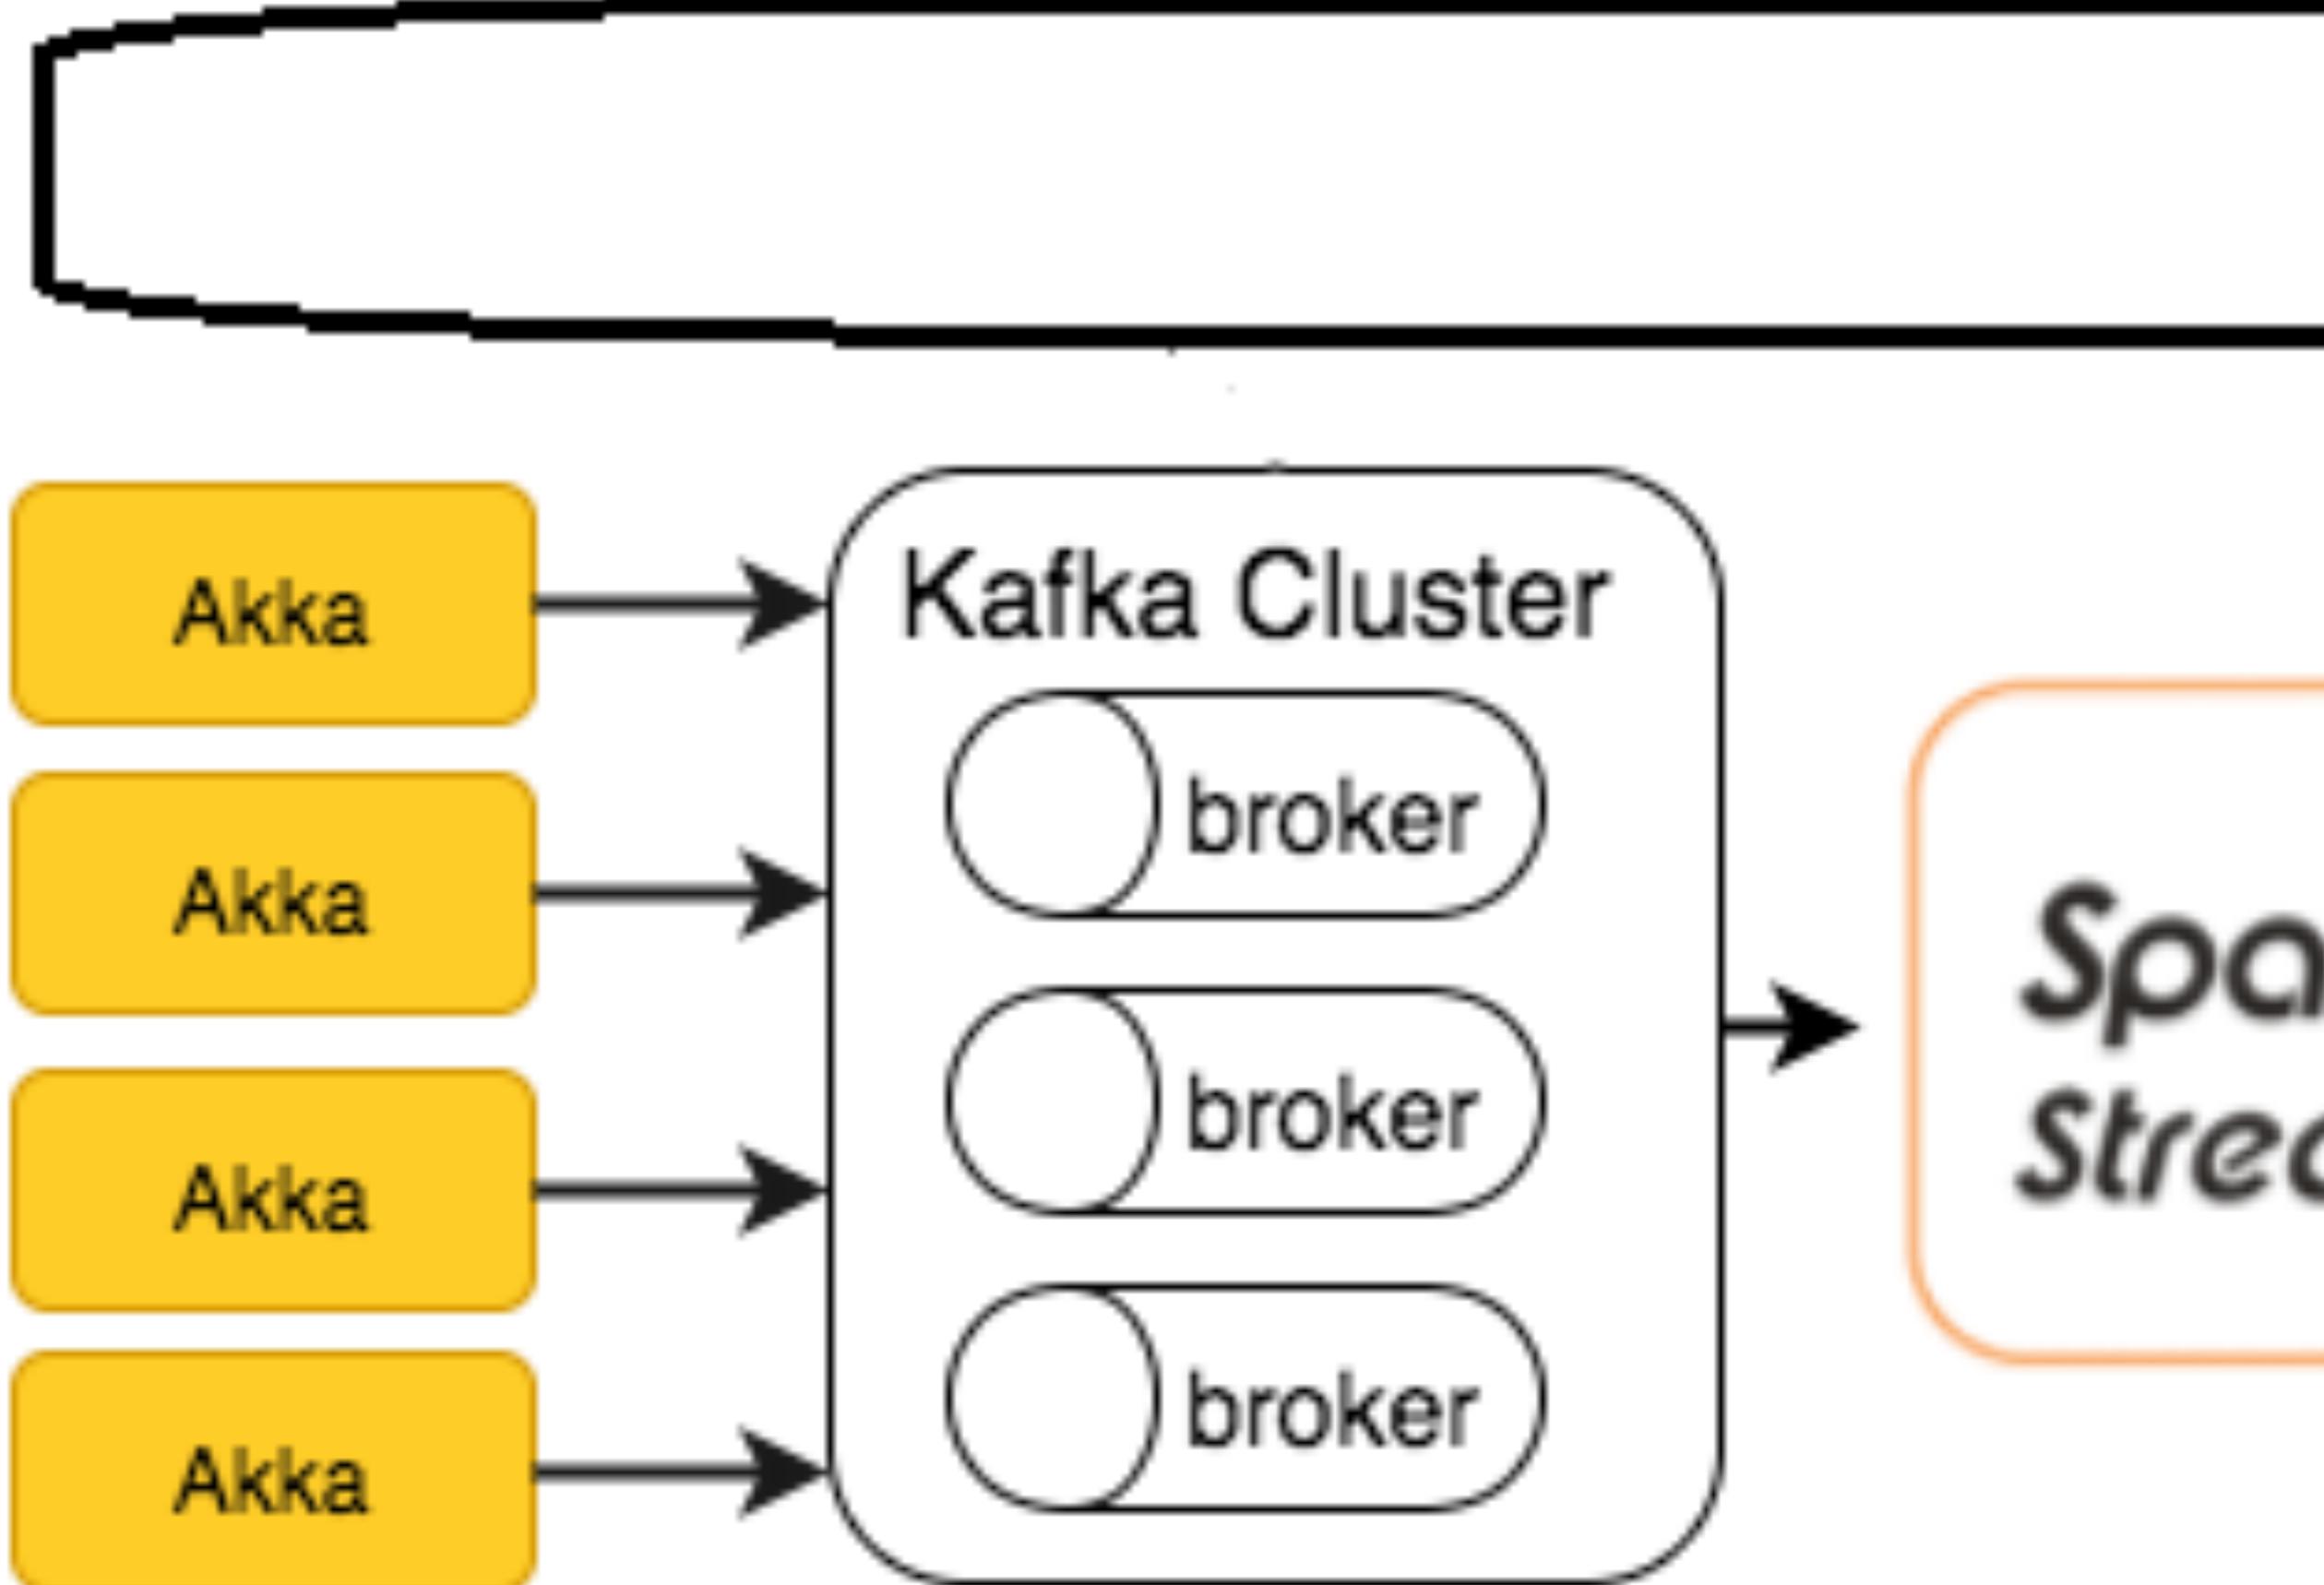
CP

Big Table	Mongo DB	Berkeley DB
Hypertable	Terrastore	Memcache DB
Hbase	Scalarmis	Redis

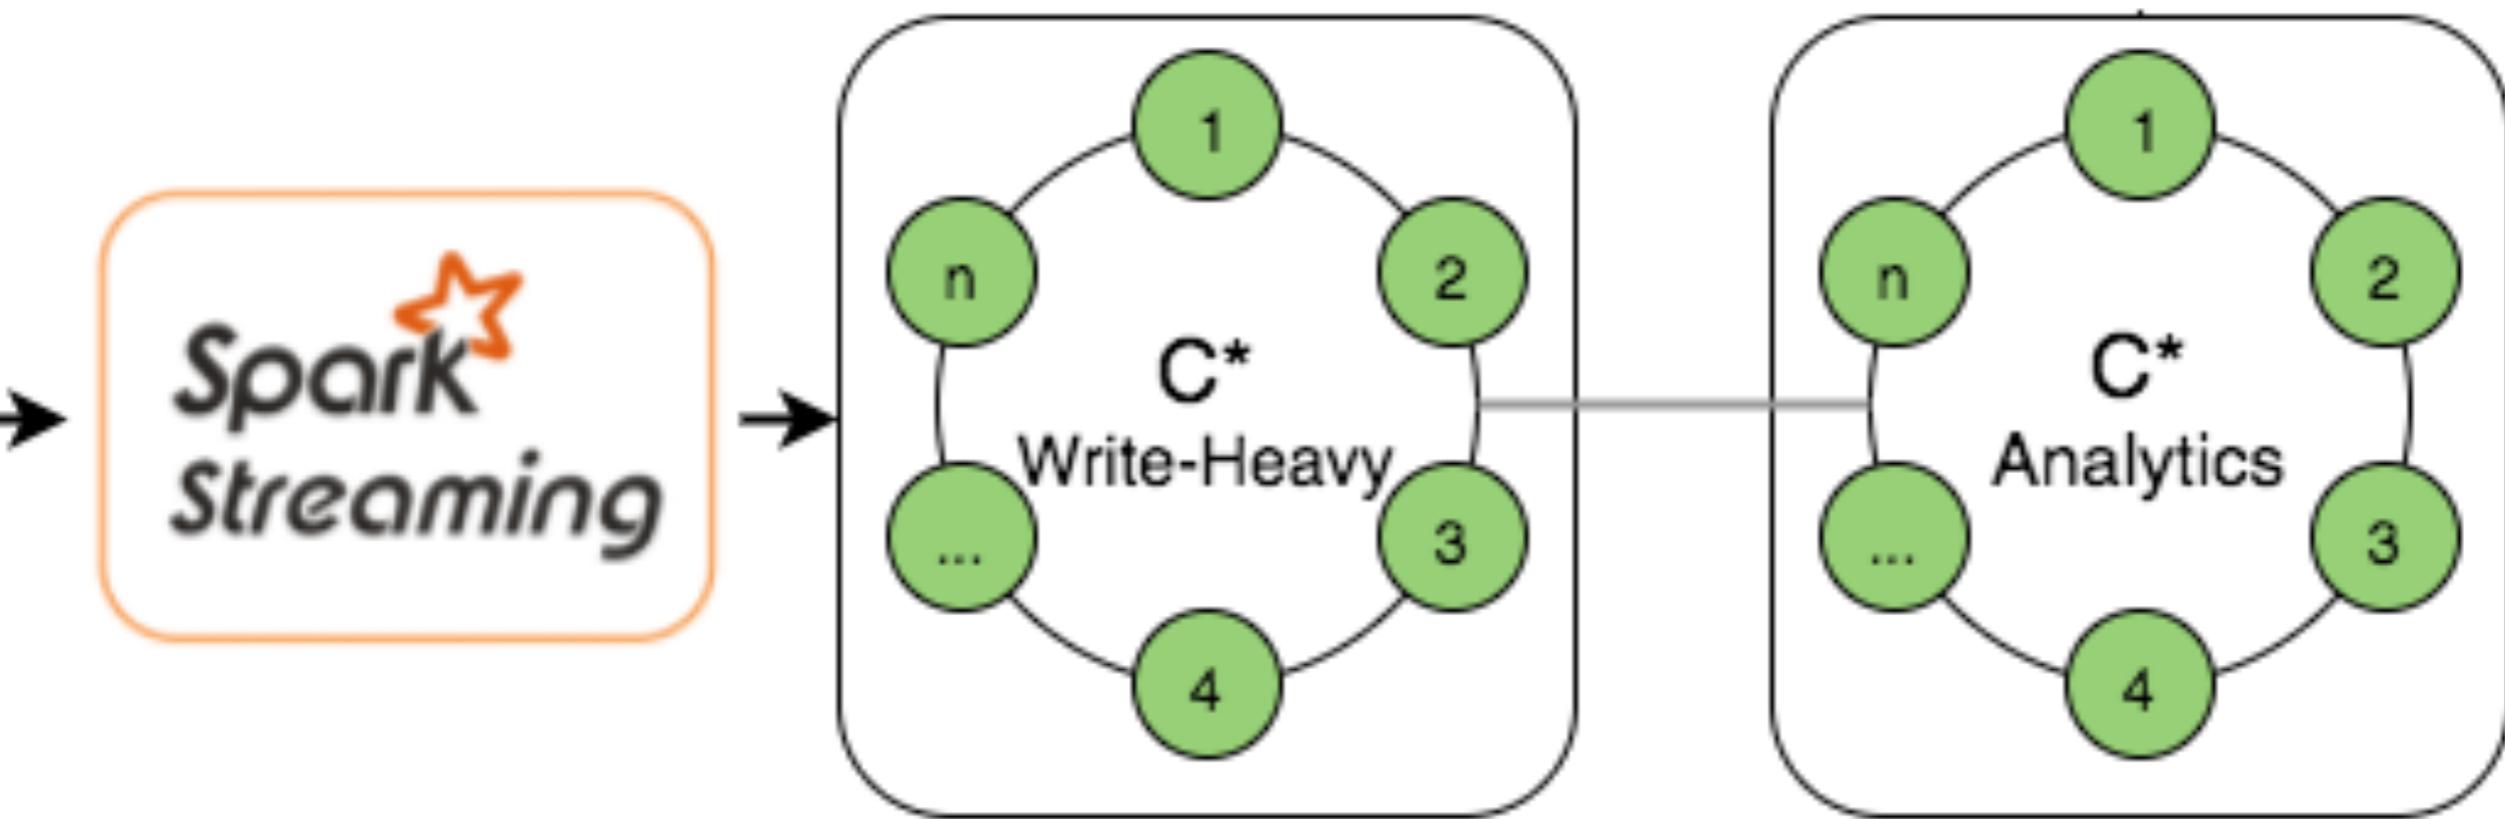
Partition Tolerance:
The system works well despite physical network partitions

MESOS





MESOS





Cassandra

A highly scalable, eventually consistent, distributed, structured key-value store.

Bigtable



Dynamo



Cassandra



Cassandra

[Home](#) [Download](#) [Getting Started](#) [Planet Cassandra](#) [Contribute](#)

[Welcome](#) [Video](#) [Slides](#)

Welcome to Apache Cassandra™

The Apache Cassandra database is the right choice when you need scalability and high availability without compromising performance. [Linear scalability](#) and proven fault-tolerance on commodity hardware or cloud infrastructure make it the perfect platform for mission-critical data. Cassandra's support for replicating across multiple datacenters is best-in-class, providing lower latency for your users and the peace of mind of knowing that you can survive regional outages.

Cassandra's data model offers the convenience of [column indexes](#) with the performance of log-structured updates, strong support for [denormalization](#) and [materialized views](#), and powerful built-in caching.

Download

[Tick-Tock release 3.4 \(Changes\)](#)

[3.0.x release 3.0.4 \(Changes\)](#)

[2.2.x release 2.2.5 \(Changes\)](#)



[Download options](#)

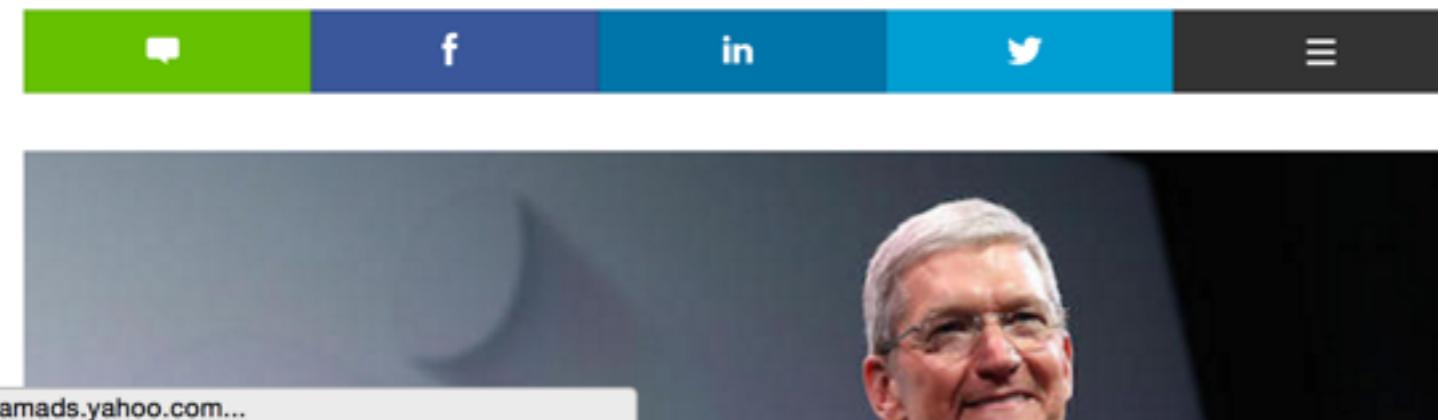
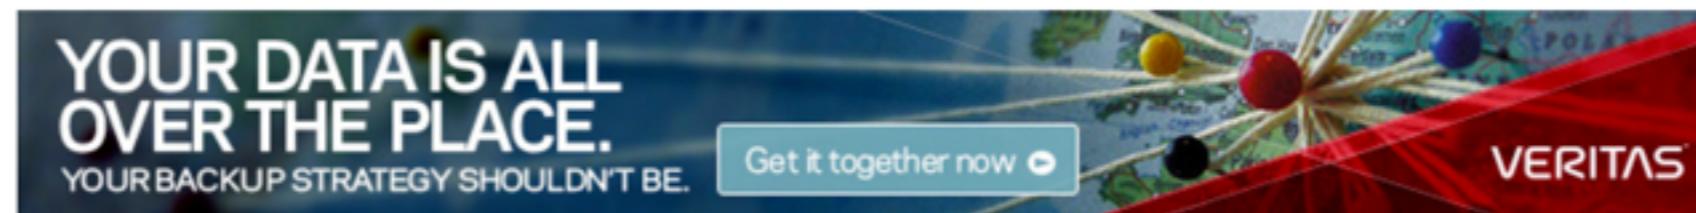
SOFTWARE



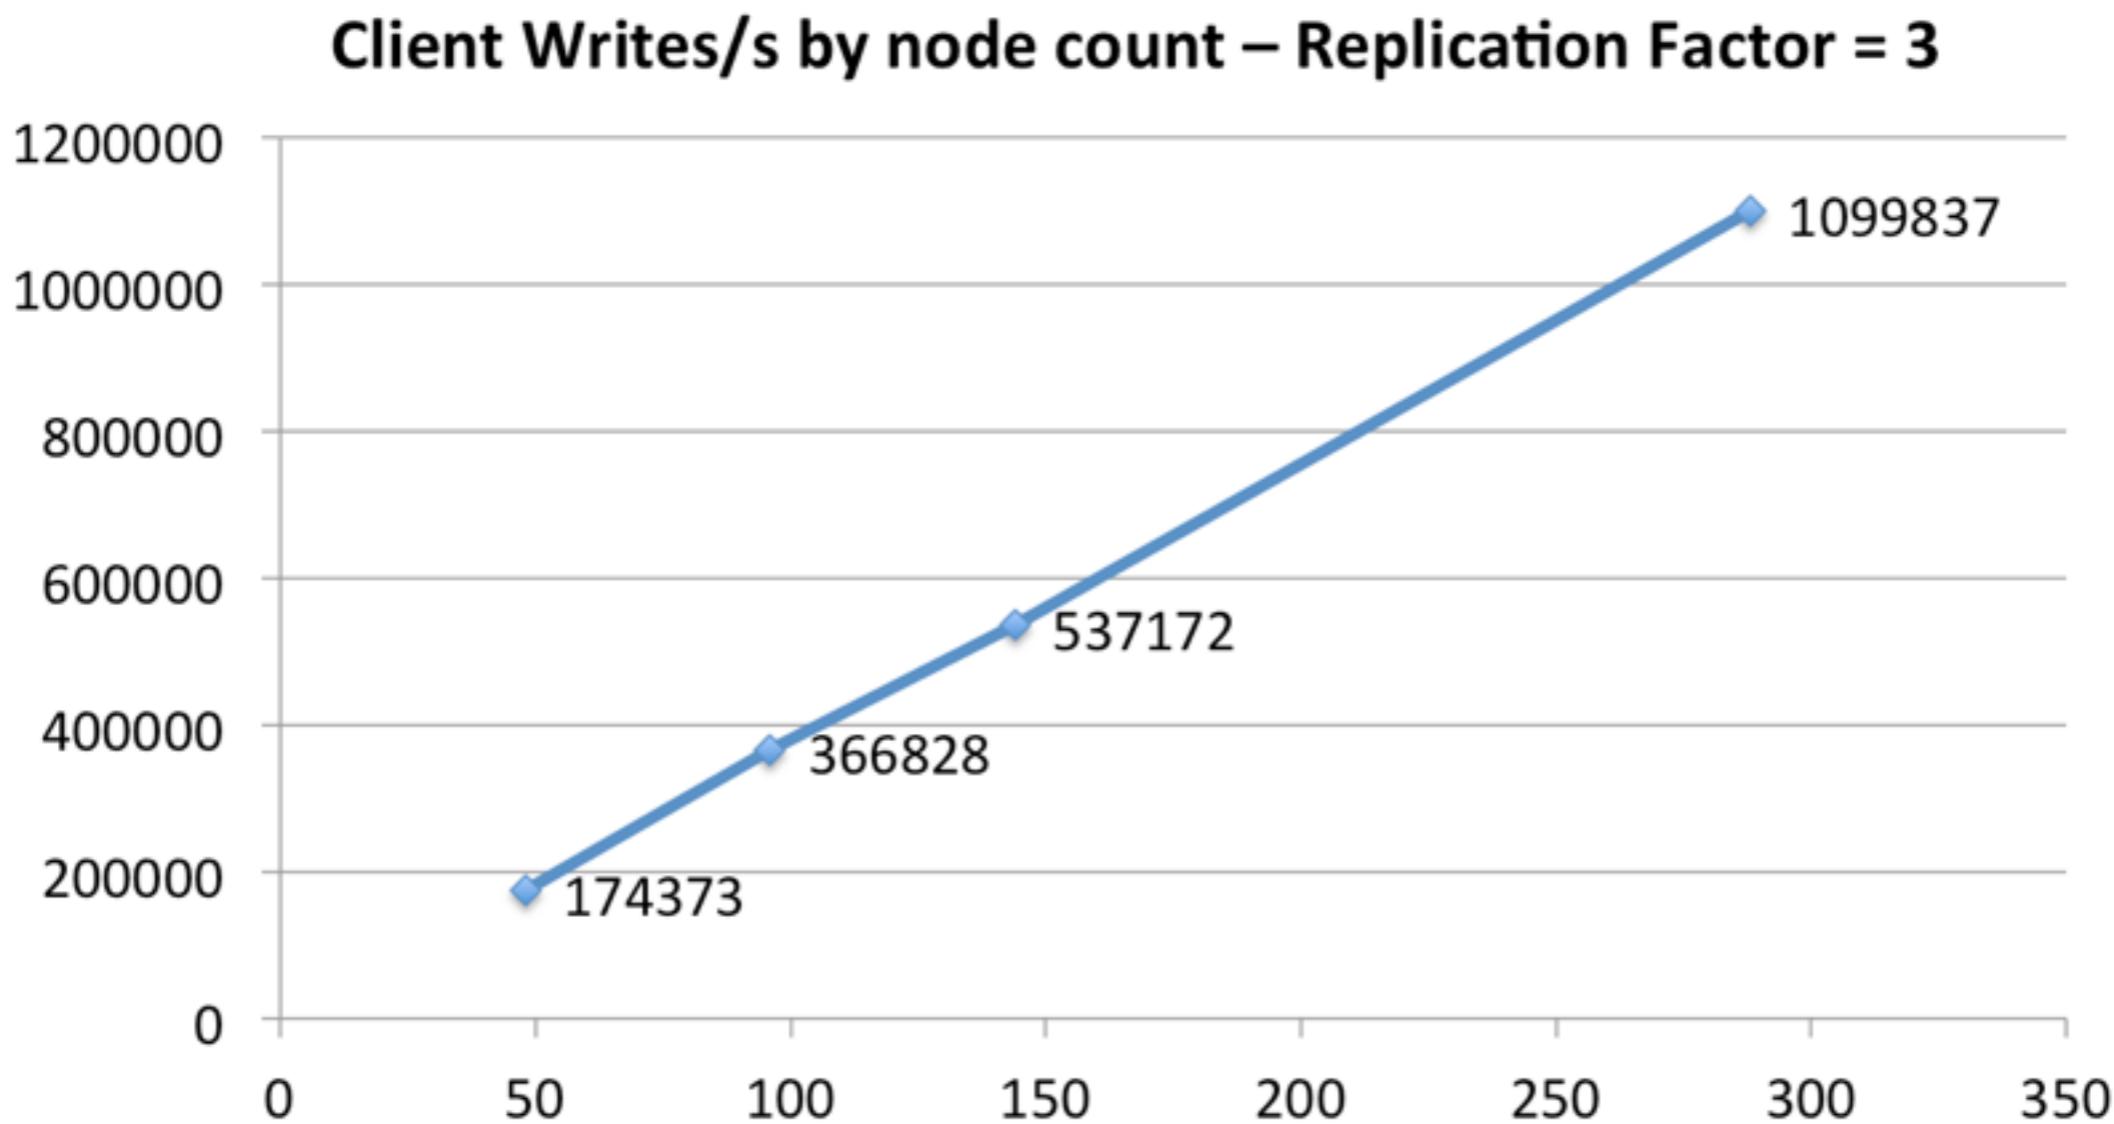
Apple's secret NoSQL sauce includes a hefty dose of Cassandra

If you want to scale like Apple, you might need to consider Cassandra. Matt Asay explains.

By Matt Asay | September 16, 2015, 6:49 AM PST



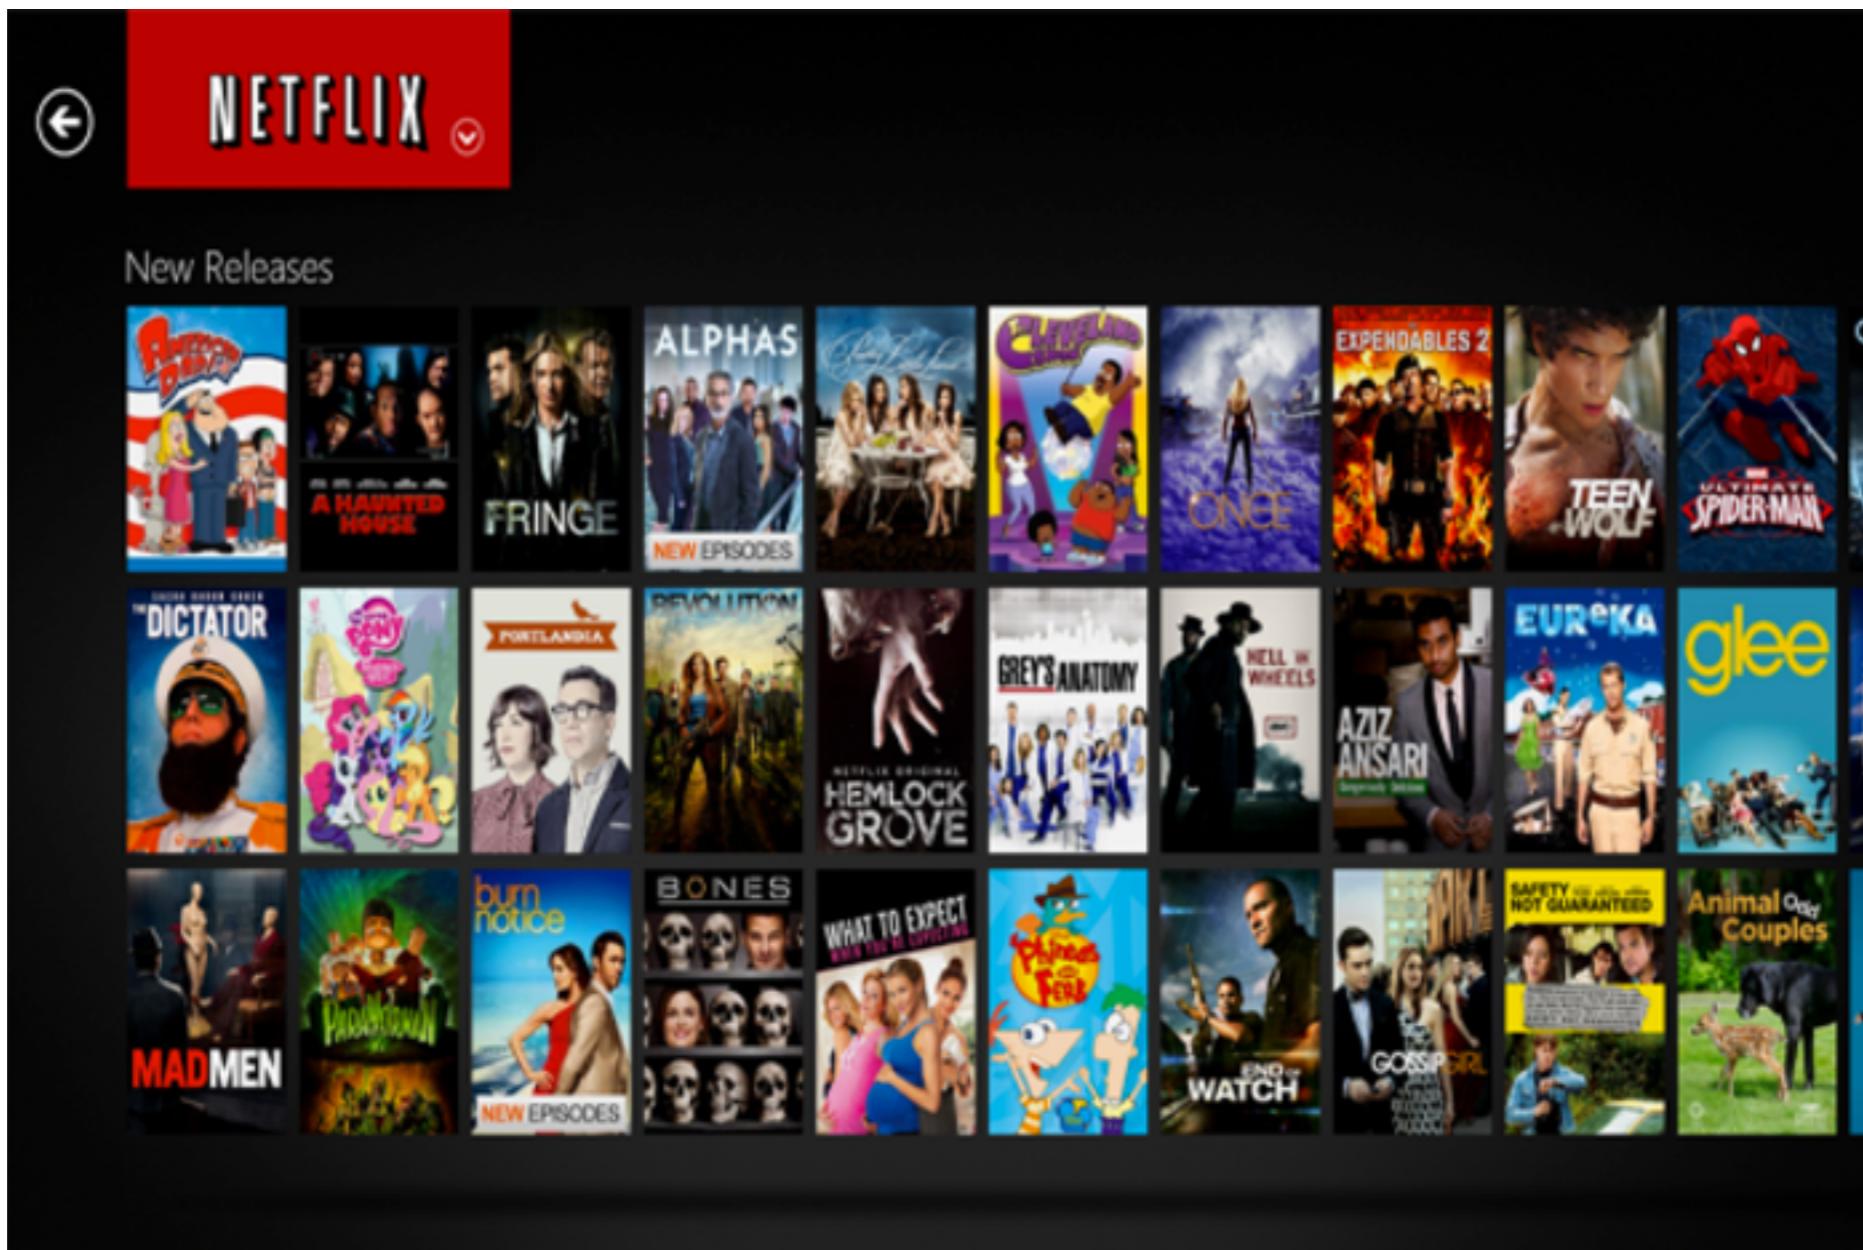
Scale-Up Linearity

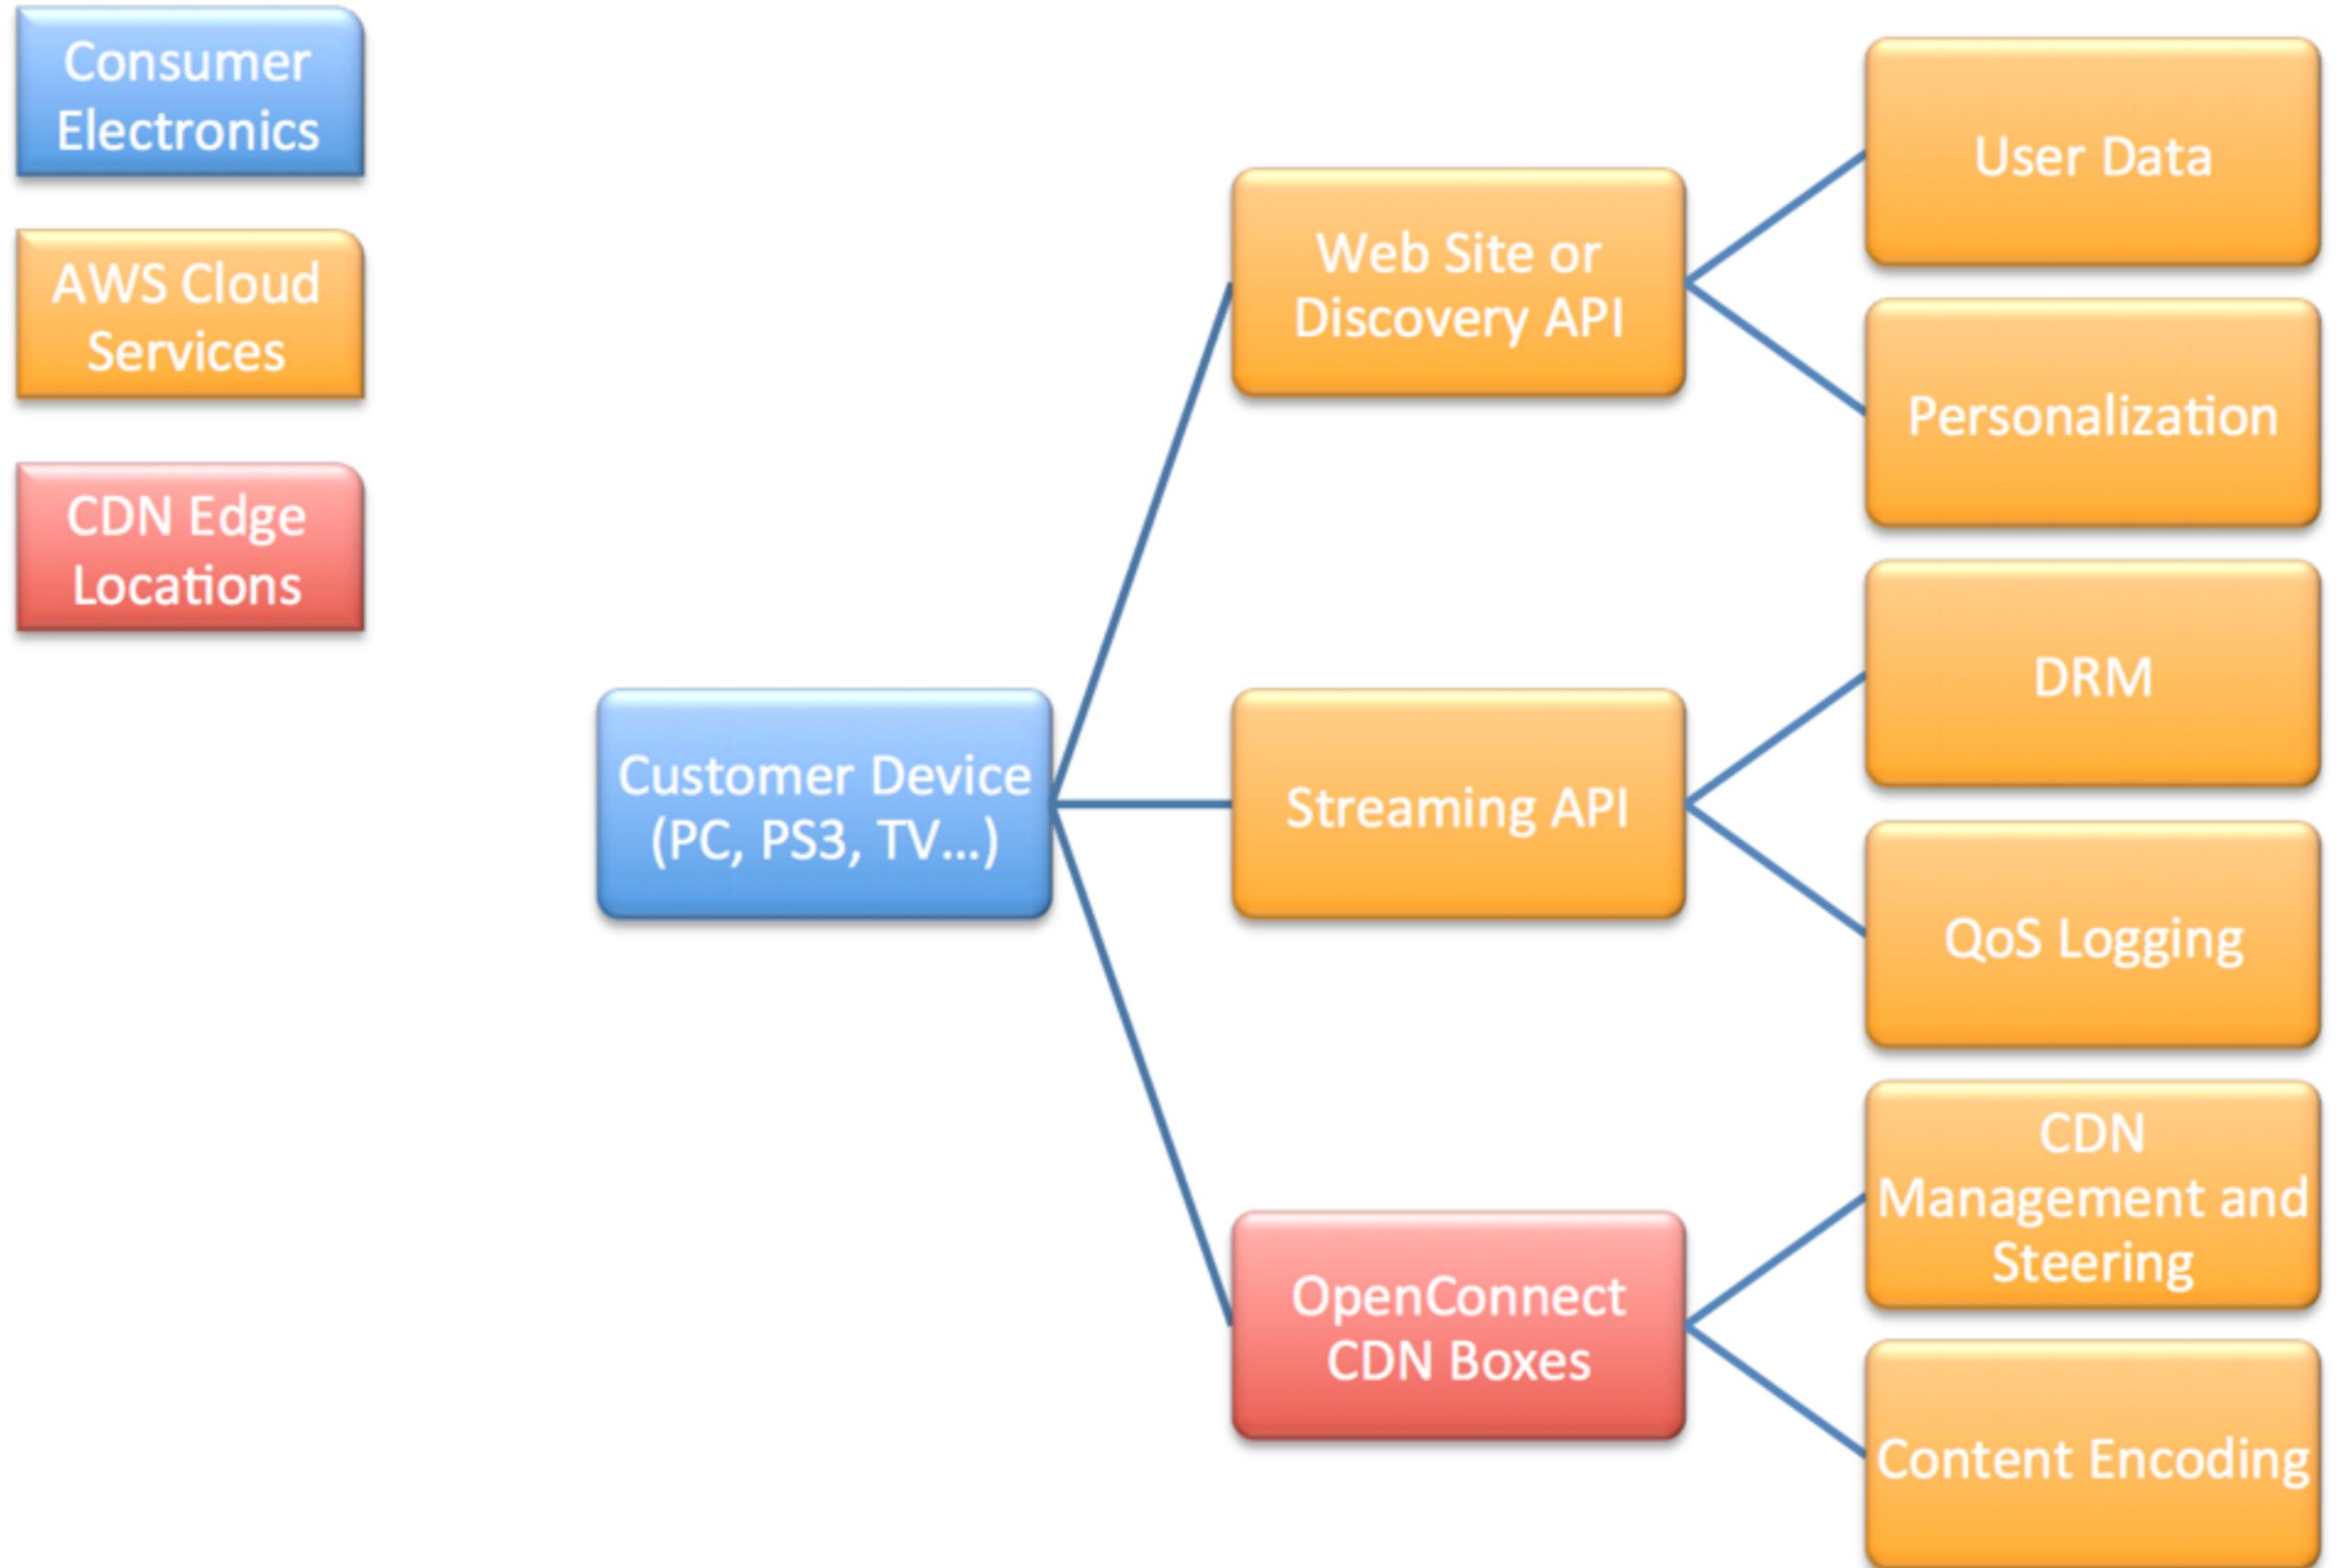


The Netflix logo, featuring the word "NETFLIX" in its signature white, blocky font against a solid orange background.The Intuit logo, with the word "intuit." in blue lowercase letters, followed by a registered trademark symbol.The Call of Duty logo, with the words "CALL OF DUTY" in a stylized, metallic font.The Twitter logo, with the word "twitter" in its signature blue, lowercase, sans-serif font.The eBay logo, with the word "ebay" in a stylized font where each letter has a different color: red, blue, yellow, and green.The Cisco logo, featuring a series of seven vertical bars of increasing height followed by the word "CISCO" in a bold, red, sans-serif font.

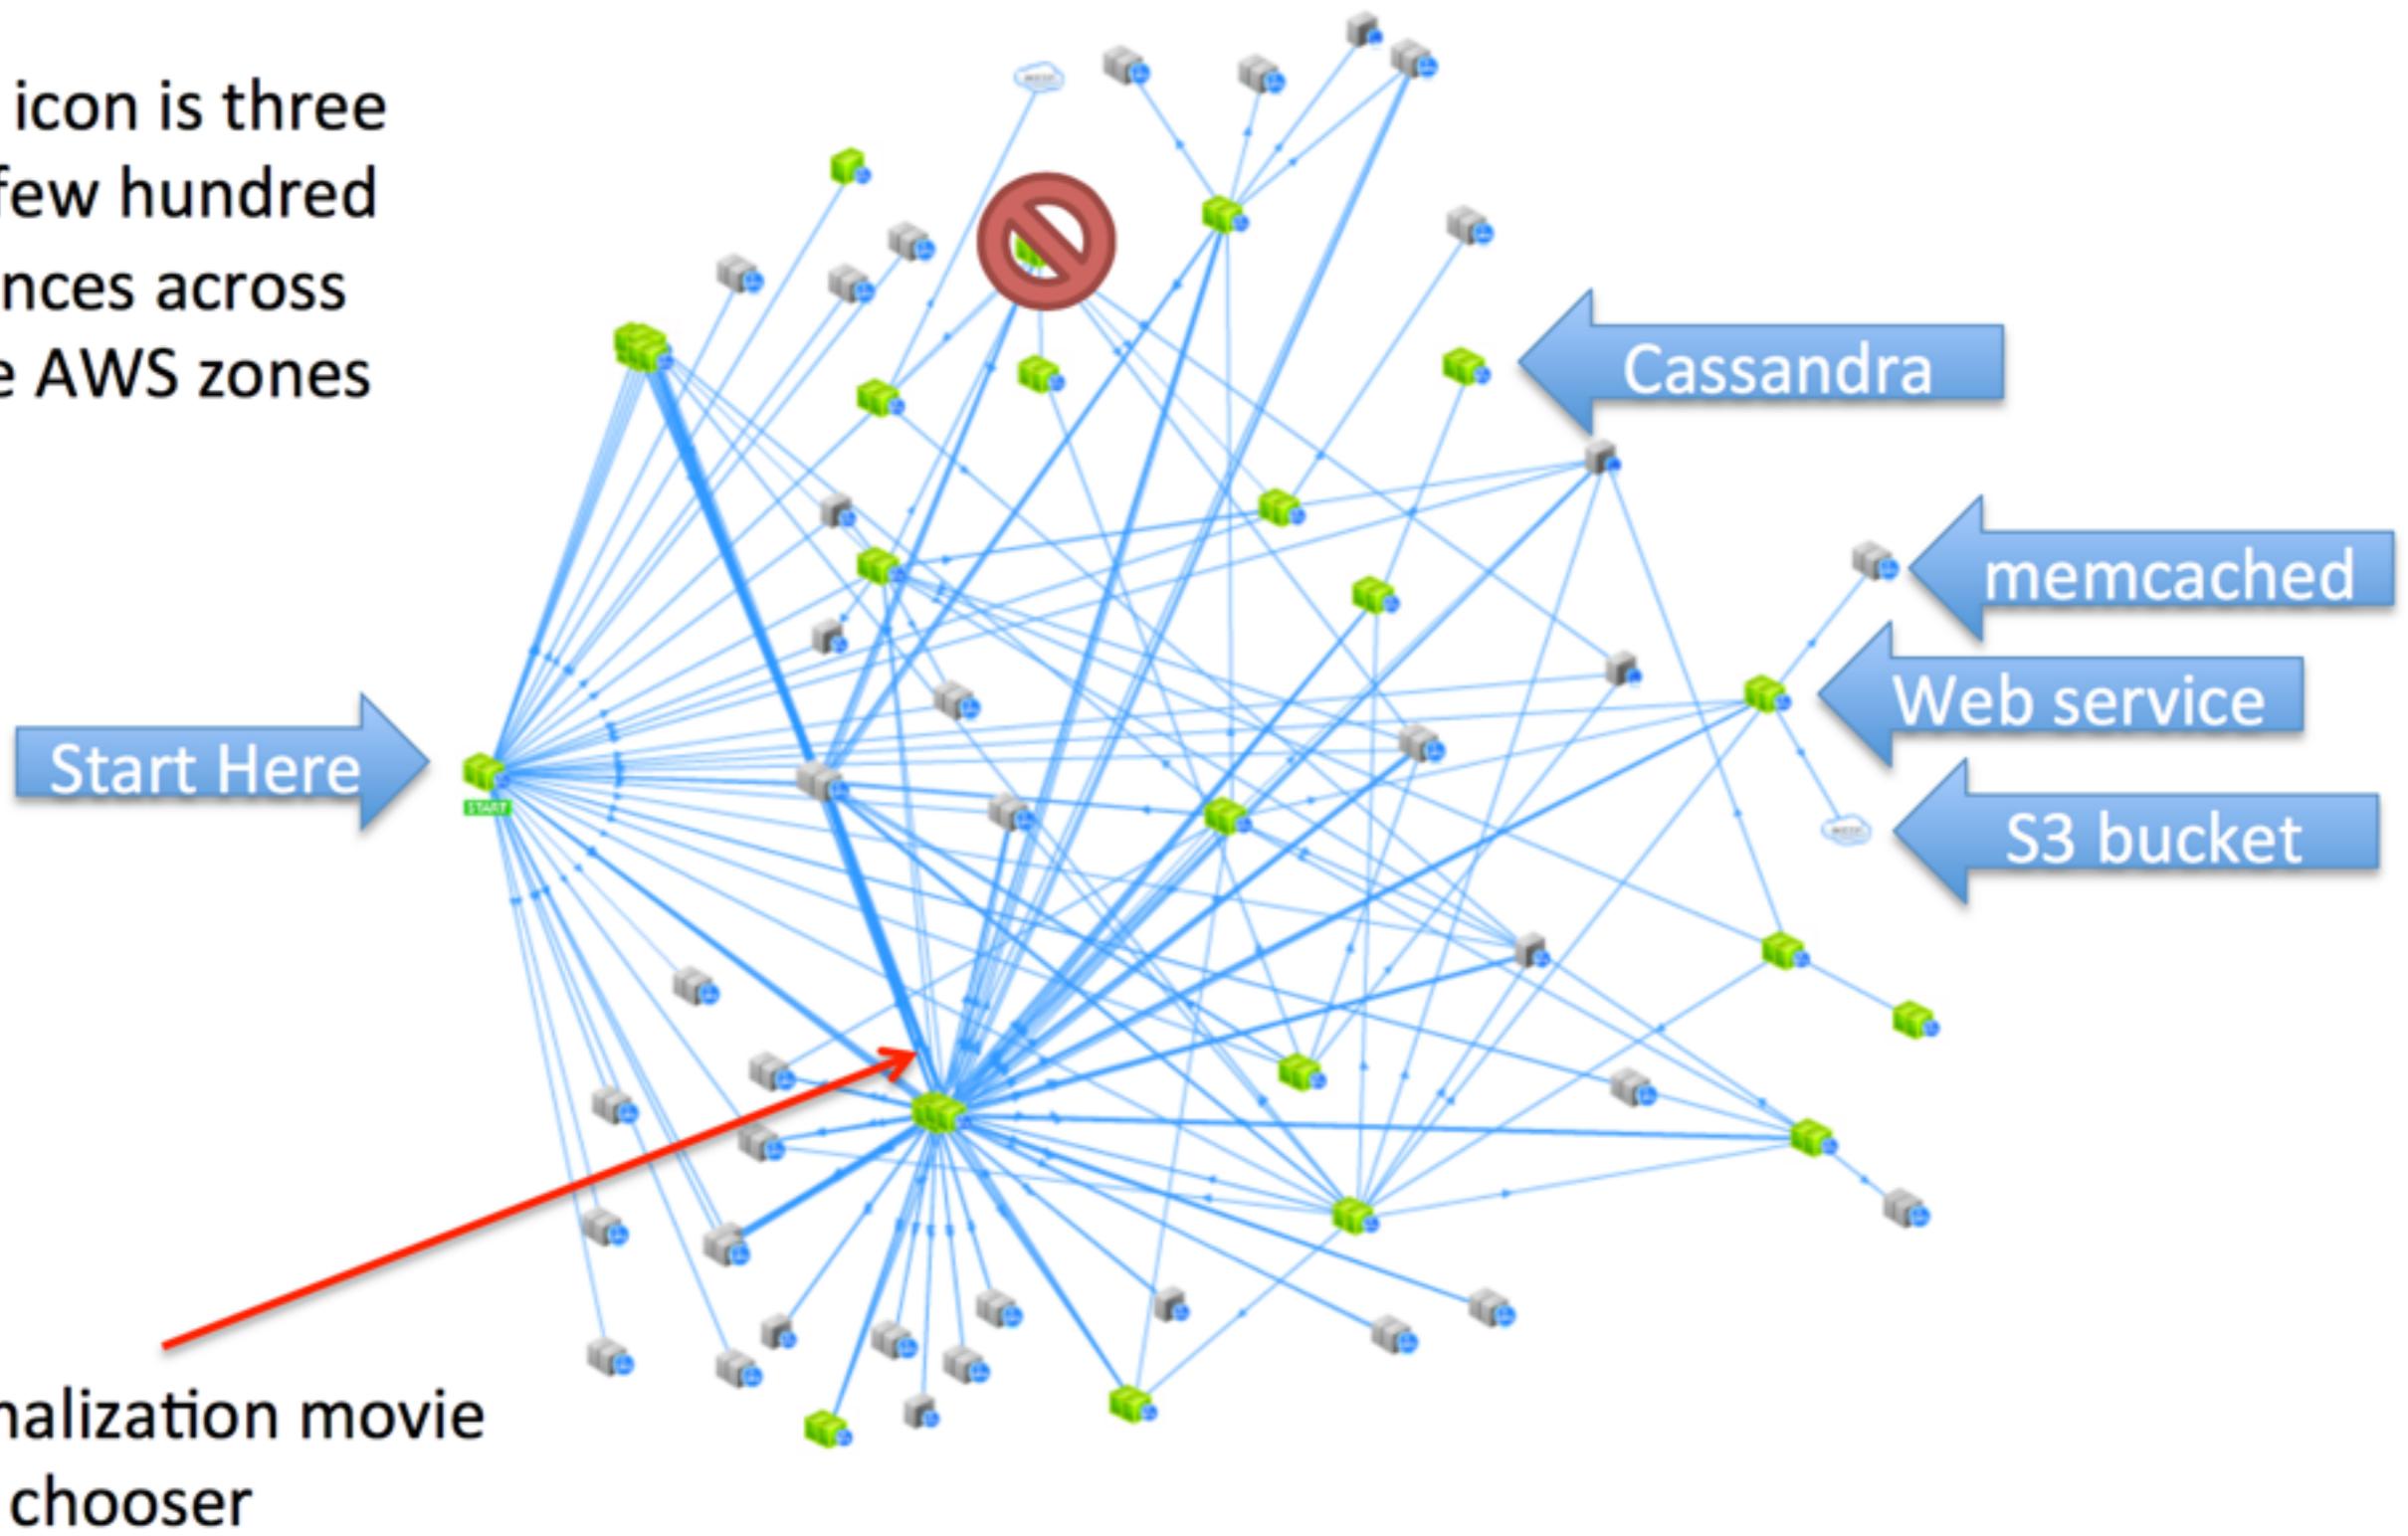
for more, see
[PlanetCassandra.org/
companies/](http://PlanetCassandra.org/companies/)

Netflix case study

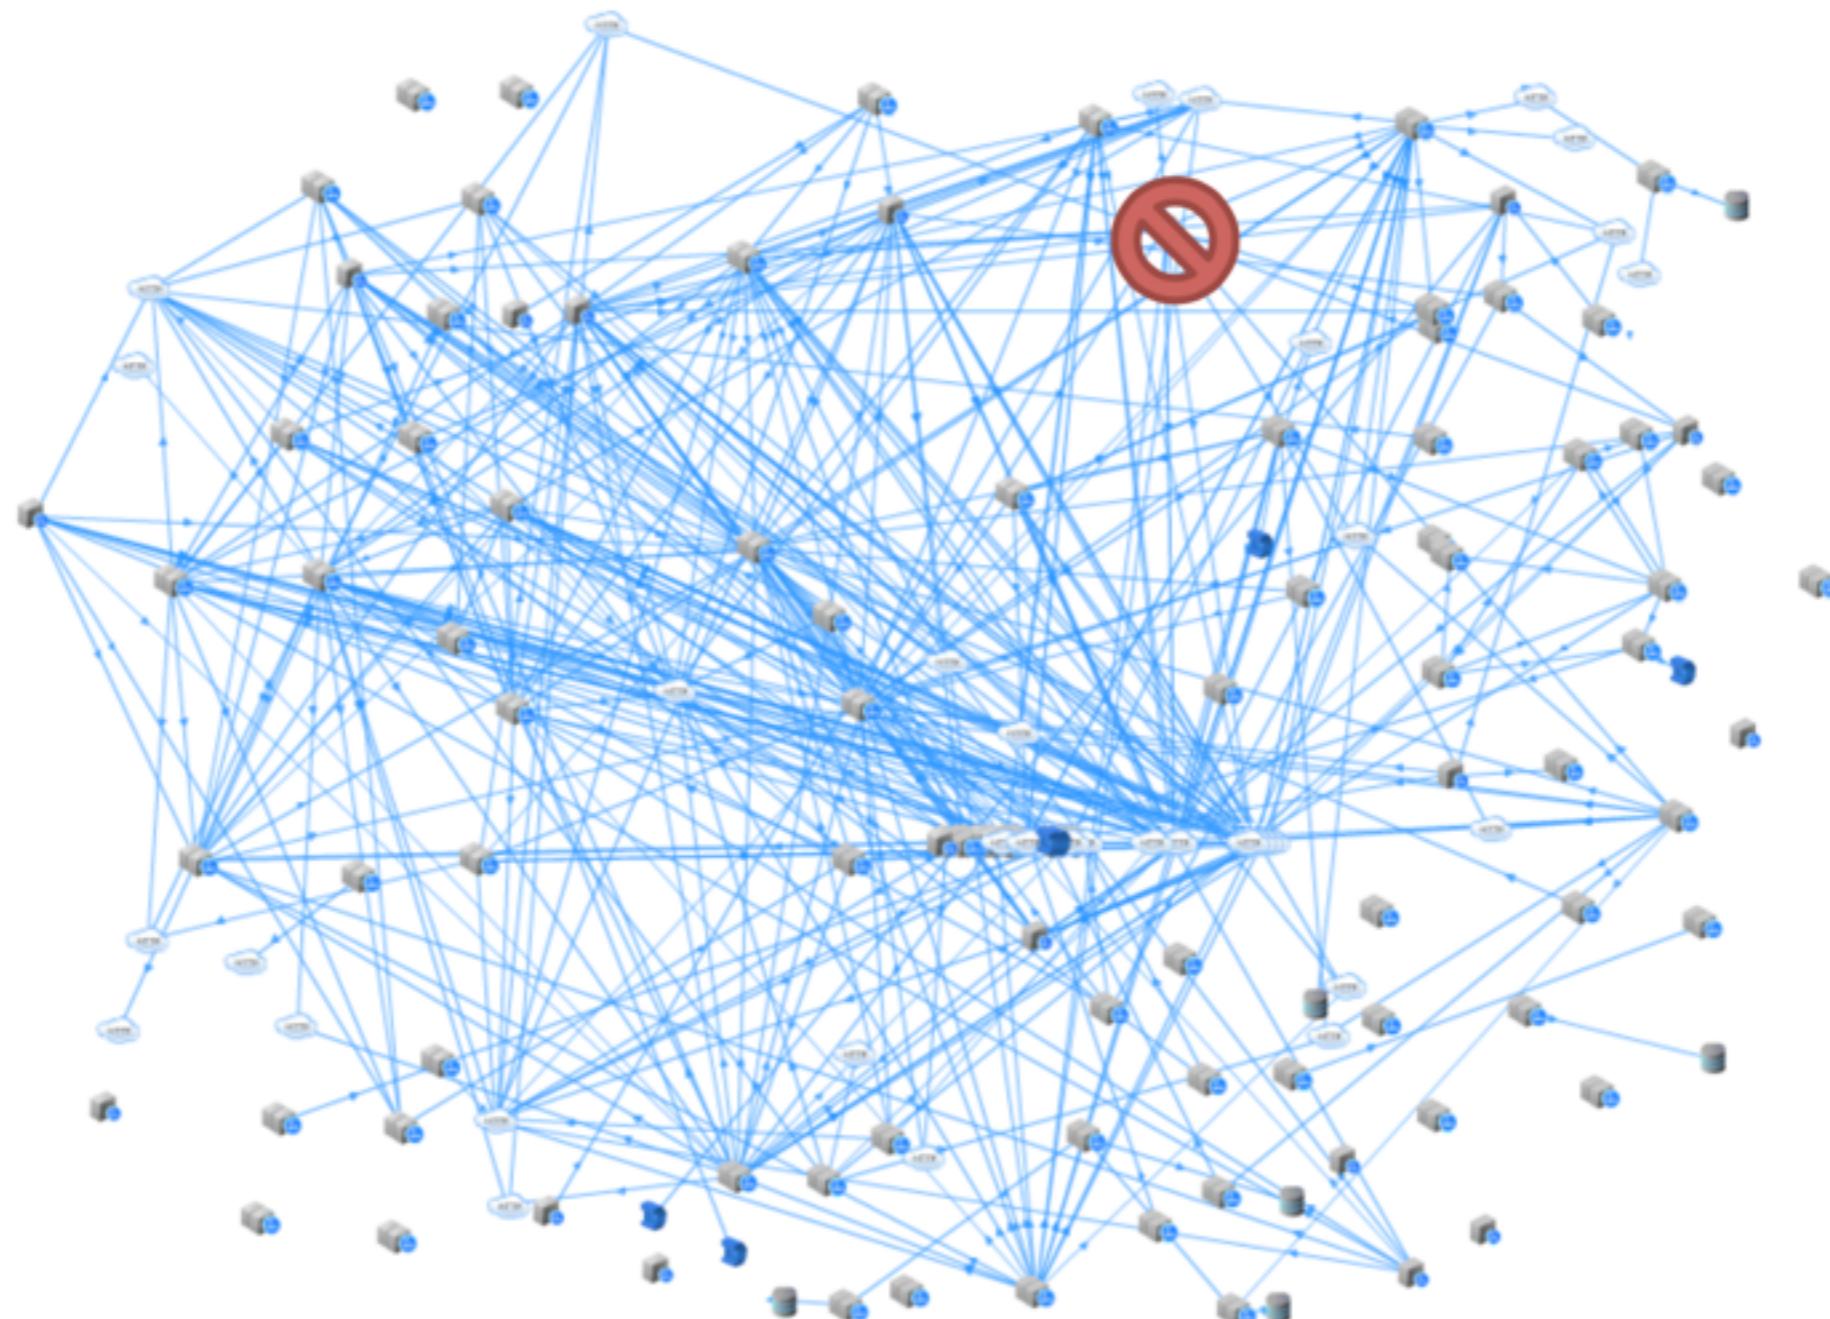


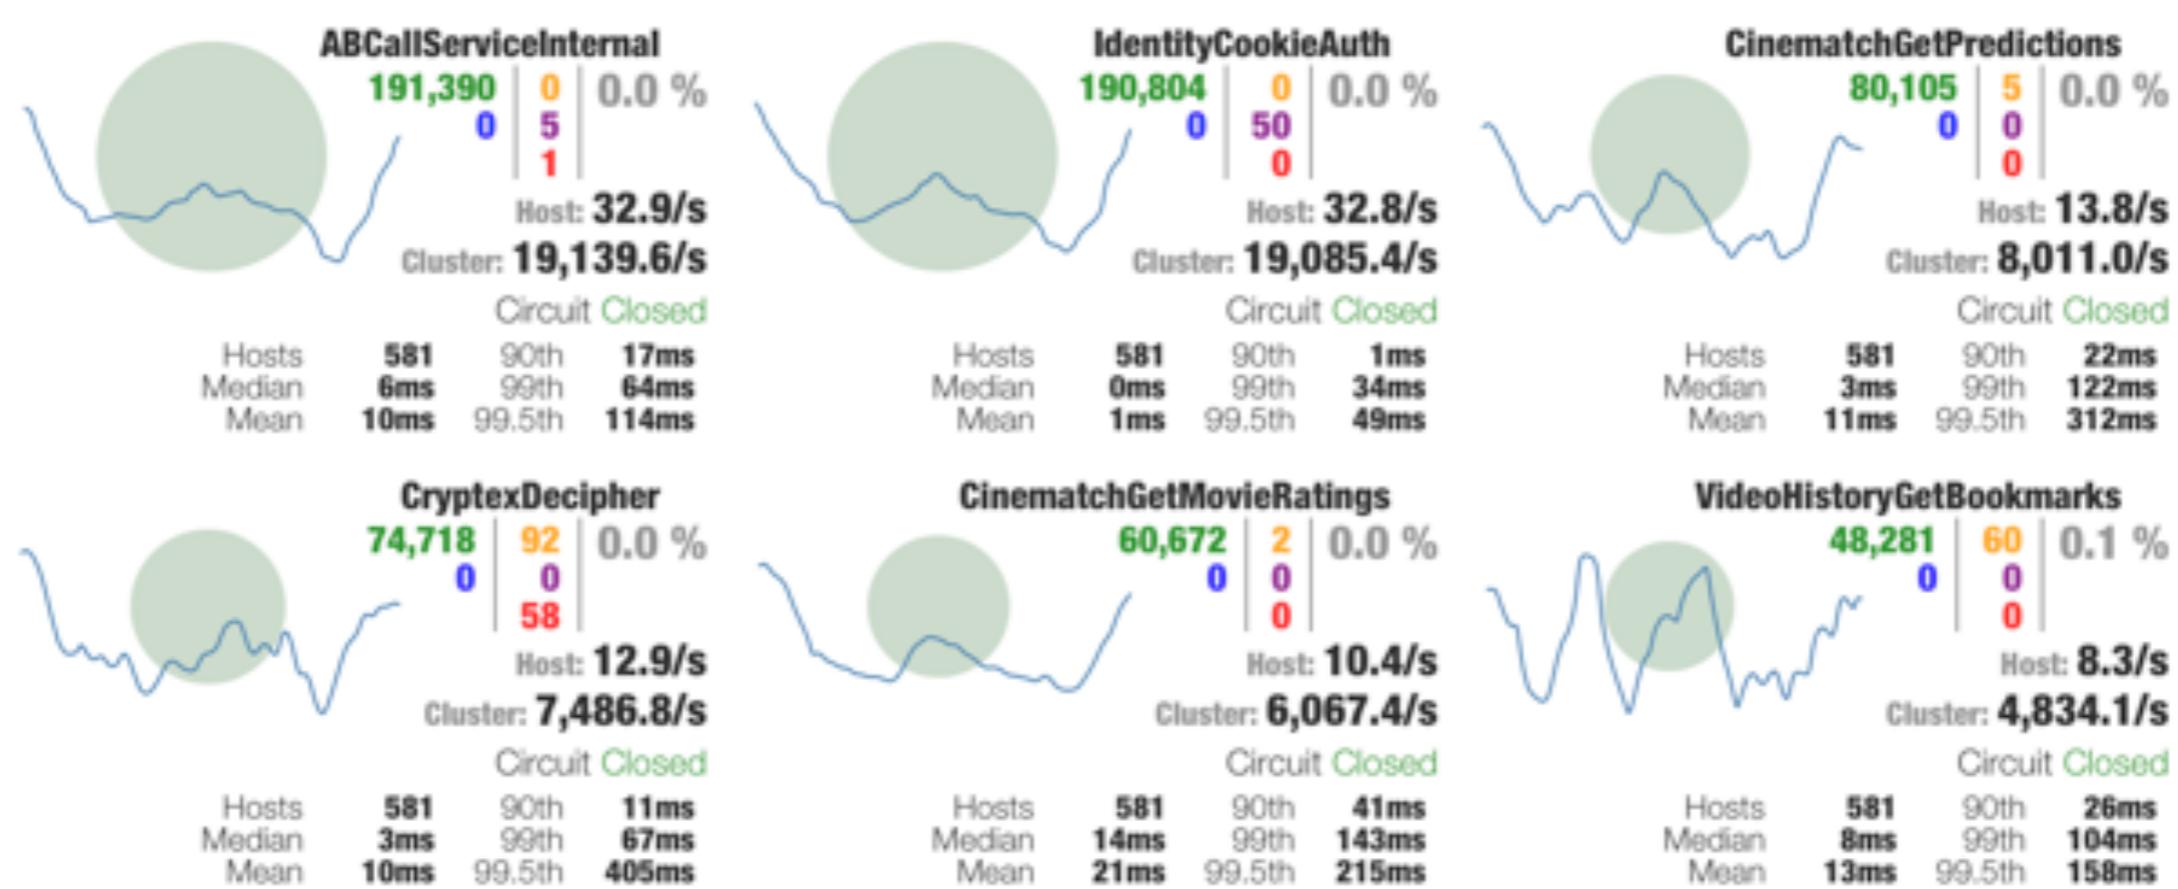


Each icon is three
to a few hundred
instances across
three AWS zones

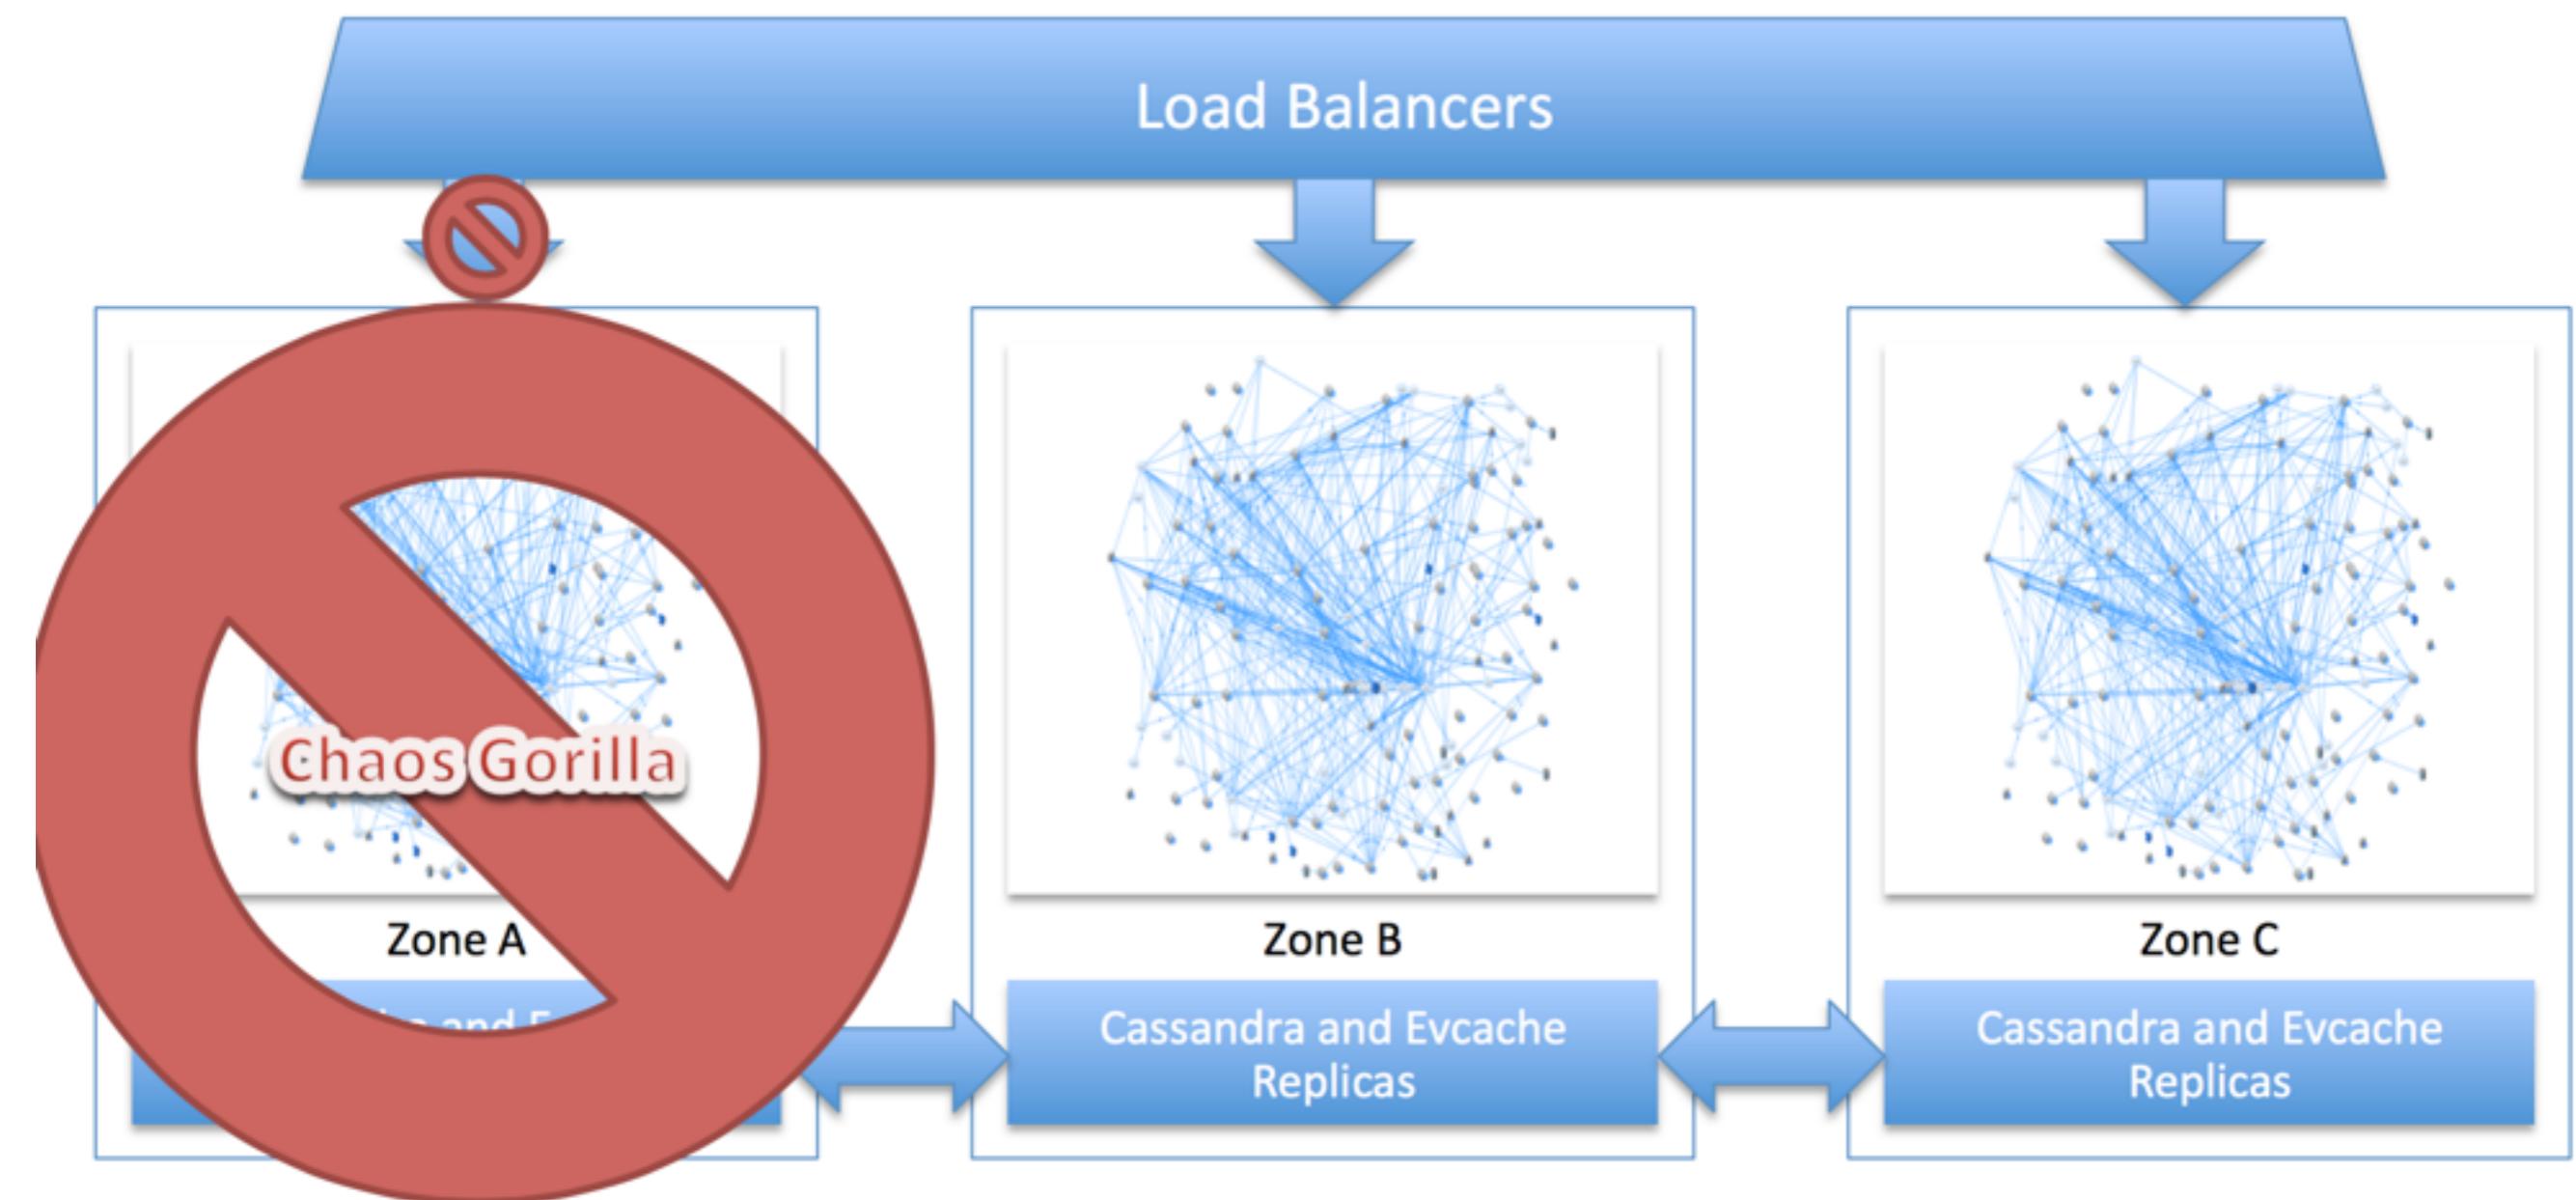


Component microservices

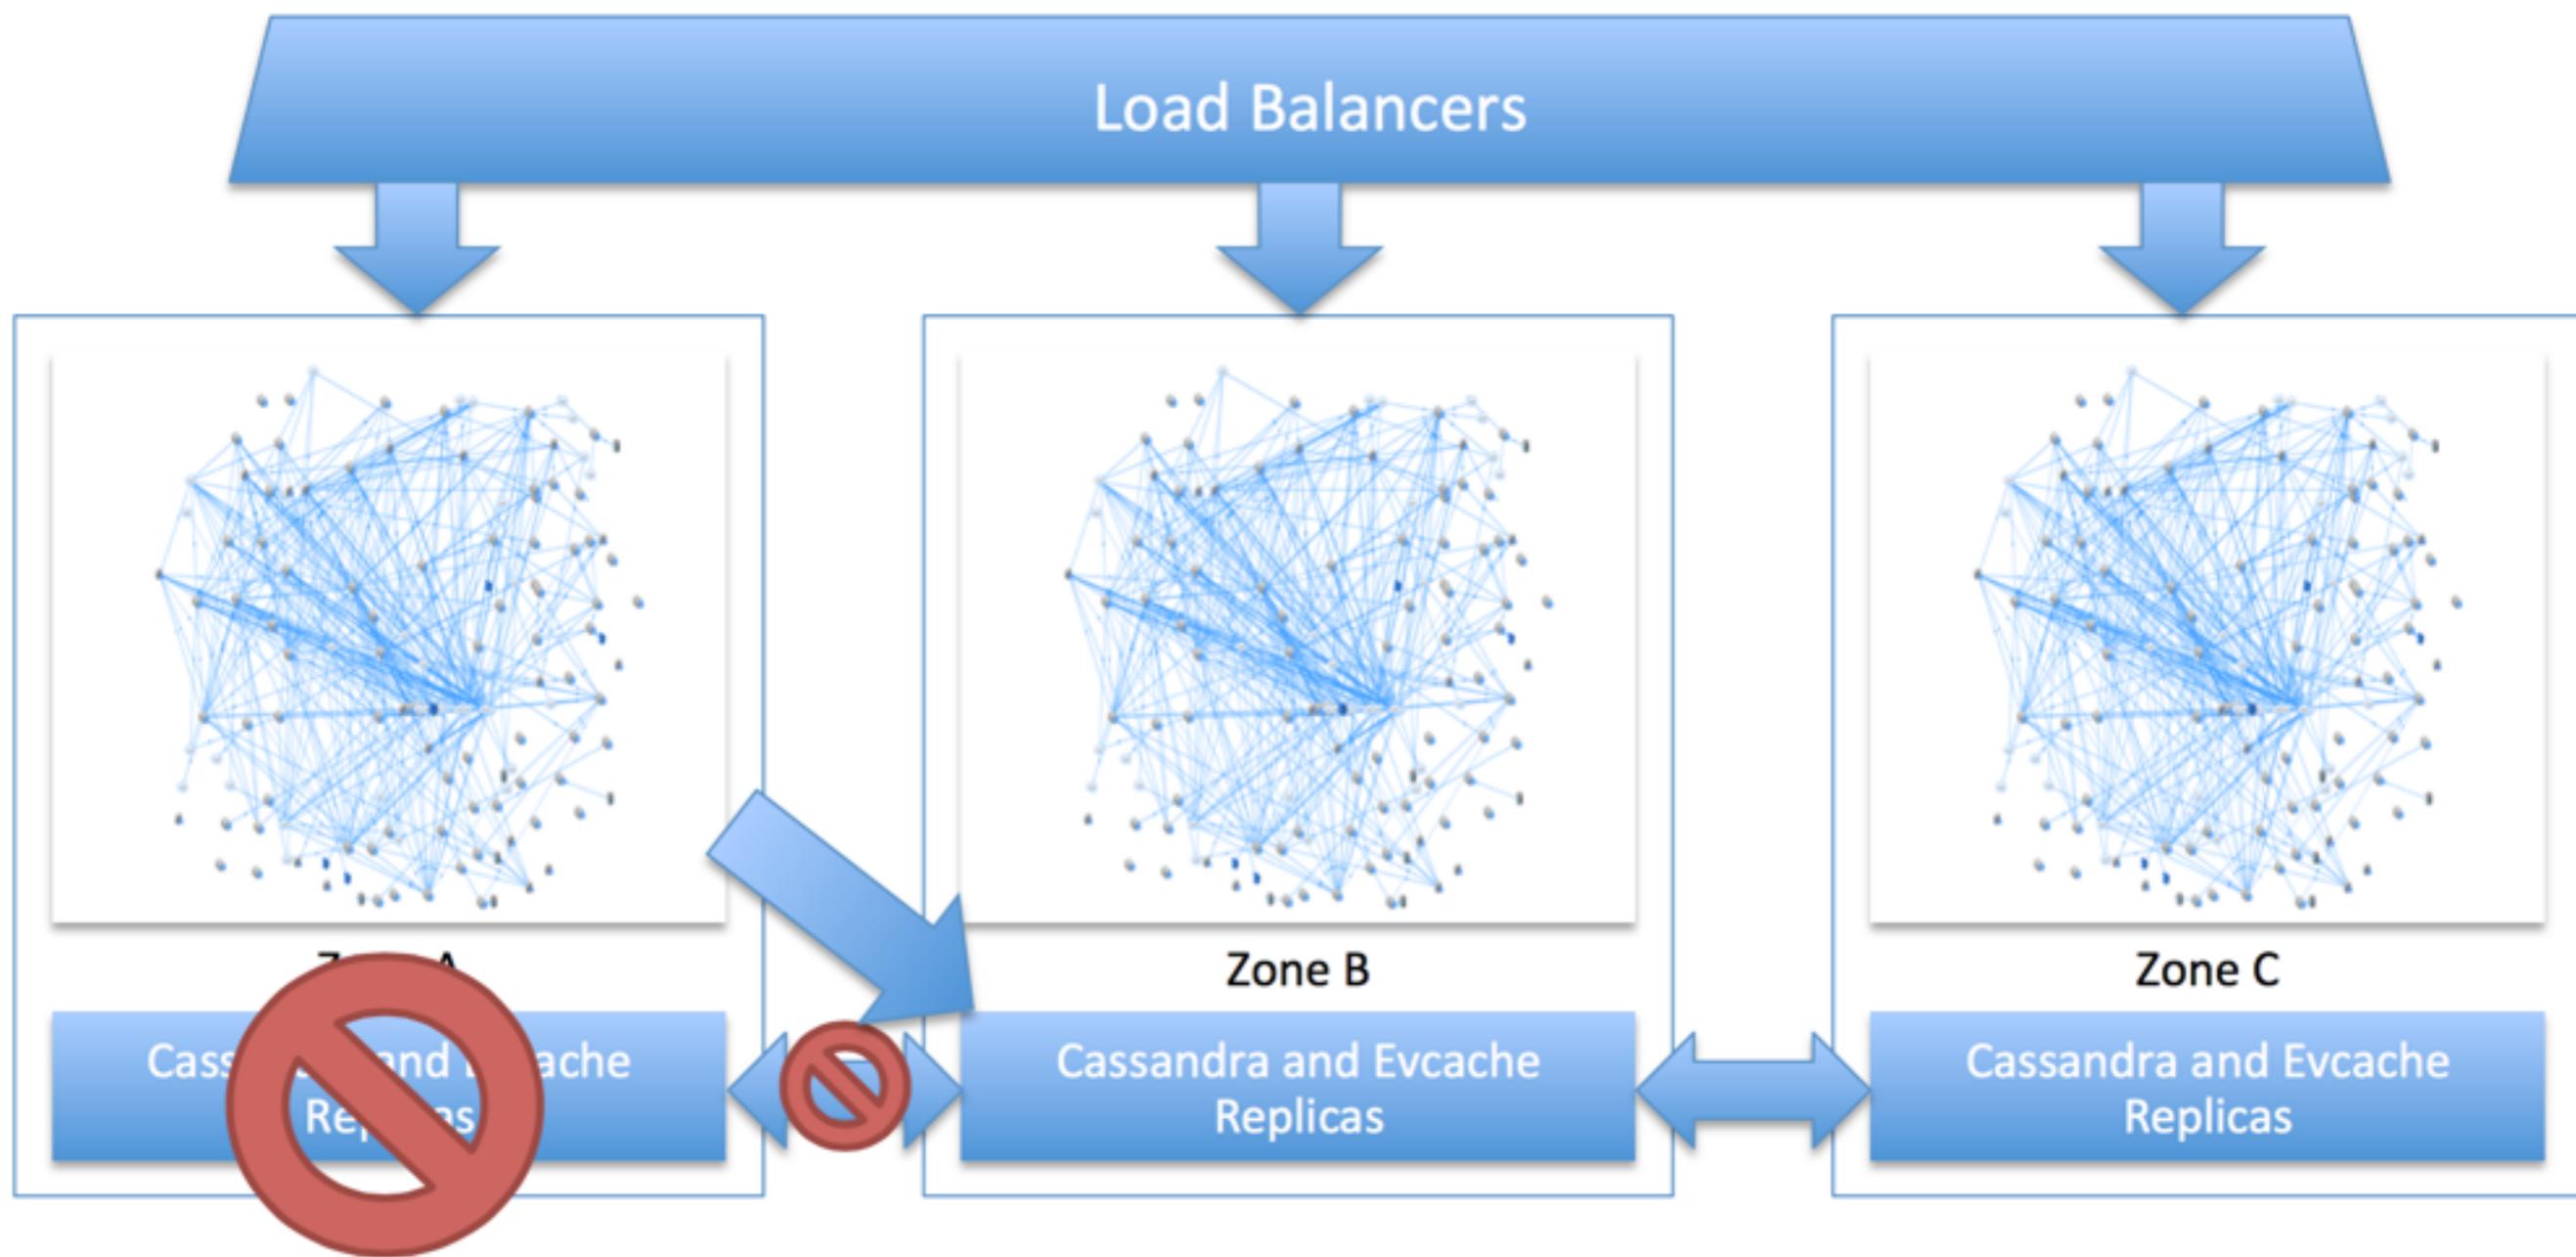




Chaos Gorilla



Cassandra maintenance

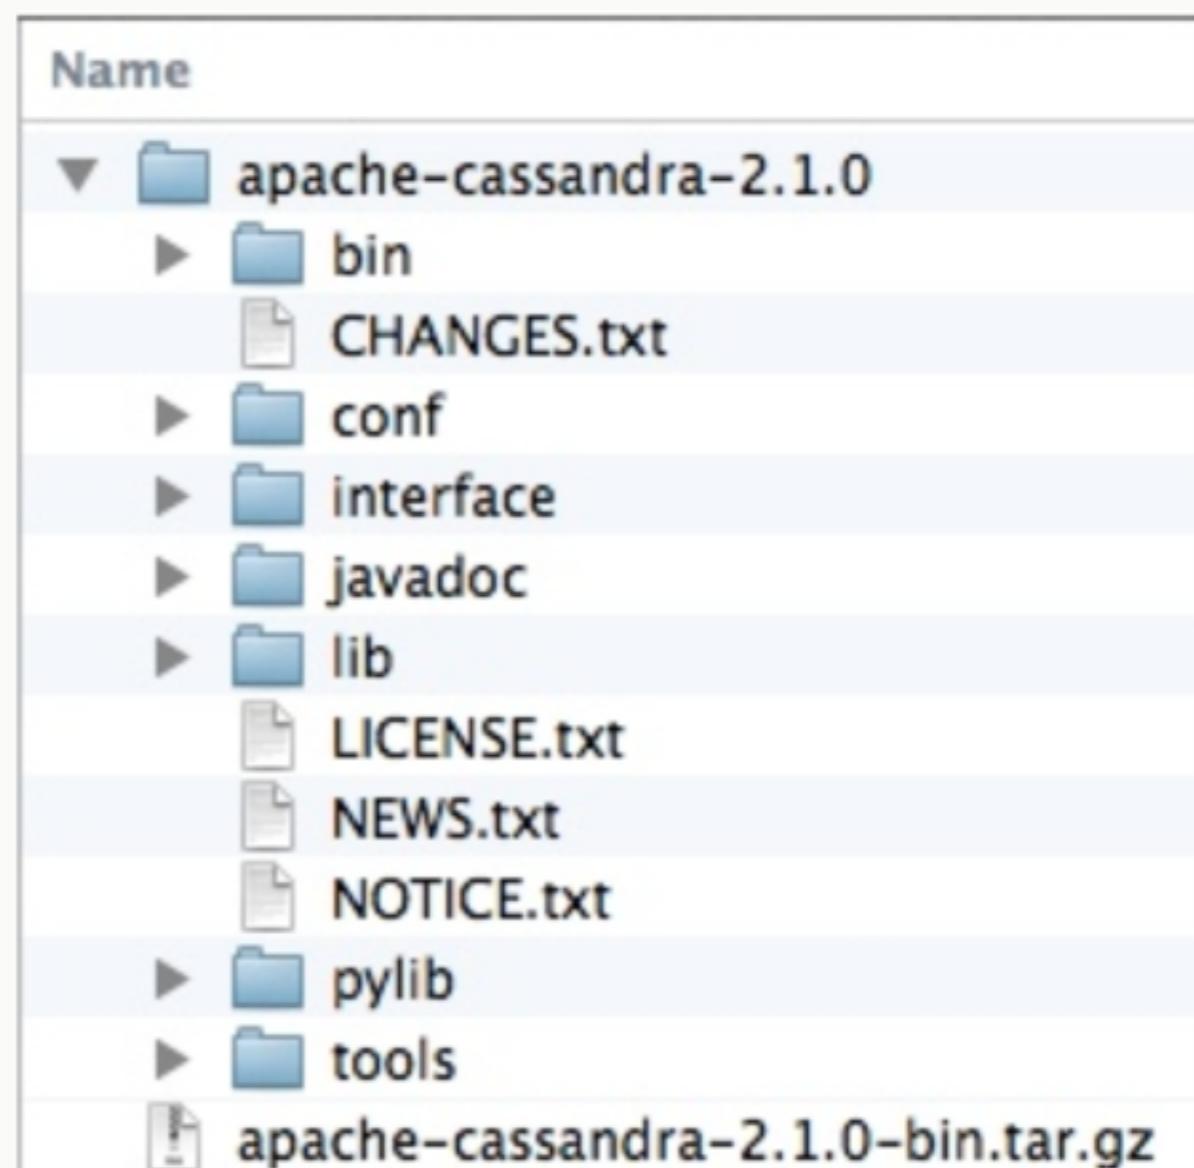


Cassandra

- 10x more read throughput
- 8x faster read latency (up to 100x faster)
- 8x more write throughput
- 10x slower write latency (with the default configuration; that is, no write durability for HBase)
- 8x faster scan latency
- 4x more scan throughput

INSTALLATION

- Tarball installations create these folders in a single location:
 - **bin**: executables (*cassandra*, *cqlsh*, *nodetool*, etc)
 - **conf**: Configuration files - *cassandra.yaml*
 - **javadoc**: C* source documentation
 - **lib**: Library dependencies (jars)
 - **pylib**: Python libraries (e.g. for *cqlsh*, which is written in Python)
 - **tools**: Additional tools (e.g. *cassandra-stress* which stresses a C* cluster)



configuration files

- **cassandra.yaml**

- One file per node--must agree with other node's files
- Parameters defined throughout

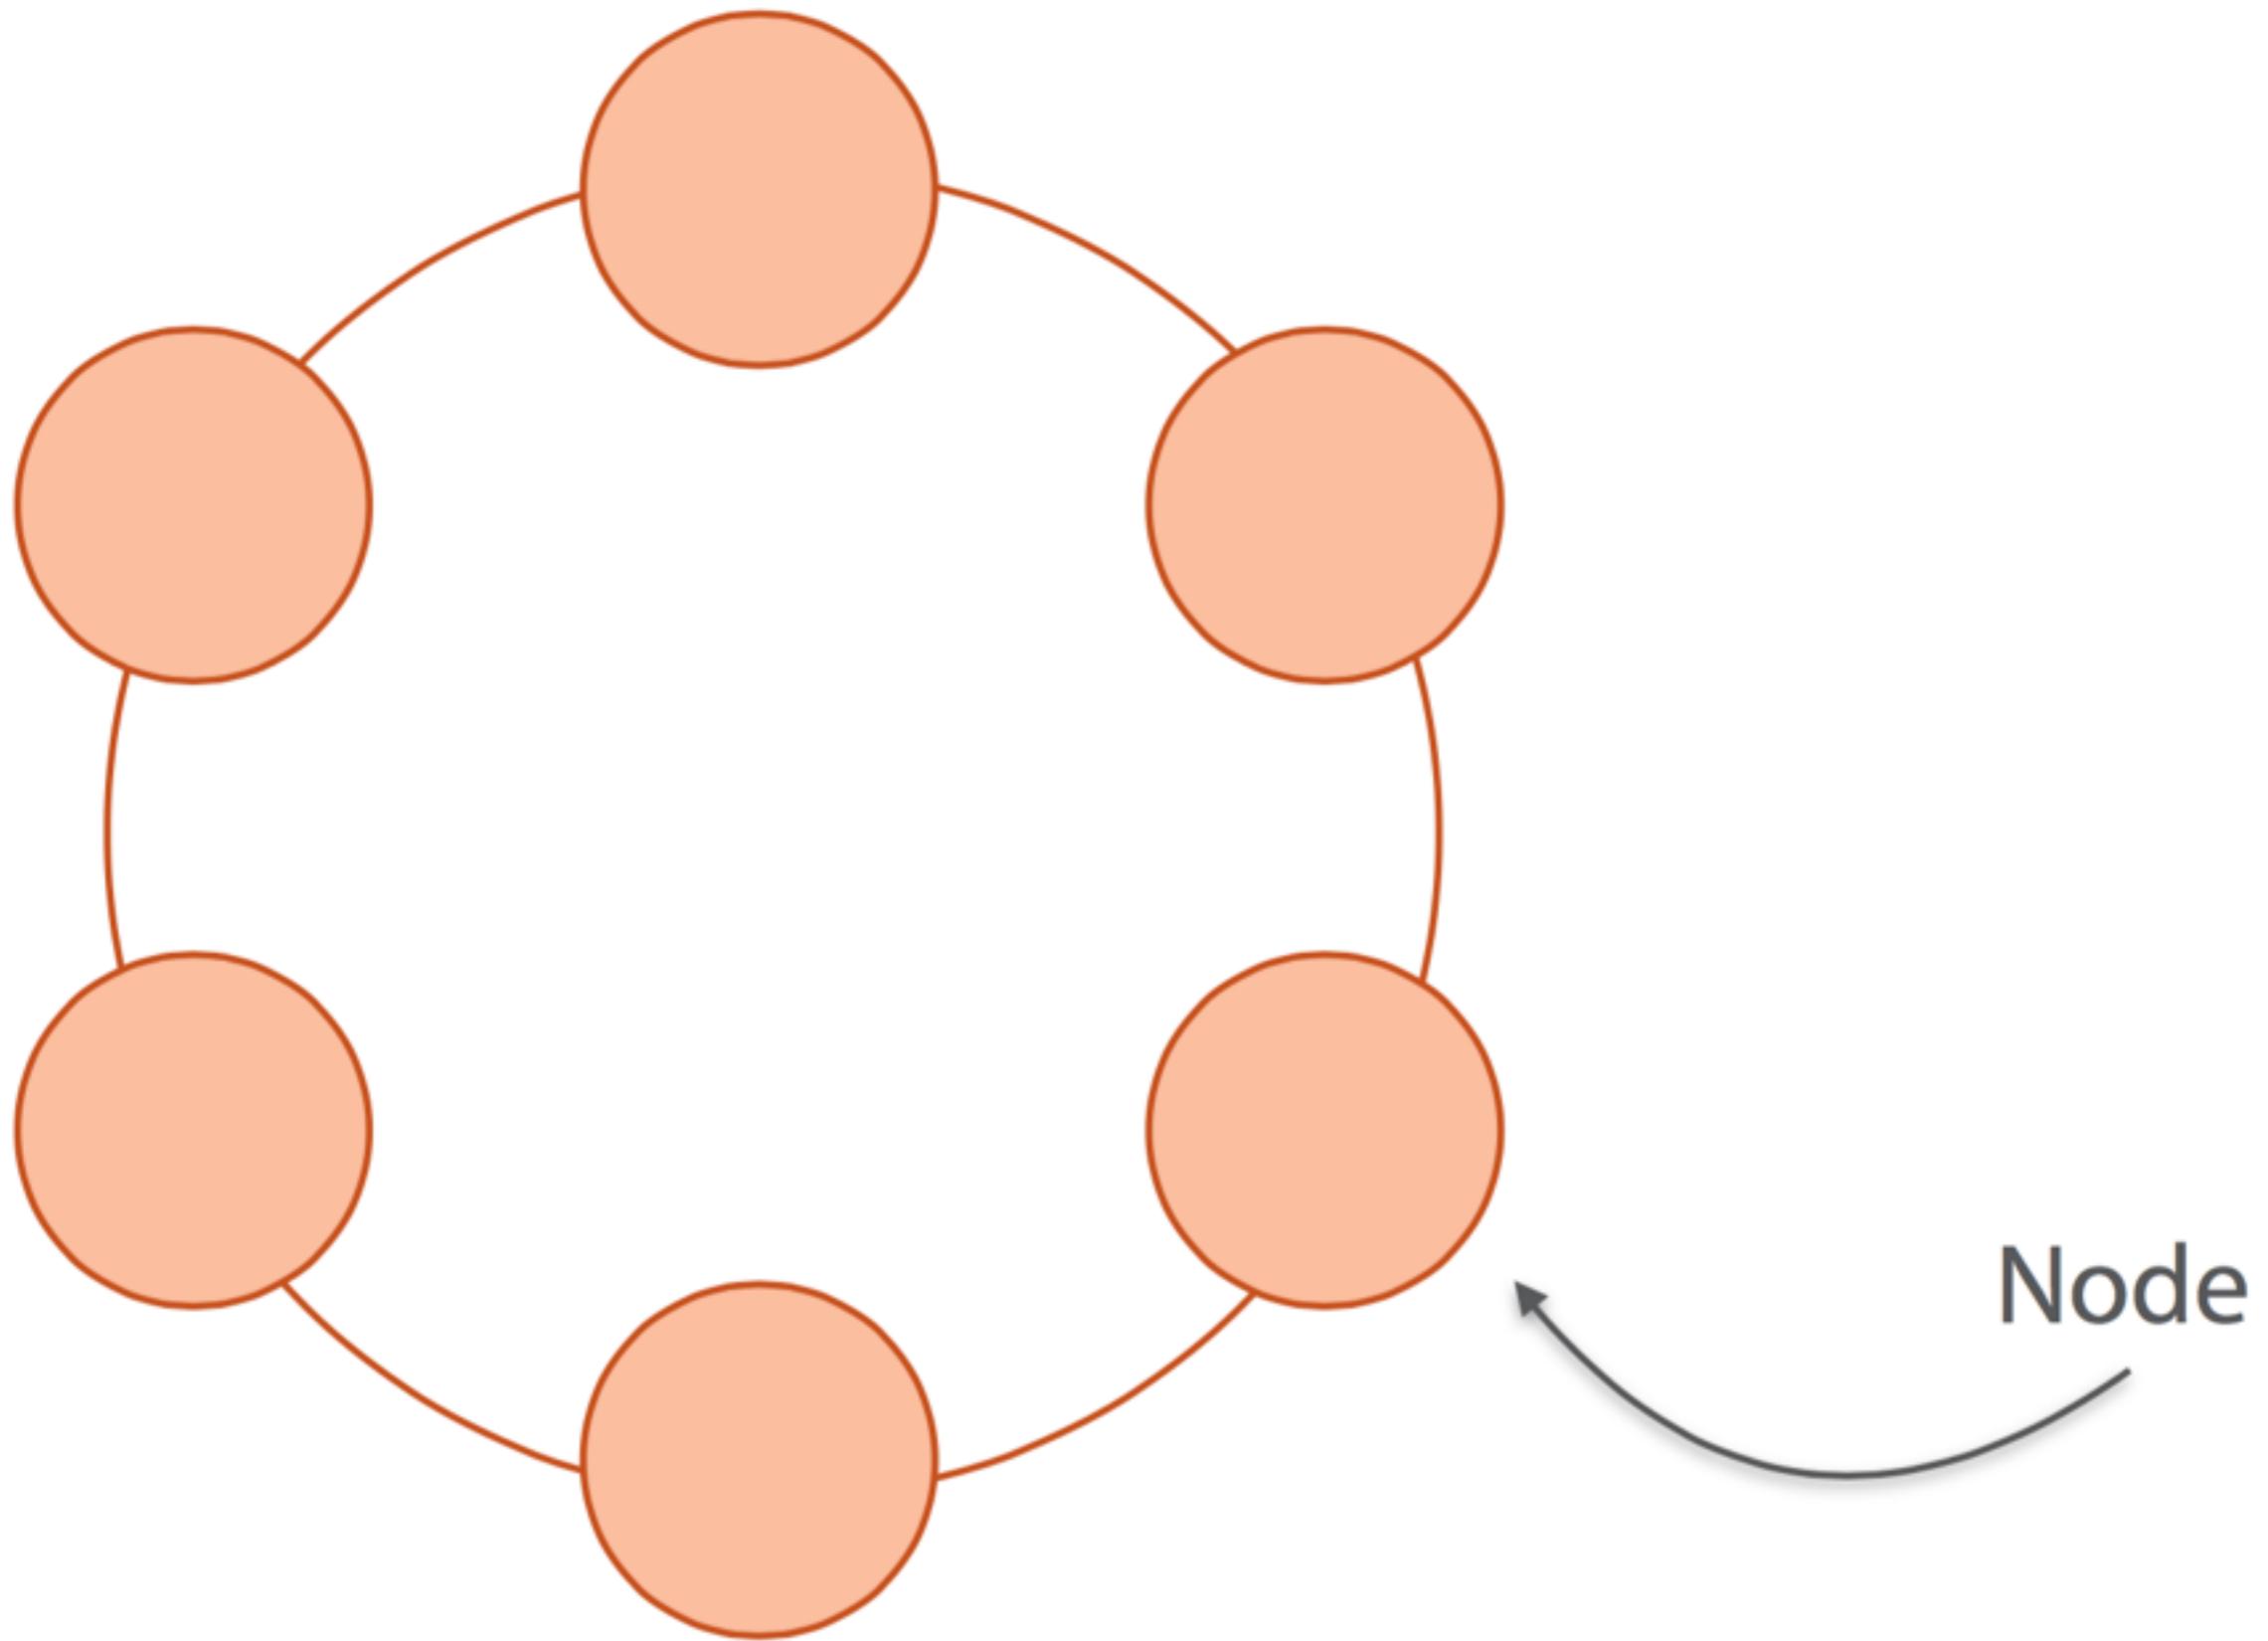
- **cassandra-env.sh**

- Memory settings
- JMX settings

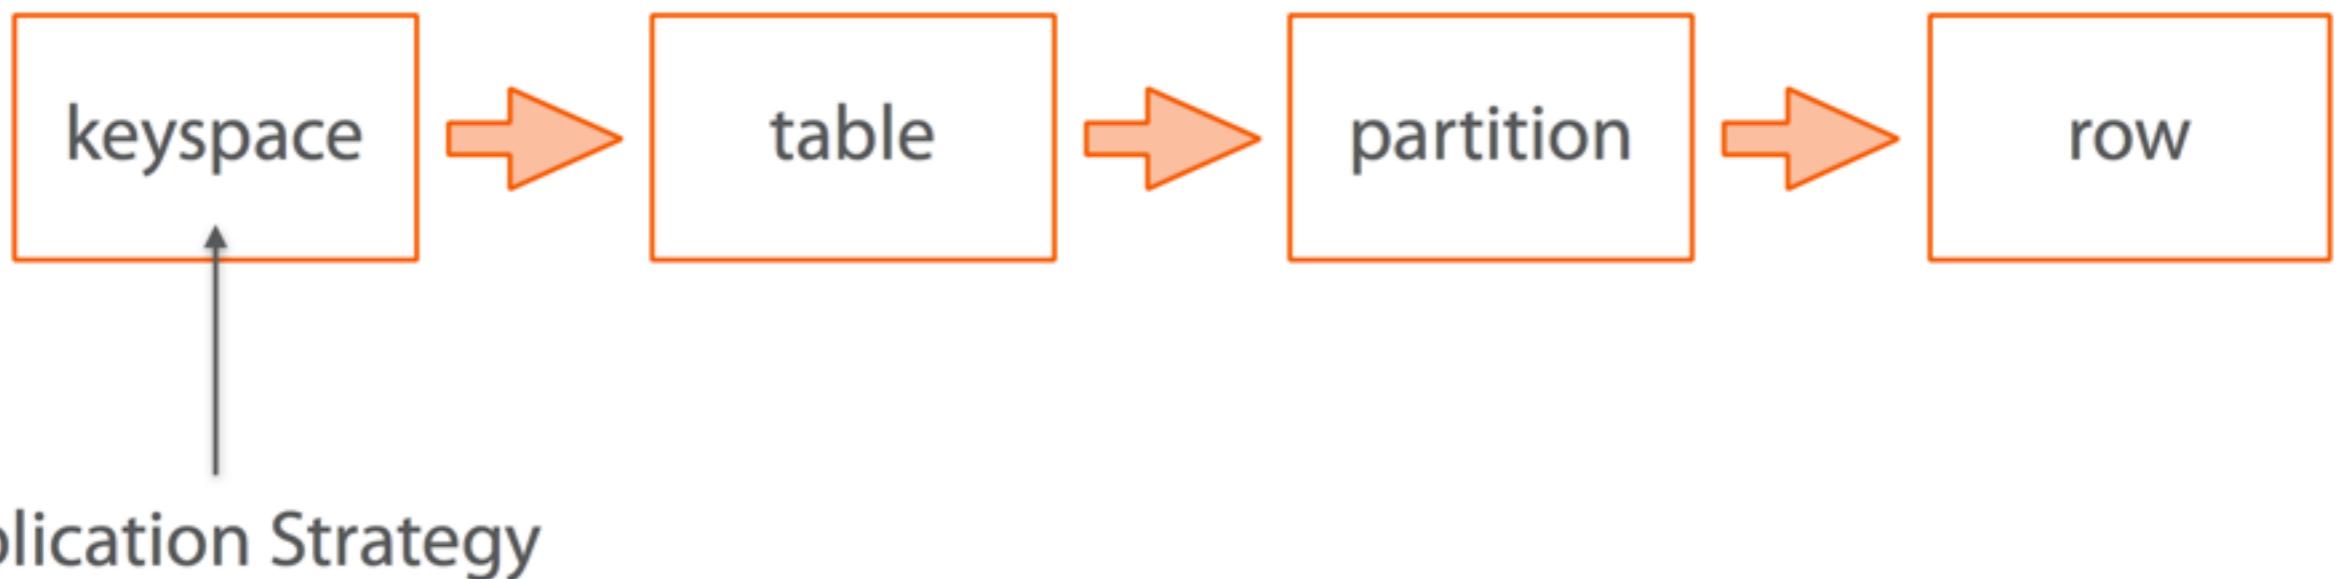
- **logback.xml**

- System log location
- Logging level

Name
apache-cassandra-2.1.0
bin
CHANGES.txt
conf
cassandra-env.ps1
cassandra-env.sh
cassandra-rackdc.properties
cassandra-topology.properties
cassandra-topology.yaml
cassandra.yaml
commitlog_archiving.properties
cqlshrc.sample
logback-tools.xml
logback.xml
metrics-reporter-config-sample.yaml
README.txt
triggers
interface
javadoc

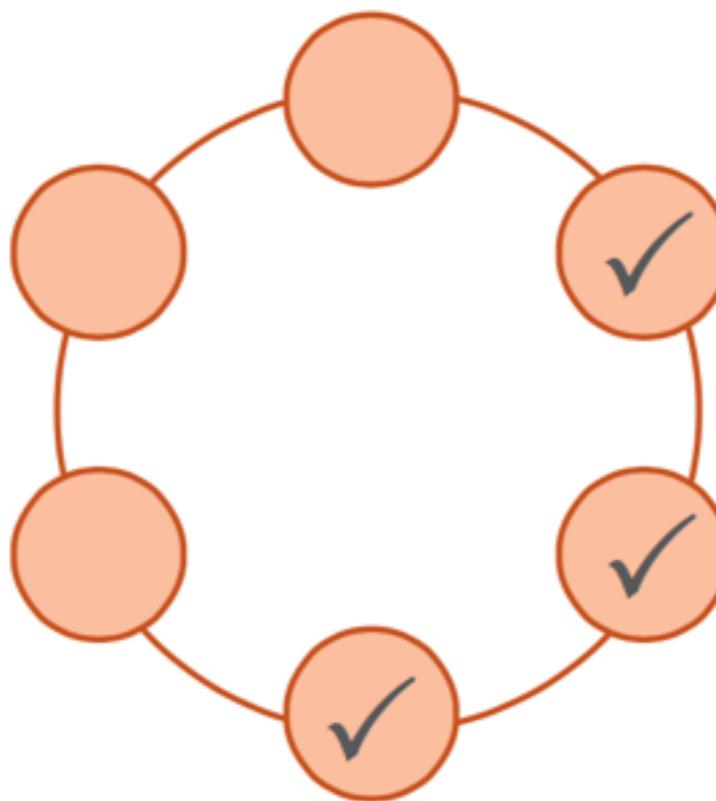


cassandra terminology



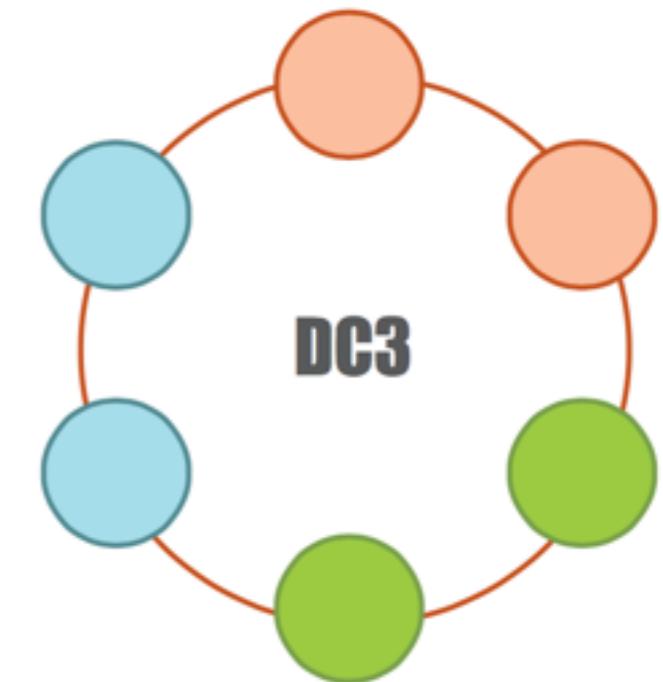
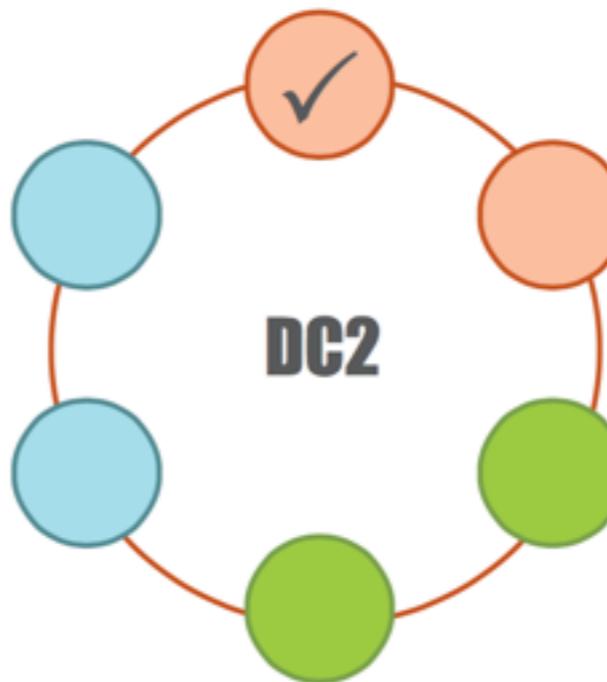
simple strategy

- CREATE KEYSPACE nfjs WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 3 };

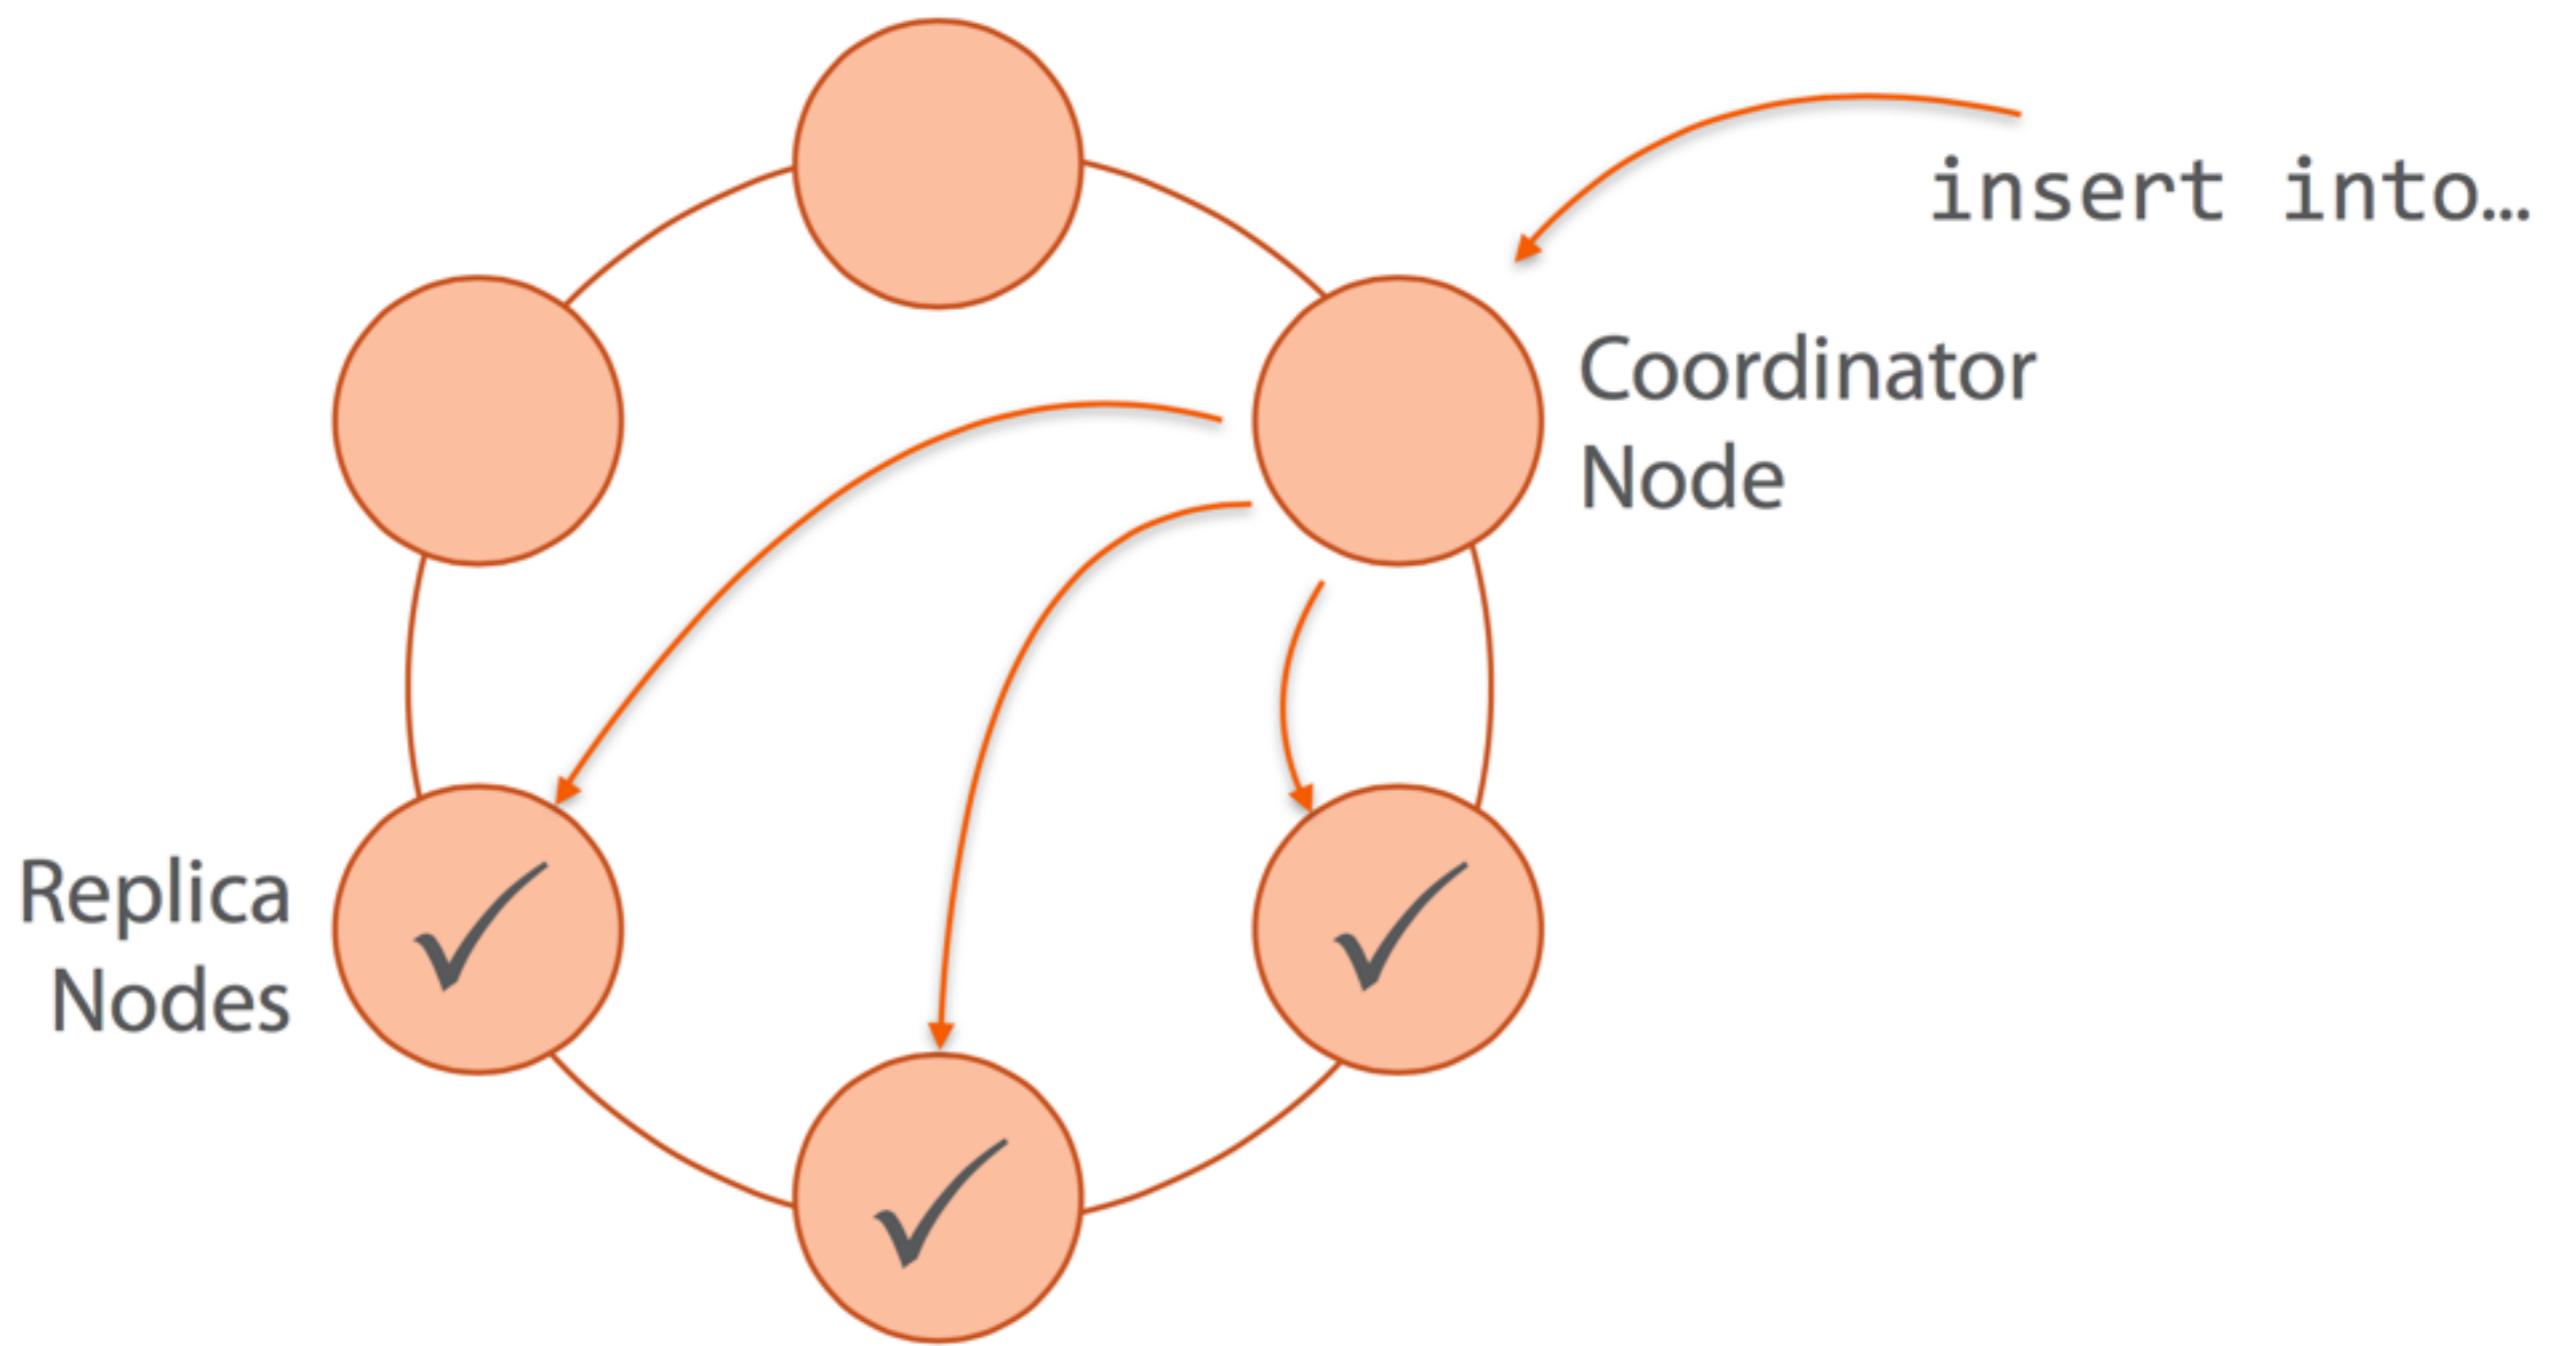


network topology strategy

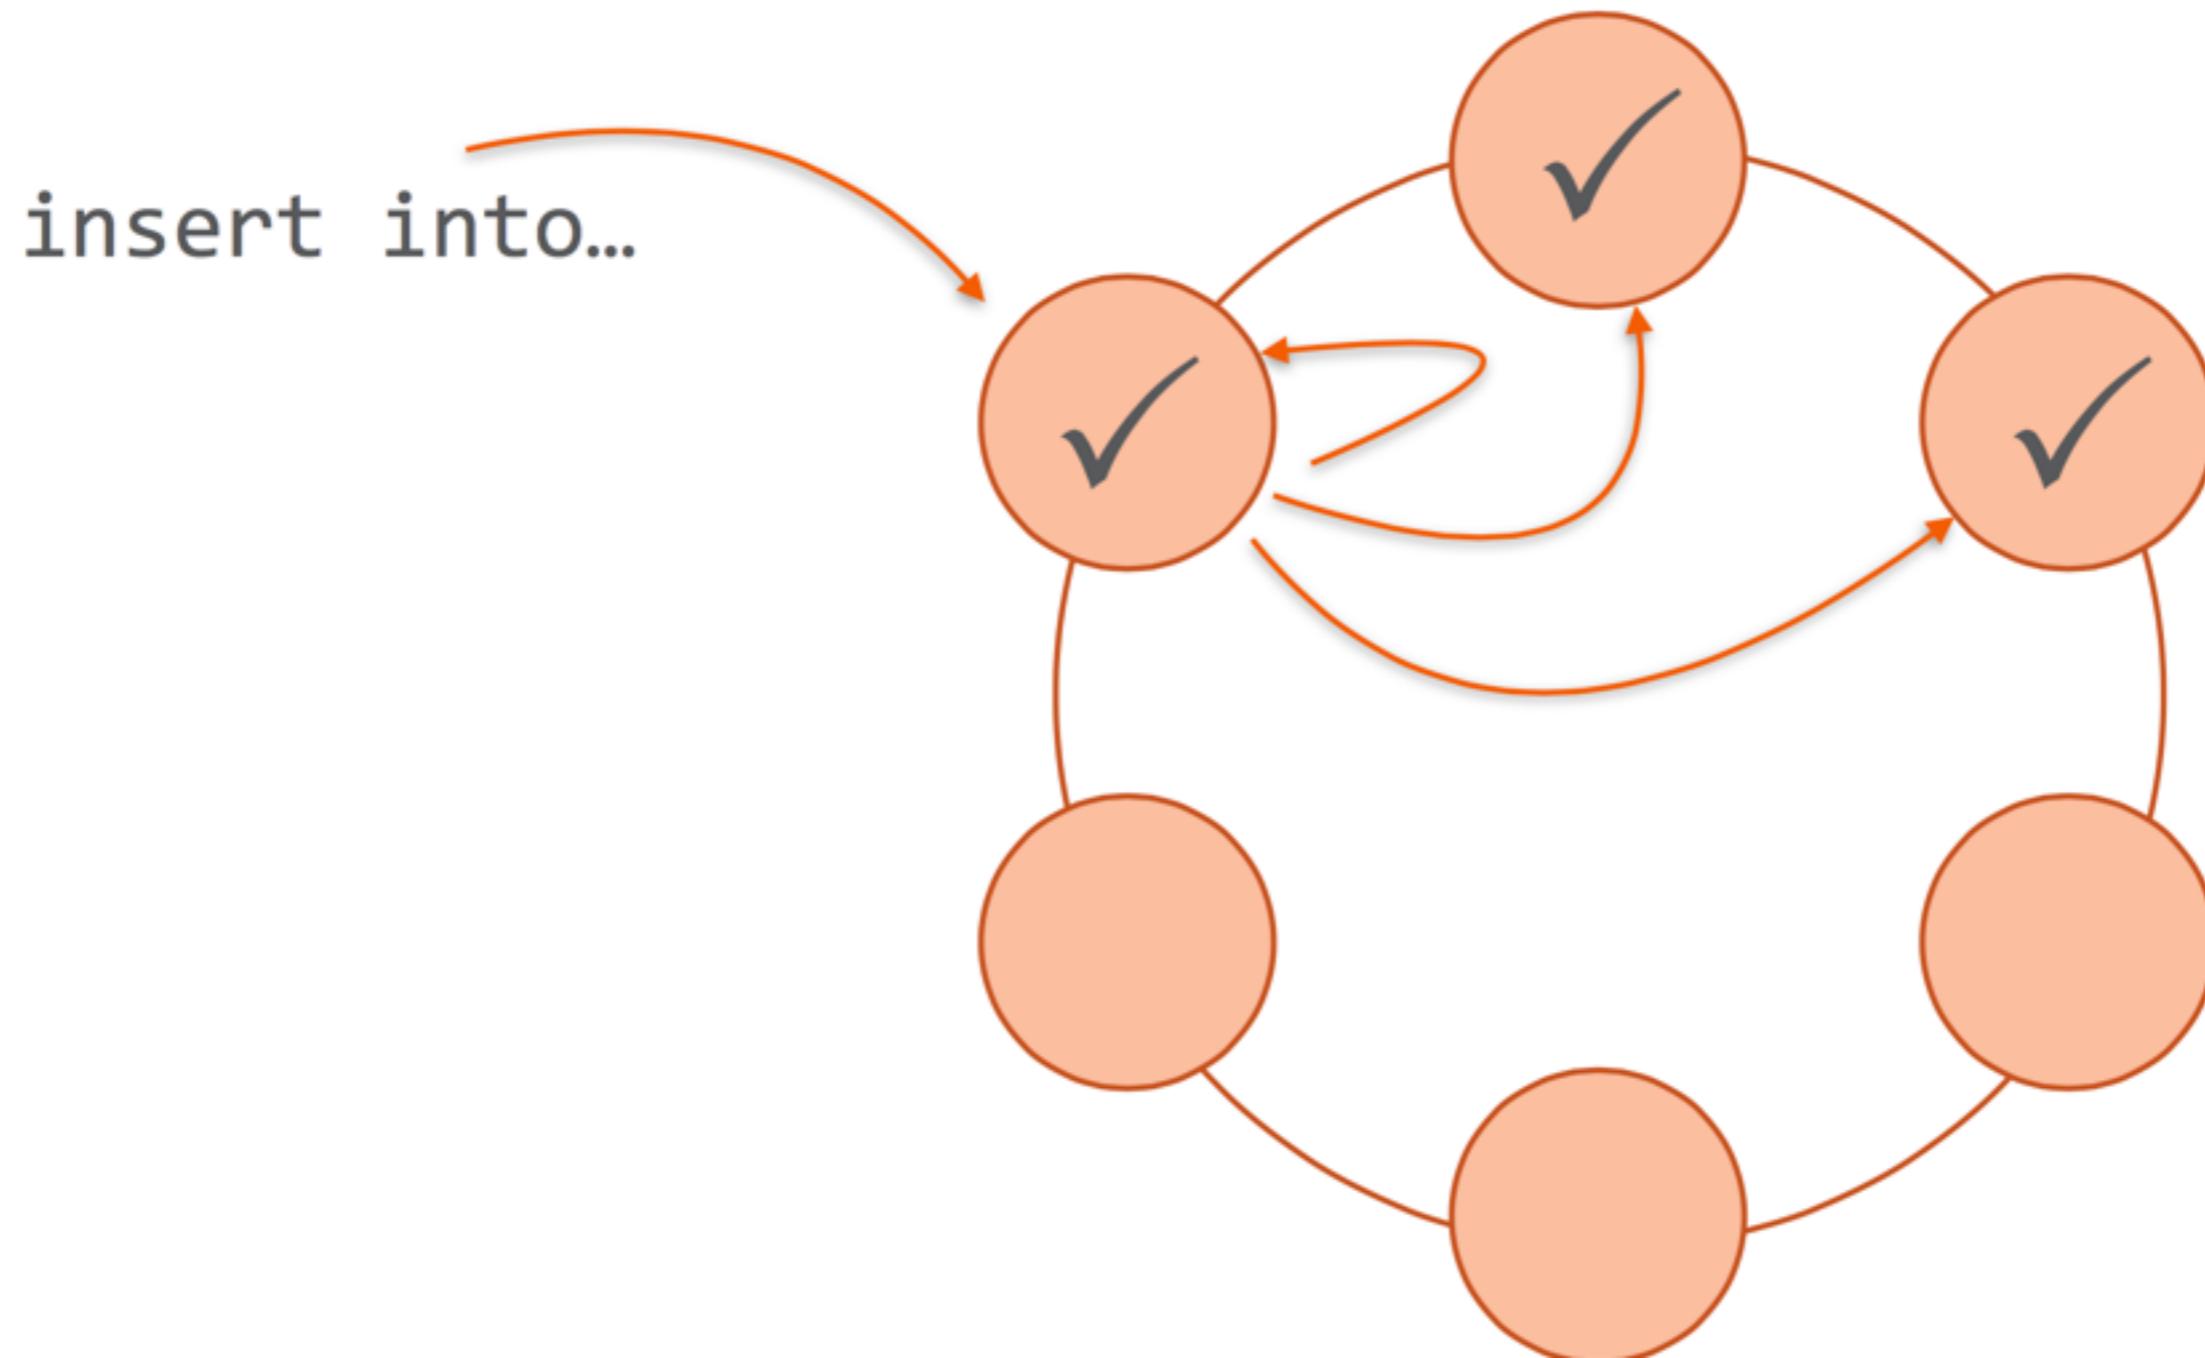
- create KEYSPACE nfjs WITH REPLICATION =
{'class': 'NetworkTopologyStrategy', 'DC1': 3, 'DC2':
1};



Reads and Writes in Cassandra



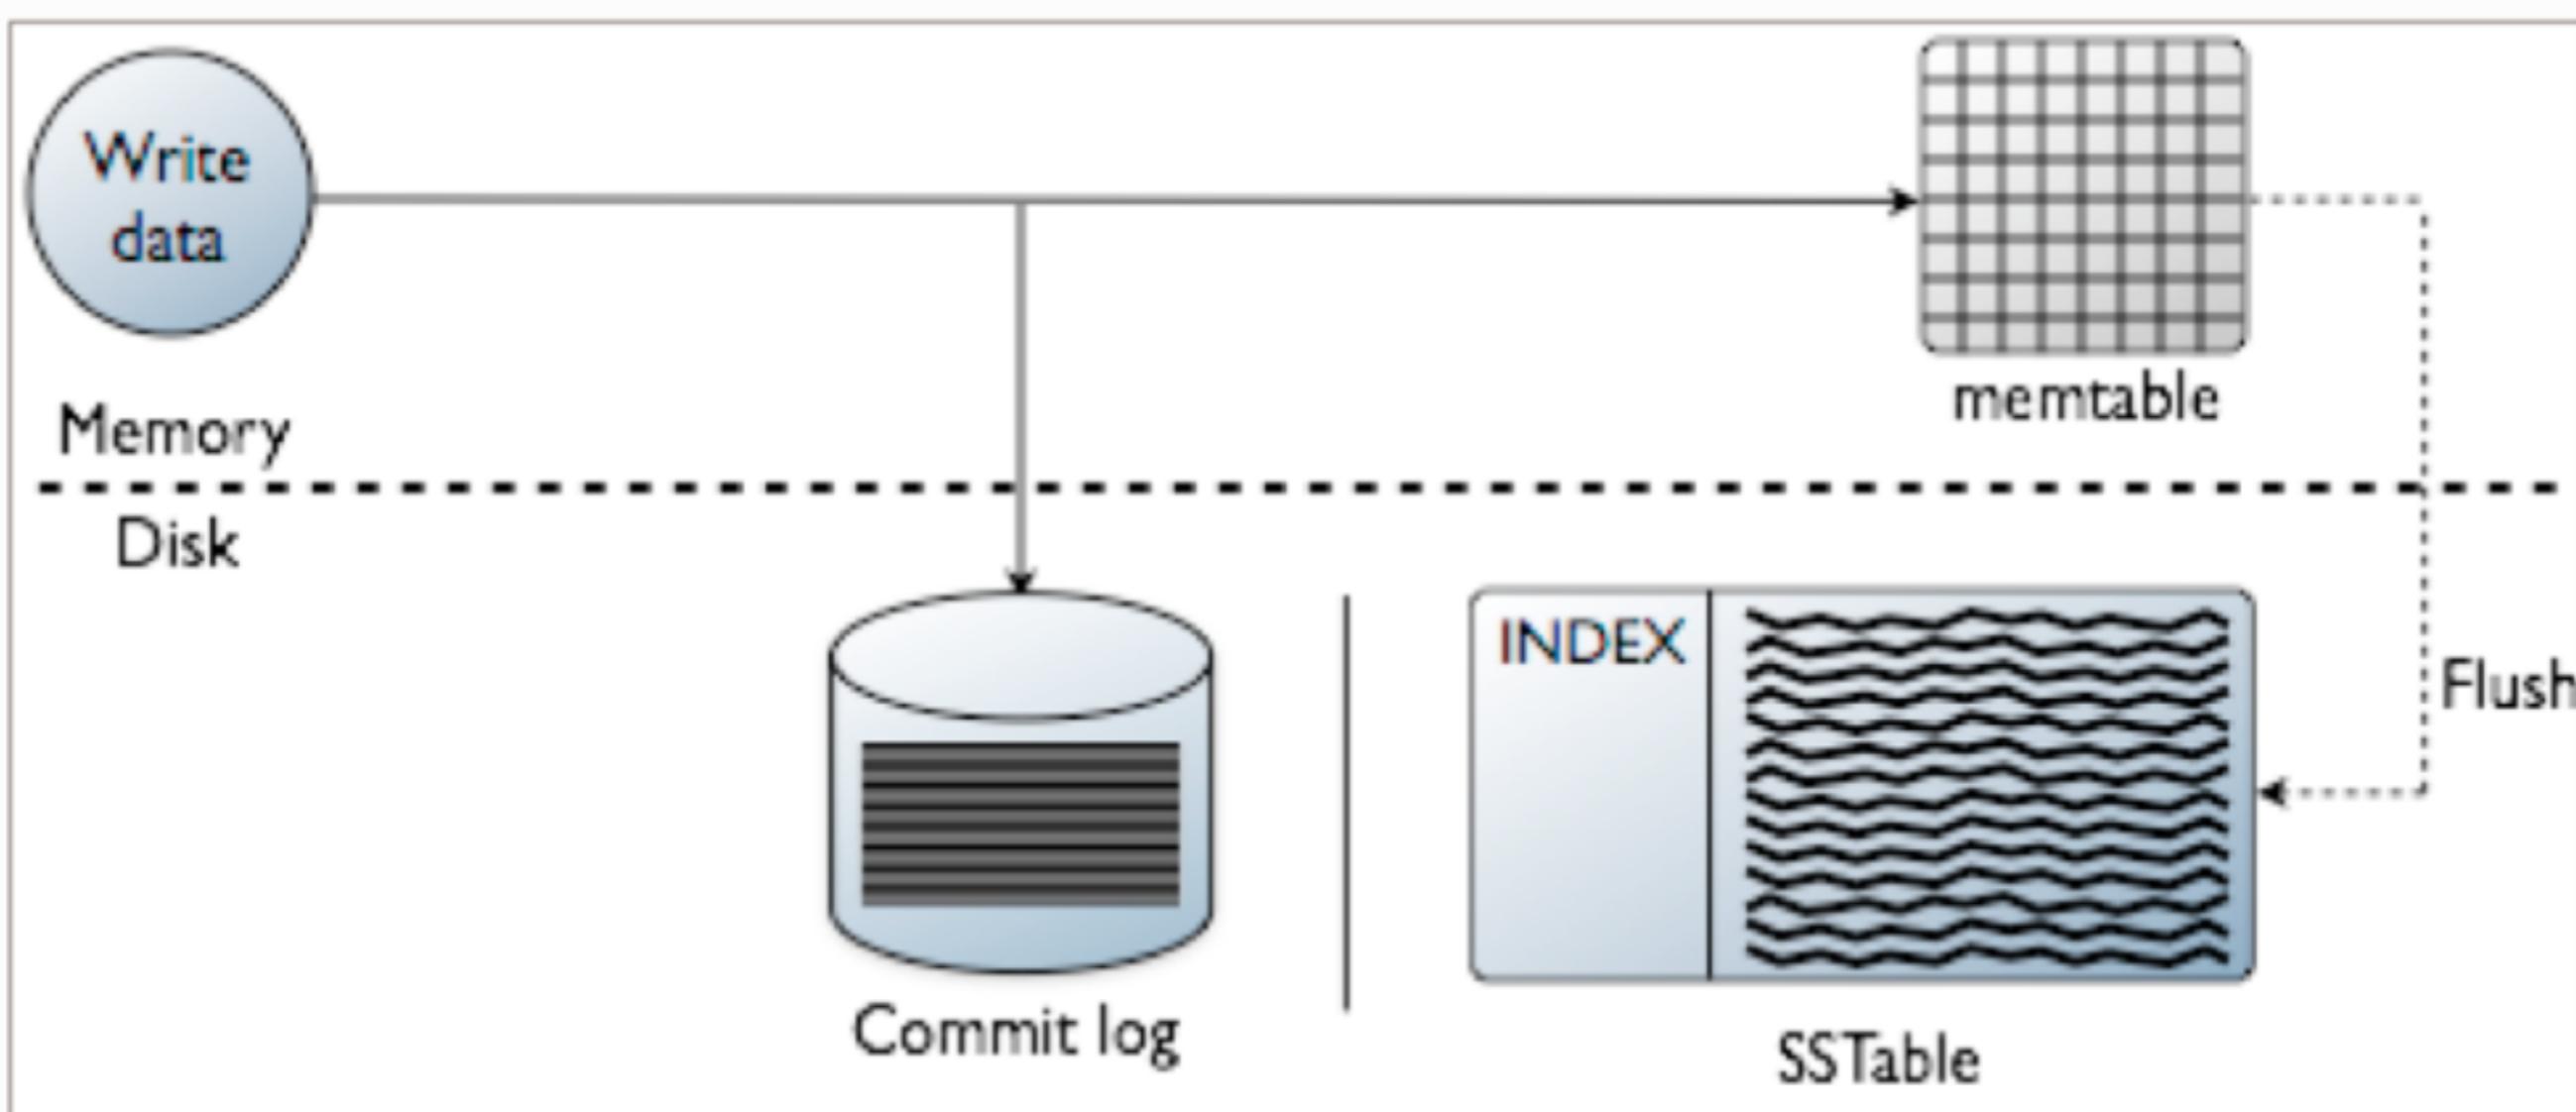
Reads and Writes in Cassandra



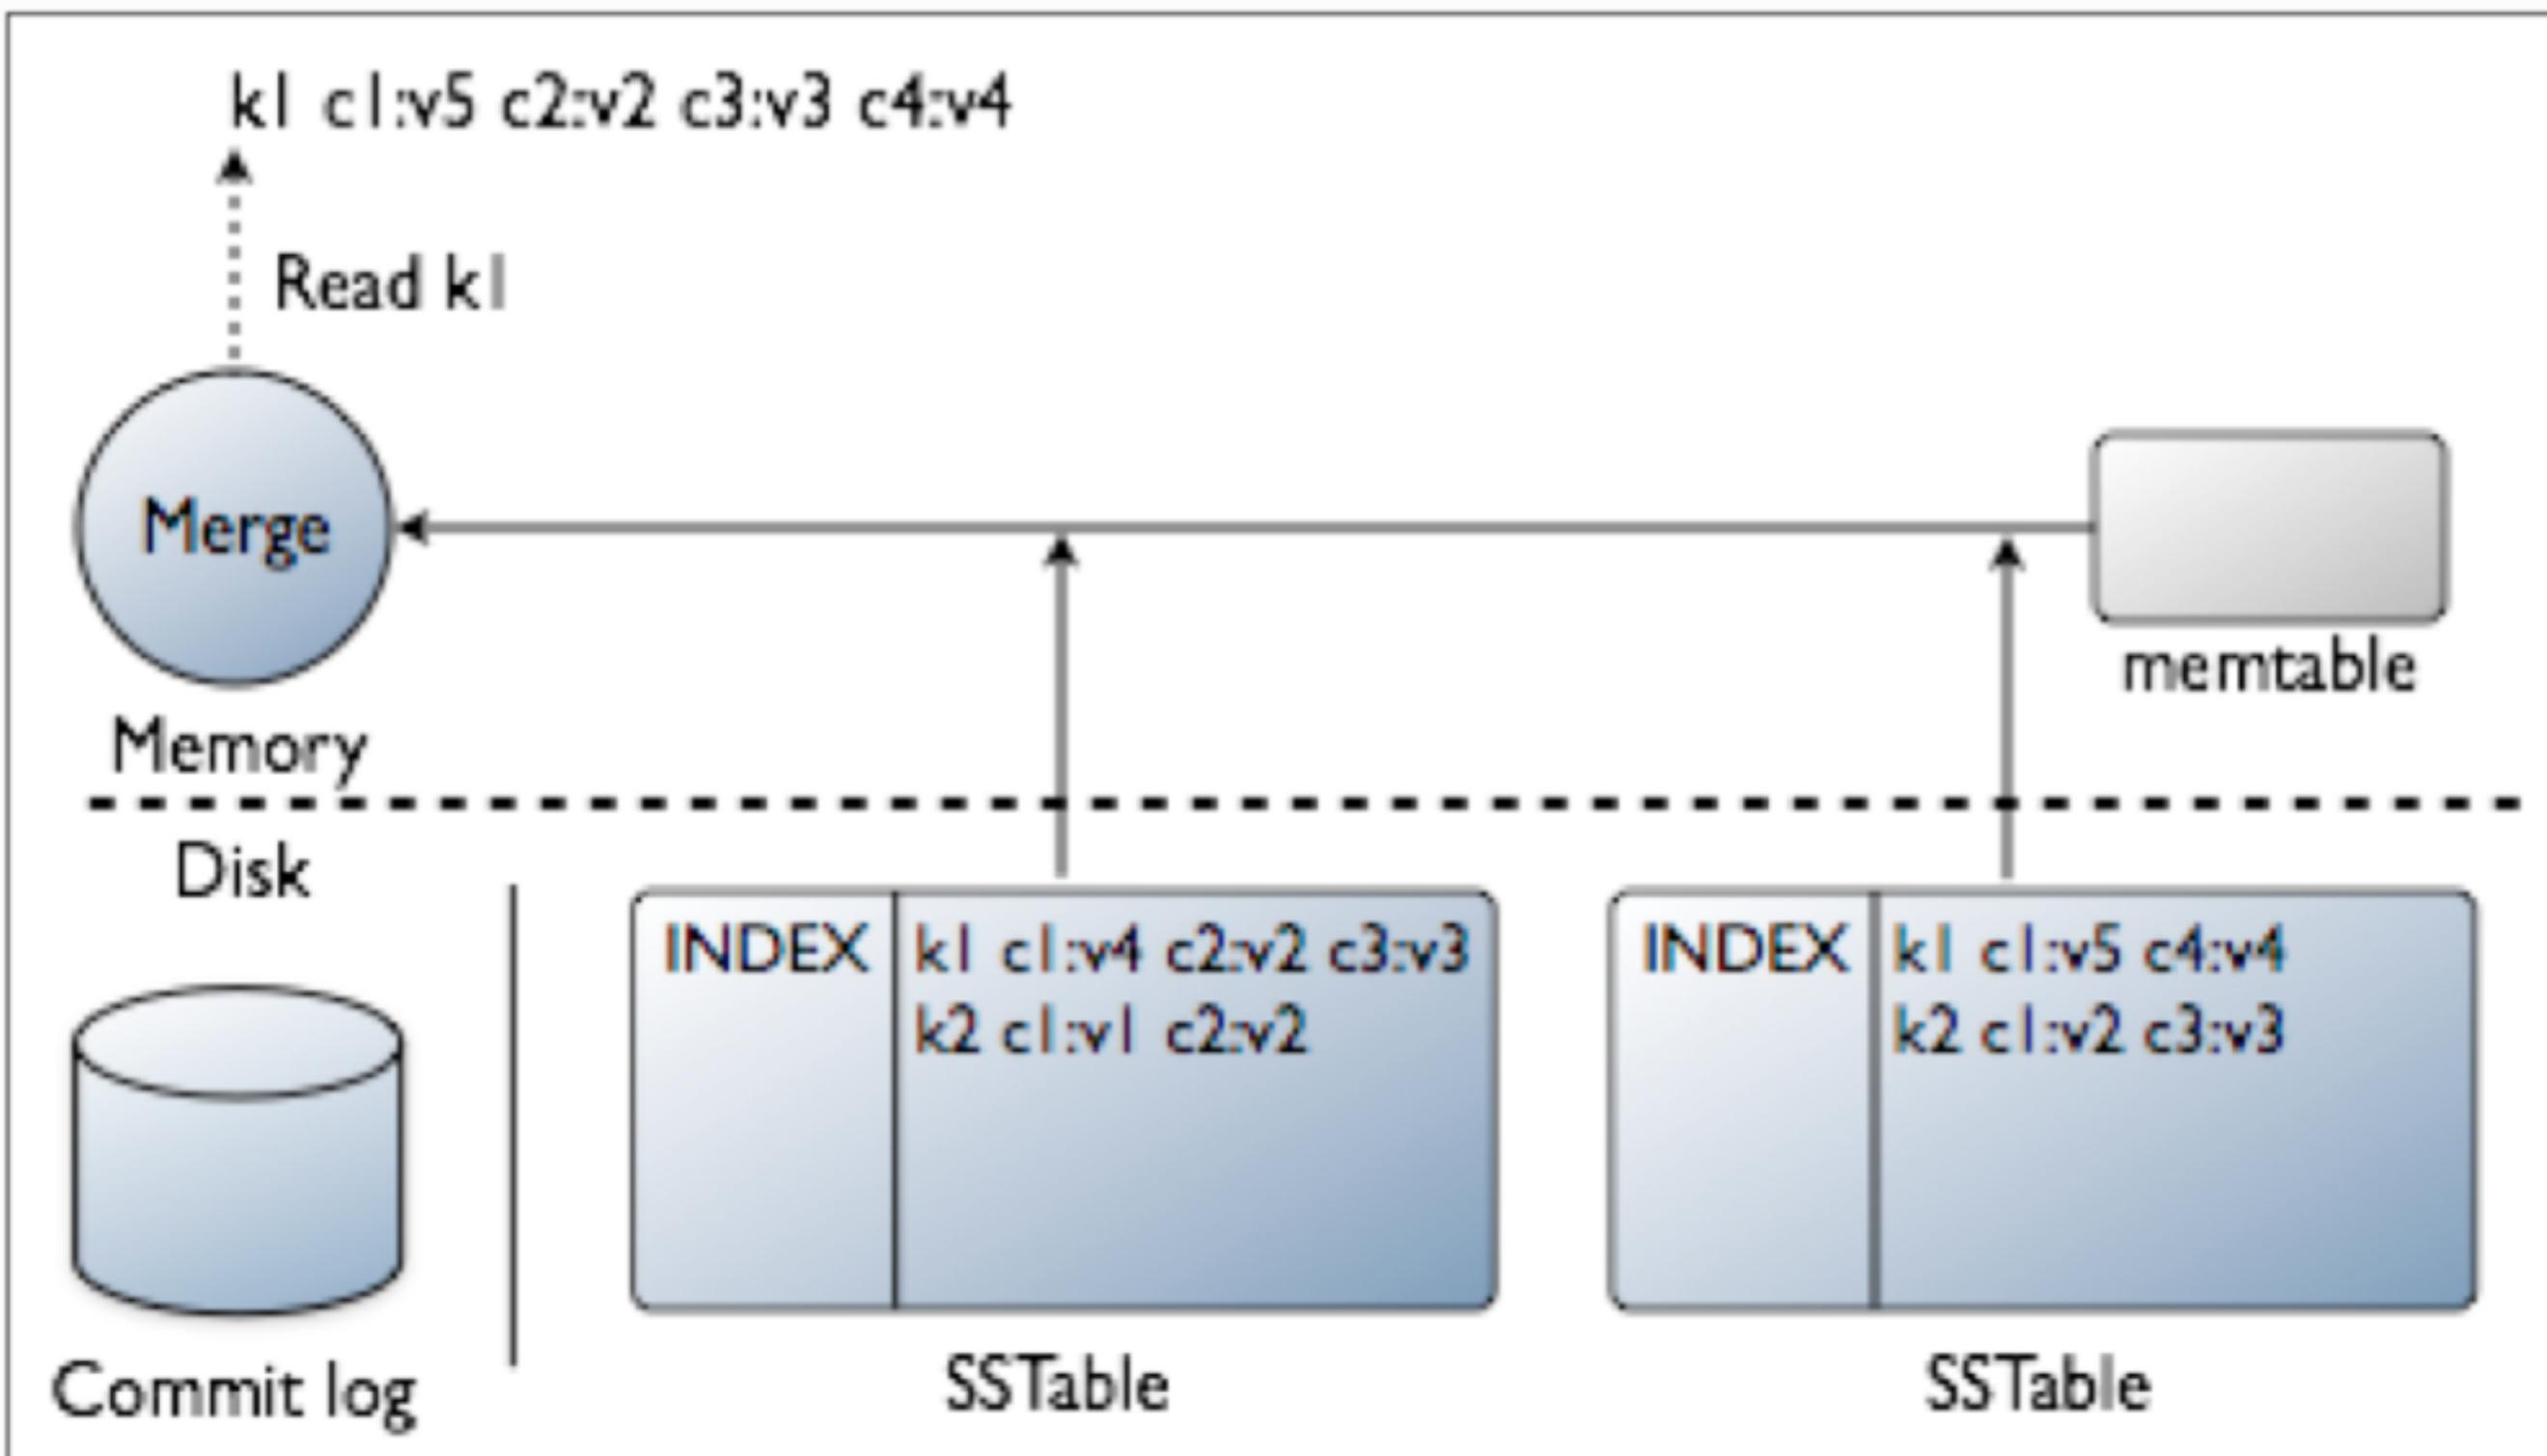
creating tables

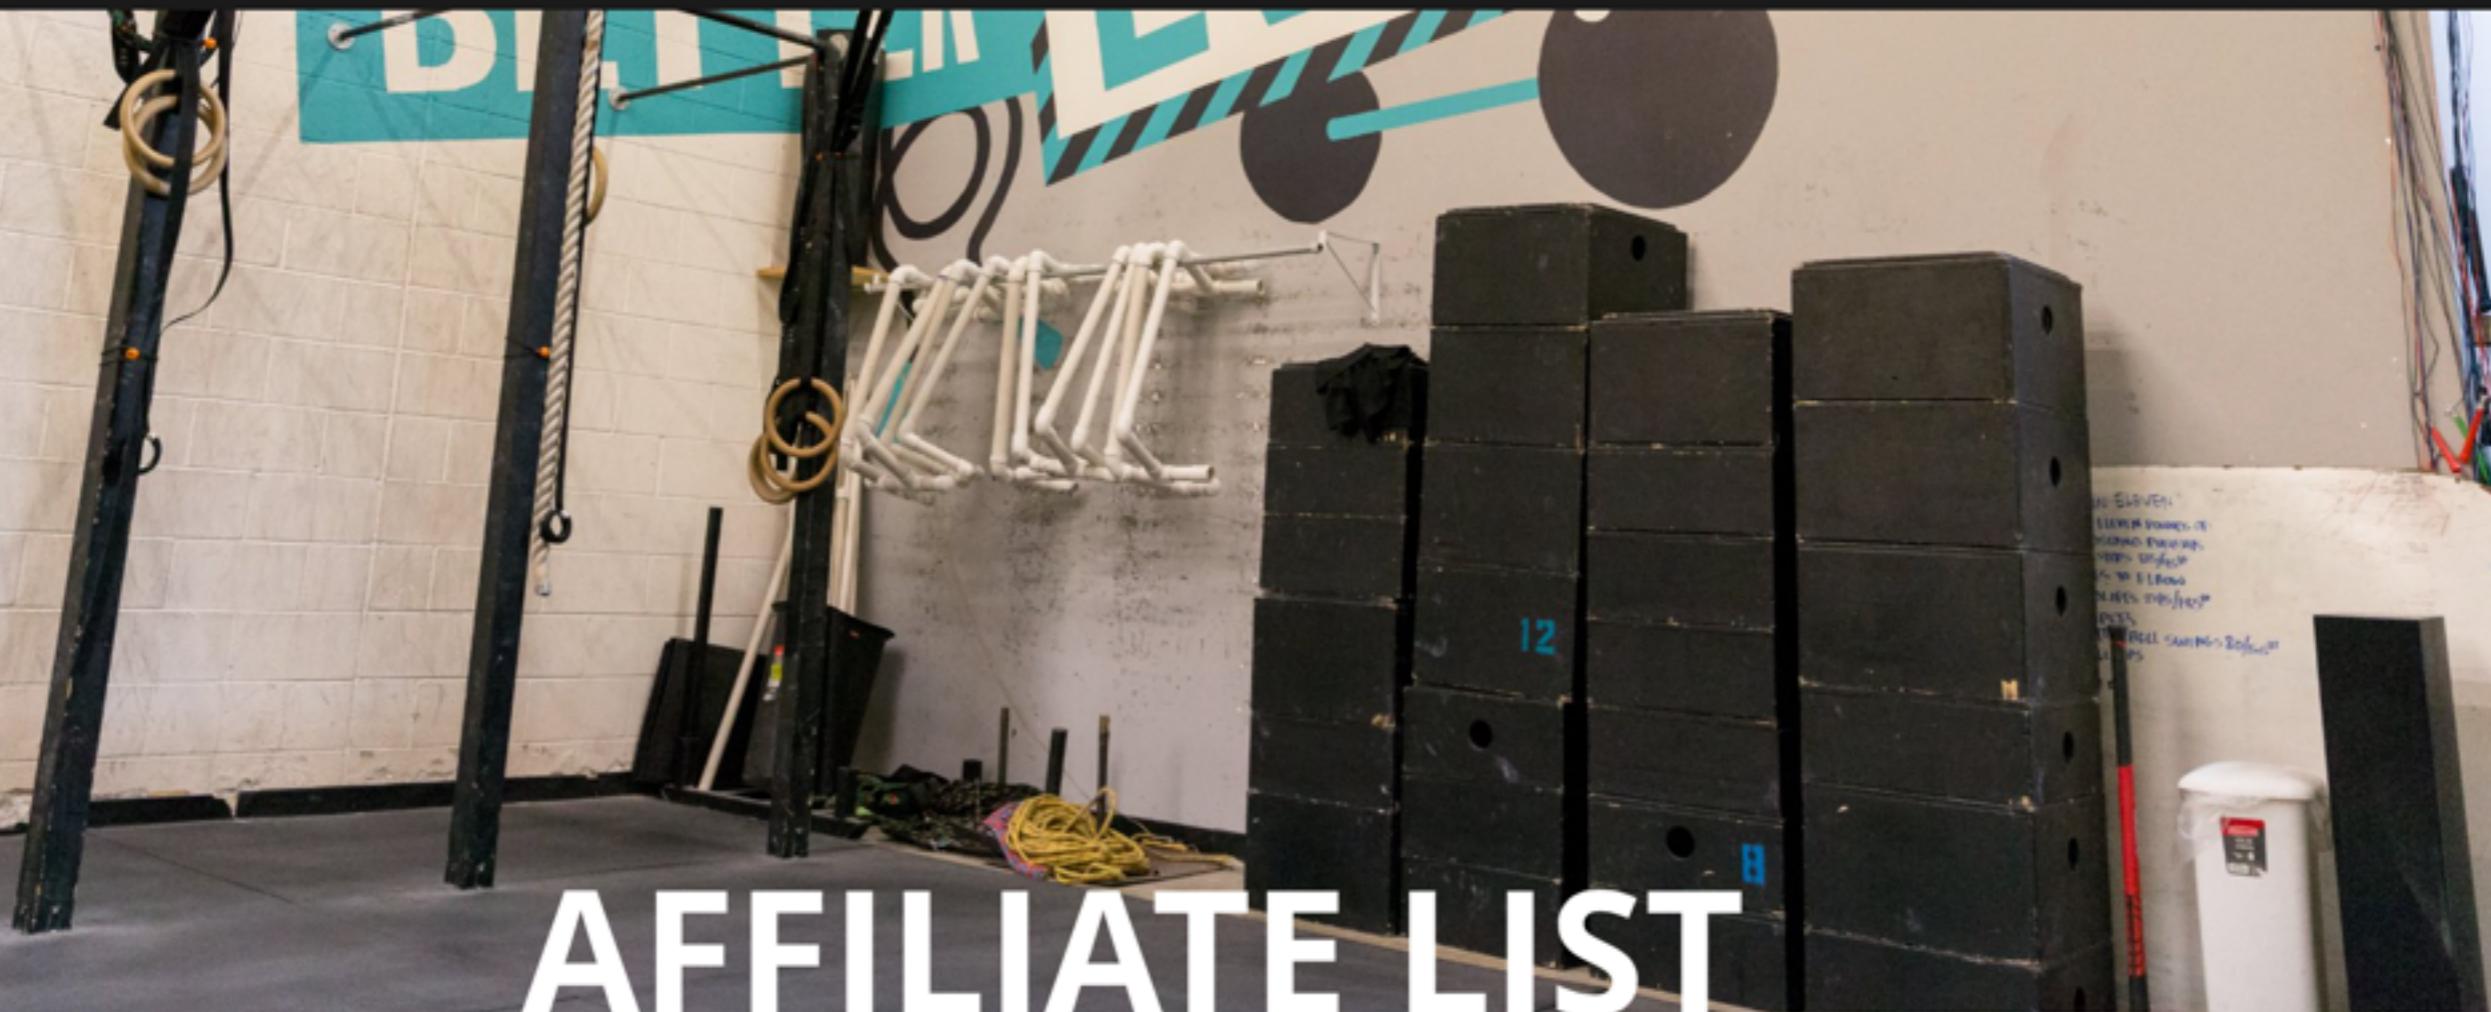
Writing the CQL statements

```
CREATE TABLE comments_by_user (
    user_id UUID,
    posted_timestamp TIMESTAMP,
    video_id TIMEUUID,
    comment TEXT,
    title TEXT,
    type TEXT,
    tags SET<TEXT>,
    preview_thumbnails MAP<INT, BLOB>,
    PRIMARY KEY ((user_id), posted_timestamp, video_id)
) WITH CLUSTERING ORDER BY (posted_timestamp DESC, video_id ASC);
```



read processing in node



[SIGN IN](#)**CrossFit®**[GET STARTED](#) [WODs](#) [COURSES](#) [CERTIFICATIONS](#) [EXERCISE & DEMOS](#) [GAMES](#) [JOURNAL](#) [AFFILIATES](#) [FOUNDATION](#) [SHOP](#)

AFFILIATE LIST[®]

[CLEAR RESULTS](#)

Partition keys

- CREATE TABLE crossfit_gyms (
- gym_name text,
- city text,
- state_province text,
- country_code text,
- PRIMARY KEY (gym_name)
-);

compound keys

```
CREATE TABLE crossfit_gyms (
    gym_name text,
    city text,
    state_province text,
    country_code text,
    PRIMARY KEY (gym_name)
);
```

Clustering keys

```
CREATE TABLE crossfit_gyms_by_location (
    country_code text,
    state_province text,
    city text,
    gym_name text,
    PRIMARY KEY (country_code, state_province, city,
    gym_name)
);
```

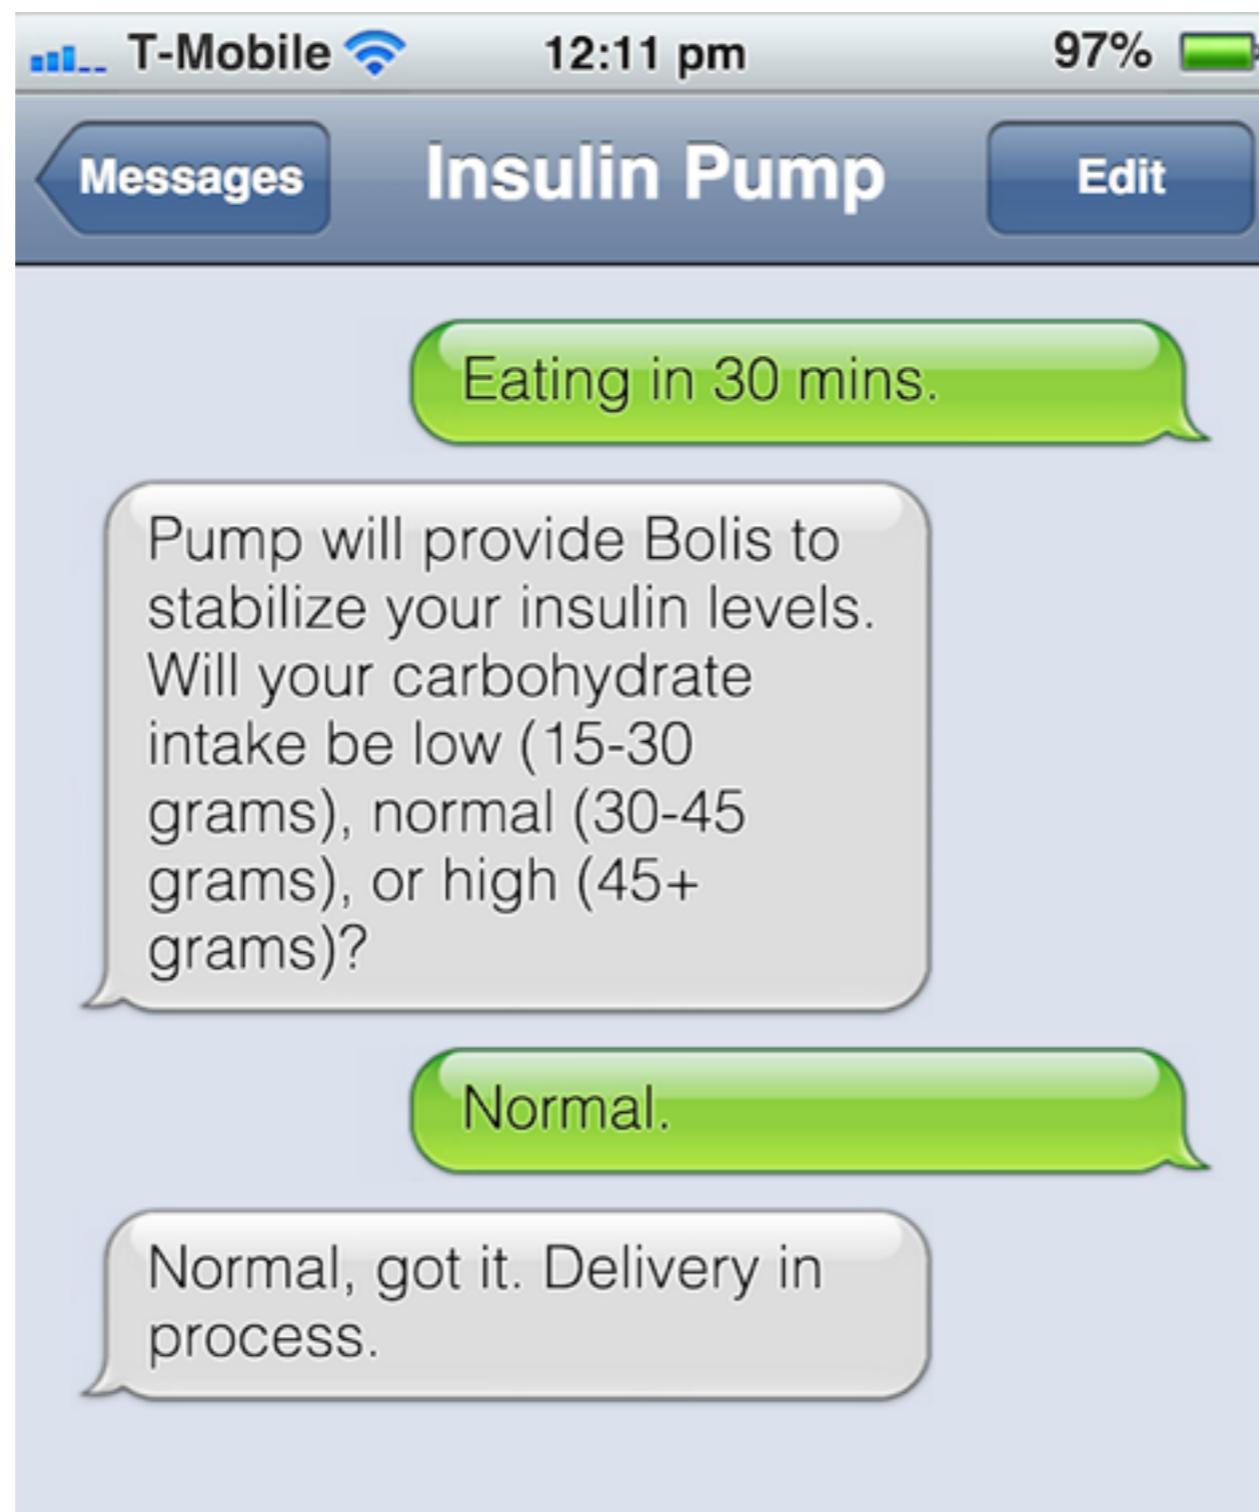
order by

```
CREATE TABLE crossfit_gyms_by_location (
    country_code text,
    state_province text,
    city text,
    gym_name text,
    PRIMARY KEY (country_code, state_province,
    city, gym_name)
) WITH CLUSTERING ORDER BY (state_province DESC,
    city ASC, gym_name ASC);
```

hashed values

- **country_code | state_province | city | gym_name**
- -----+-----+-----+-----
- CAN | ON | Toronto | CrossFit Leslieville
- CAN | ON | Toronto | CrossFit Toronto
- CAN | BC | Vancouver | CrossFit BC
- CAN | BC | Vancouver | CrossFit Vancouver
- USA | NY | New York | CrossFit Metropolis
- USA | NY | New York | CrossFit NYC
- USA | NV | Las Vegas | CrossFit Las Vegas
- USA | NV | Las Vegas | Kaizen CrossFit
- USA | CA | San Francisco | LaLanne Fitness CrossFit
- USA | CA | San Francisco | San Francisco CrossFit

Medical Devices as IoT



insulin

- CREATE KEYSPACE insulin_pump
- WITH REPLICATION = { 'class' : 'SimpleStrategy',
'replication_factor' : 3 }
- USE insulin_pump;

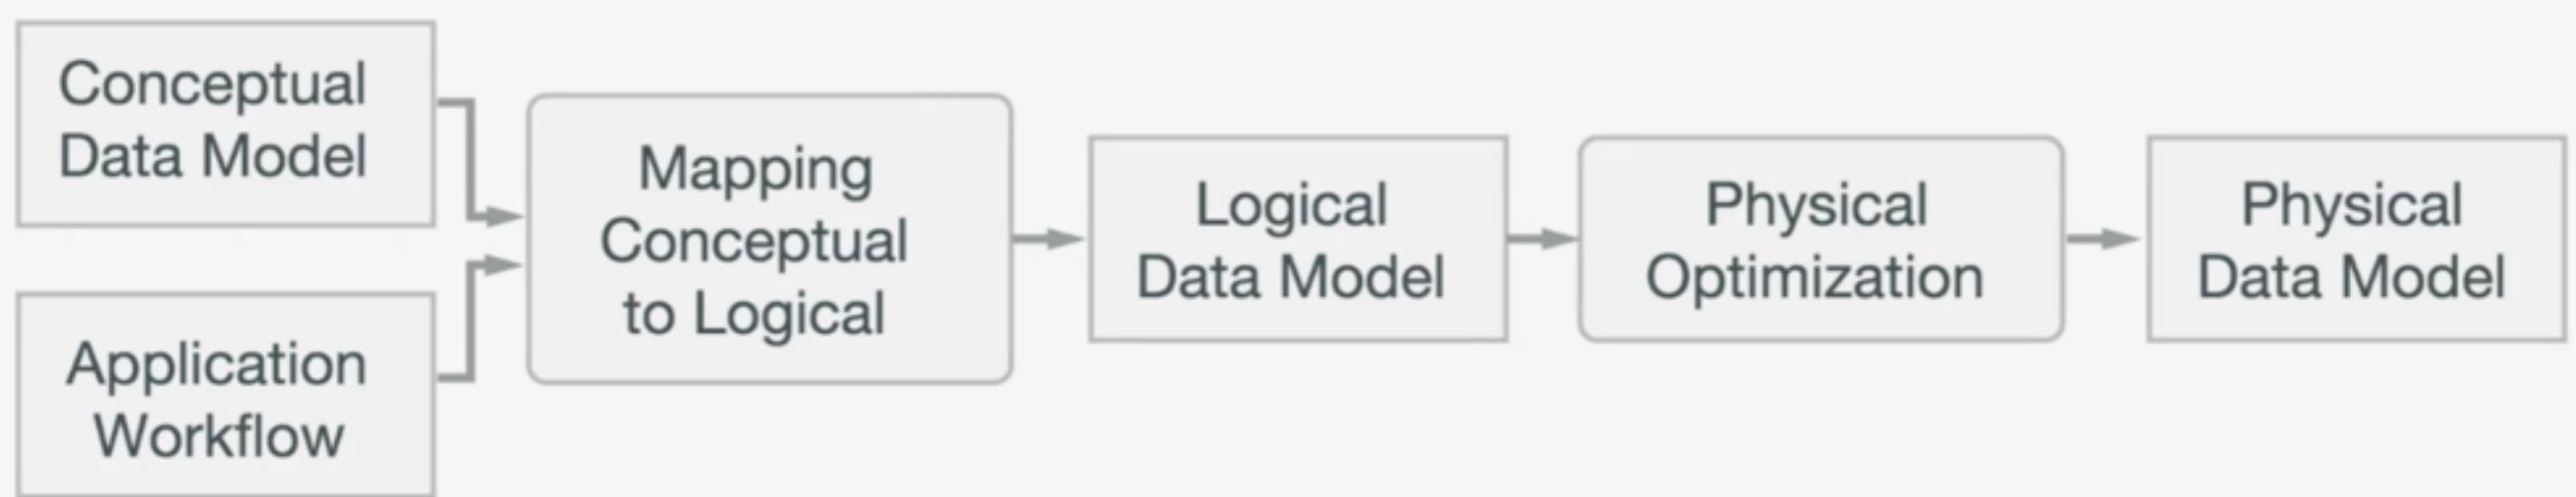
insulin_level table

- CREATE TABLE insulin_level (
- insulin_pump_id text,
- event_time timestamp,
- insulin_level_mcU_ml int,
- **PRIMARY KEY (insulin_pump_id,event_time)**
-);

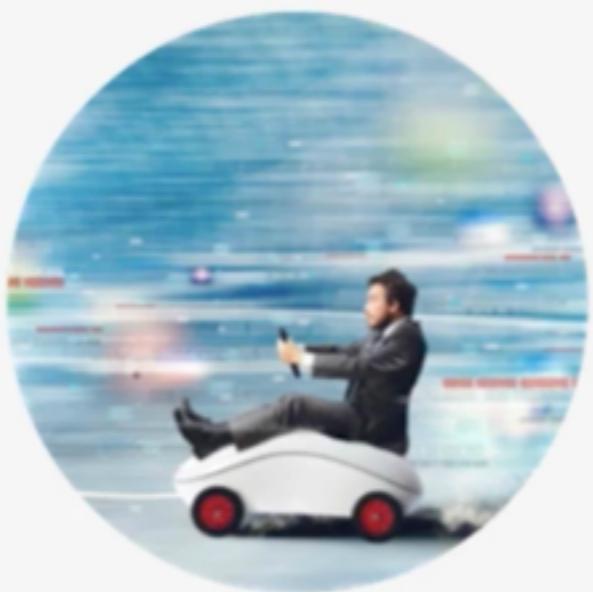
insulin_level_by_day

- CREATE TABLE insulin_level_by_day (
- insulin_pump_id text,
- date text,
- event_time timestamp,
- insulin_level_mcU_ml int,
- PRIMARY KEY ((insulin_pump_id,date),event_time)
-);

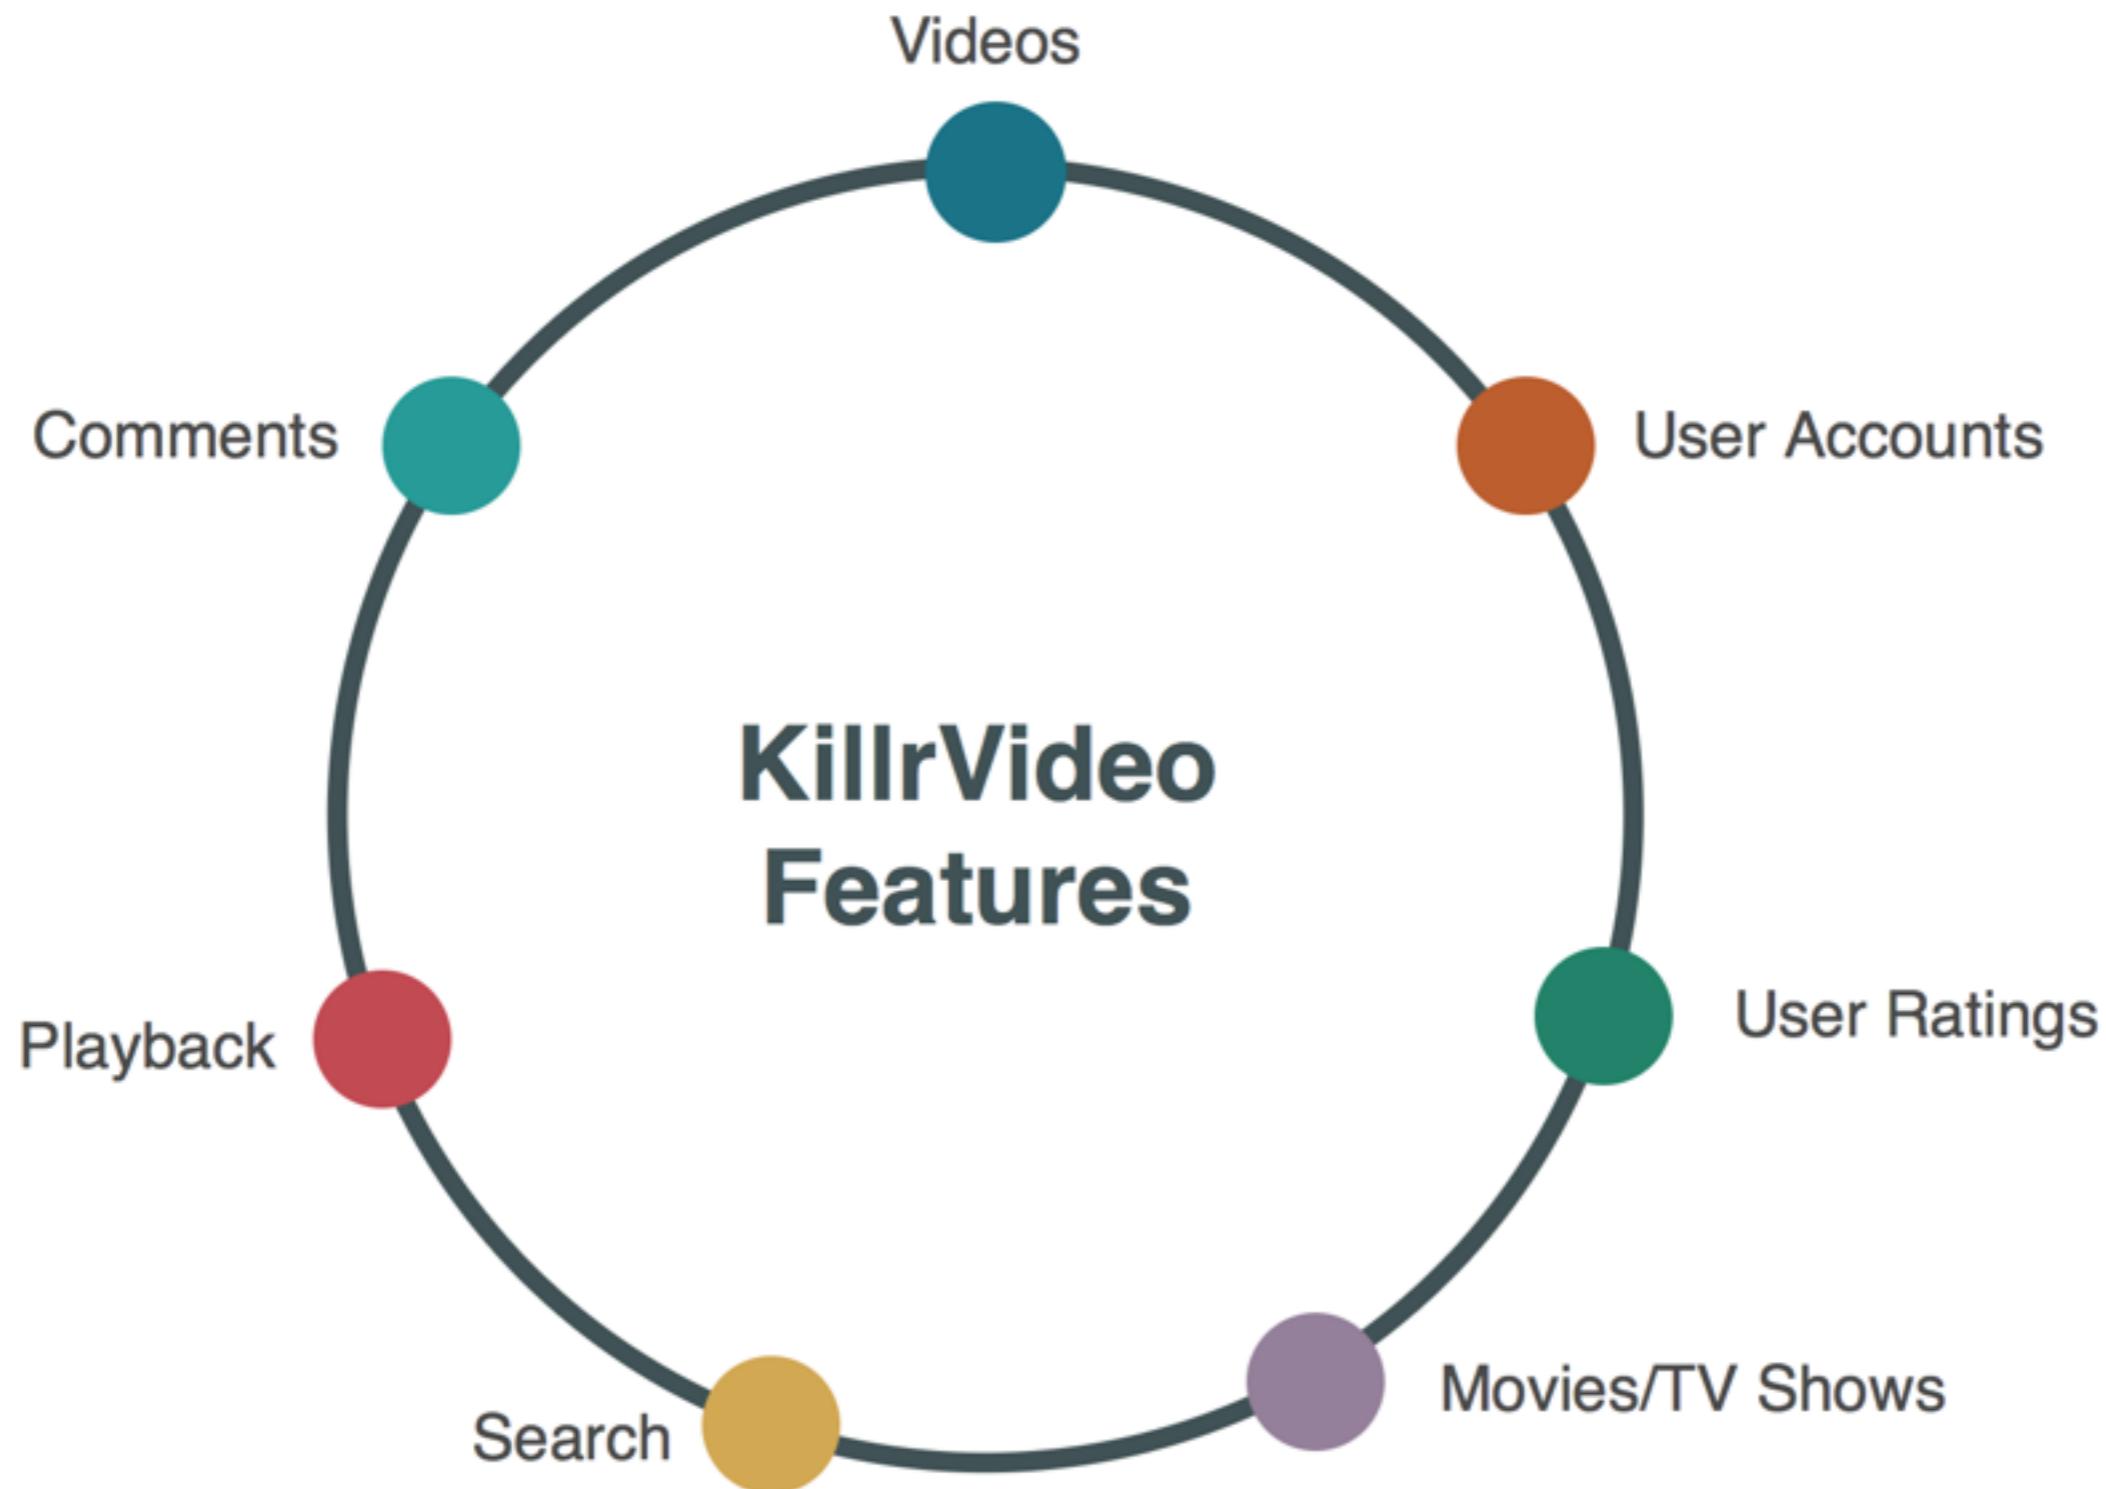
Cassandra data model



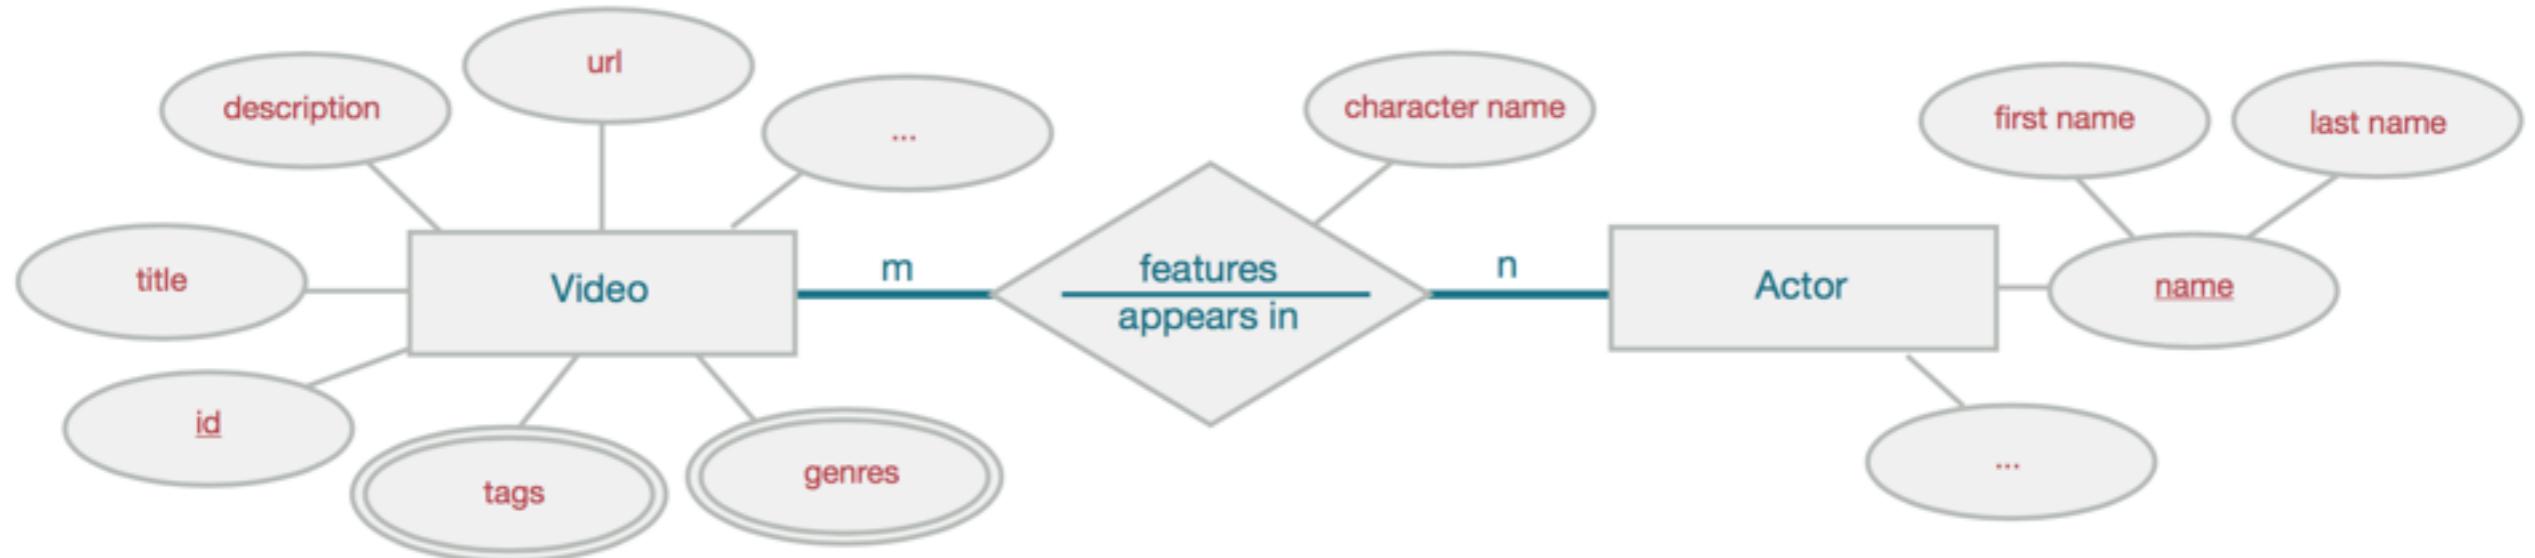
queries



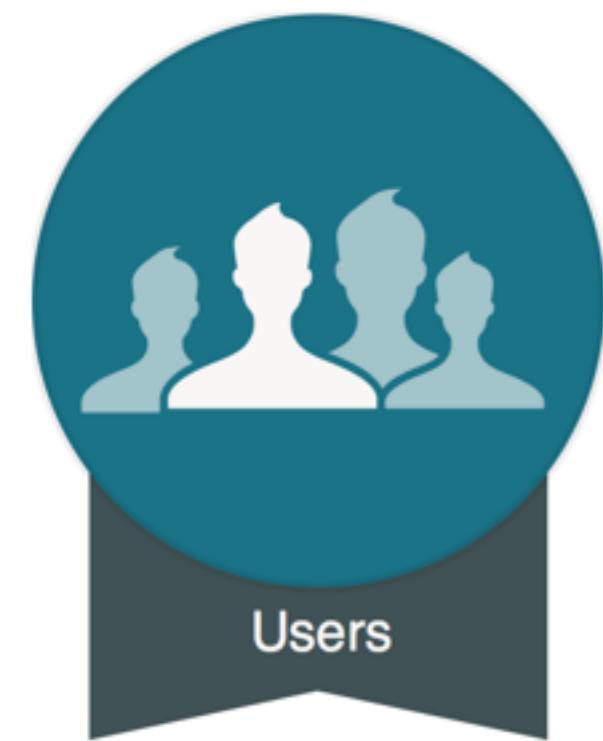
killrVideo.com



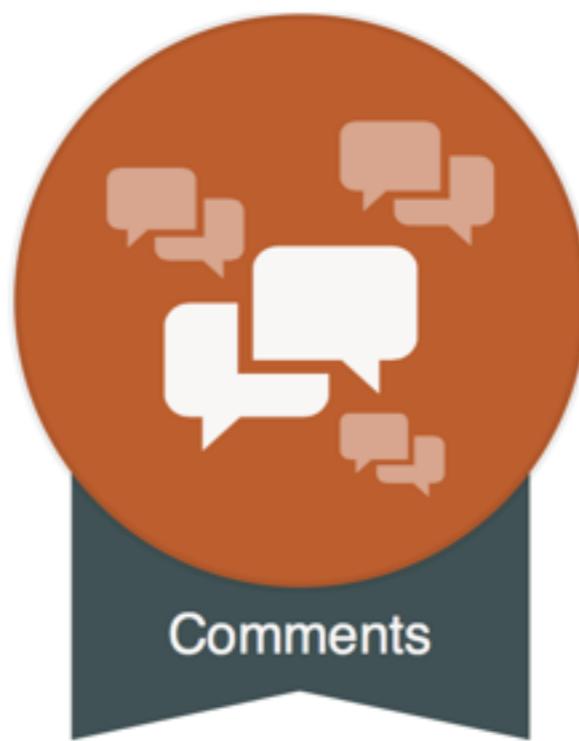
Conceptual Data Modeling



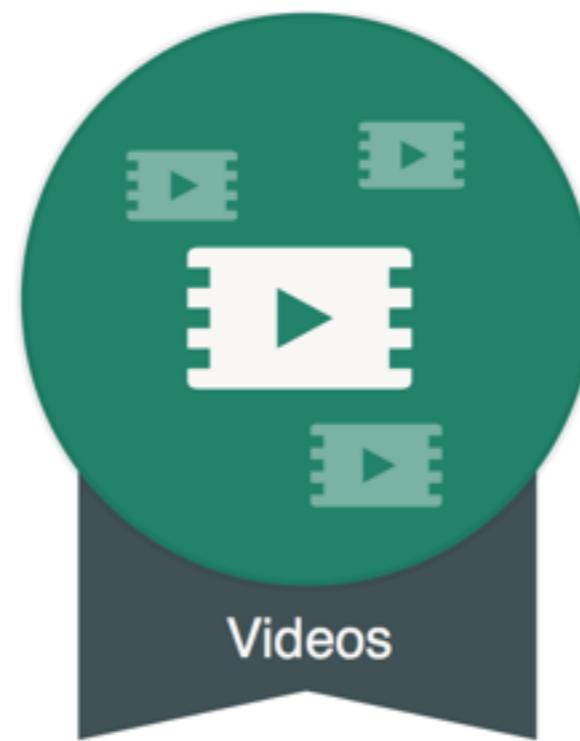
purpose



Users



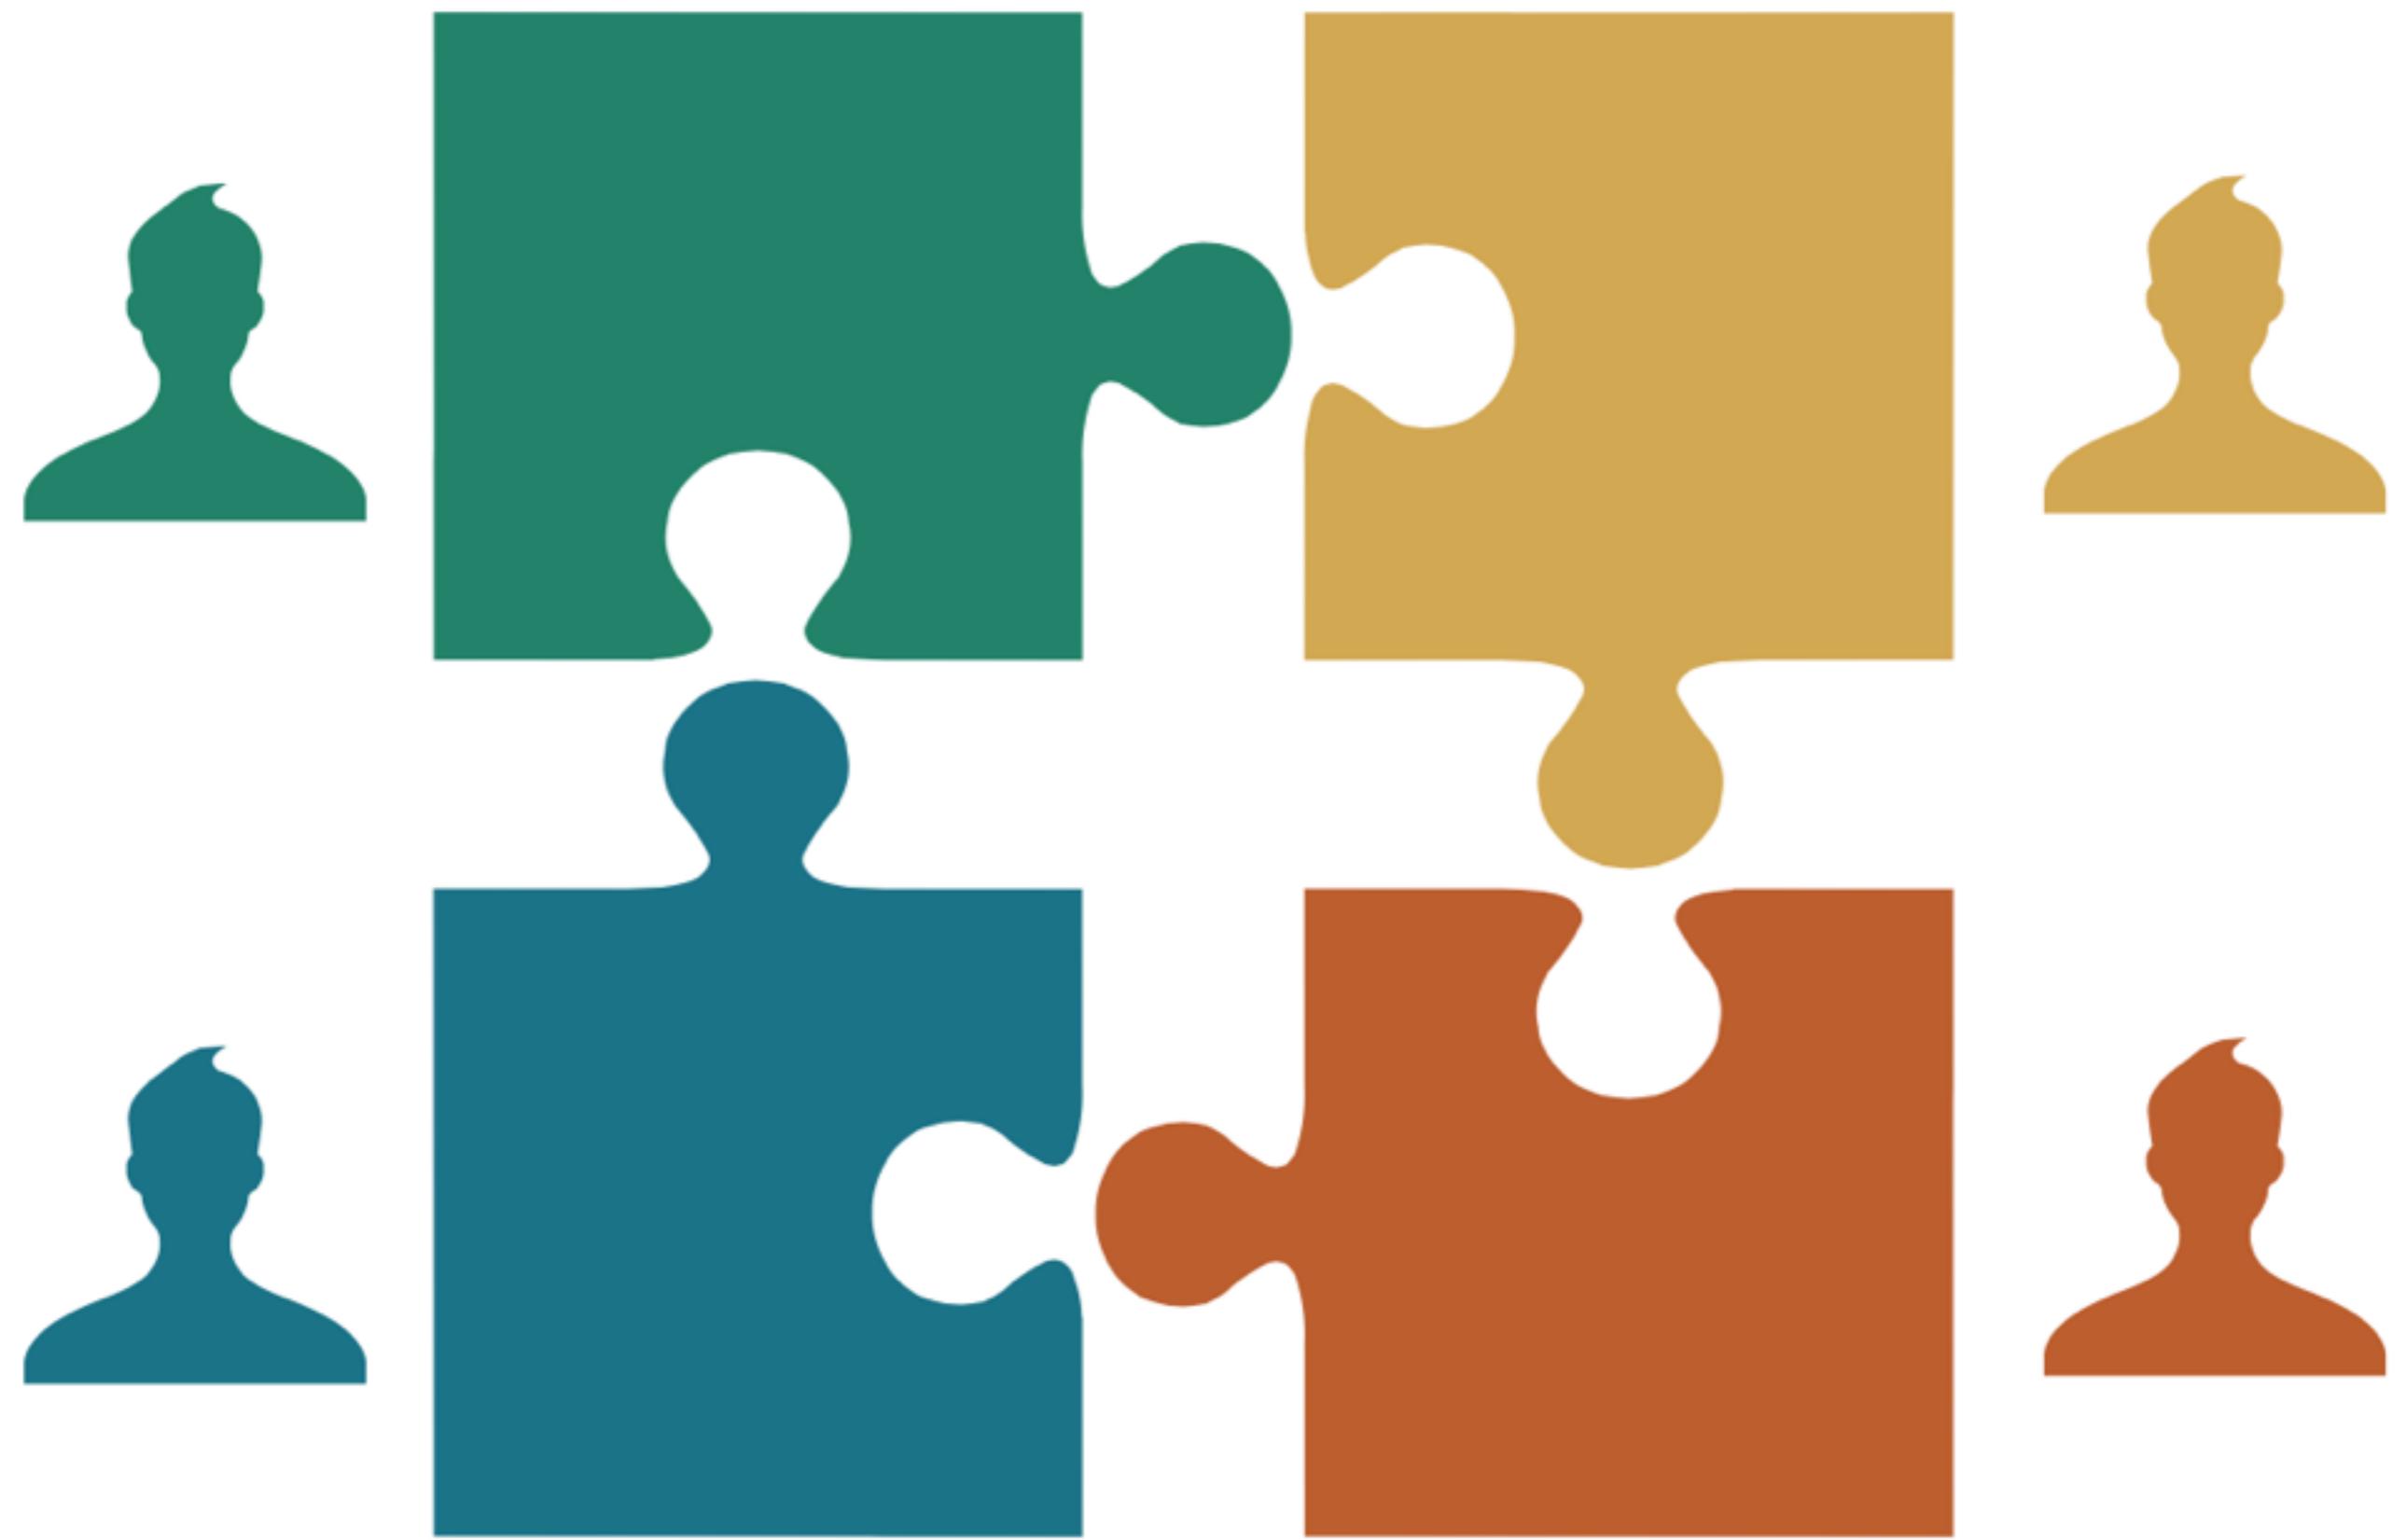
Comments



Videos



Ratings



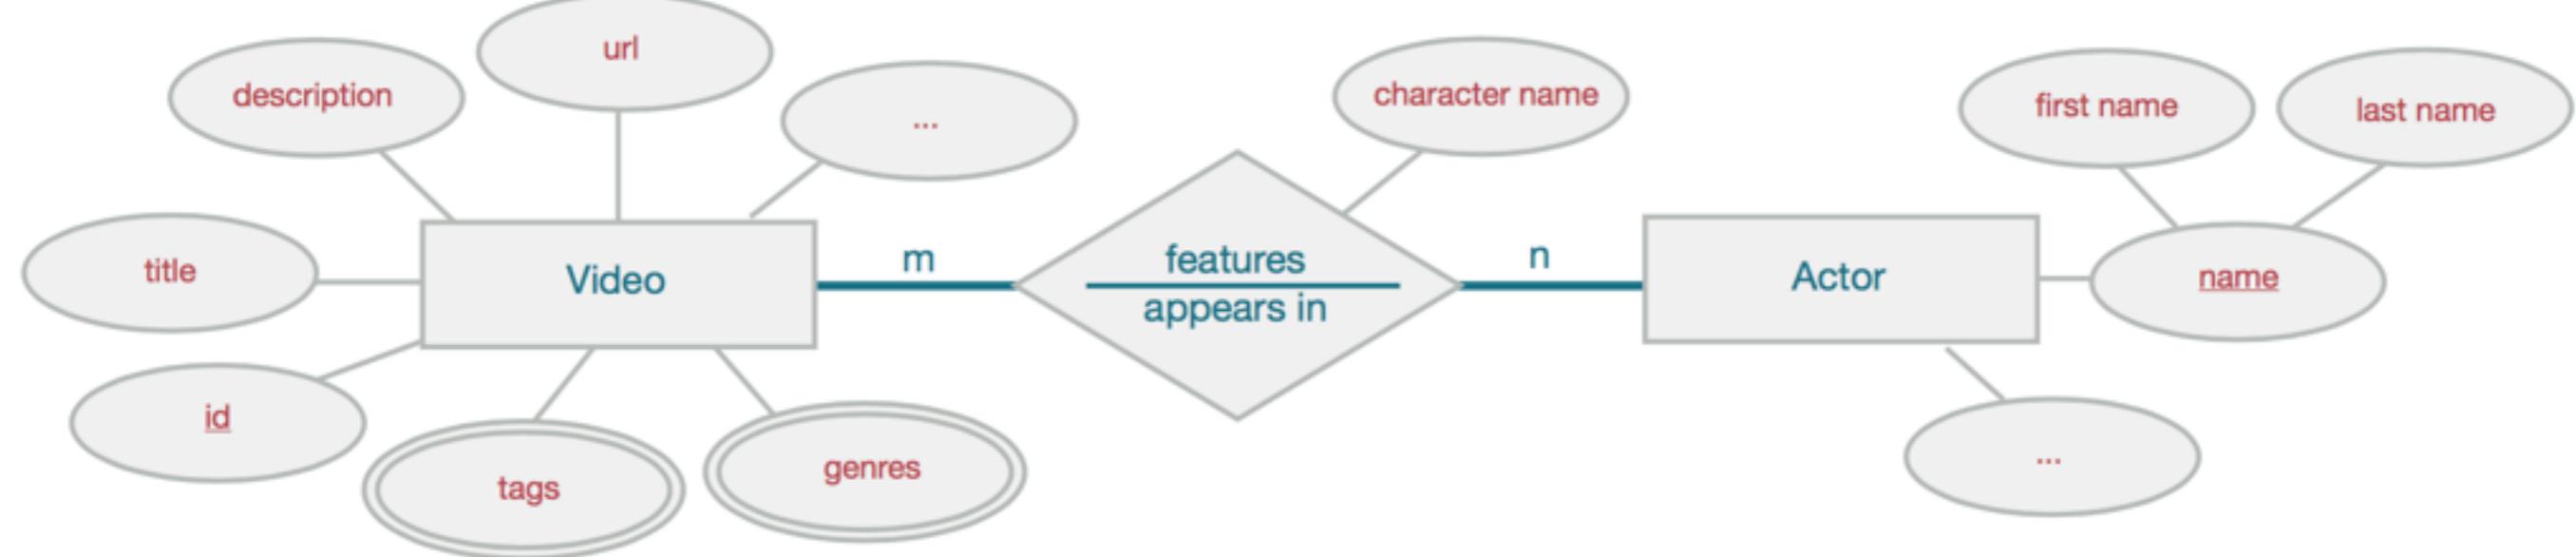
abstraction



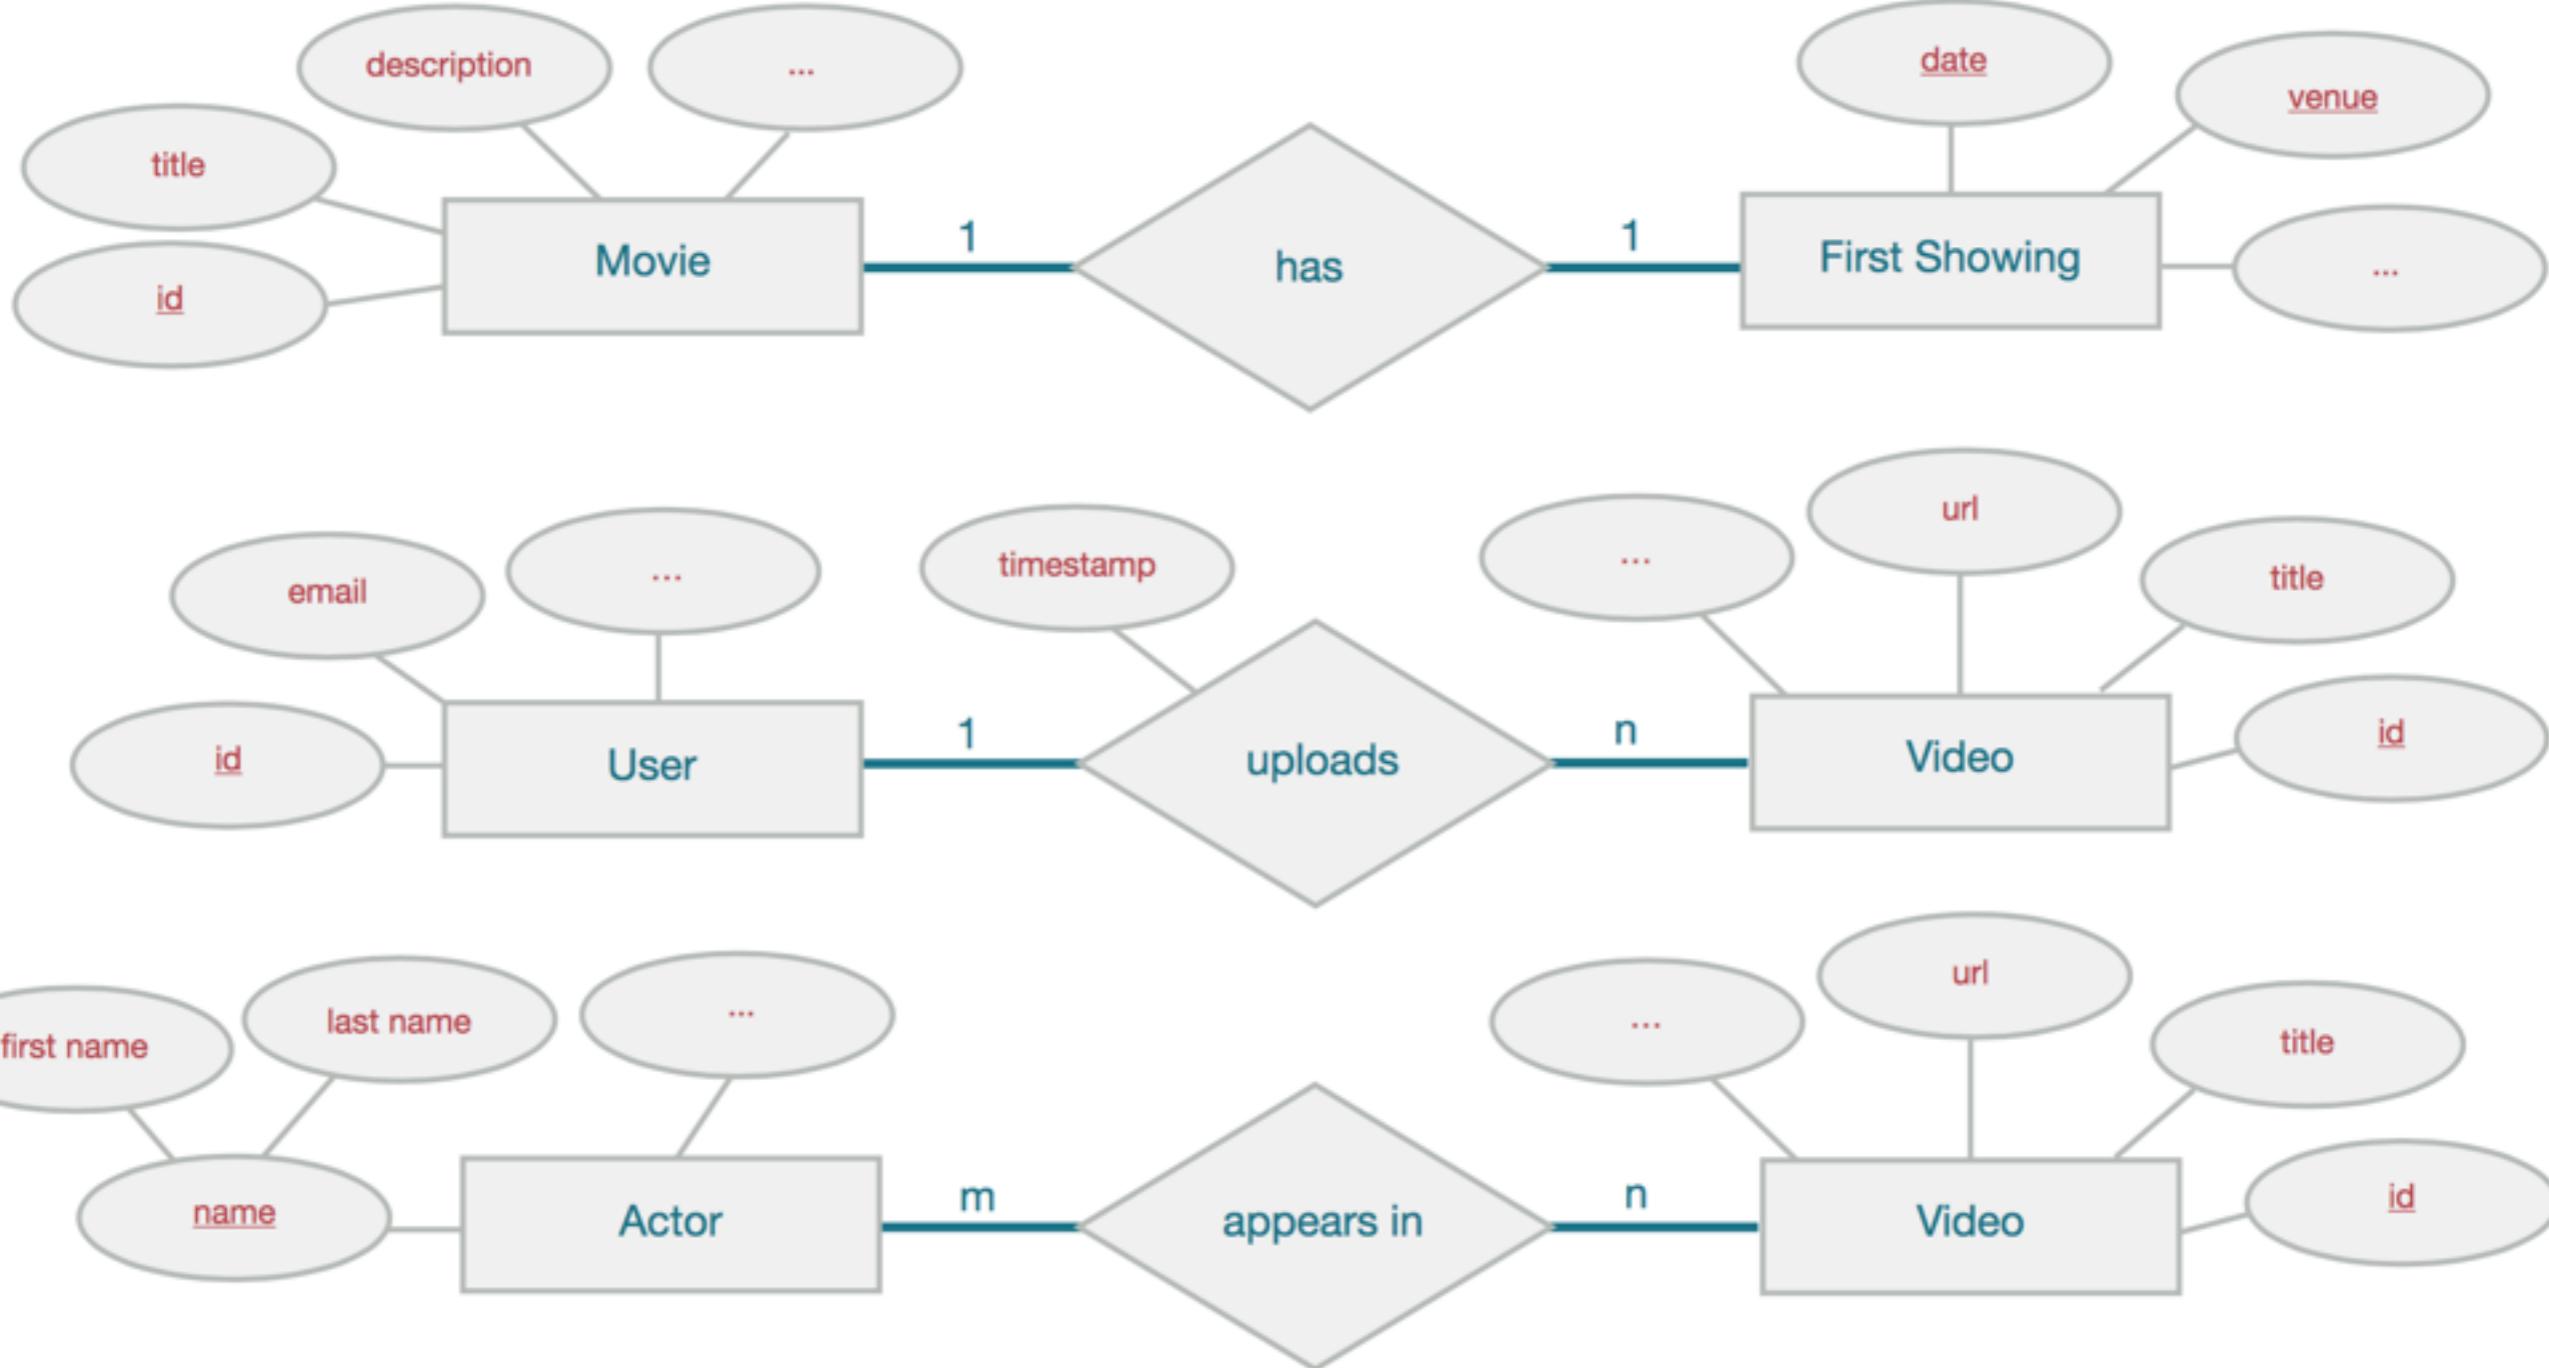
abstraction



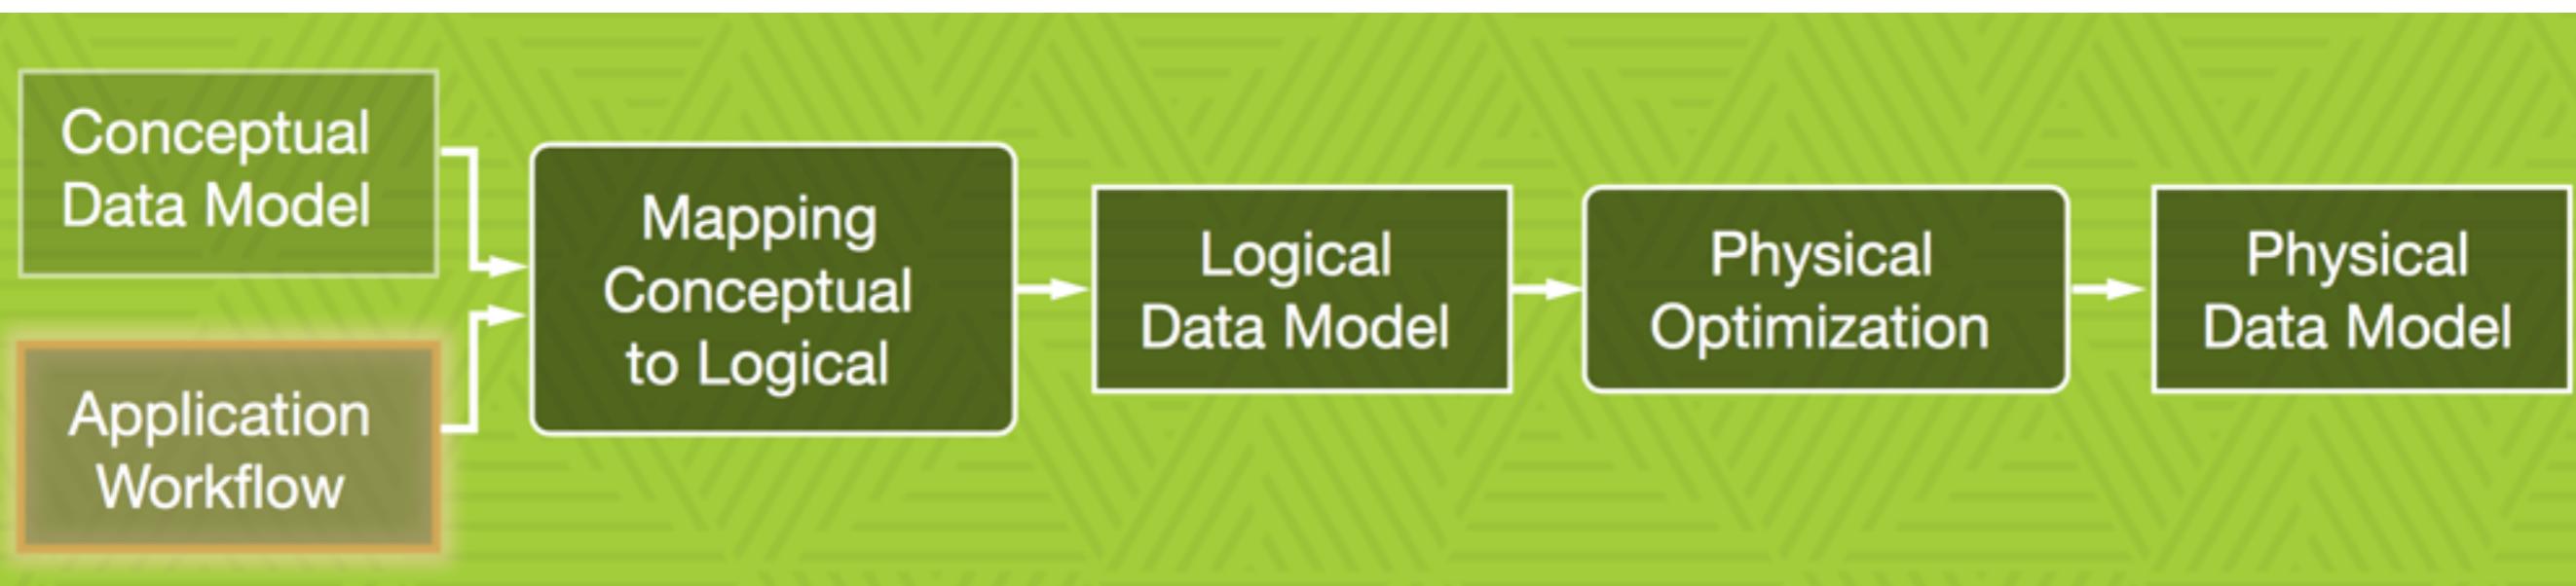
Entity-Relationship (ER) Model



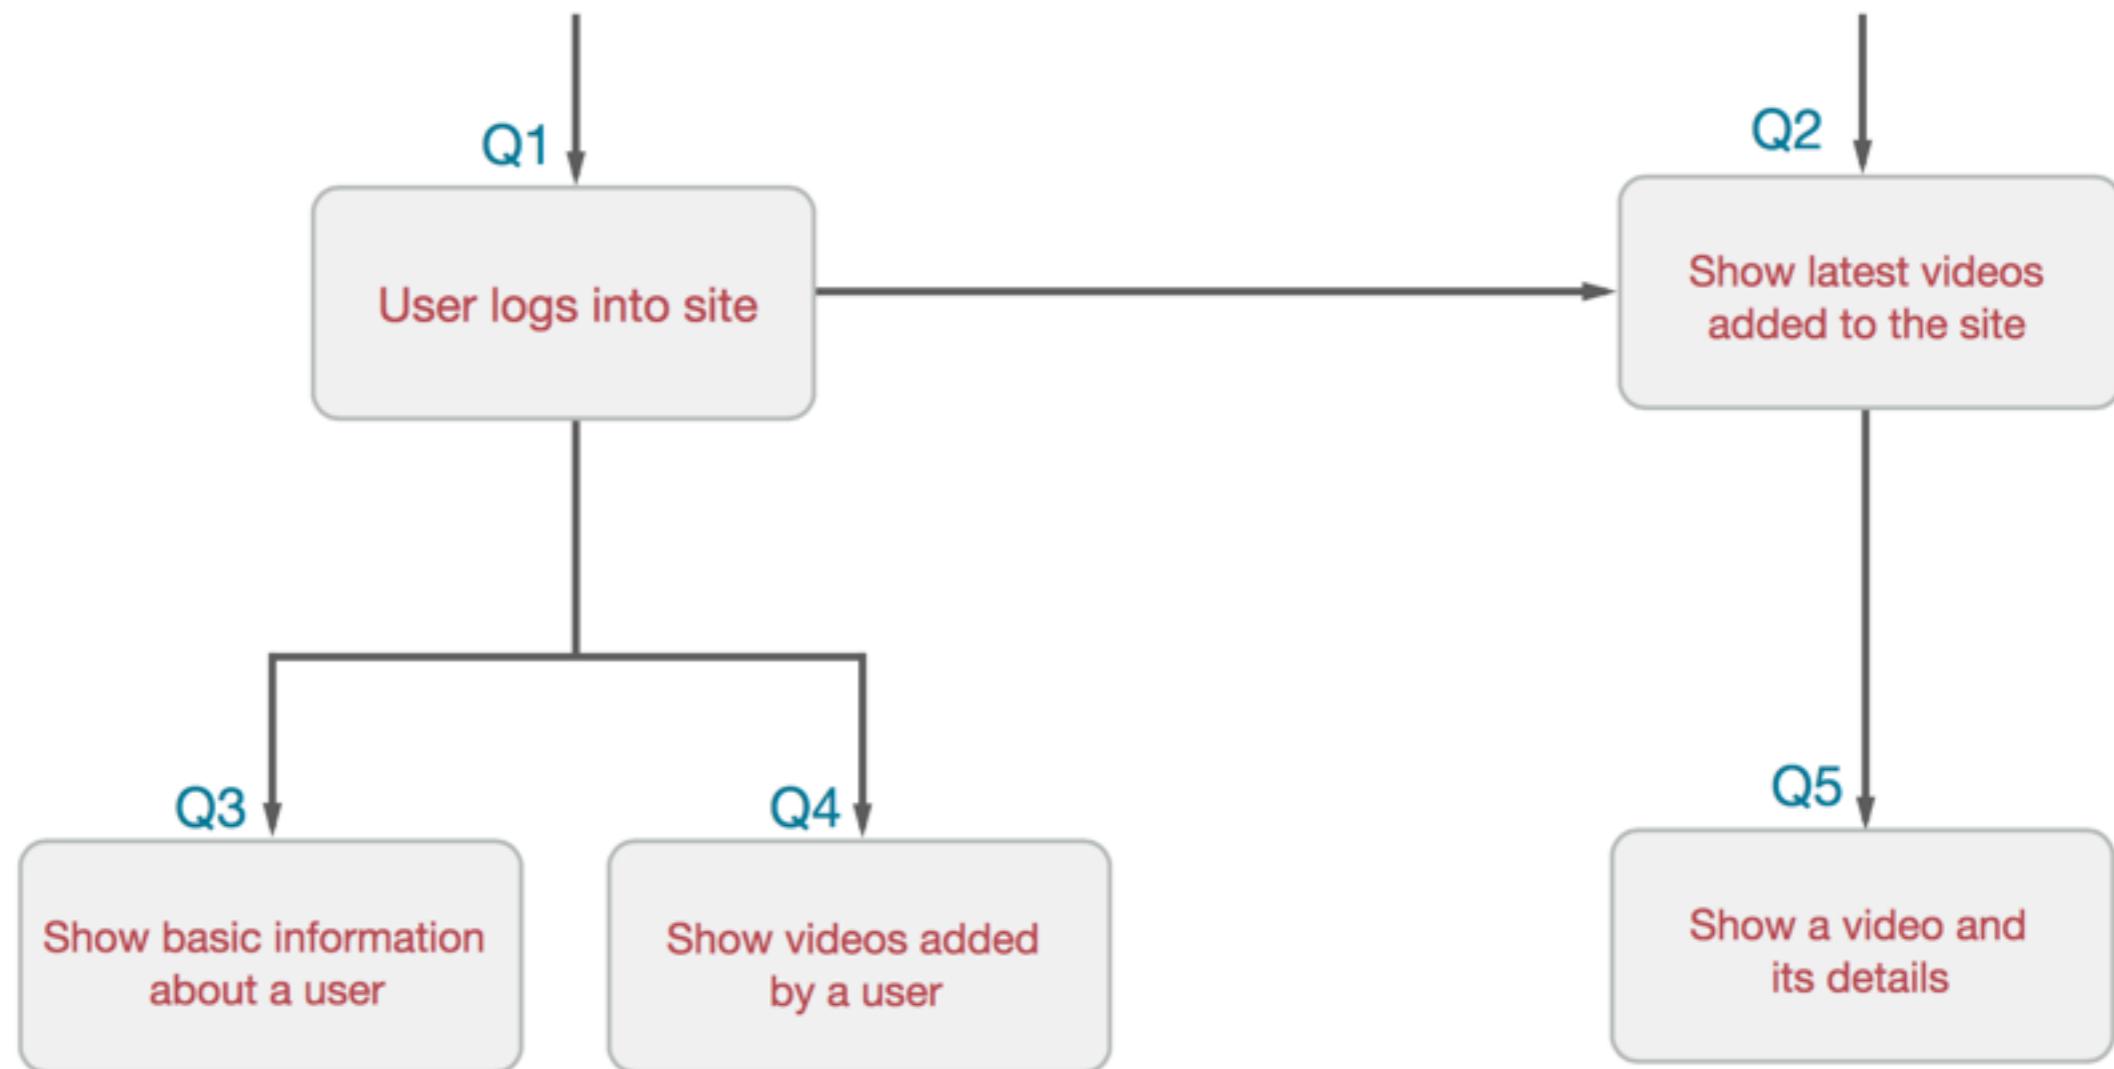
relationship keys



application workflow model



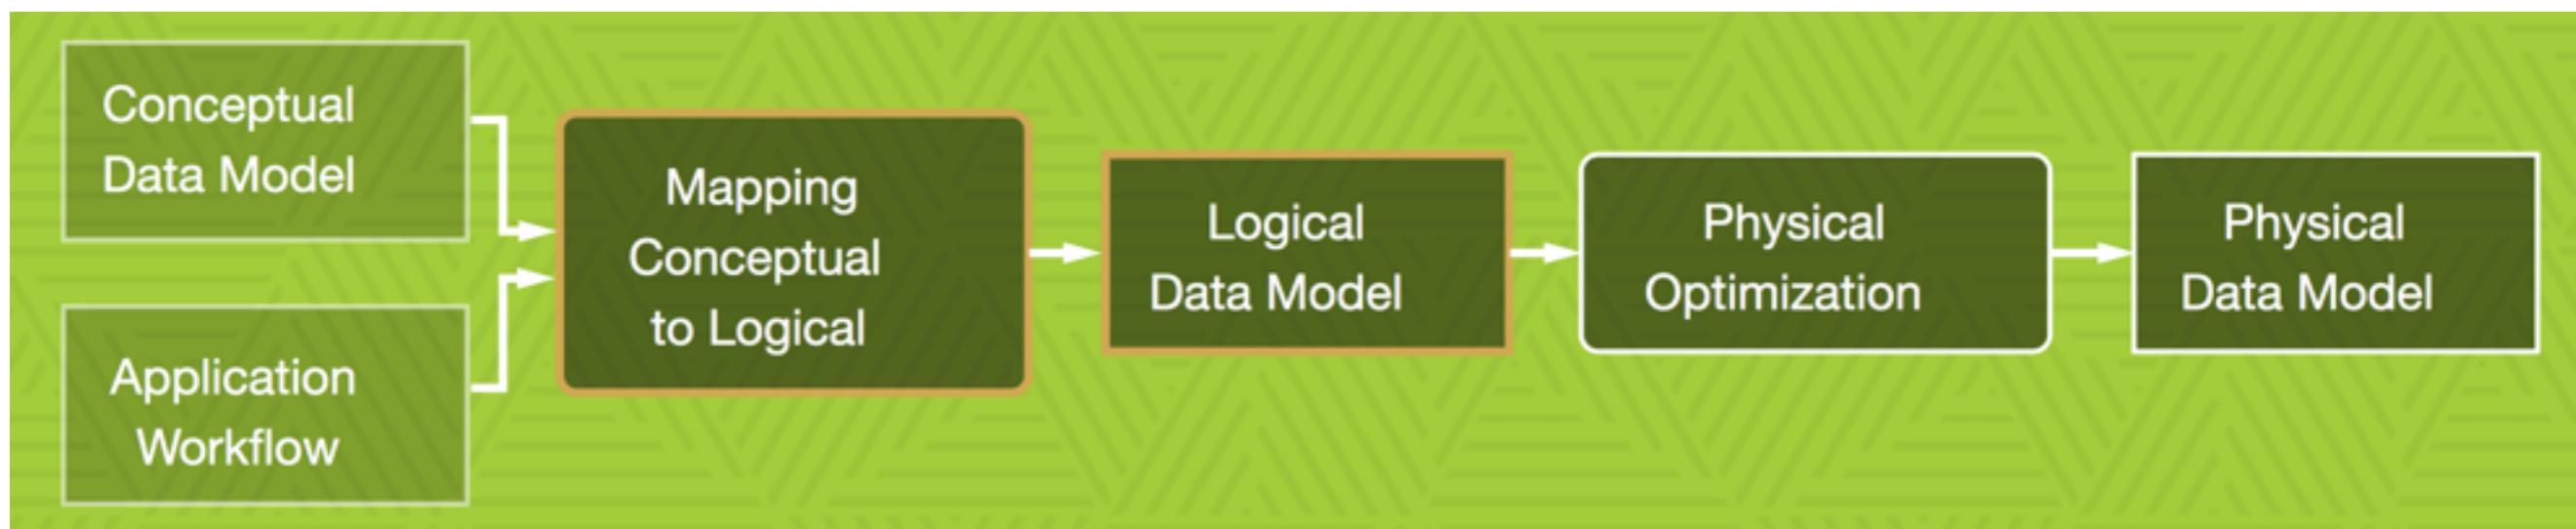
Task: Find a video that has a particular video id



ACCESS PATTERNS

- Q1: Find a user with a **specified email**
- Q2: Find most recently uploaded **videos**
- Q3: Find a **user with a specified id**
- Q4: Find videos uploaded by a **user with a known id**(show most recently uploaded videos first)
- Q5: Find a video with a **specified video id**

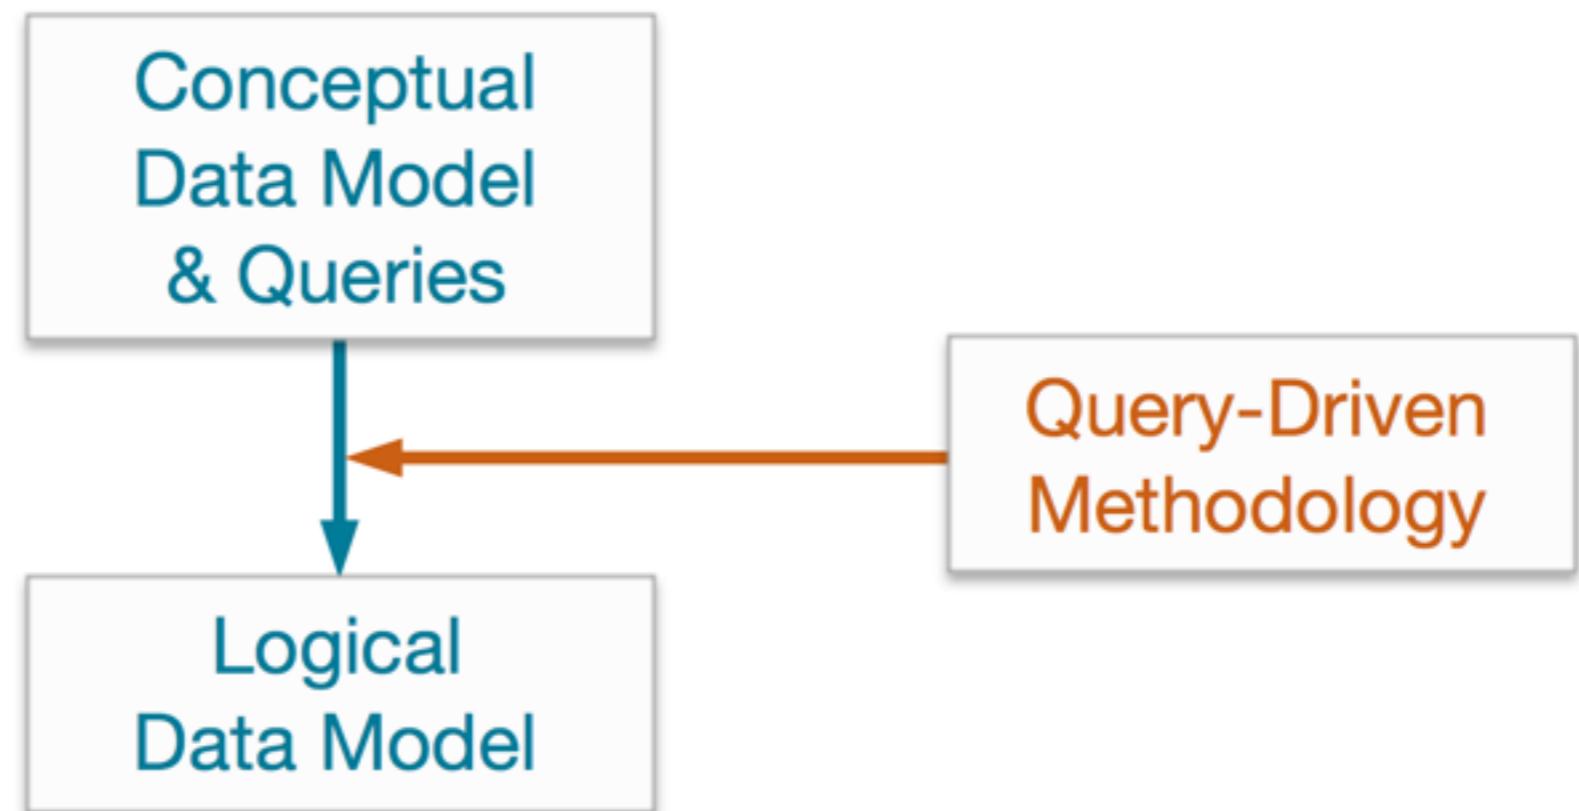
Mapping Conceptual to Logical



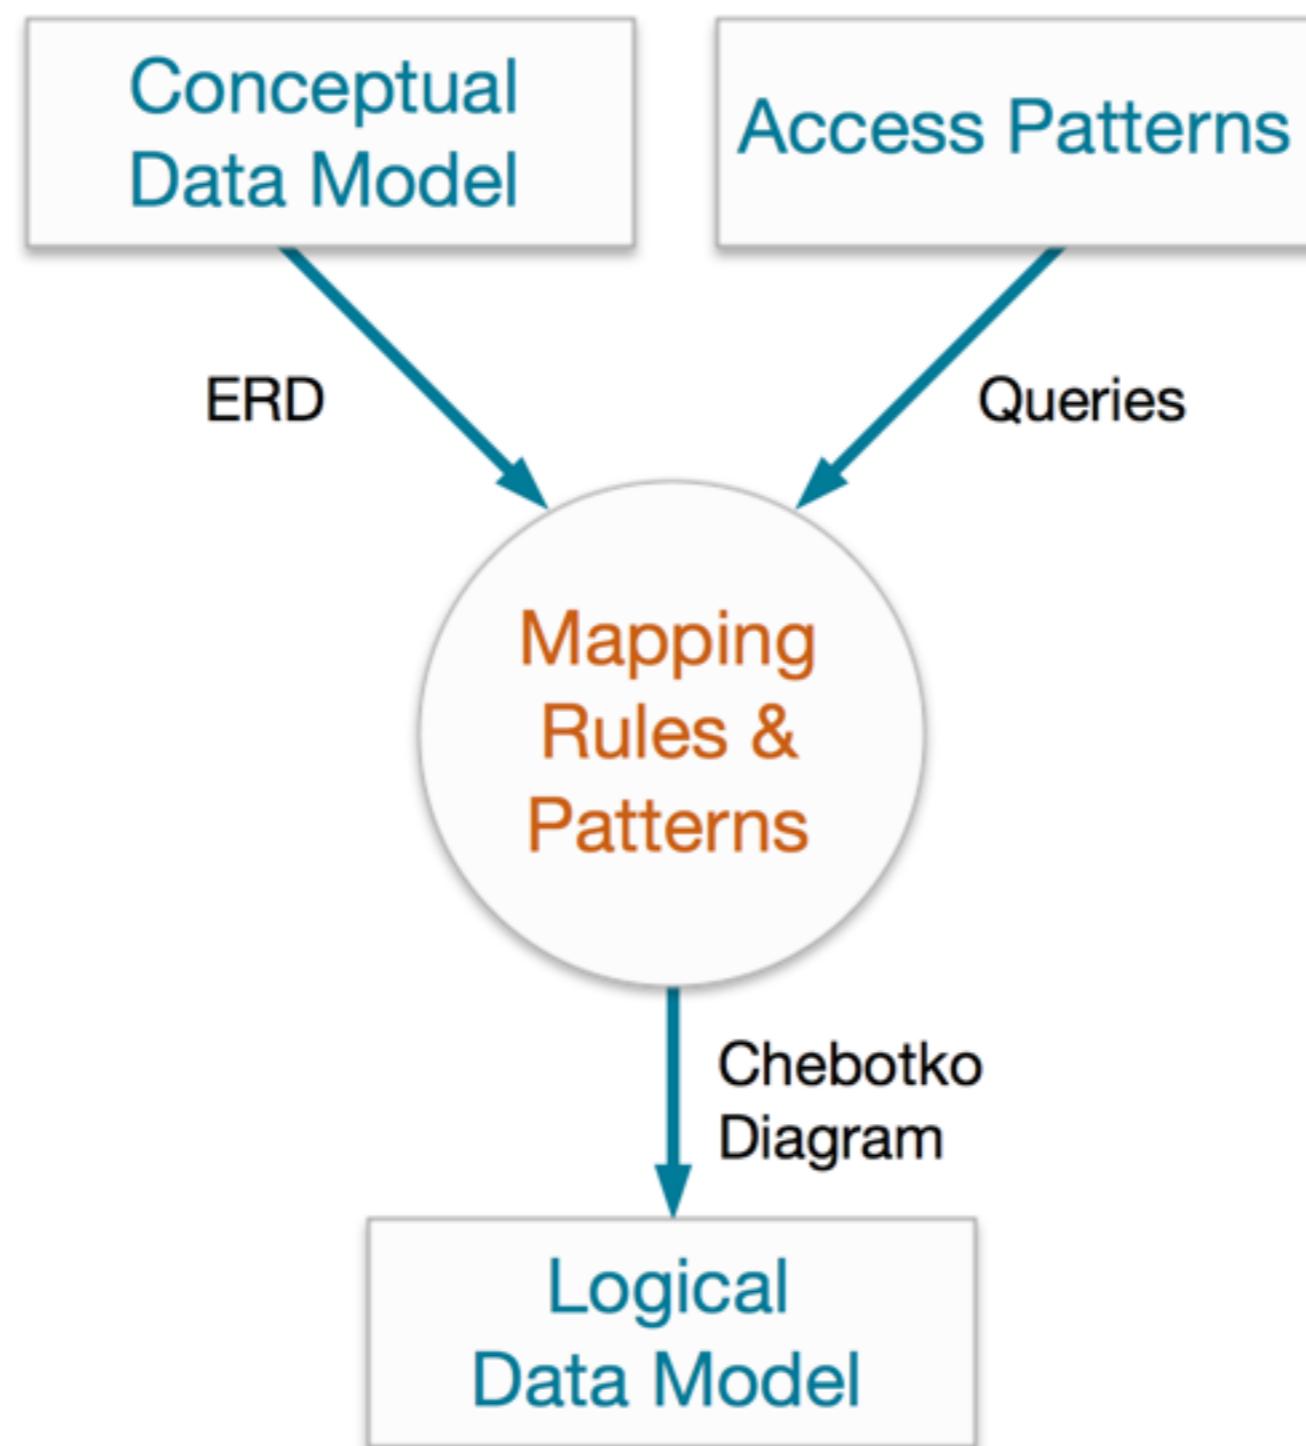
Data Modeling Methodology

Process to design a logical model

- Uses a top-down approach
- Can be algorithmically defined
- Effective in the long run—vs. dark art

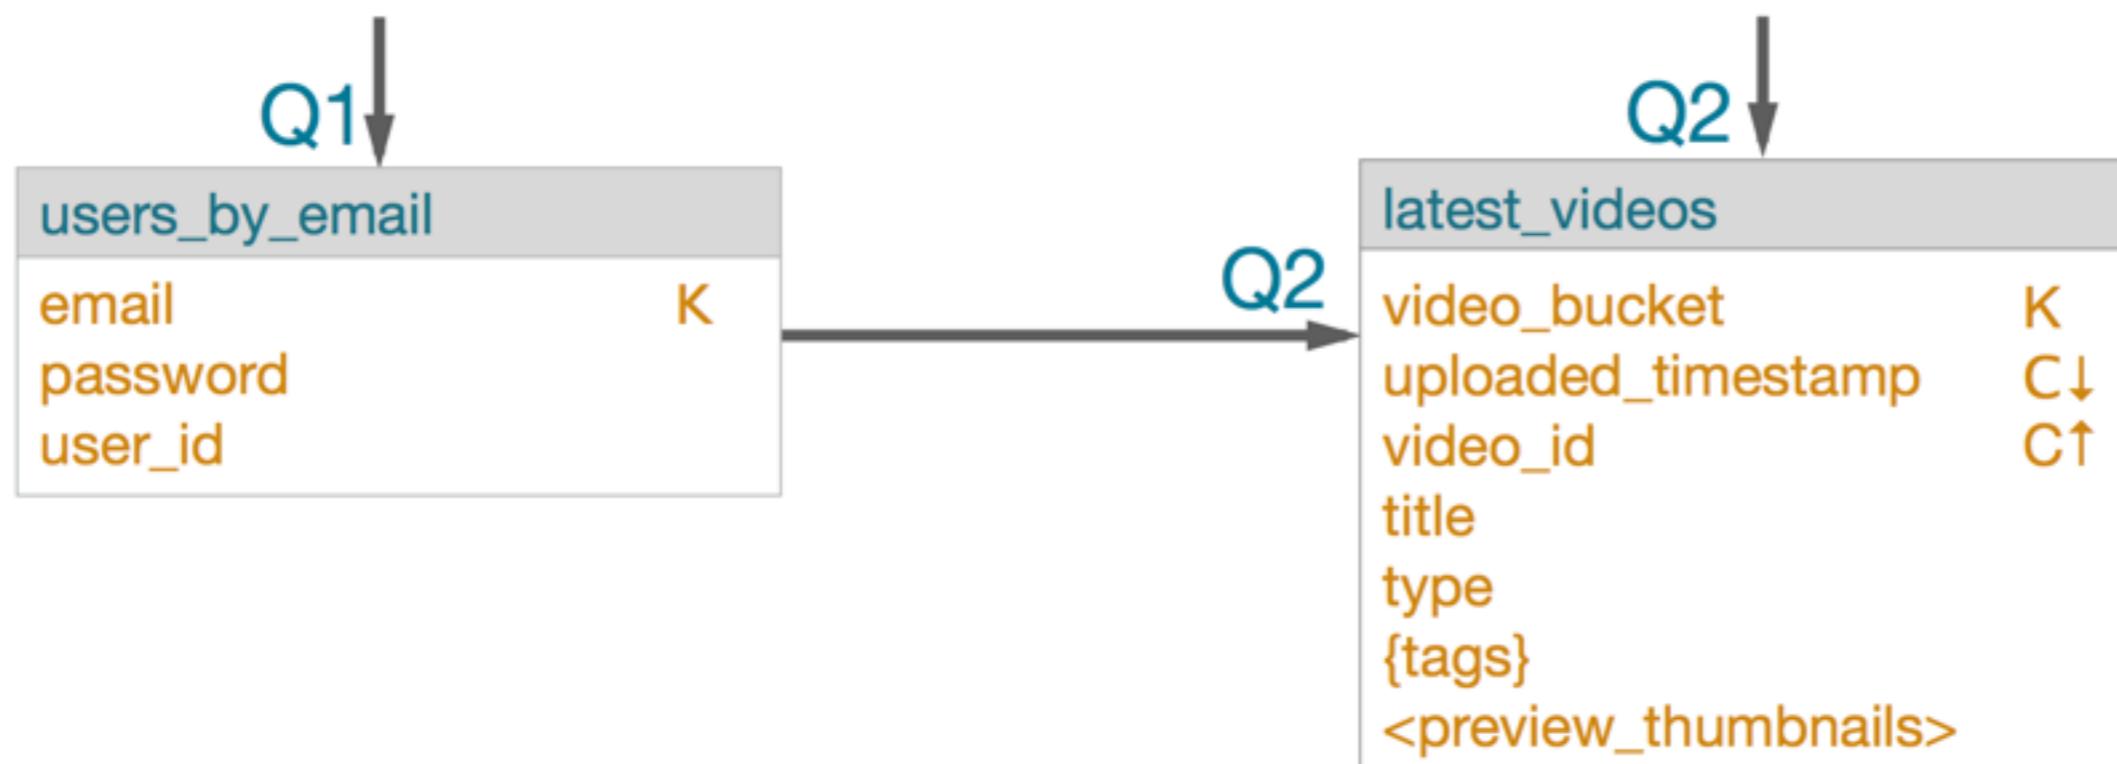


Query-Driven Data Modeling



Chebotko Diagram

- Visual diagram for Cassandra tables and access patterns



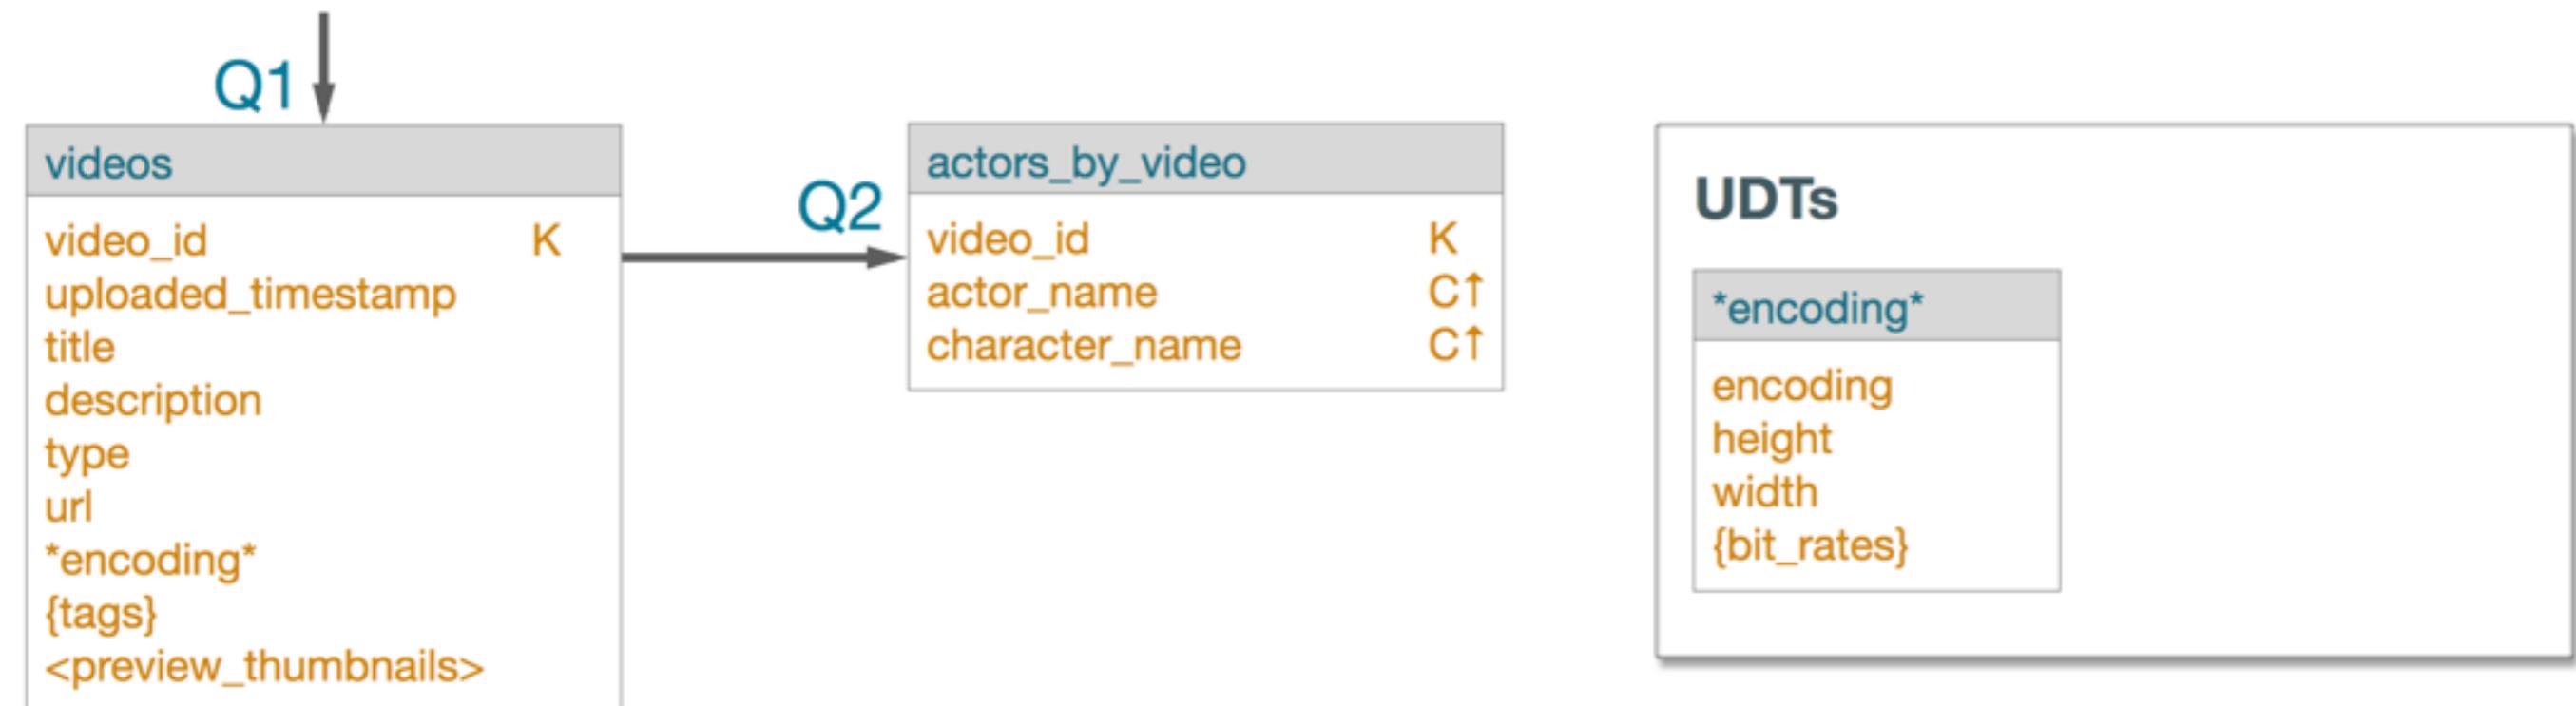
ACCESS PATTERNS

Q1: Find a user with a **specified email**

Q2: Find most recently uploaded videos

Chebotko Diagrams

- Graphical representation of Cassandra database schema design
- Documents the logical and physical data model

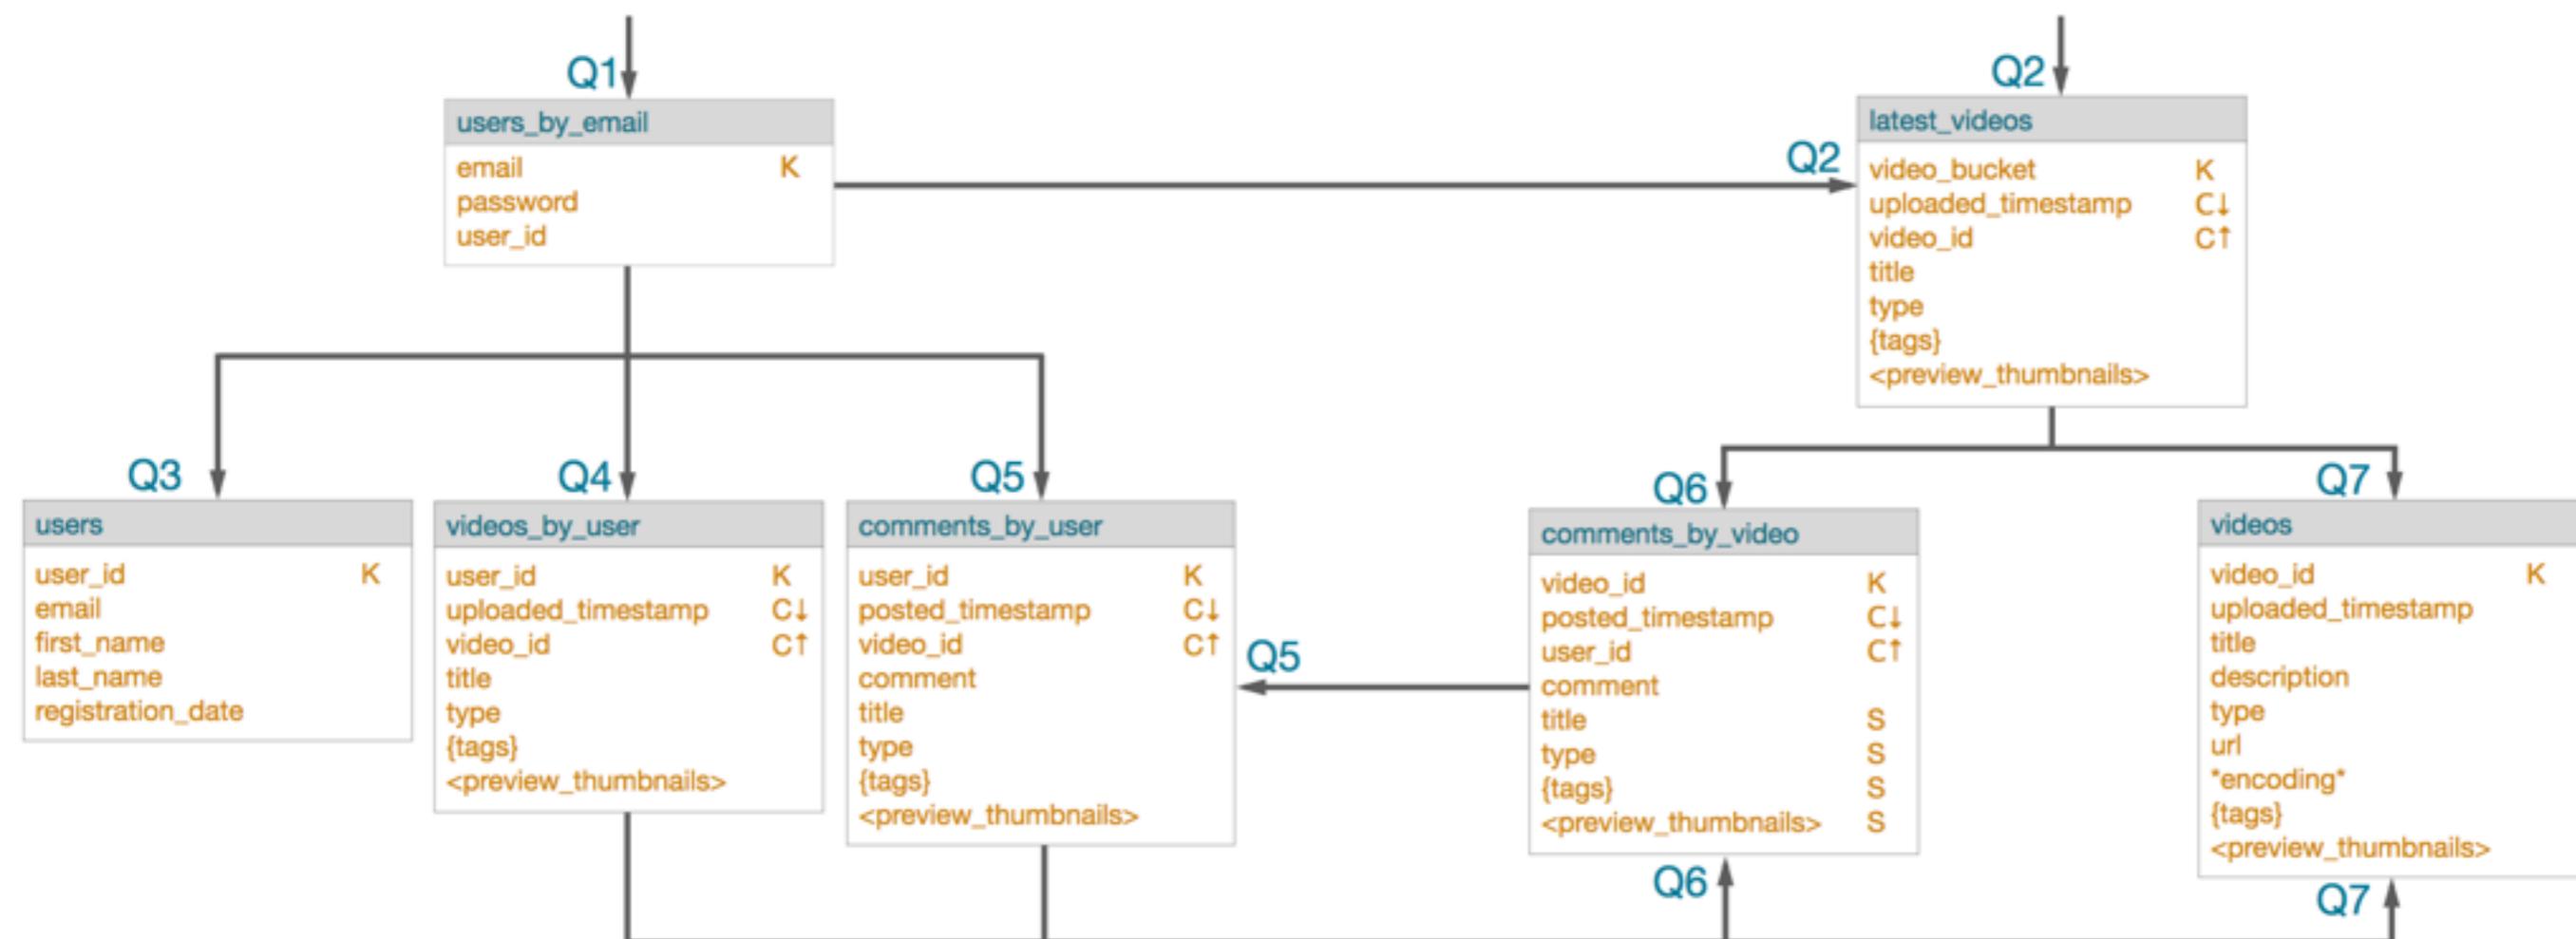


ACCESS PATTERNS

Q1: Find a video with a **specified video id**

Q2: Find actors for a **video with a known id** (show actor names in ascending order)

Example Chebotko Diagram



ACCESS PATTERNS

- Q1: Find a user with a **specified email**
- Q2: Find most recently uploaded videos
- Q3: Find a **user with a specified id**
- Q4: Find videos uploaded by a **user with a known id** (show most recently uploaded videos first)
- Q5: Find comments posted by a **user with a known id** (show most recently posted comments first)
- Q6: Find comments posted for a **video with a known id** (show most recent comments first)
- Q7: Find a video with a **specified video id**

UDTs

encoding
encoding
height
width
{bit_rates}

Duplicate Data

Better to duplicate than to join data

- Partition per query and data nesting may result in data duplication
 - Query results are pre-computed and materialized
 - Data can be duplicated across tables, partitions, and / or rows

videos_by_actor	
actor	K
release_date	C↓
video_id	C↑
title	
type	
{tags}	
<preview_thumbnails>	

videos_by_genre	
genre	K
release_date	C↓
video_id	C↑
title	
type	
{tags}	
<preview_thumbnails>	

videos_by_tag	
tag	K
release_date	C↓
video_id	C↑
title	
type	
{tags}	
<preview_thumbnails>	

Duplicate Data

Data duplication can scale, joins cannot

videos_by_actor		
actor	K	
video_id	C↑	
character_name	C↑	

Q1 →

videos		
video_id	K	
uploaded_timestamp	C↓	
title		
description		
type		
release_date		
{tags}		
<preview_thumbnails>		
{genres}		

Q2 ↓

videos_by_actor		
actor	K	
video_id	C↑	
character_name	C↑	
uploaded_timestamp		
title		
description		
type		
release_date		
{tags}		
<preview_thumbnails>		
{genres}		

Q1 →

VS.

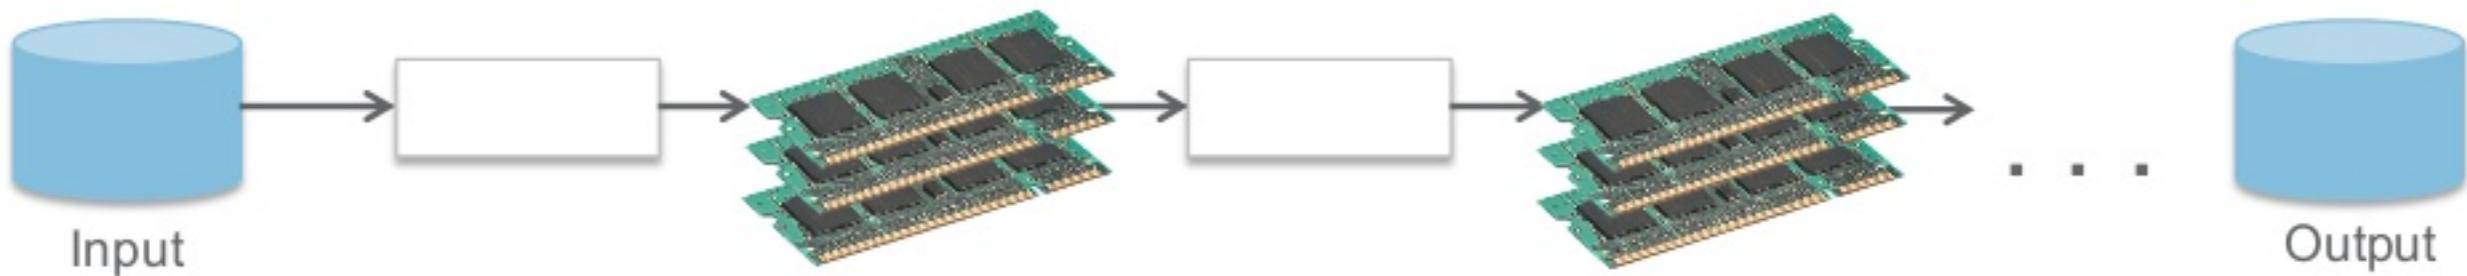
spark

Apache Spark for data in motion

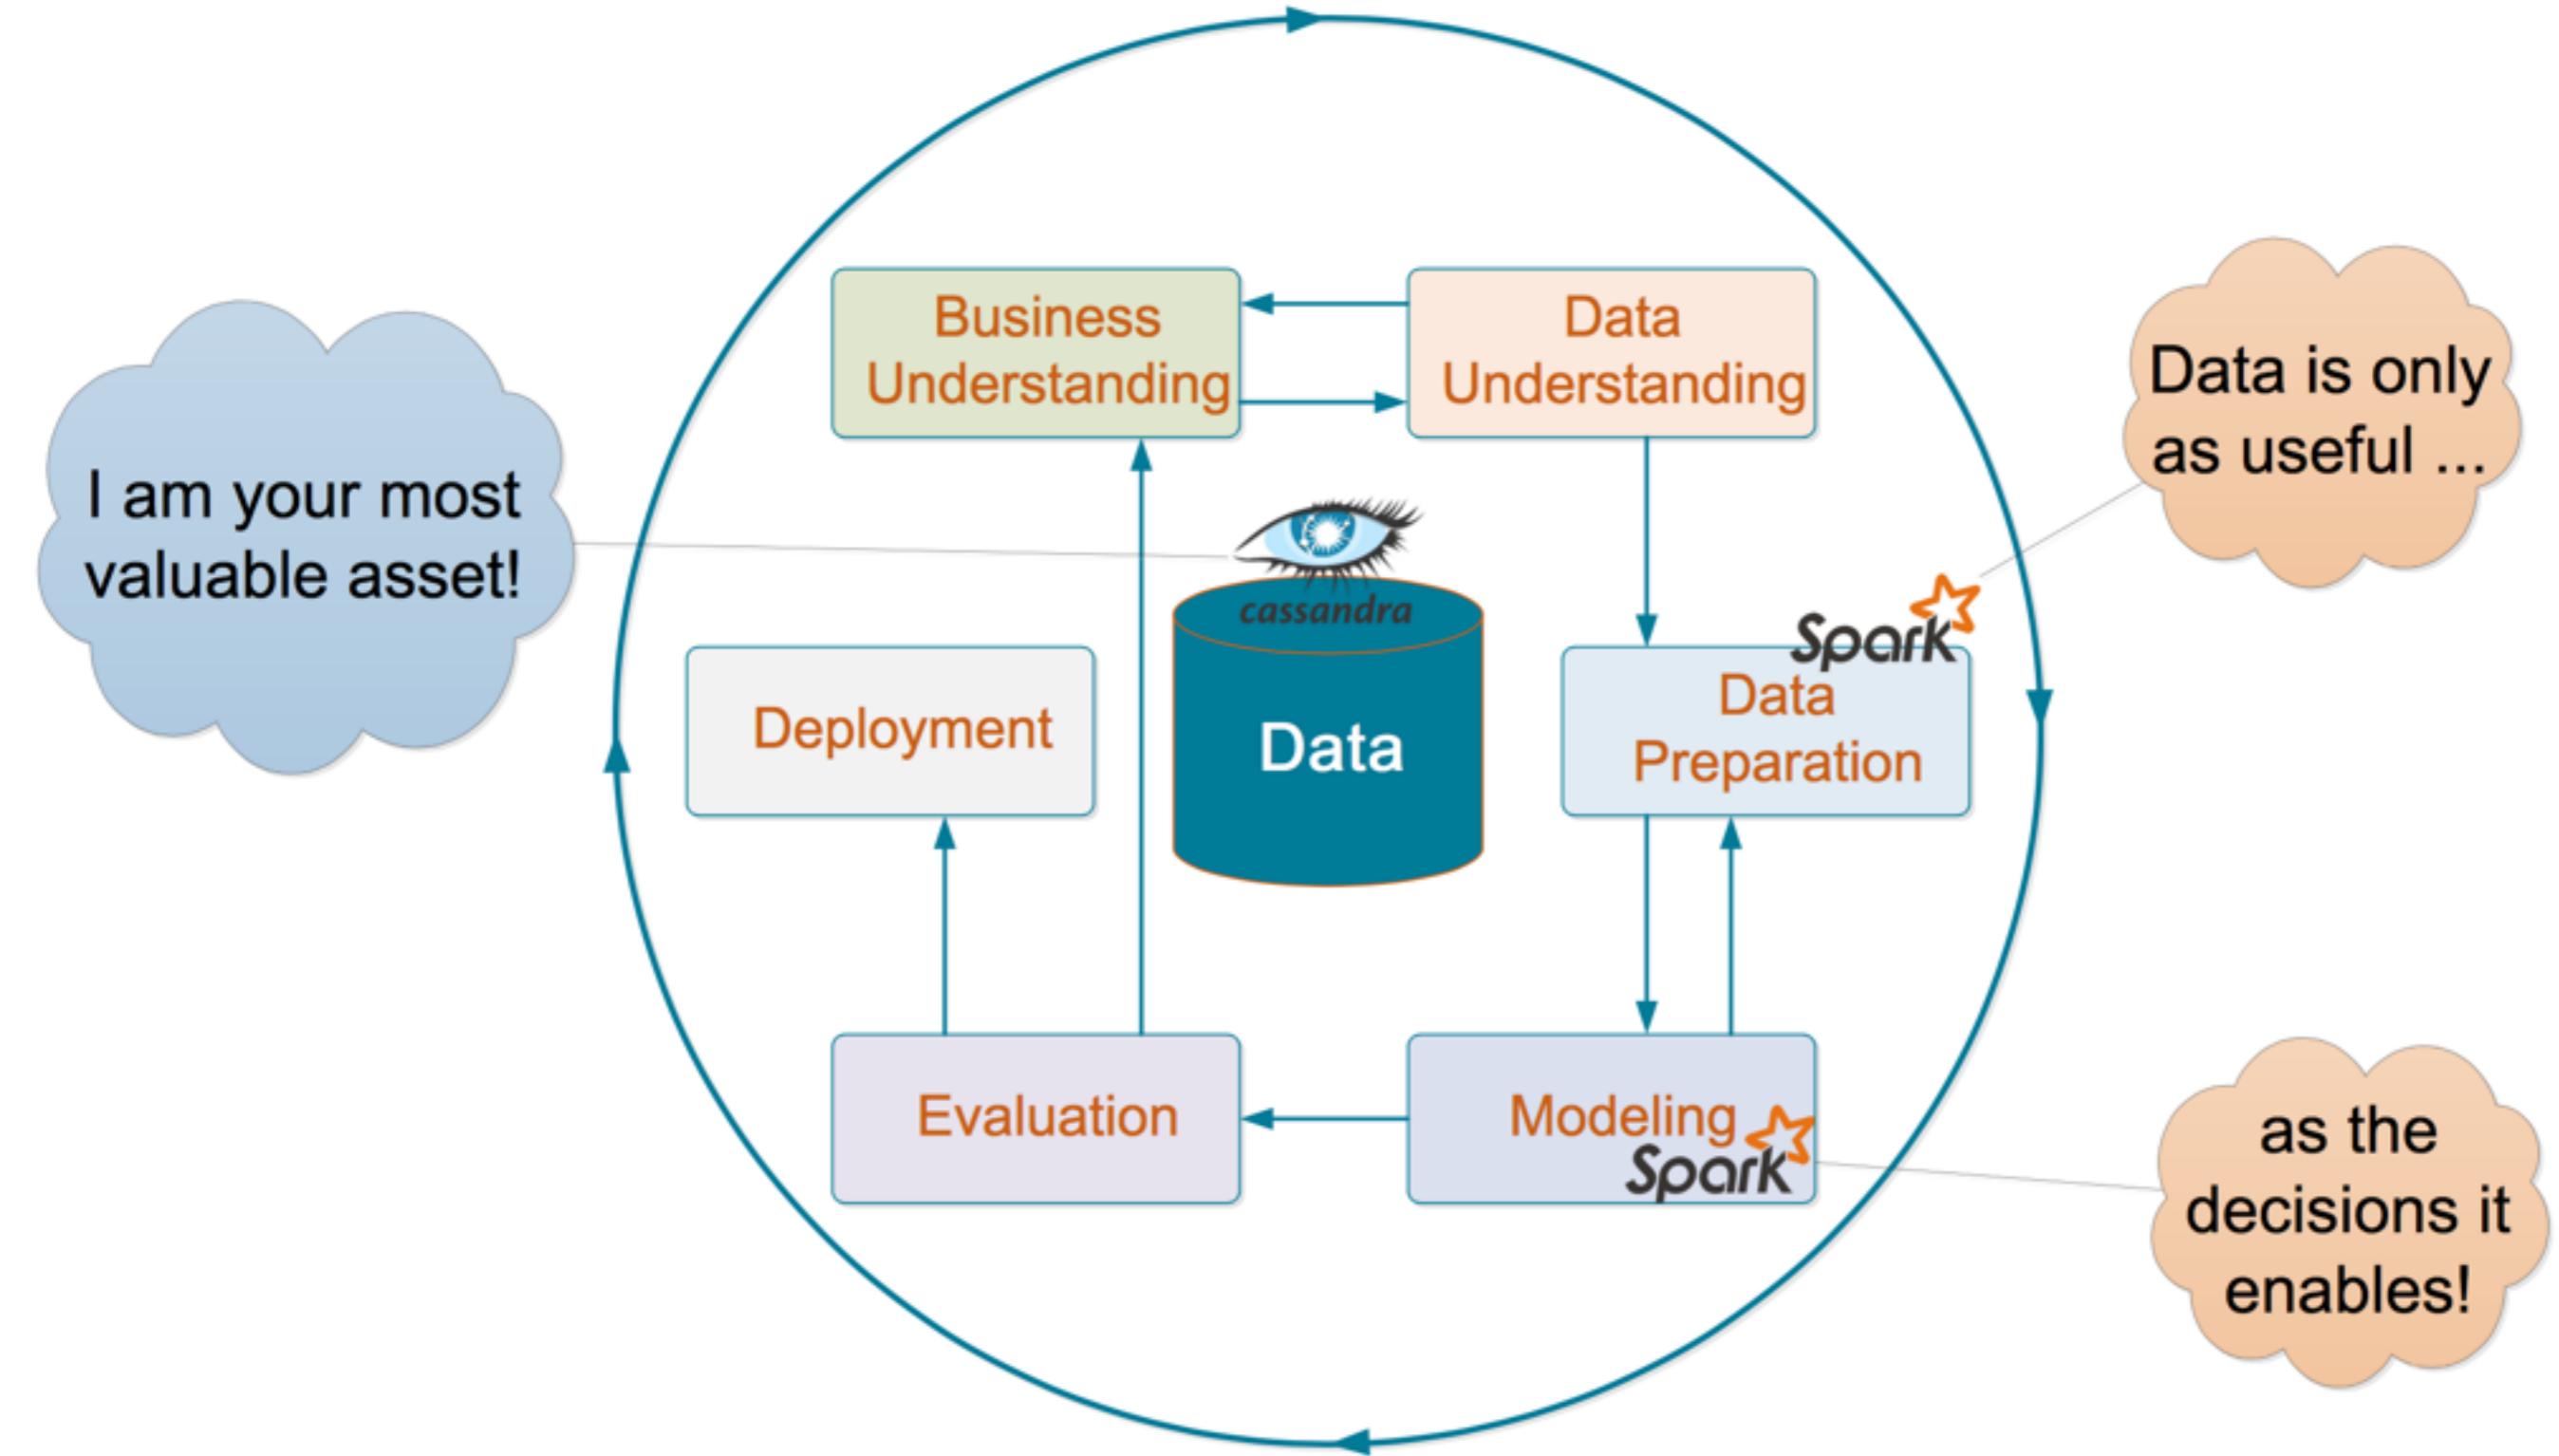
Hadoop MapReduce: Data Sharing on Disk



Spark: Speed up processing by using Memory instead of Disks



spark / cassandra role



Distributed computation engine designed for big data and in-memory processing

- Interactive and batch analytics
- Up to 100x faster than Hadoop
- 5-10x less code than Hadoop

Spark
SQL

Spark
Streaming

MLlib

GraphX

SparkR





COMPANY

[All Posts](#)[Partners](#)

Spark officially sets a new record in large-scale sorting

November 5, 2014 | by Reynold Xin



	Hadoop MR Record	Spark Record	Spark 1 PB
Data Size	102.5 TB	100 TB	1000 TB
Elapsed Time	72 mins	23 mins	234 mins
# Nodes	2100	206	190
# Cores	50400 physical	6592 virtualized	6080 virtualized
Cluster disk throughput	3150 GB/s (est.)	618 GB/s	570 GB/s
Sort Benchmark Daytona Rules	Yes	Yes	No
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network	virtualized (EC2) 10Gbps network
Sort rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Sort rate/node	0.67 GB/min	20.7 GB/min	22.5 GB/min

big data

```
public class WordCount {  
    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(LongWritable key, Text value, Context context) throws IOException,  
IOException, InterruptedException {  
            String line = value.toString();  
            StringTokenizer tokenizer = new StringTokenizer(line);  
            while (tokenizer.hasMoreTokens()) {  
                word.set(tokenizer.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
  
    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
        public void reduce(Text key, Iterable<IntWritable> values, Context context)  
throws IOException, InterruptedException {  
            int sum = 0;  
            for (IntWritable val : values) {  
                sum += val.get();  
            }  
            context.write(key, new IntWritable(sum));  
        }  
    }  
  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
  
        Job job = new Job(conf, "wordcount");  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        job.setMapperClass(Map.class);  
        job.setReducerClass(Reduce.class);  
  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.waitForCompletion(true);  
    }  
}
```

big data

```
public class WordCount {  
    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
        private final static IntWritable one = new IntWritable(1);  
        private Text word = new Text();  
  
        public void map(LongWritable key, Text value, Context context) throws IOException,  
IOException, InterruptedException {  
            String line = value.toString();  
            StringTokenizer tokenizer = new StringTokenizer(line);  
            while (tokenizer.hasMoreTokens()) {  
                word.set(tokenizer.nextToken());  
                context.write(word, one);  
            }  
        }  
    }  
  
    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
        public void reduce(Text key, Iterable<IntWritable> values, Context context)  
throws IOException, InterruptedException {  
            int sum = 0;  
            for (IntWritable val : values) {  
                sum += val.get();  
            }  
            context.write(key, new IntWritable(sum));  
        }  
    }  
  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
  
        Job job = new Job(conf, "wordcount");  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        job.setMapperClass(Map.class);  
        job.setReducerClass(Reduce.class);  
  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.waitForCompletion(true);  
    }  
}
```

```
object WordCount{  
  
    def main(def main(args: Array[String])){  
  
        val sparkConf = new SparkConf()  
  
            .setAppName("wordcount")  
  
        val sc = new SparkContext(sparkConf)  
  
        sc.textFile(args(0))  
            .flatMap(_.split(" "))+  
            .countByValue  
            .saveAsTextFile(args(1))  
    }  
}
```

Why spark?

- Readability
- Expressiveness
- Fast
- Testability
- Interactive
- Fault Tolerant
- Unify Big Data





DataFrames

Spark
SQL

Spark
Streaming

MLlib
(machine
learning)

GraphX
(graph)

Spark Core

Who is using spark?



Yahoo!

CONVIVA

Goldman
Sachs

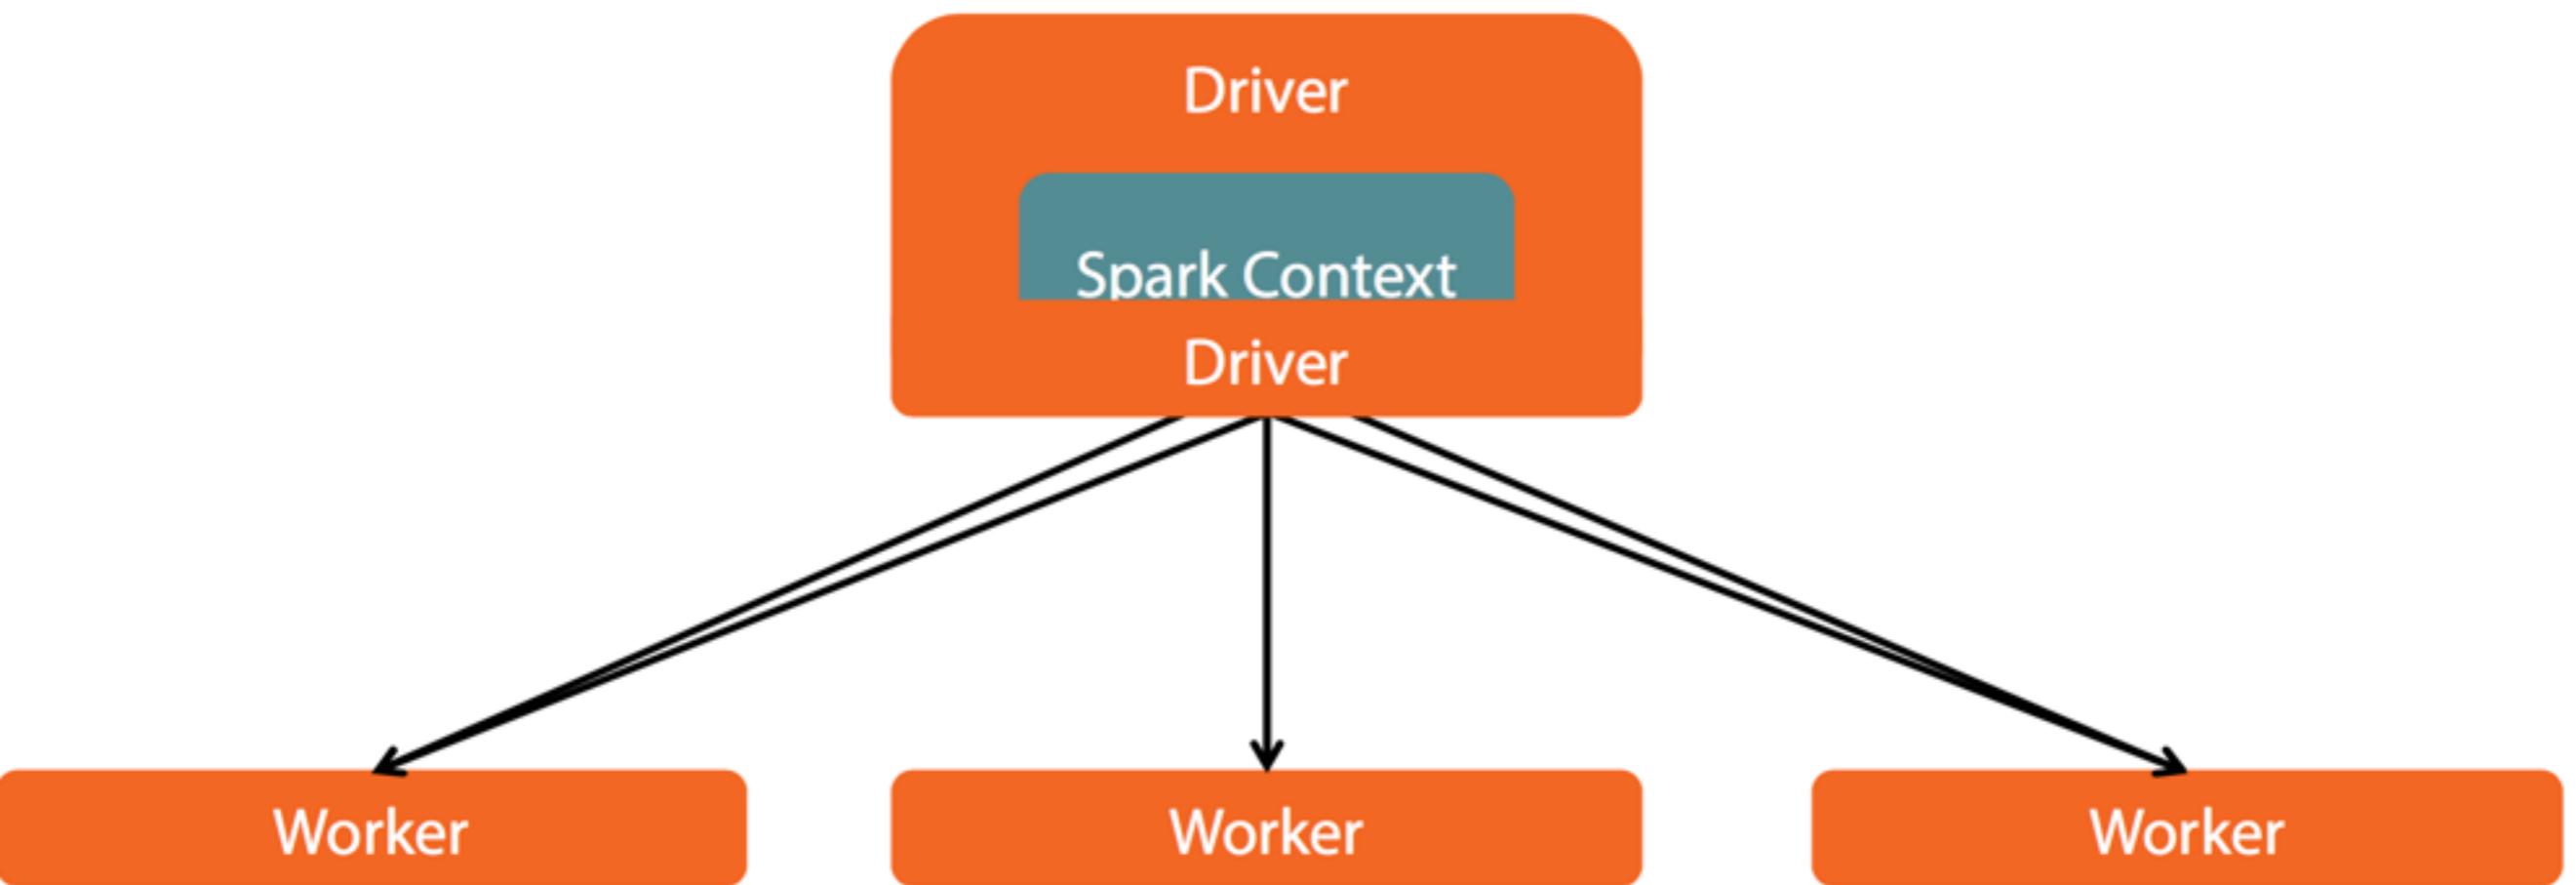
PANDORA

NETFLIX

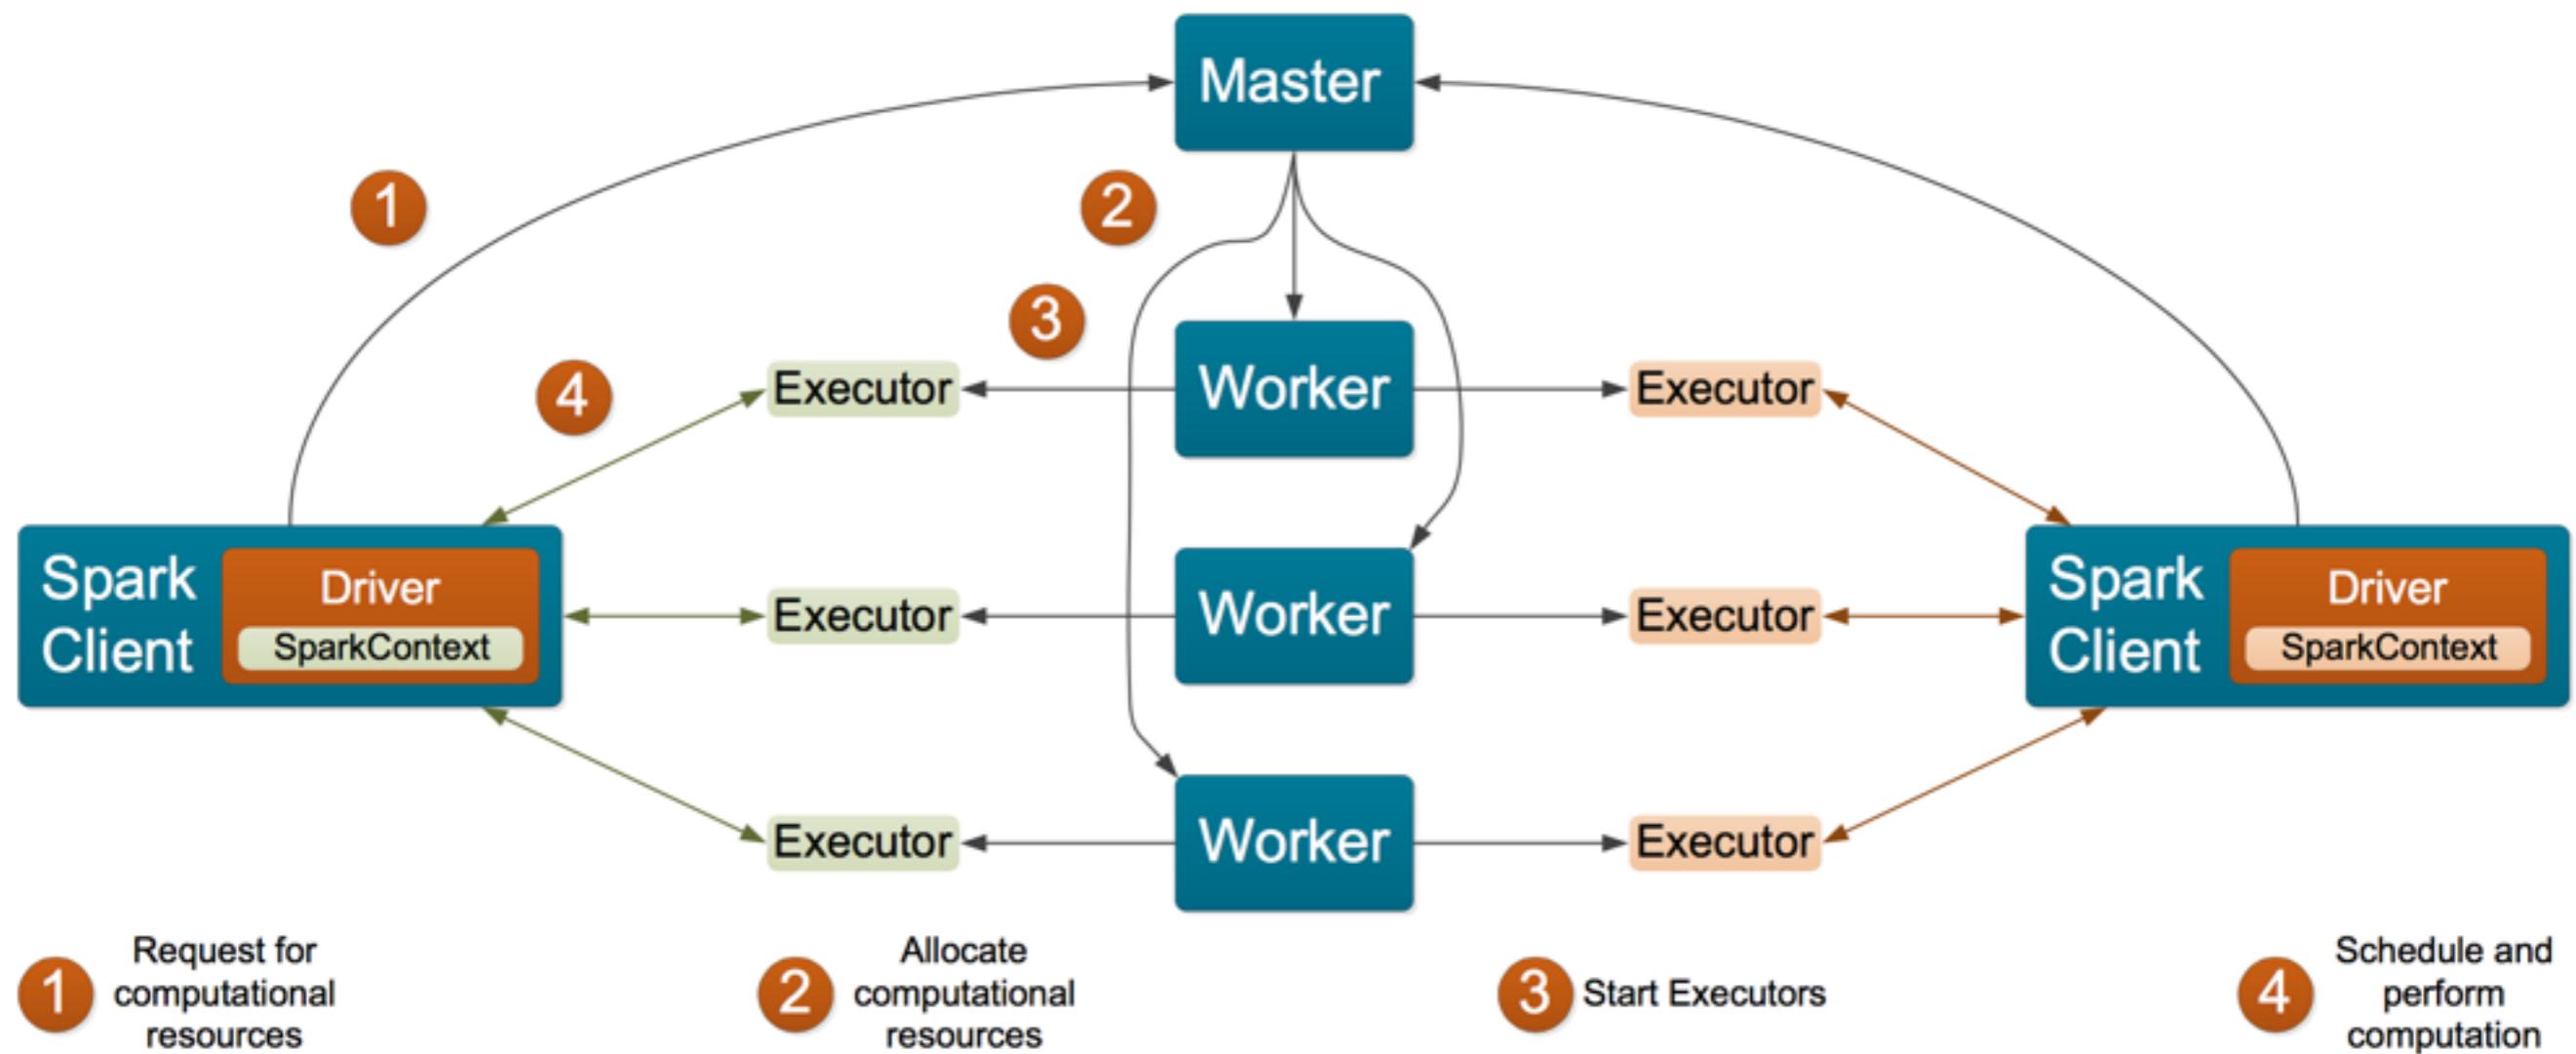
comcast

ebay

hhmi | janelia
Research Campus



spark architecture



Spark download



Download Libraries ▾ Documentation ▾ Examples Community ▾ FAQ

Download Apache Spark™

Our latest stable version is Apache Spark 1.6.2, released on June 25, 2016 ([release notes](#)) ([git tag](#))

1. Choose a Spark release:
2. Choose a package type:
3. Choose a download type:
4. Download Spark: [spark-1.6.2-bin-hadoop2.6.tgz](#)
5. Verify this release using the [1.6.2 signatures and checksums](#) and [project release KEYS](#).

Note: Scala 2.11 users should download the Spark source package and build [with Scala 2.11 support](#).

<http://www-us.apache.org/dist/spark/spark-1.6.2/spark-1.6.2-bin-hadoop2.6.tgz>

Add `export SPARK_LOCAL_IP="127.0.0.1"` to `load-spark-env.sh`

 .DS_Store	
 bin	▶
 CHANGES.txt	▶
 conf	▶
 count.txt	▶
 count3	▶
 data	▶
 derby.log	▶
 ec2	▶
 examples	▶
 lib	▶
 LICENSE	▶
 licenses	▶
 metastore_db	▶
 NOTICE	▶
 python	▶
 R	▶
 README.md	▶
 RELEASE	▶
 sbin	▶
	 beeline
	 beeline.cmd
	 load-spark-env.cmd
	 load-spark-env.sh
	 pyspark
	 pyspark.cmd
	 pyspark2.cmd
	 run-example
	 run-example.cmd
	 run-example2.cmd
	 spark-class
	 spark-class.cmd
	 spark-class2.cmd
	 spark-shell
	 spark-shell.cmd
	 spark-shell2.cmd
	 spark-sql
	 spark-submit
	 spark-submit.cmd
	 spark-submit2.cmd
	 sparkR
	 sparkR.cmd
	 sparkR2.cmd

Spark shell

- `./bin/spark-shell --master local[2]`
- Interactive Scala-REPL-based Spark client

Hello spark

- scala> `println("Hello NFJS Spark!")`
- Hello NFJS Spark!
- scala> `val seven = 2 + 5`
- `seven: Int = 7`

spark context : sc

```
scala> sc.  
accumulable  
accumulator  
addJar  
appName  
applicationId  
binaryFiles  
broadcast  
cancelJobGroup  
clearFiles  
clearJobGroup  
defaultMinSplits  
emptyRDD  
files  
getCheckpointDir  
getExecutorMemoryStatus  
getLocalProperty  
getPoolForName  
getSchedulingMode  
hadoopFile  
initLocalProperties  
isLocal  
jars  
killExecutors  
master  
newAPIHadoopFile  
objectFile  
range  
runApproximateJob  
sequenceFile  
setCheckpointDir  
setJobGroup  
setLogLevel  
startTime  
stop  
tachyonFolderName  
toString  
version  
  
accumulableCollection  
addFile  
addSparkListener  
applicationAttemptId  
asInstanceOf  
binaryRecords  
cancelAllJobs  
clearCallSite  
clearJars  
defaultMinPartitions  
defaultParallelism  
externalBlockStoreFolderName  
getAllPools  
getConf  
getExecutorStorageStatus  
getPersistentRDDs  
getRDDStorageInfo  
hadoopConfiguration  
hadoopRDD  
isInstanceOf  
isStopped  
killExecutor  
makeRDD  
metricsSystem  
newAPIHadoopRDD  
parallelize  
requestExecutors  
runJob  
setCallSite  
setJobDescription  
setLocalProperty  
sparkUser  
statusTracker  
submitJob  
textFile  
union  
wholeTextFiles
```

Word Count - to directory

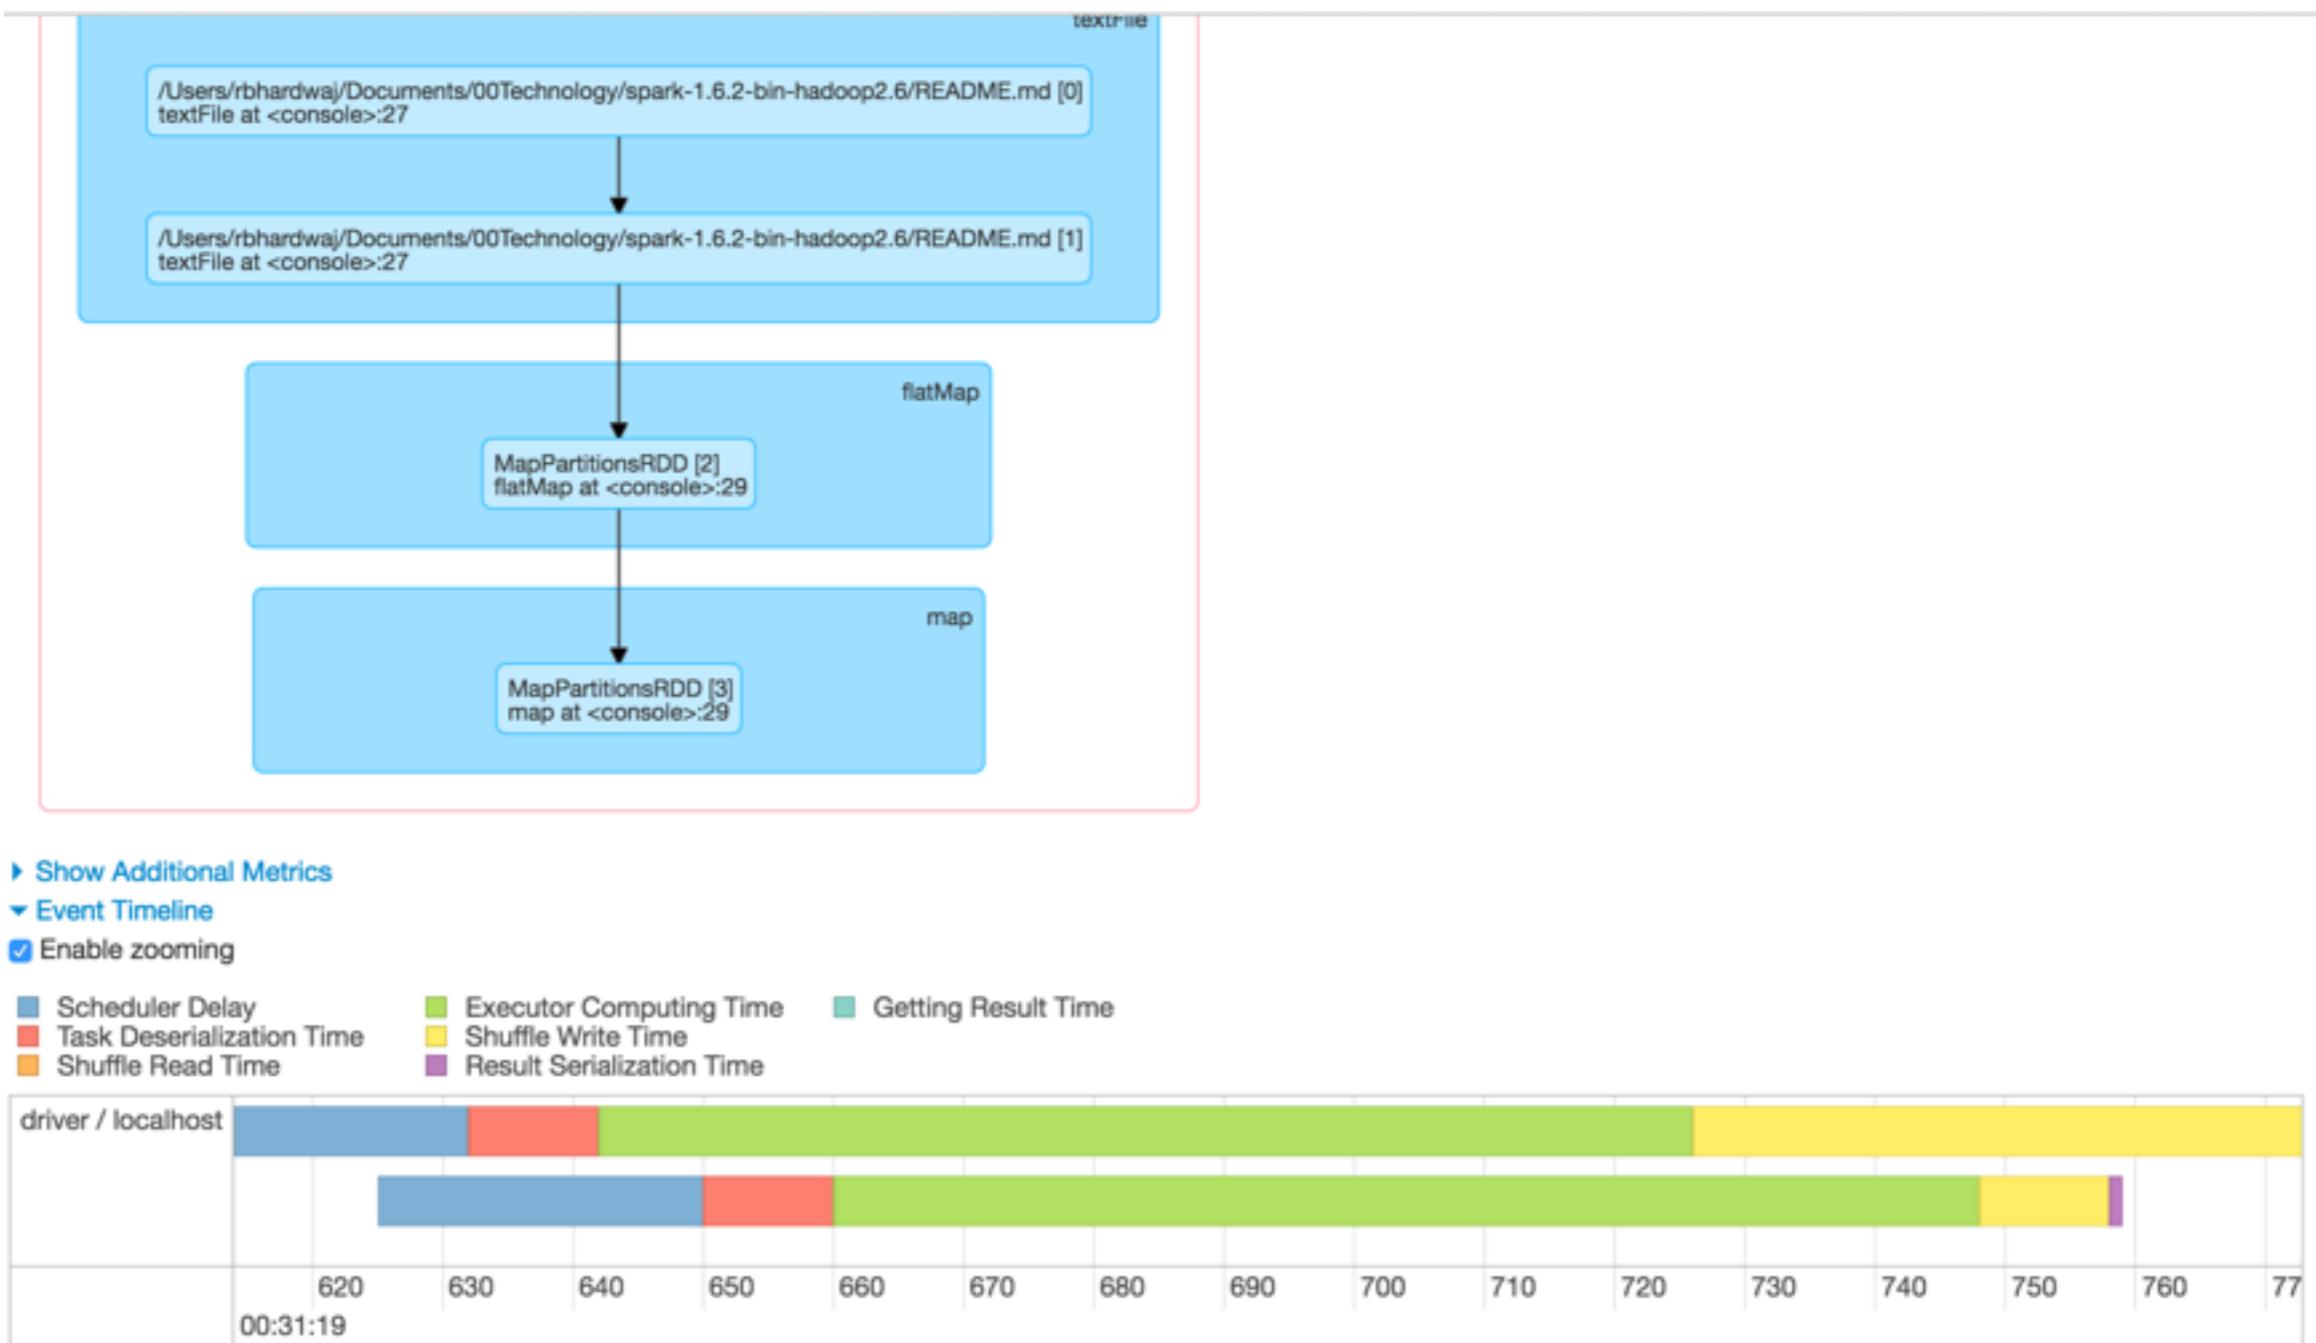
- :paste
- // Entering paste mode (ctrl-D to finish)
- val textFile = sc.textFile("README.md")
- val counts = textFile
- .flatMap(line => line.split(" "))
- .map(word => (word, 1))
- .reduceByKey(_ + _)
- counts.saveAsTextFile("count5")
- // Exiting paste mode, now interpreting.

Application web ui

The screenshot shows the Spark 1.6.2 web UI interface. At the top, there is a header bar with a back arrow, forward arrow, and refresh button. The URL is 'localhost:4040/jobs/'. On the right side of the header are various icons for navigation and settings. Below the header, the Spark logo is displayed with the version '1.6.2'. To the right of the logo are tabs for 'Jobs', 'Stages', 'Storage', 'Environment', 'Executors', and 'SQL'. The 'Jobs' tab is currently selected. On the far right of the header, it says 'Spark shell application UI'. The main content area is titled 'Spark Jobs [\(?\)](#)'. It displays system statistics: 'Total Uptime: 16 min', 'Scheduling Mode: FIFO', and 'Completed Jobs: 1'. There is a link to 'Event Timeline'. Below this, a section titled 'Completed Jobs (1)' lists one job. The table has columns: 'Job Id', 'Description', 'Submitted', 'Duration', 'Stages: Succeeded/Total', and 'Tasks (for all stages): Succeeded/Total'. The job listed is '0' with the description 'saveAsTextFile at <console>:32', submitted on '2016/07/17 00:31:19', duration '0.6 s', stages '2/2', and tasks '4/4'.

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	saveAsTextFile at <console>:32	2016/07/17 00:31:19	0.6 s	2/2	4/4

stages



Wordcount comma separated

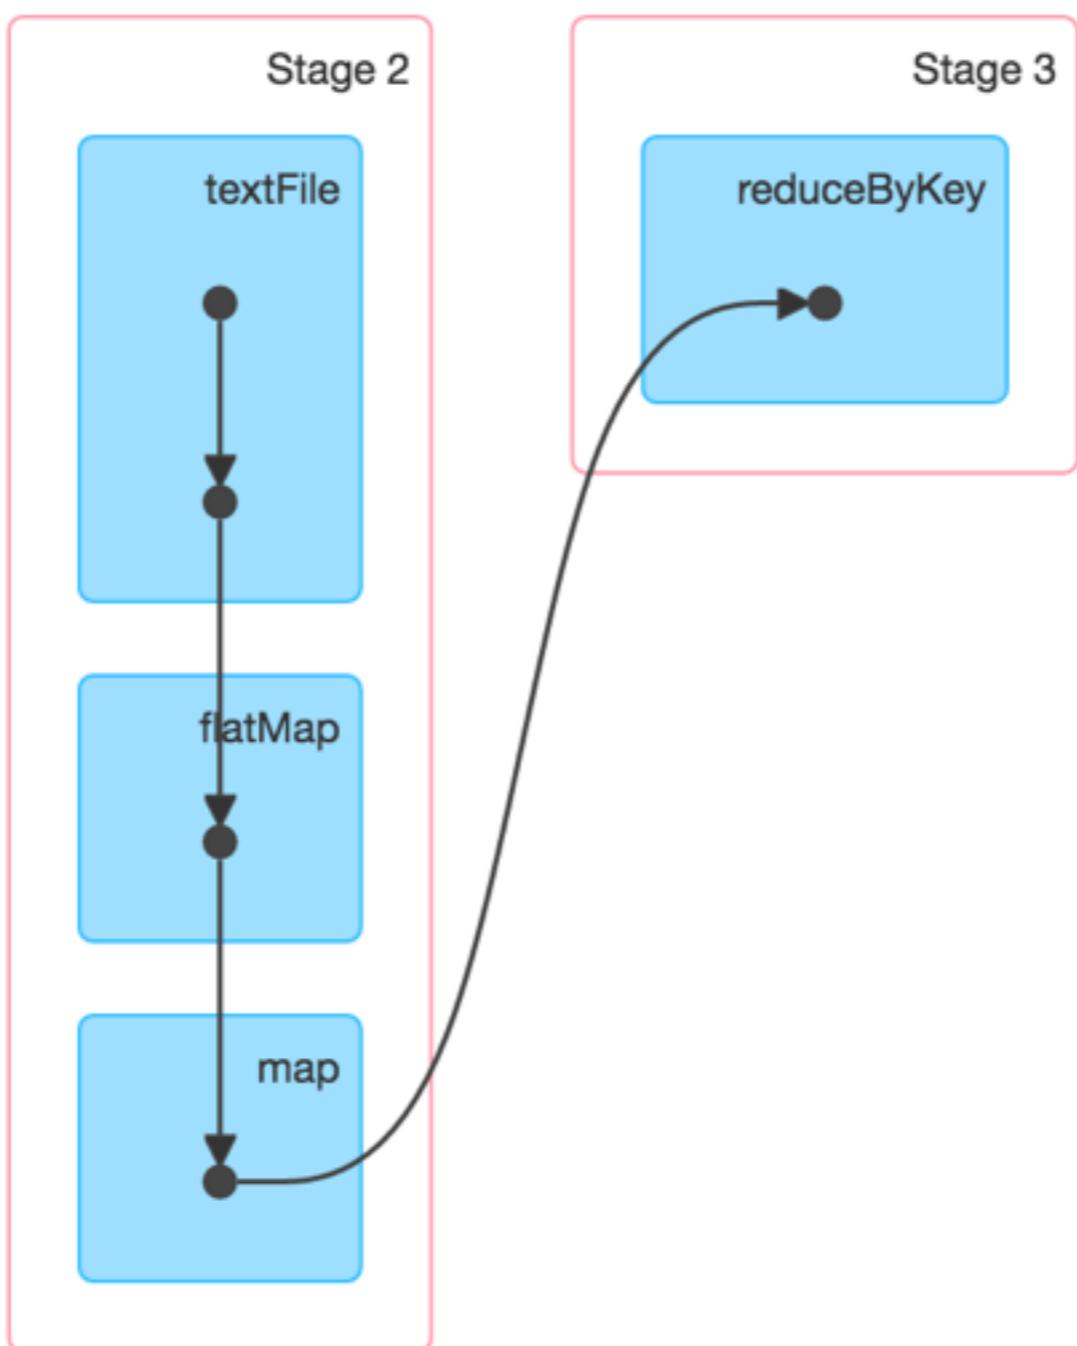
- :paste
- // Entering paste mode (ctrl-D to finish)
- sc.textFile("README.md")
- .flatMap(line => line.split(",").drop(1))
- .map(word => (word,1))
- .reduceByKey{case (x,y) => x + y}
- .collect()
- .foreach(println)
- // Exiting paste mode, now interpreting.

Details for Job 1

Status: SUCCEEDED

Completed Stages: 2

- ▶ Event Timeline
- ▼ DAG Visualization

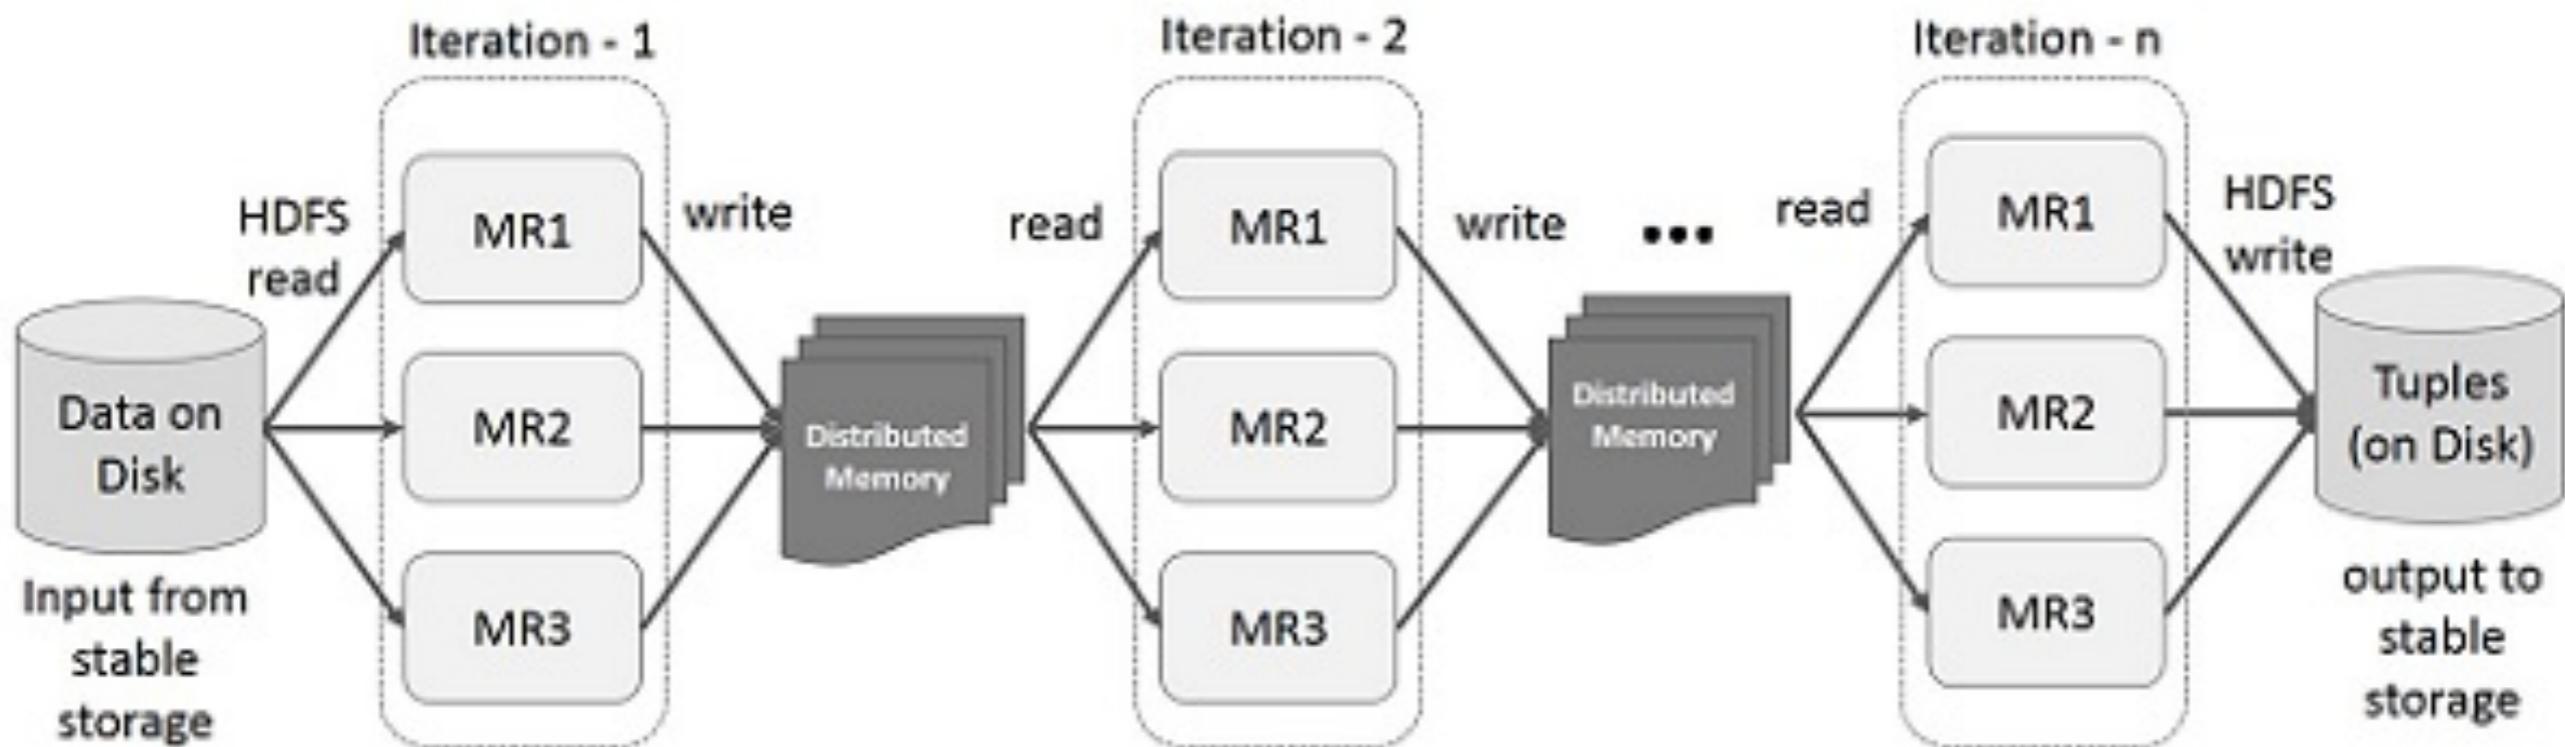


Resilient Distributed Datasets (RDD)

- Immutable
- Re-computable
- Fault tolerant
- Reusable

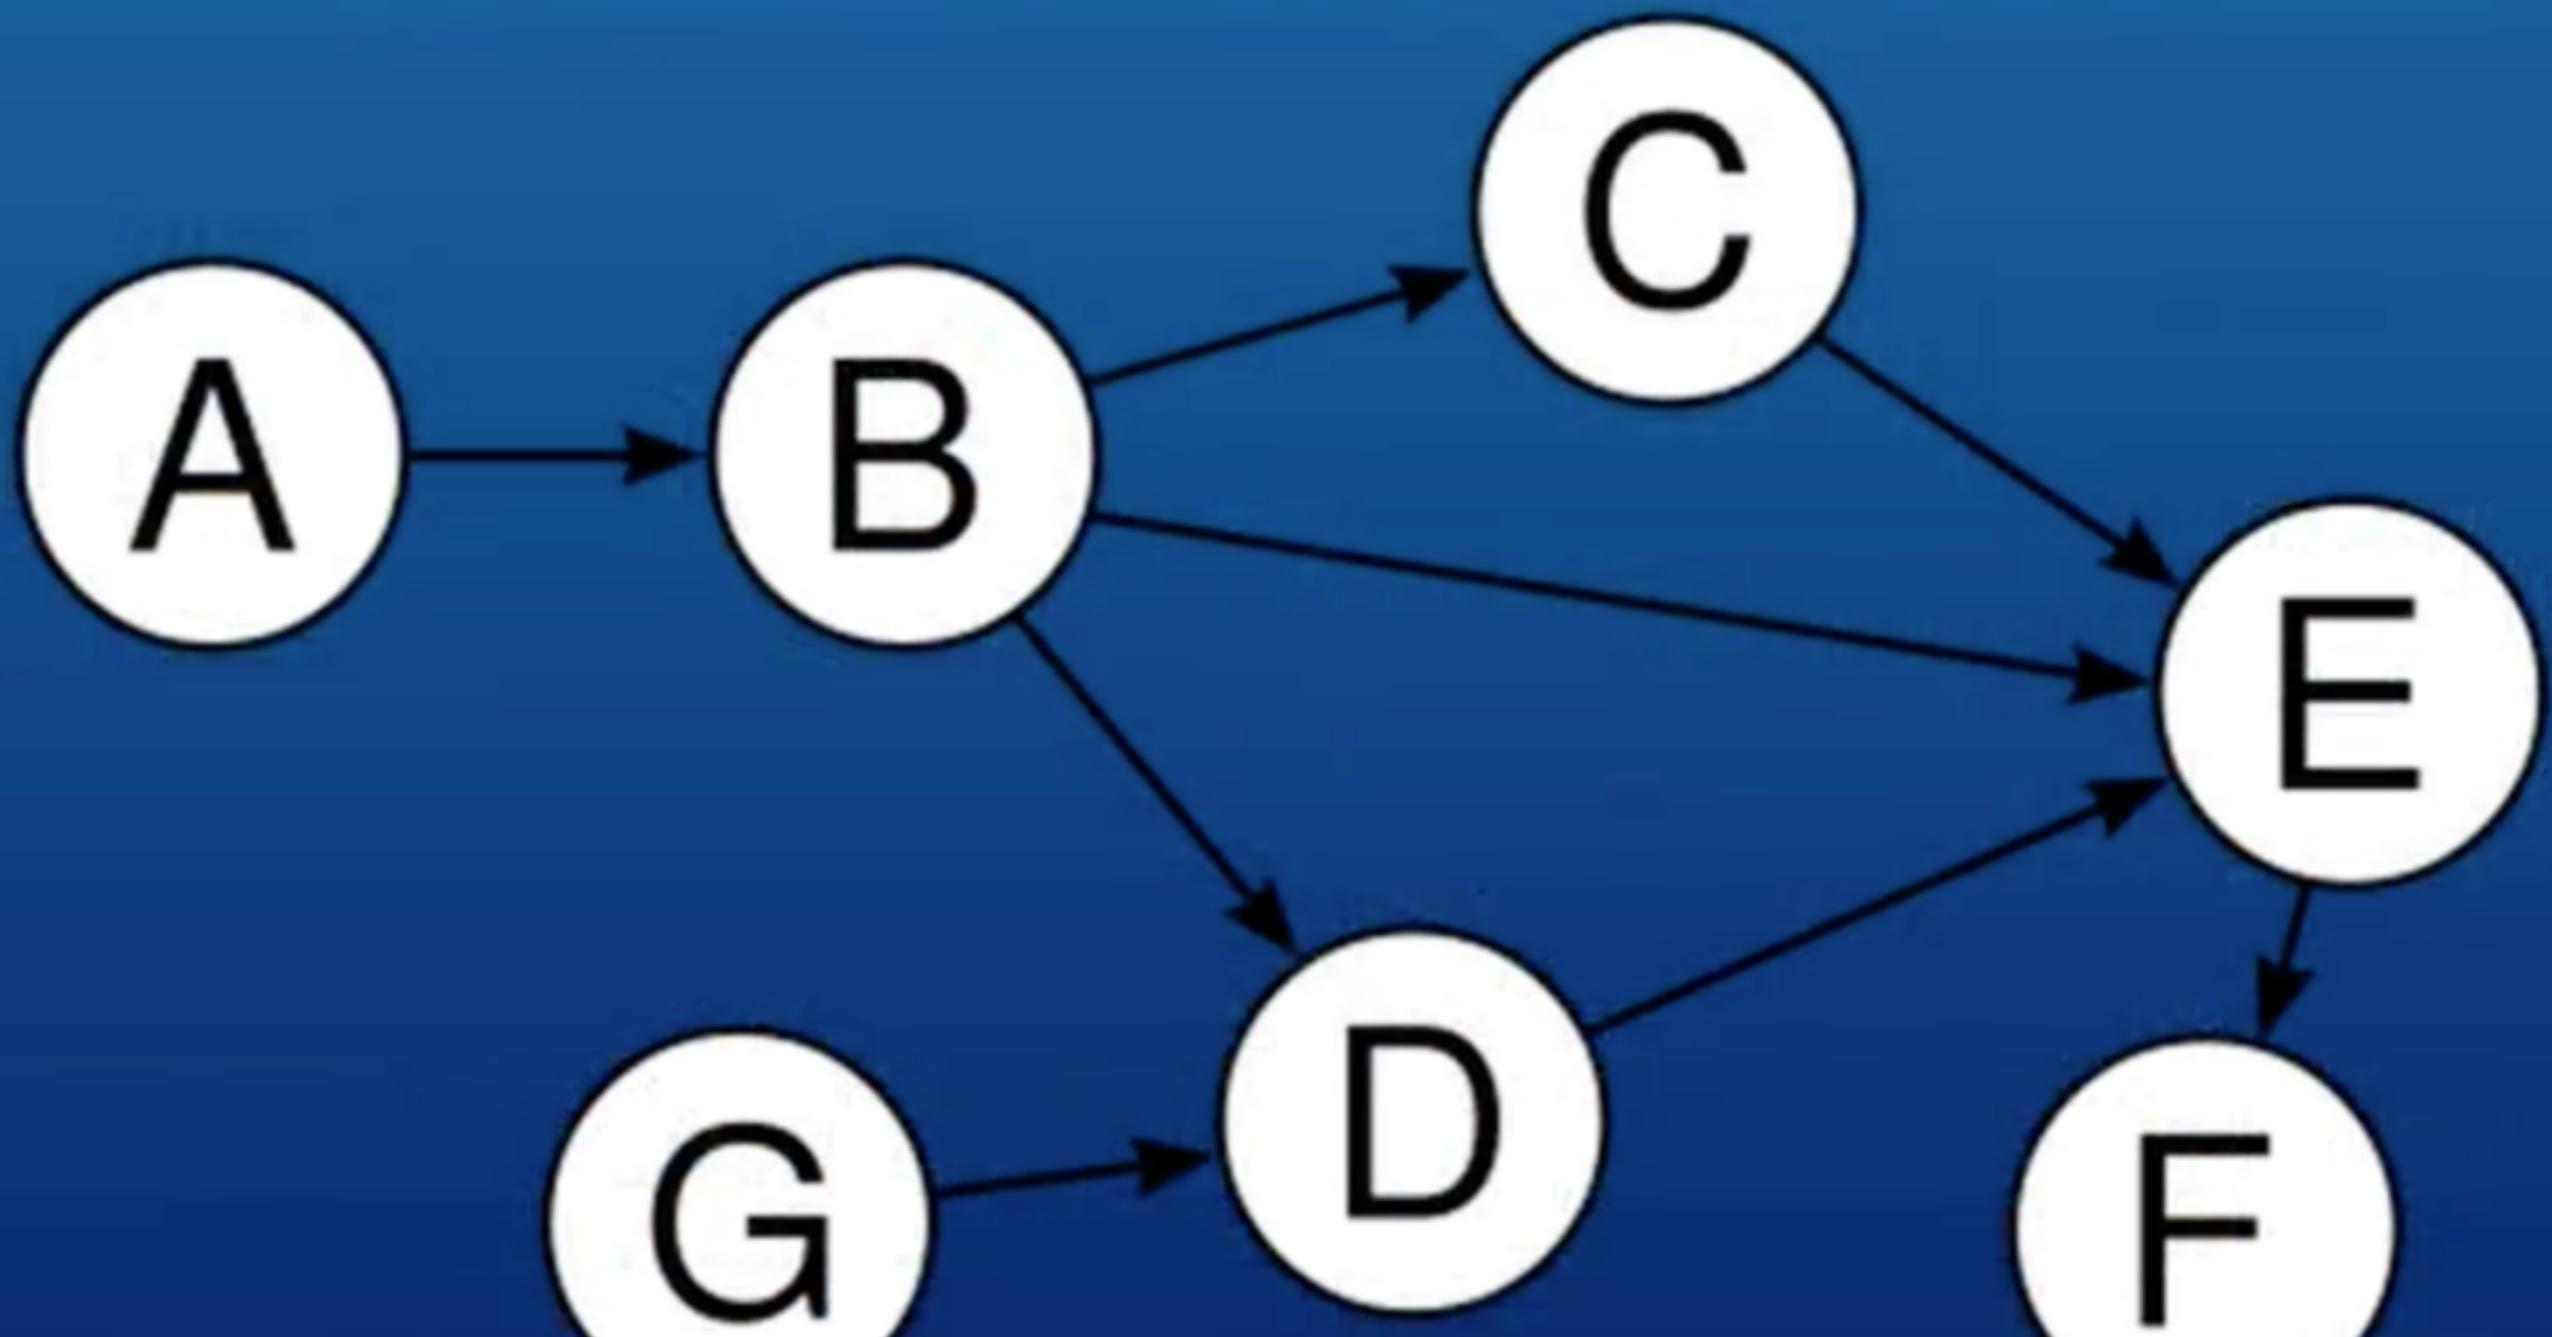
Resilient Distributed Dataset

- RDD

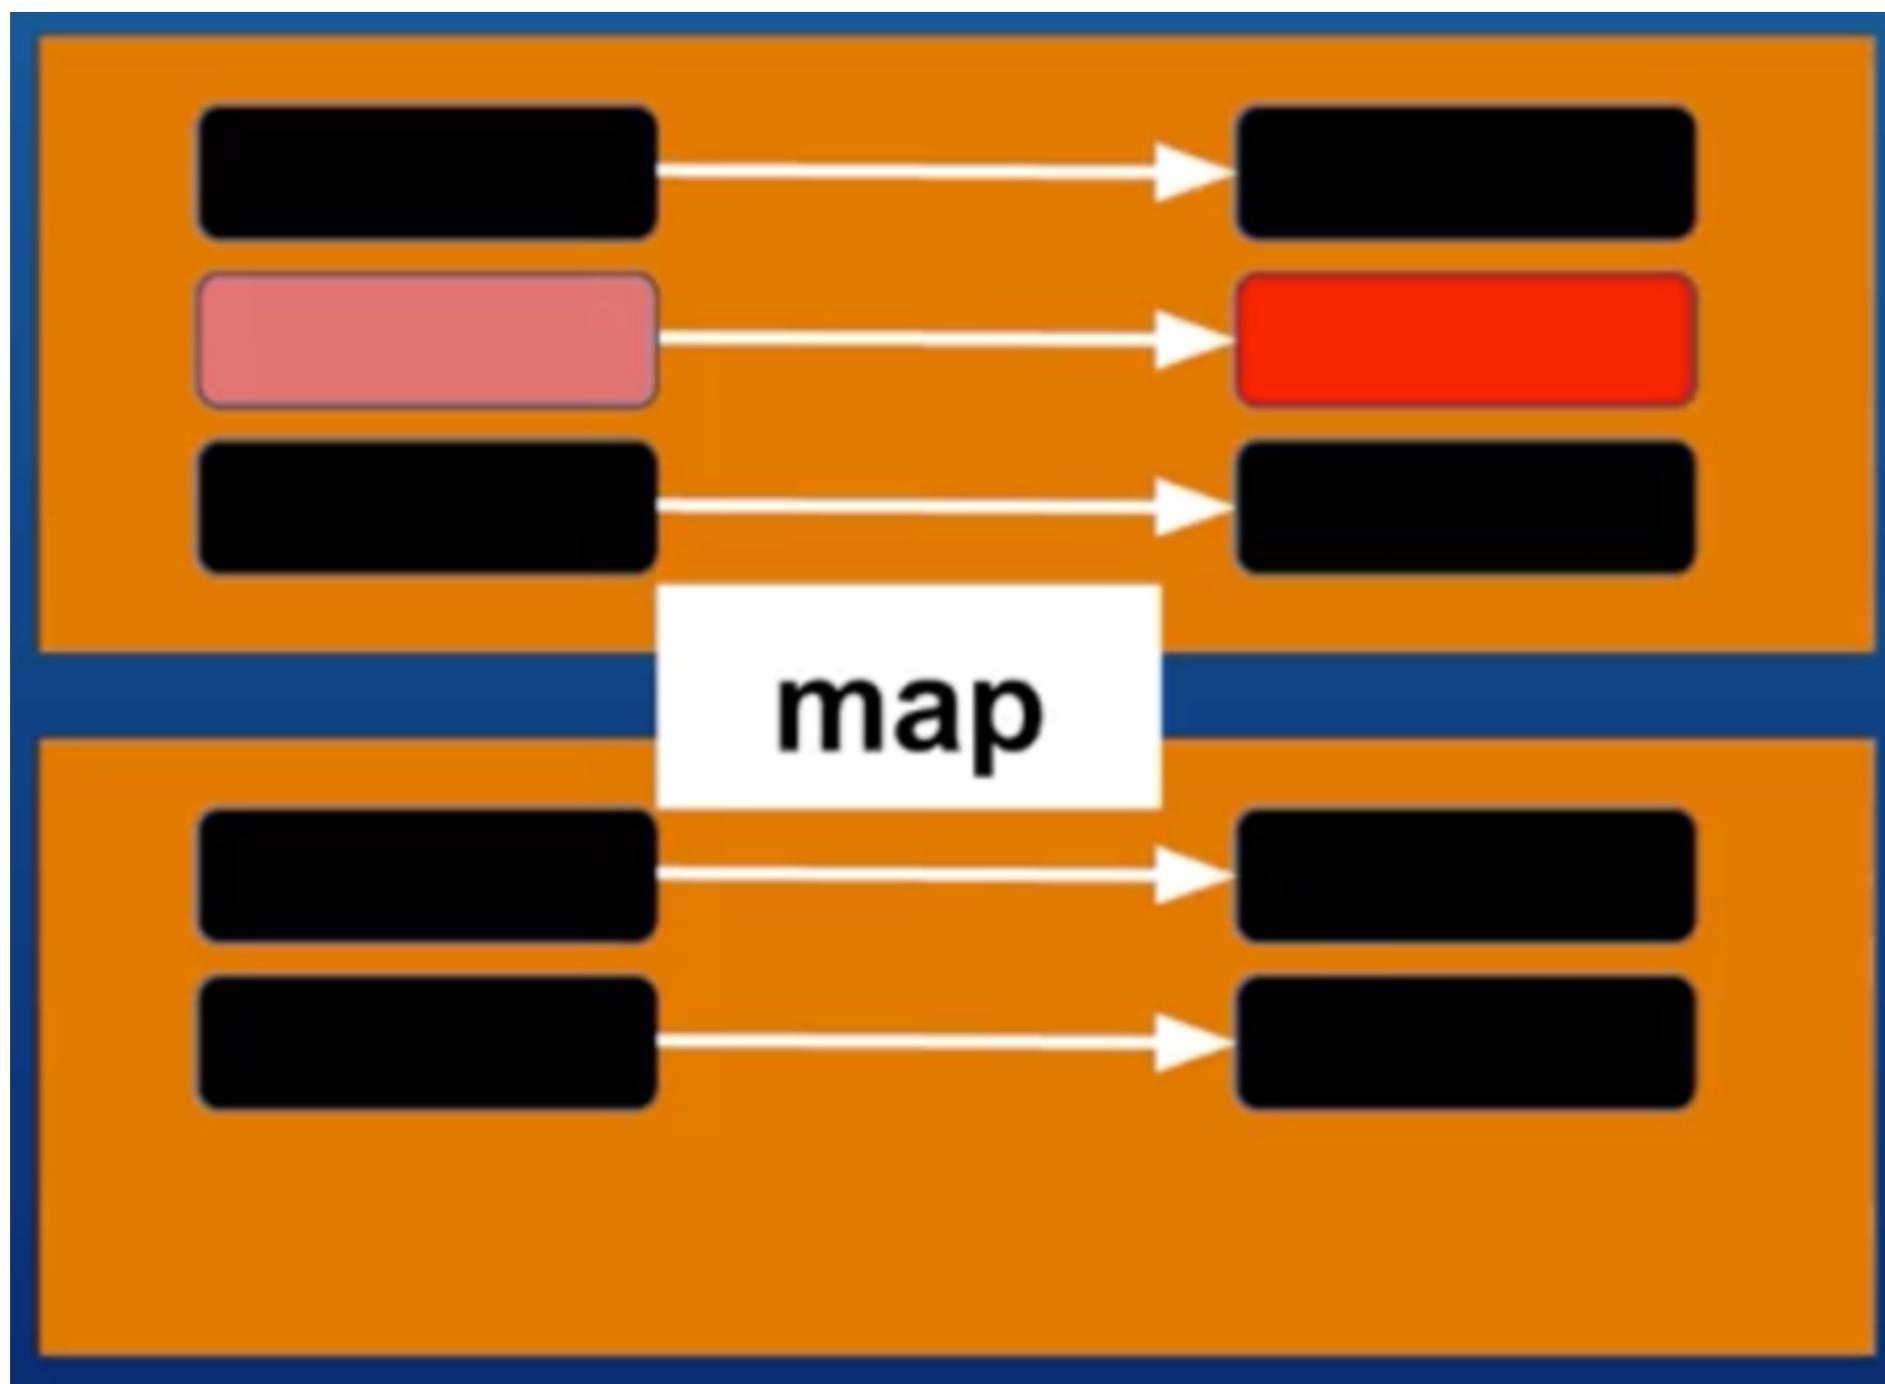


- DAG: Directed Acyclic Graph
- Transformations - map - filter -...
- Actions - collect - collect, count, reduce -...

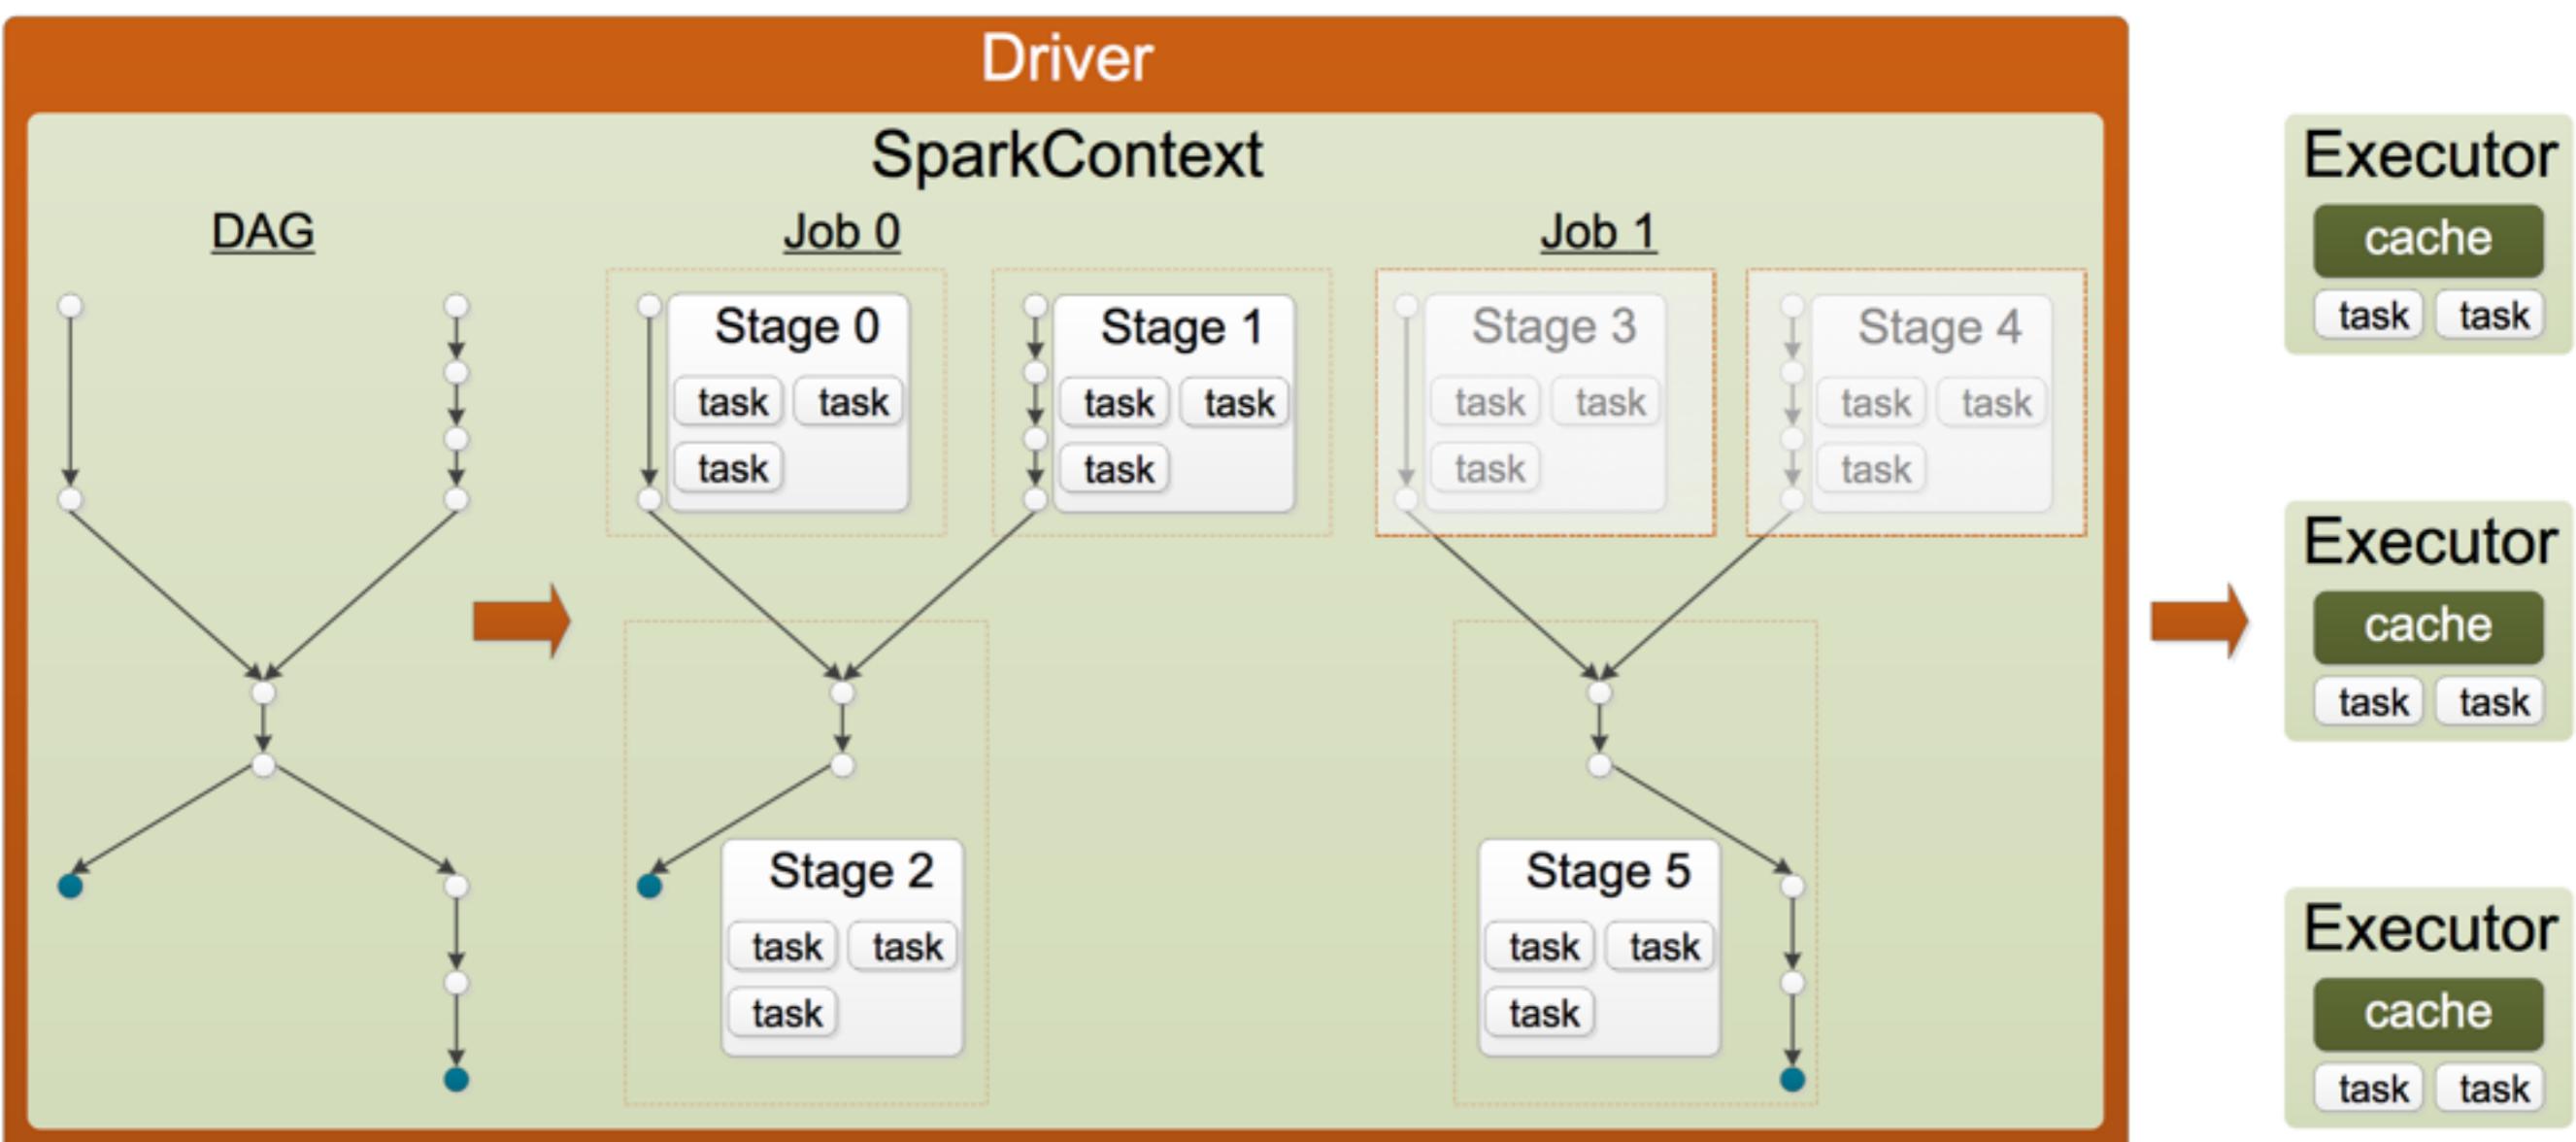
DAG: Directed Acyclic Graph



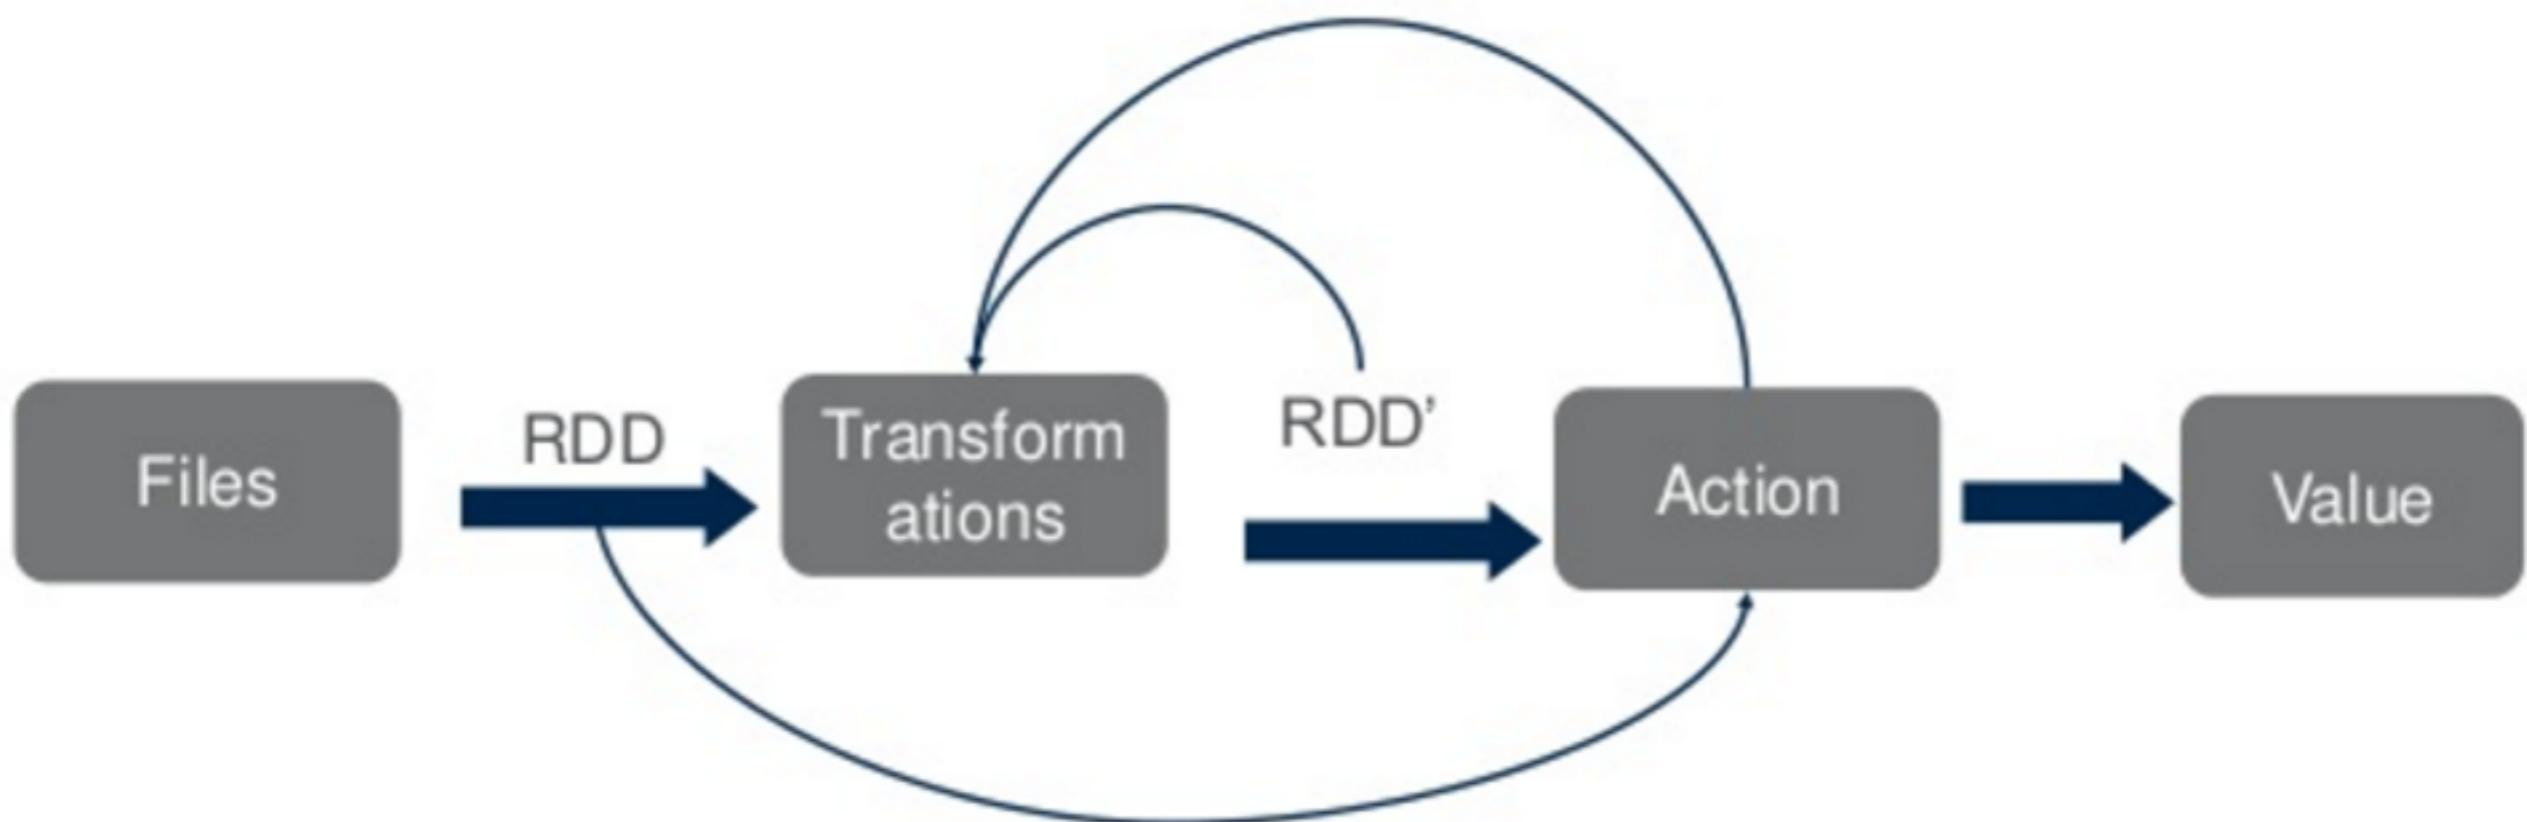
Recovery

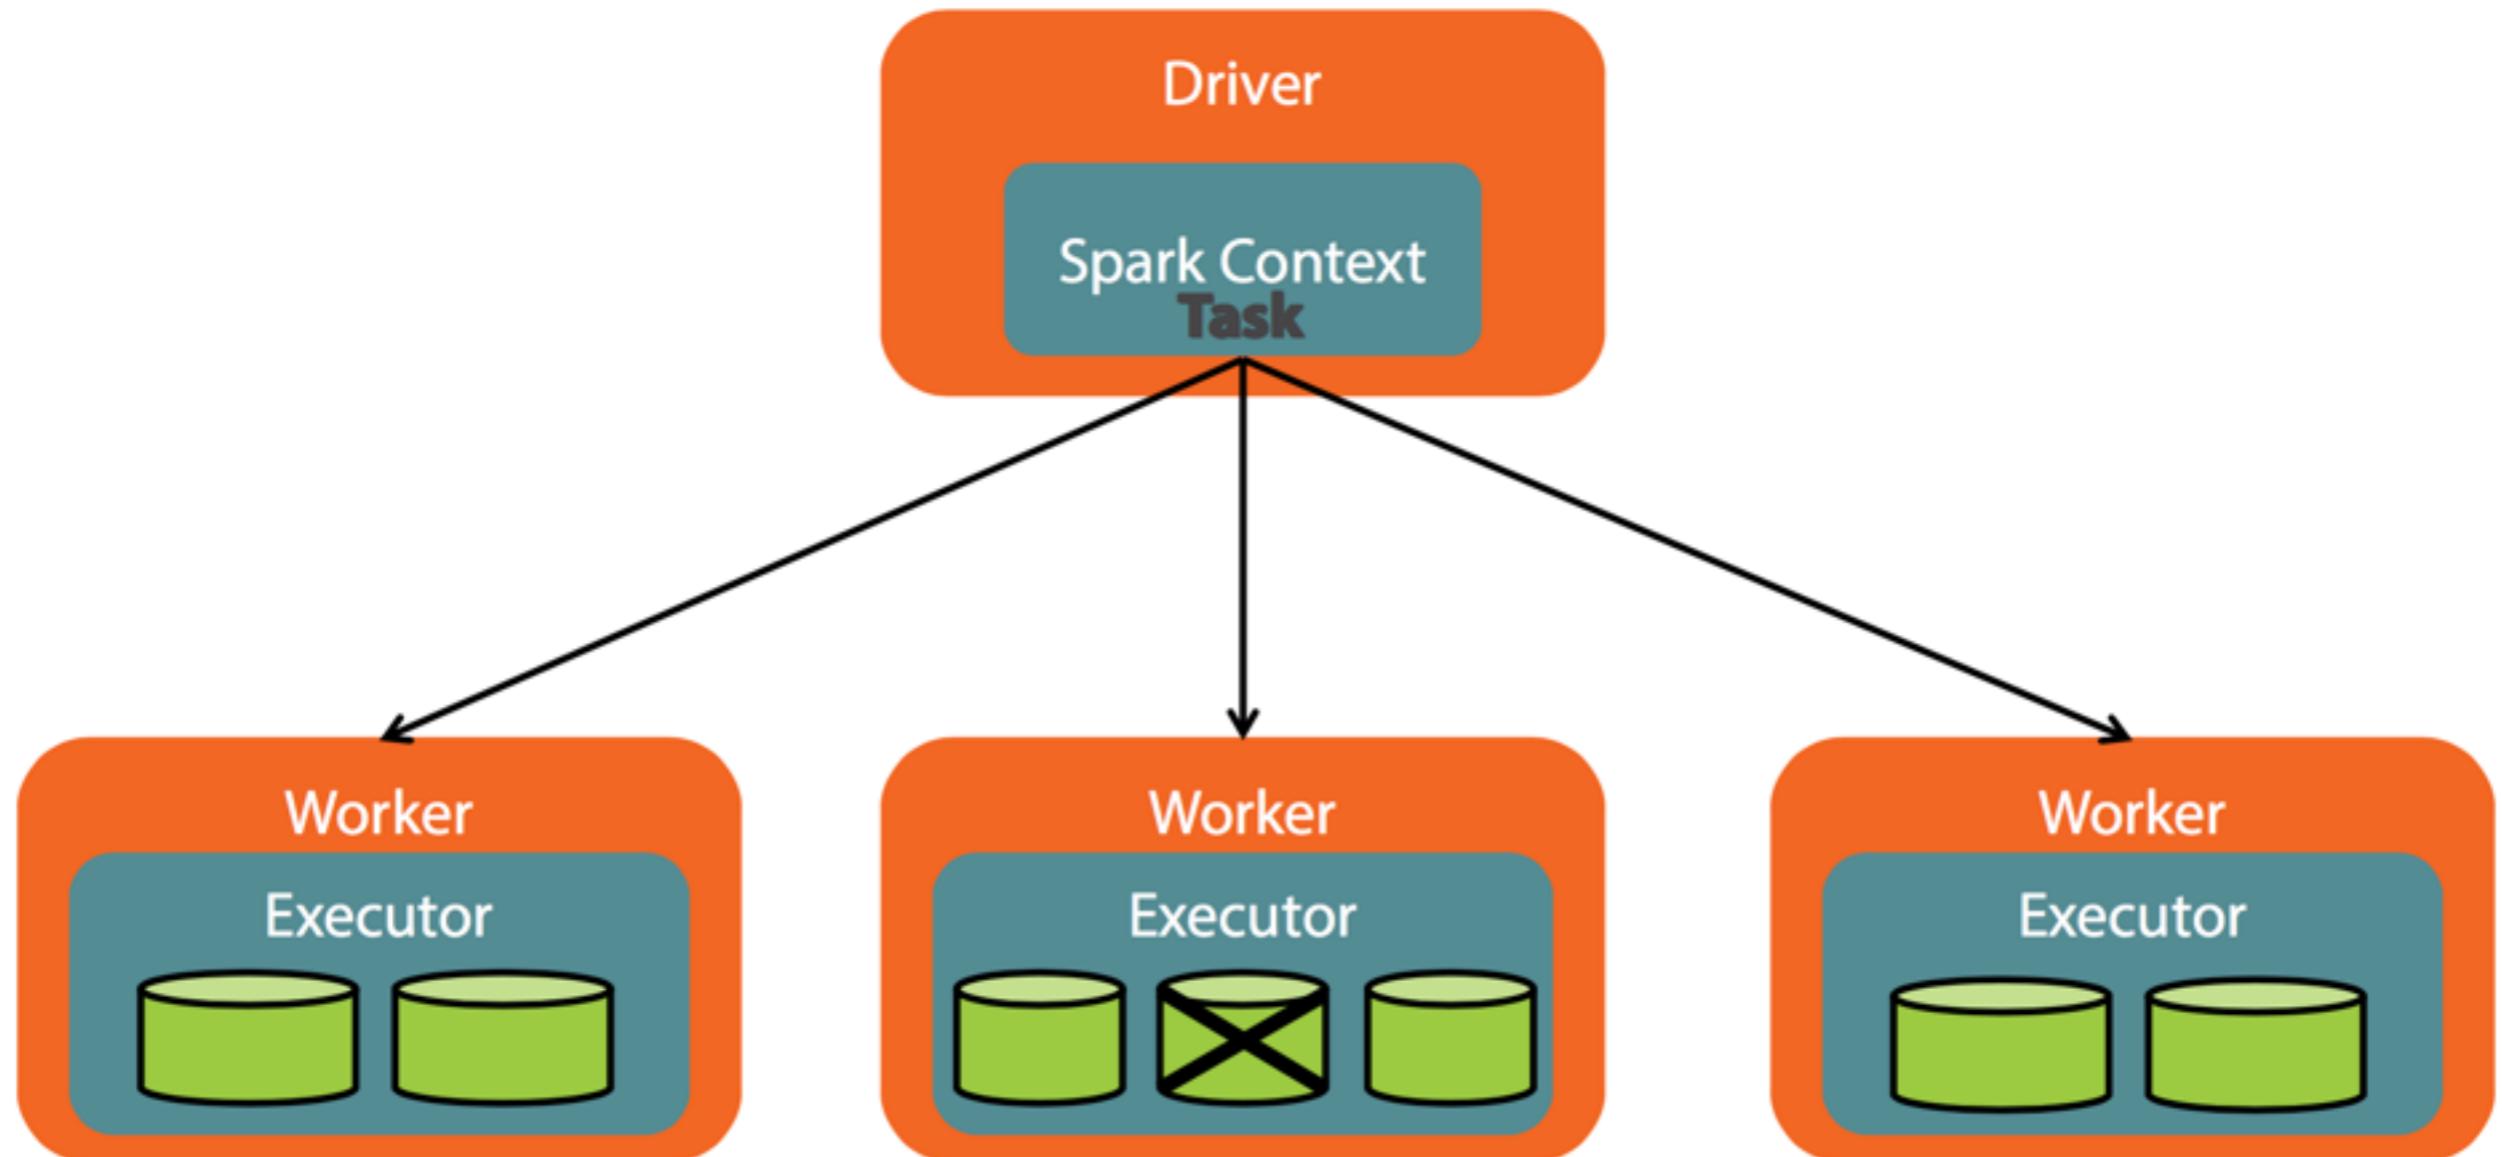


dag and stages

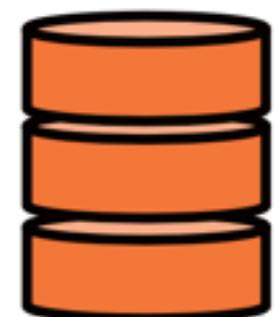


Spark general flow



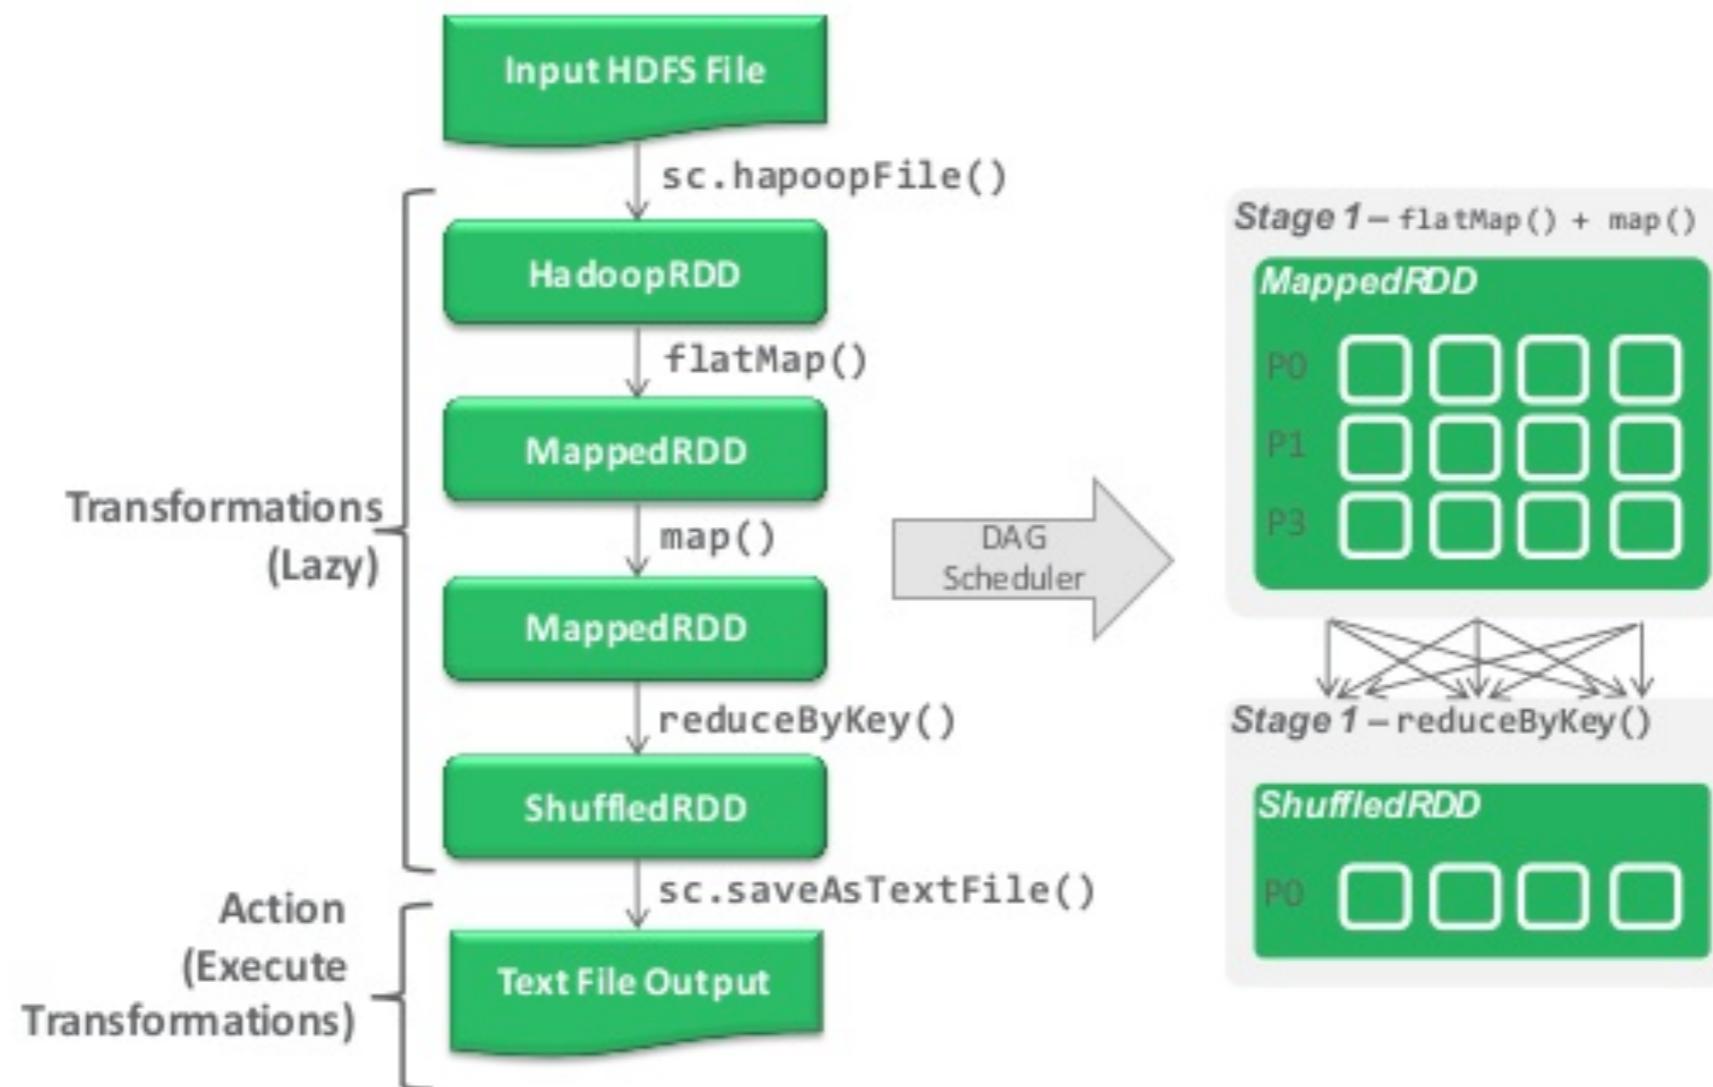


input



S3

Spark workflow



Spark SQL example

```
import com.datastax.spark.connector._

// Connect to the Spark cluster
val conf = new SparkConf(true)...
val sc = new SparkContext(conf)

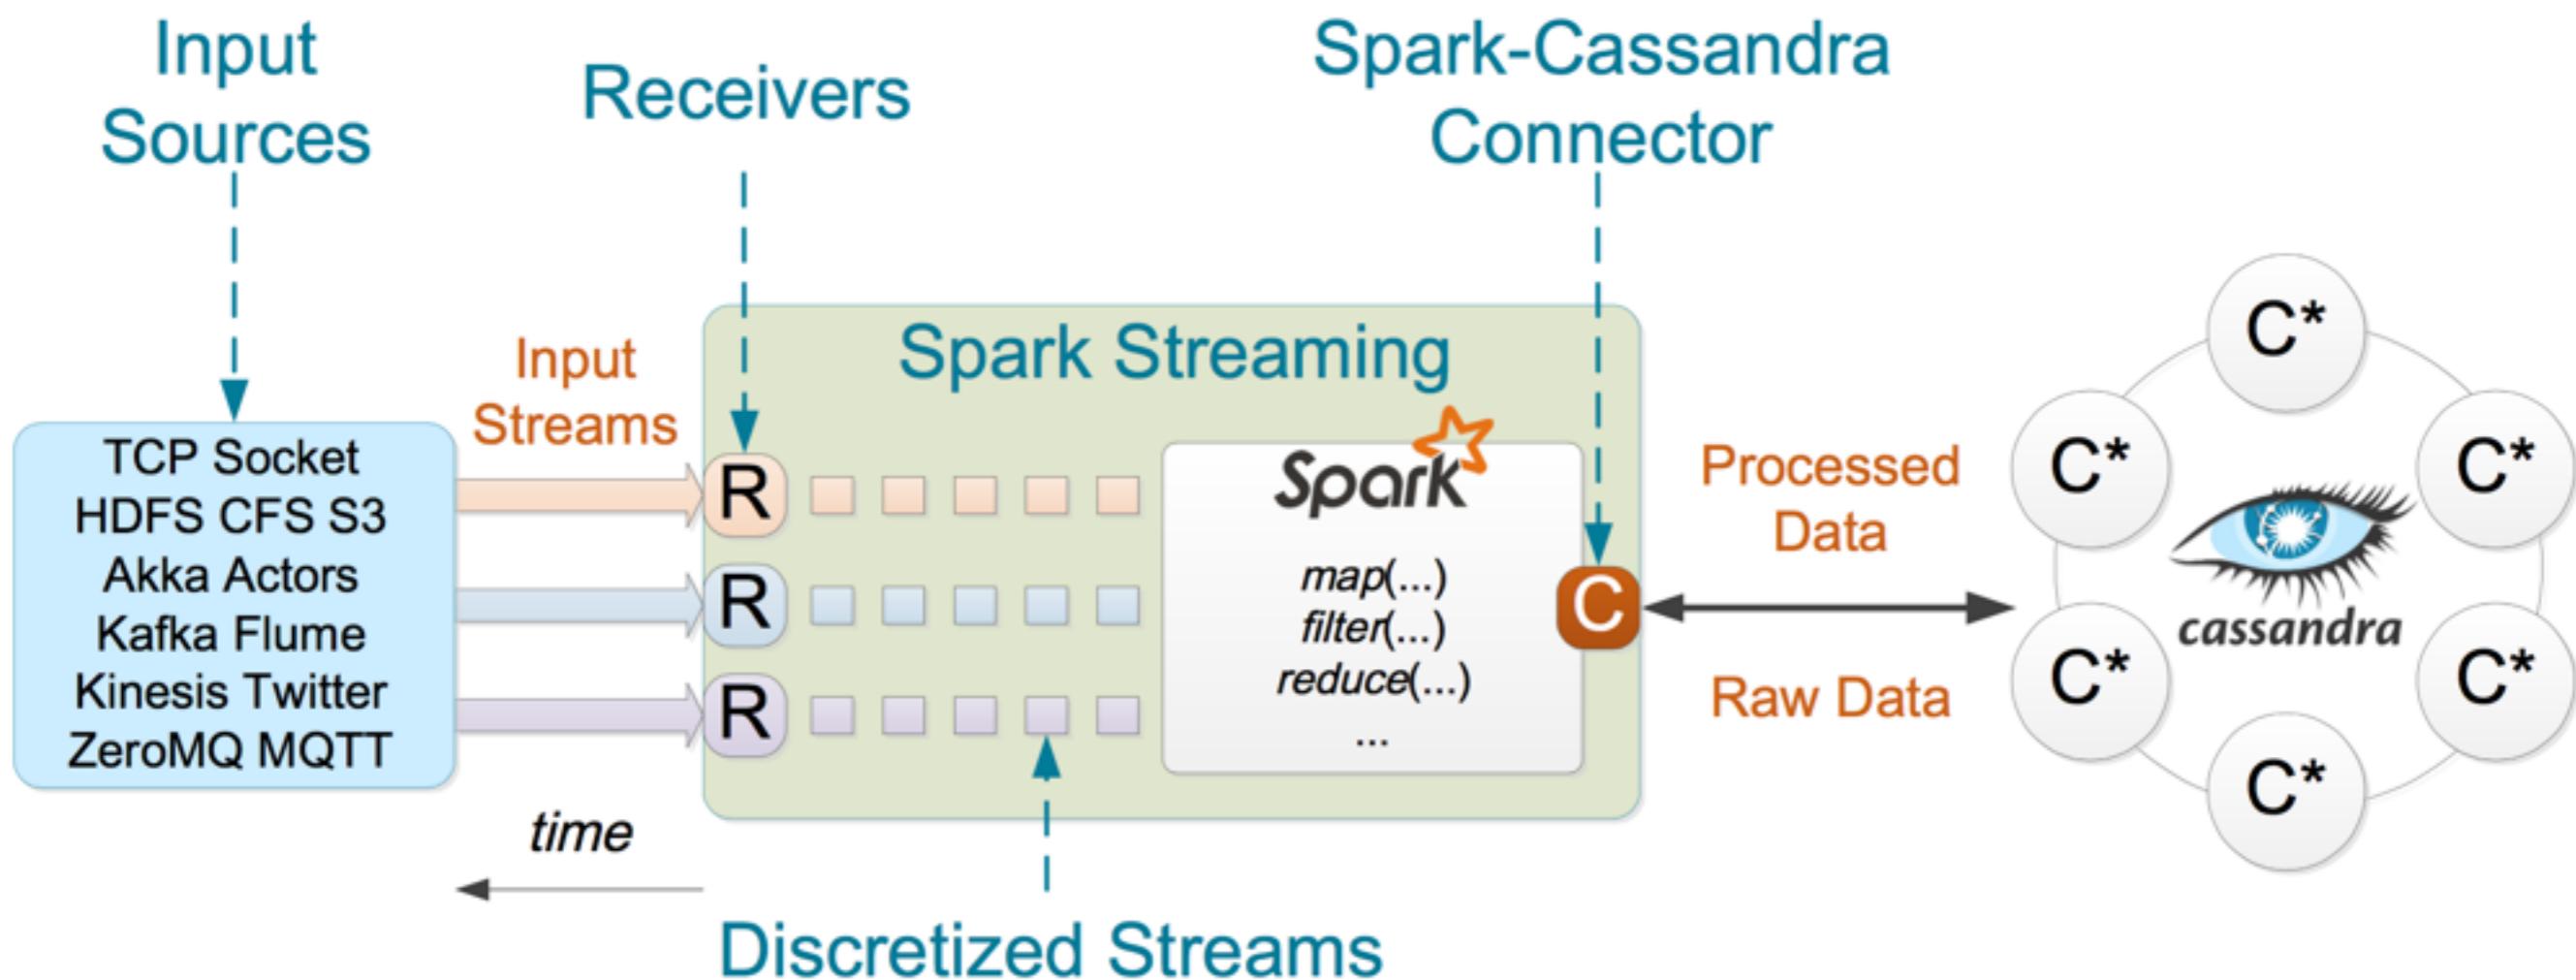
// Create Cassandra SQL context
val cc = new CassandraSQLContext(sc)

// Execute SQL query
val rdd = cc.sql("SELECT * FROM keyspace.table WHERE ...")
```

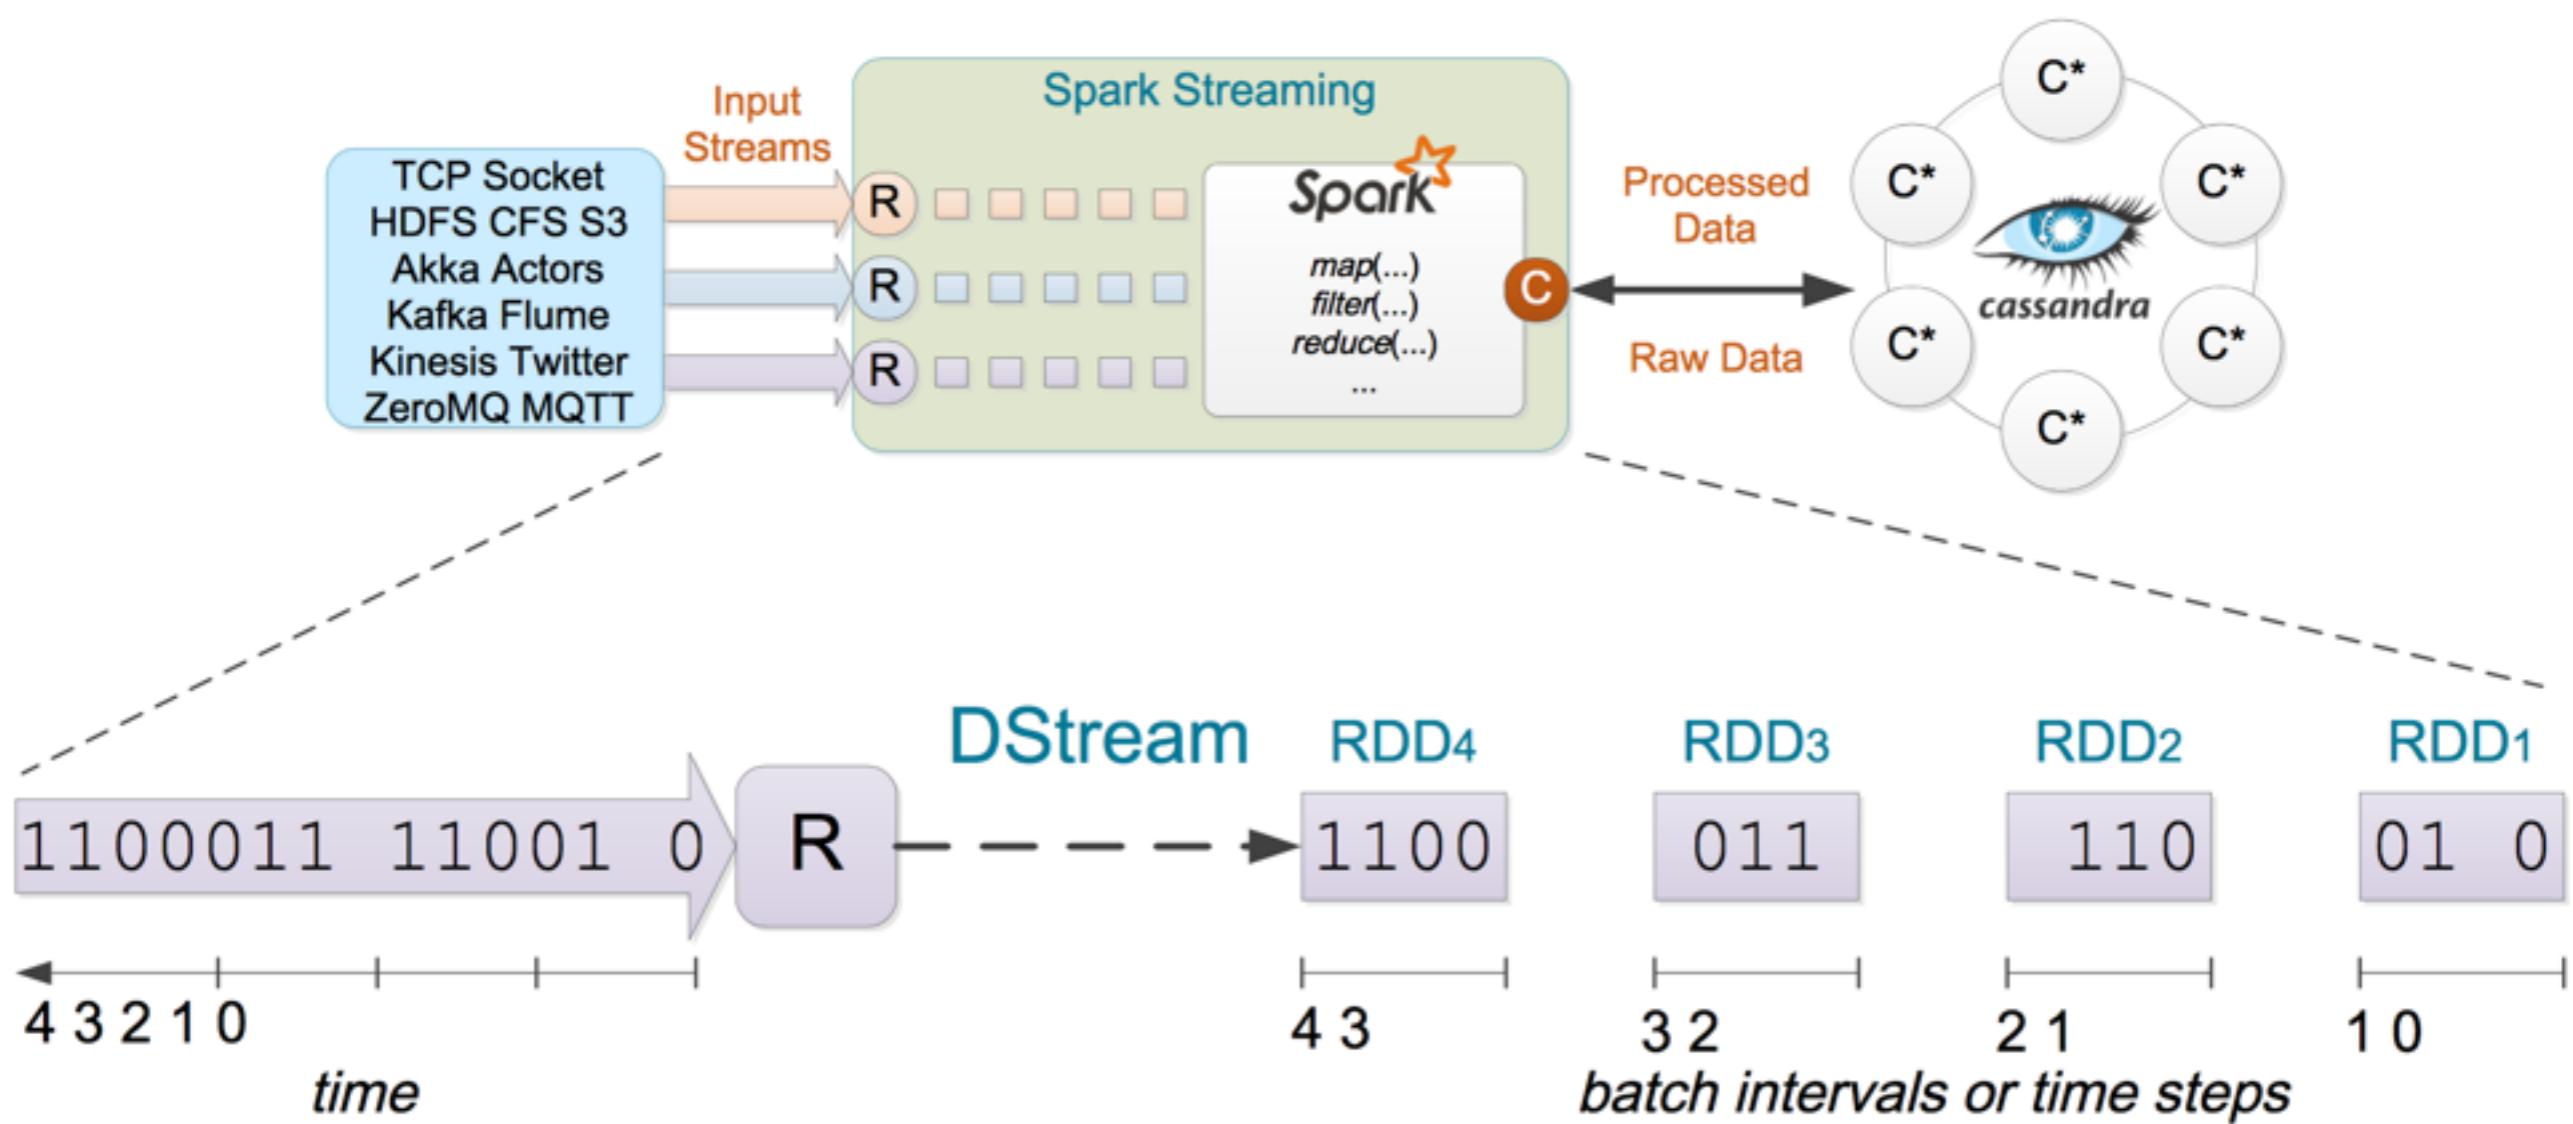
Spark Streaming



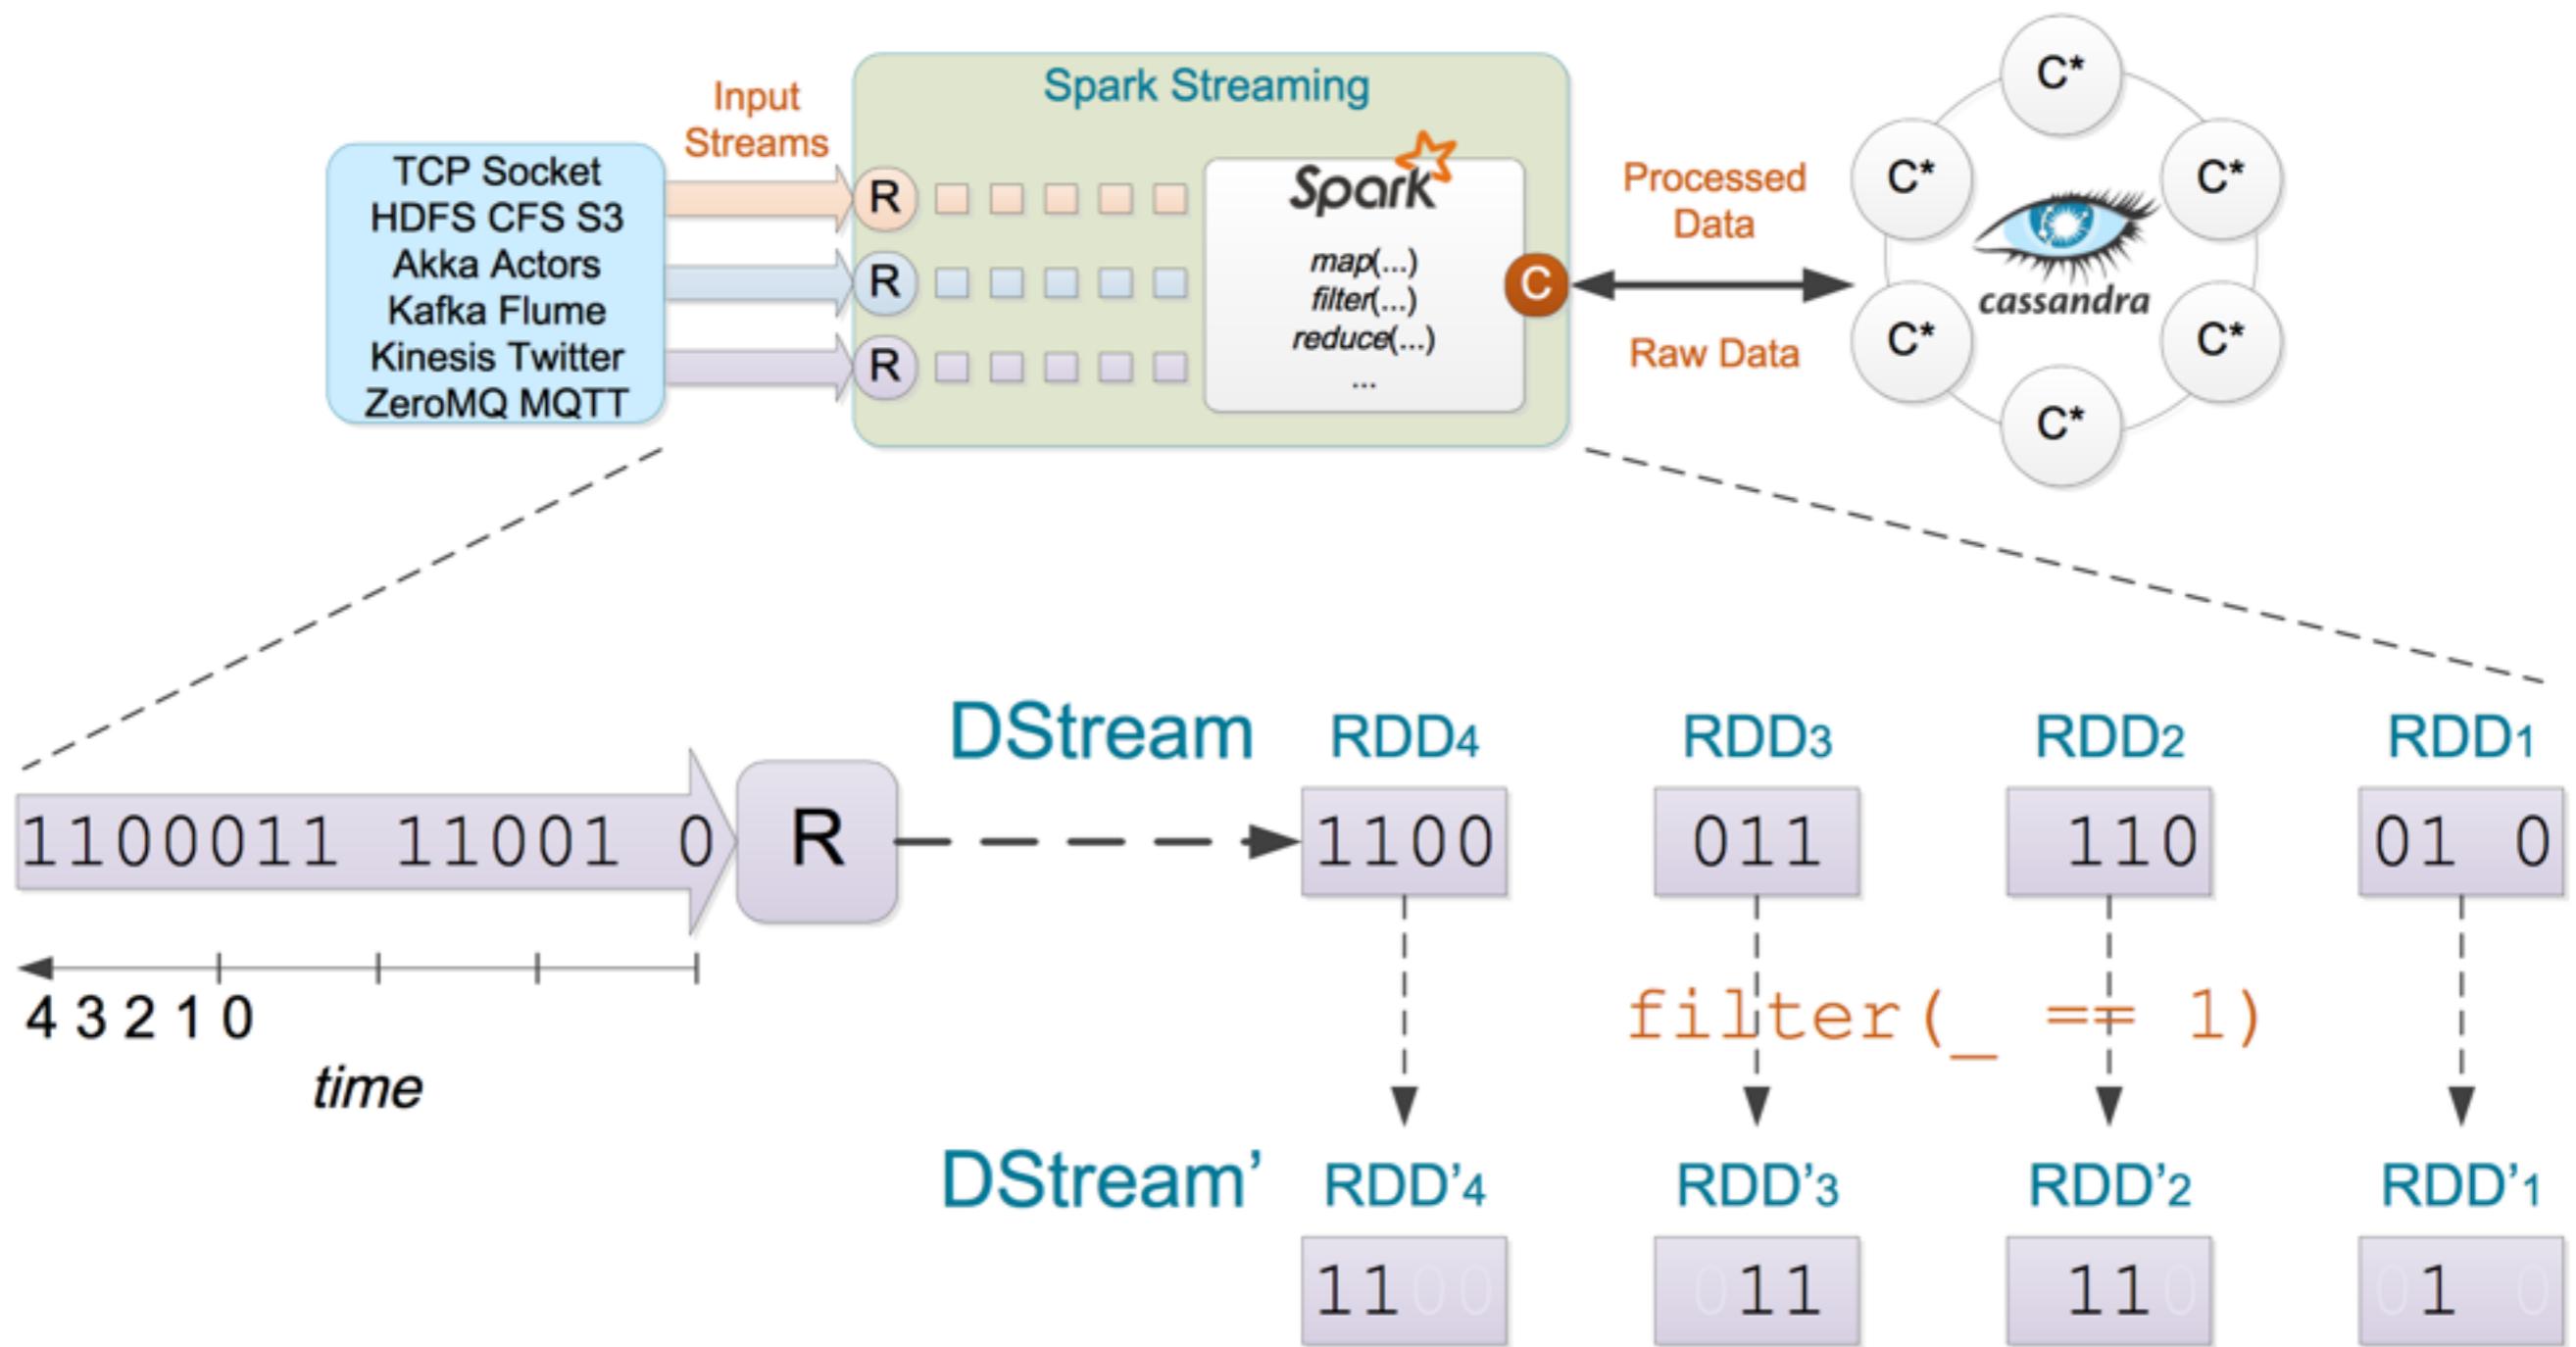
spark streaming



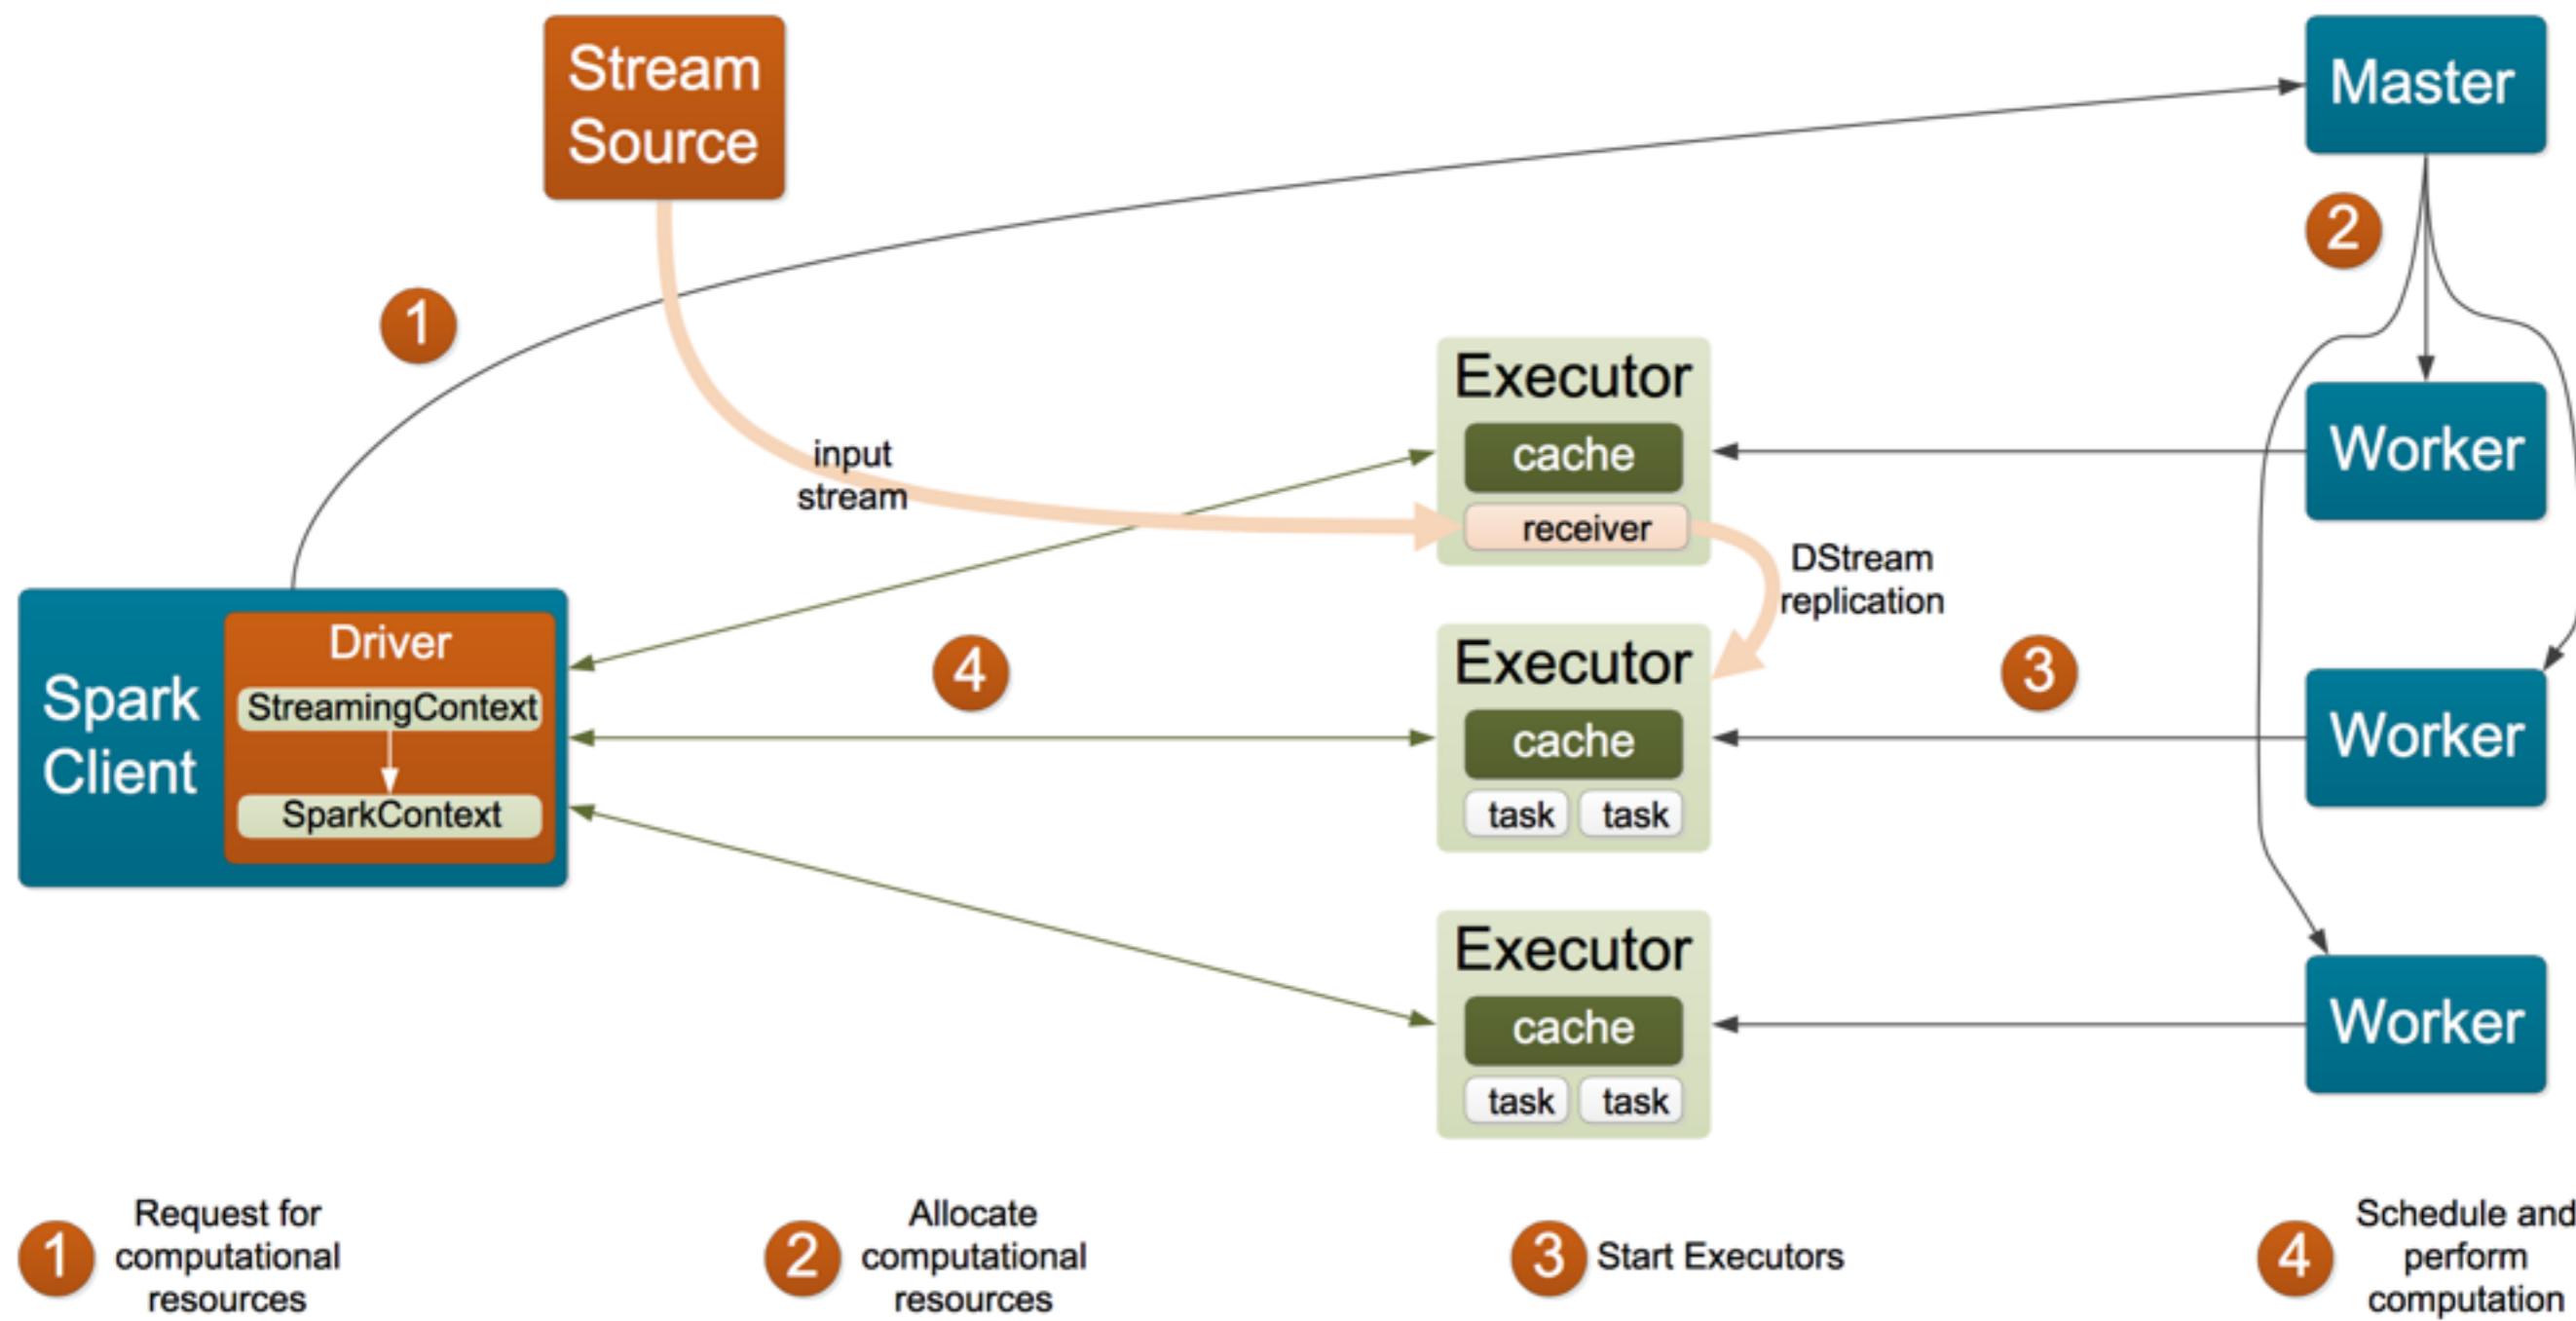
Discretized Stream — DStream



DStream Transformation



Spark Streaming Architecture



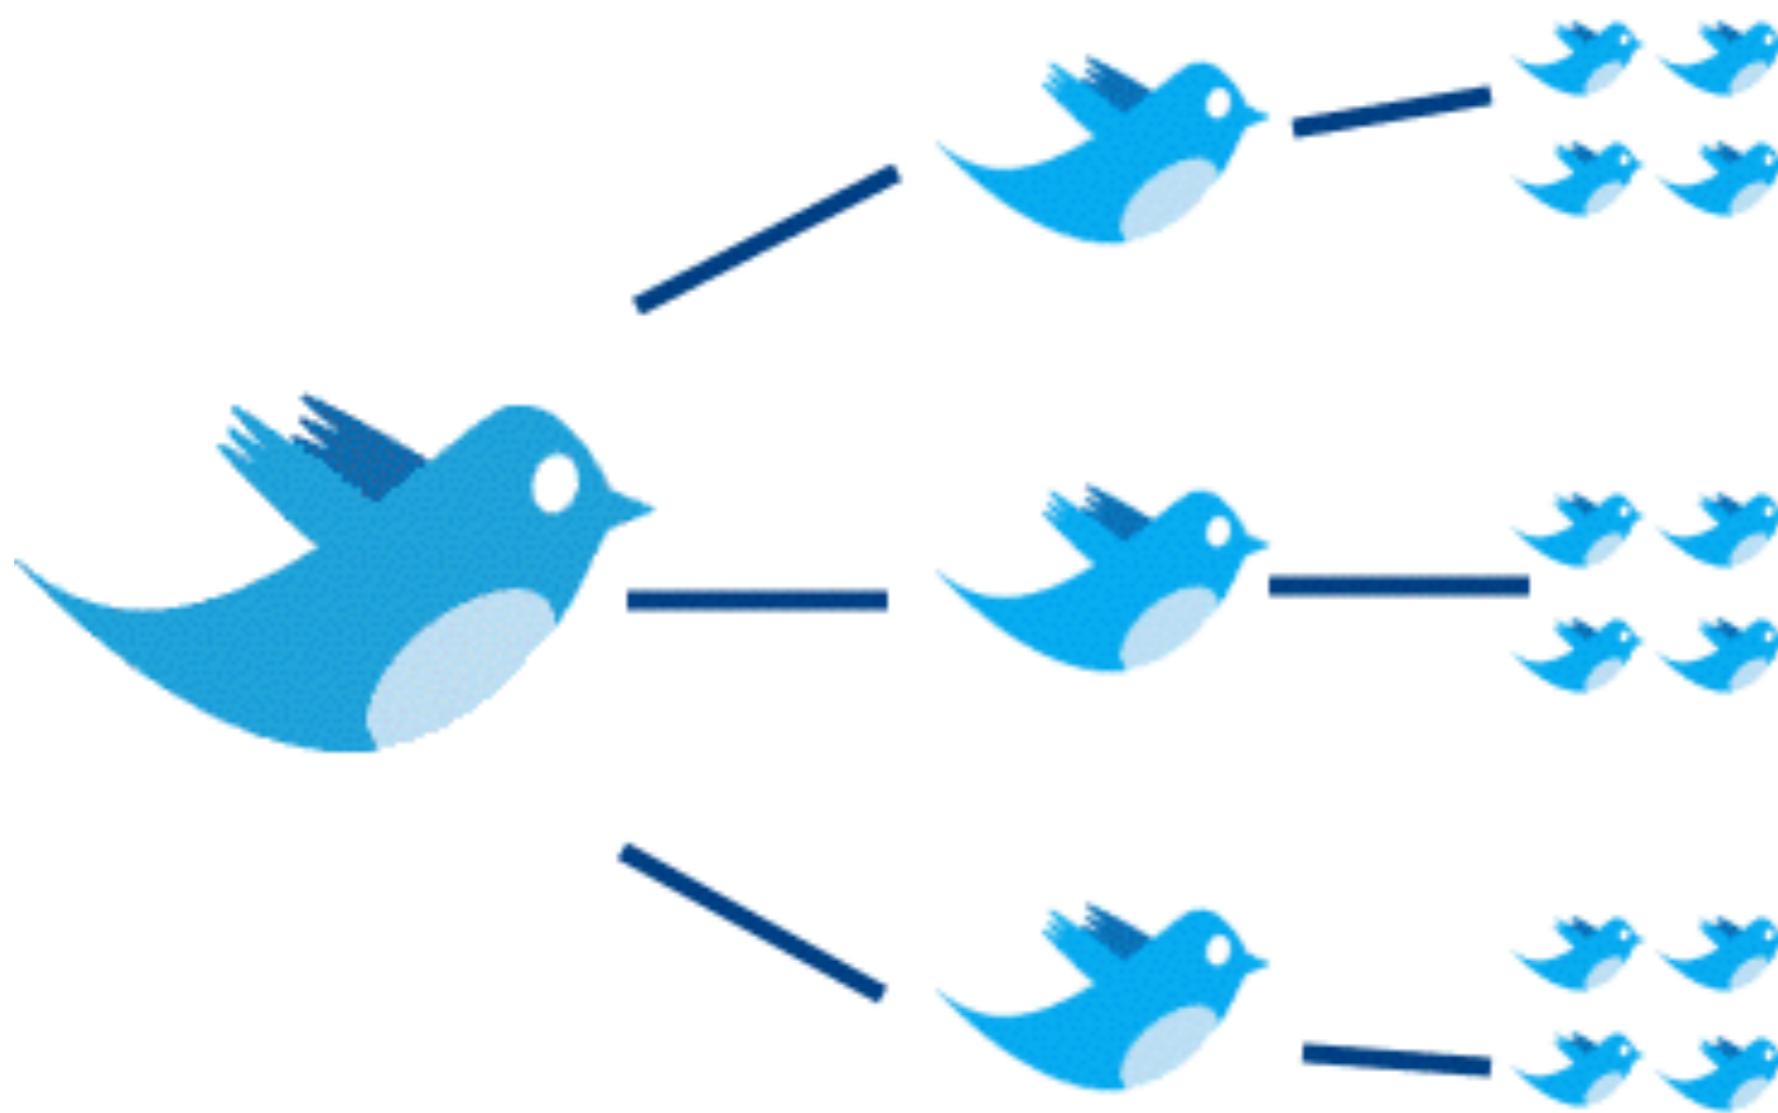
Data Stream Sources

Basic Sources	Advanced Sources	Custom Sources
<ul style="list-style-type: none">• File systems• Socket connections• Akka Actors	<ul style="list-style-type: none">• Kafka• Flume• Kinesis• Twitter• ZeroMQ• MQTT	<ul style="list-style-type: none">• User-defined receivers

QueueStream

```
object QueueStream {  
  
    def main(args: Array[String]) {  
  
        StreamingExamples.setStreamingLogLevels()  
        val sparkConf = new SparkConf().setAppName("QueueStream")  
        // Create the context  
        val ssc = new StreamingContext(sparkConf, Seconds(1))  
  
        // Create the queue through which RDDs can be pushed to  
        // a QueueInputDStream  
        val rddQueue = new SynchronizedQueue[RDD[Int]]()  
  
        // Create the QueueInputDStream and use it do some processing  
        val inputStream = ssc.queueStream(rddQueue)  
        val mappedStream = inputStream.map(x => (x % 10, 1))  
        val reducedStream = mappedStream.reduceByKey(_ + _)  
        reducedStream.print()  
        ssc.start()  
  
        // Create and push some RDDs into  
        for (i <- 1 to 30) {  
            rddQueue += ssc.sparkContext.makeRDD(1 to 1000, 10)  
            Thread.sleep(1000)  
        }  
        ssc.stop()  
    }  
}
```

twitter stream



streamingexample.scala

```
import org.apache.spark.{SparkConf, SparkContext}
import org.apache.spark.SparkContext._
import org.apache.spark.streaming._
import org.apache.spark.streaming.StreamingContext._
import com.datastax.spark.connector._
import com.datastax.spark.connector.streaming._

object StreamingExample {
    def main(args: Array[String]) {

        val conf = new SparkConf(true)
            .setAppName("Streaming Example")
            .setMaster("spark://127.0.0.1:7077")
            .set("spark.cassandra.connection.host", "127.0.0.1")
            .set("spark.cleaner.ttl", "3600")
            .setJars(Array("/.../target/scala-2.10/streaming-example_2.10-0.1.jar"))

        val ssc = new StreamingContext(conf, Seconds(4))

        val stream = ssc.socketTextStream("127.0.0.1", 9999)

        stream.flatMap(record => record.split(" "))
            .map(word => (word, 1))
            .reduceByKey(_ + _)
            .print()

        ssc.start()
        ssc.awaitTermination()
    }
}

ssc.awaitTermination()
ssc.stop()
```

results output

```
-----  
Time: 1443548140000 ms  
-----
```

```
(Drama,22)  
(Comedy,12)  
(Action,16)  
(Romance,9)  
(Dance,11)  
(Family,18)  
(Adventure,7)  
(Horror,10)  
(Thriller,19)  
(Fantasy,11)  
...
```

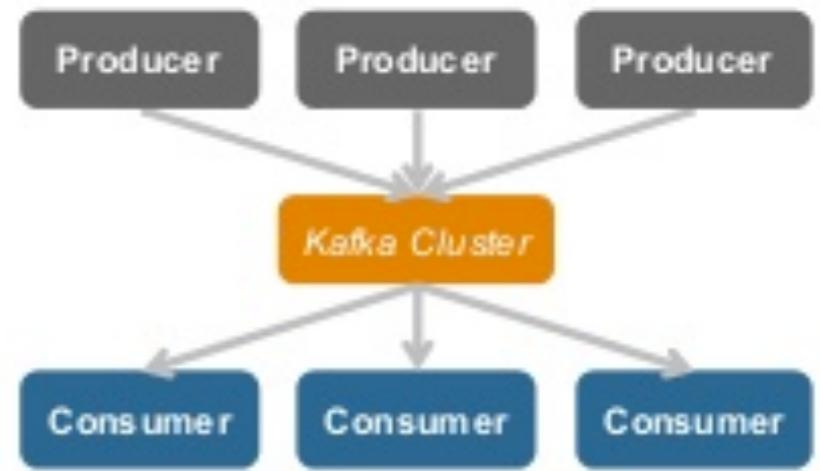
```
-----  
Time: 1443548144000 ms  
-----
```

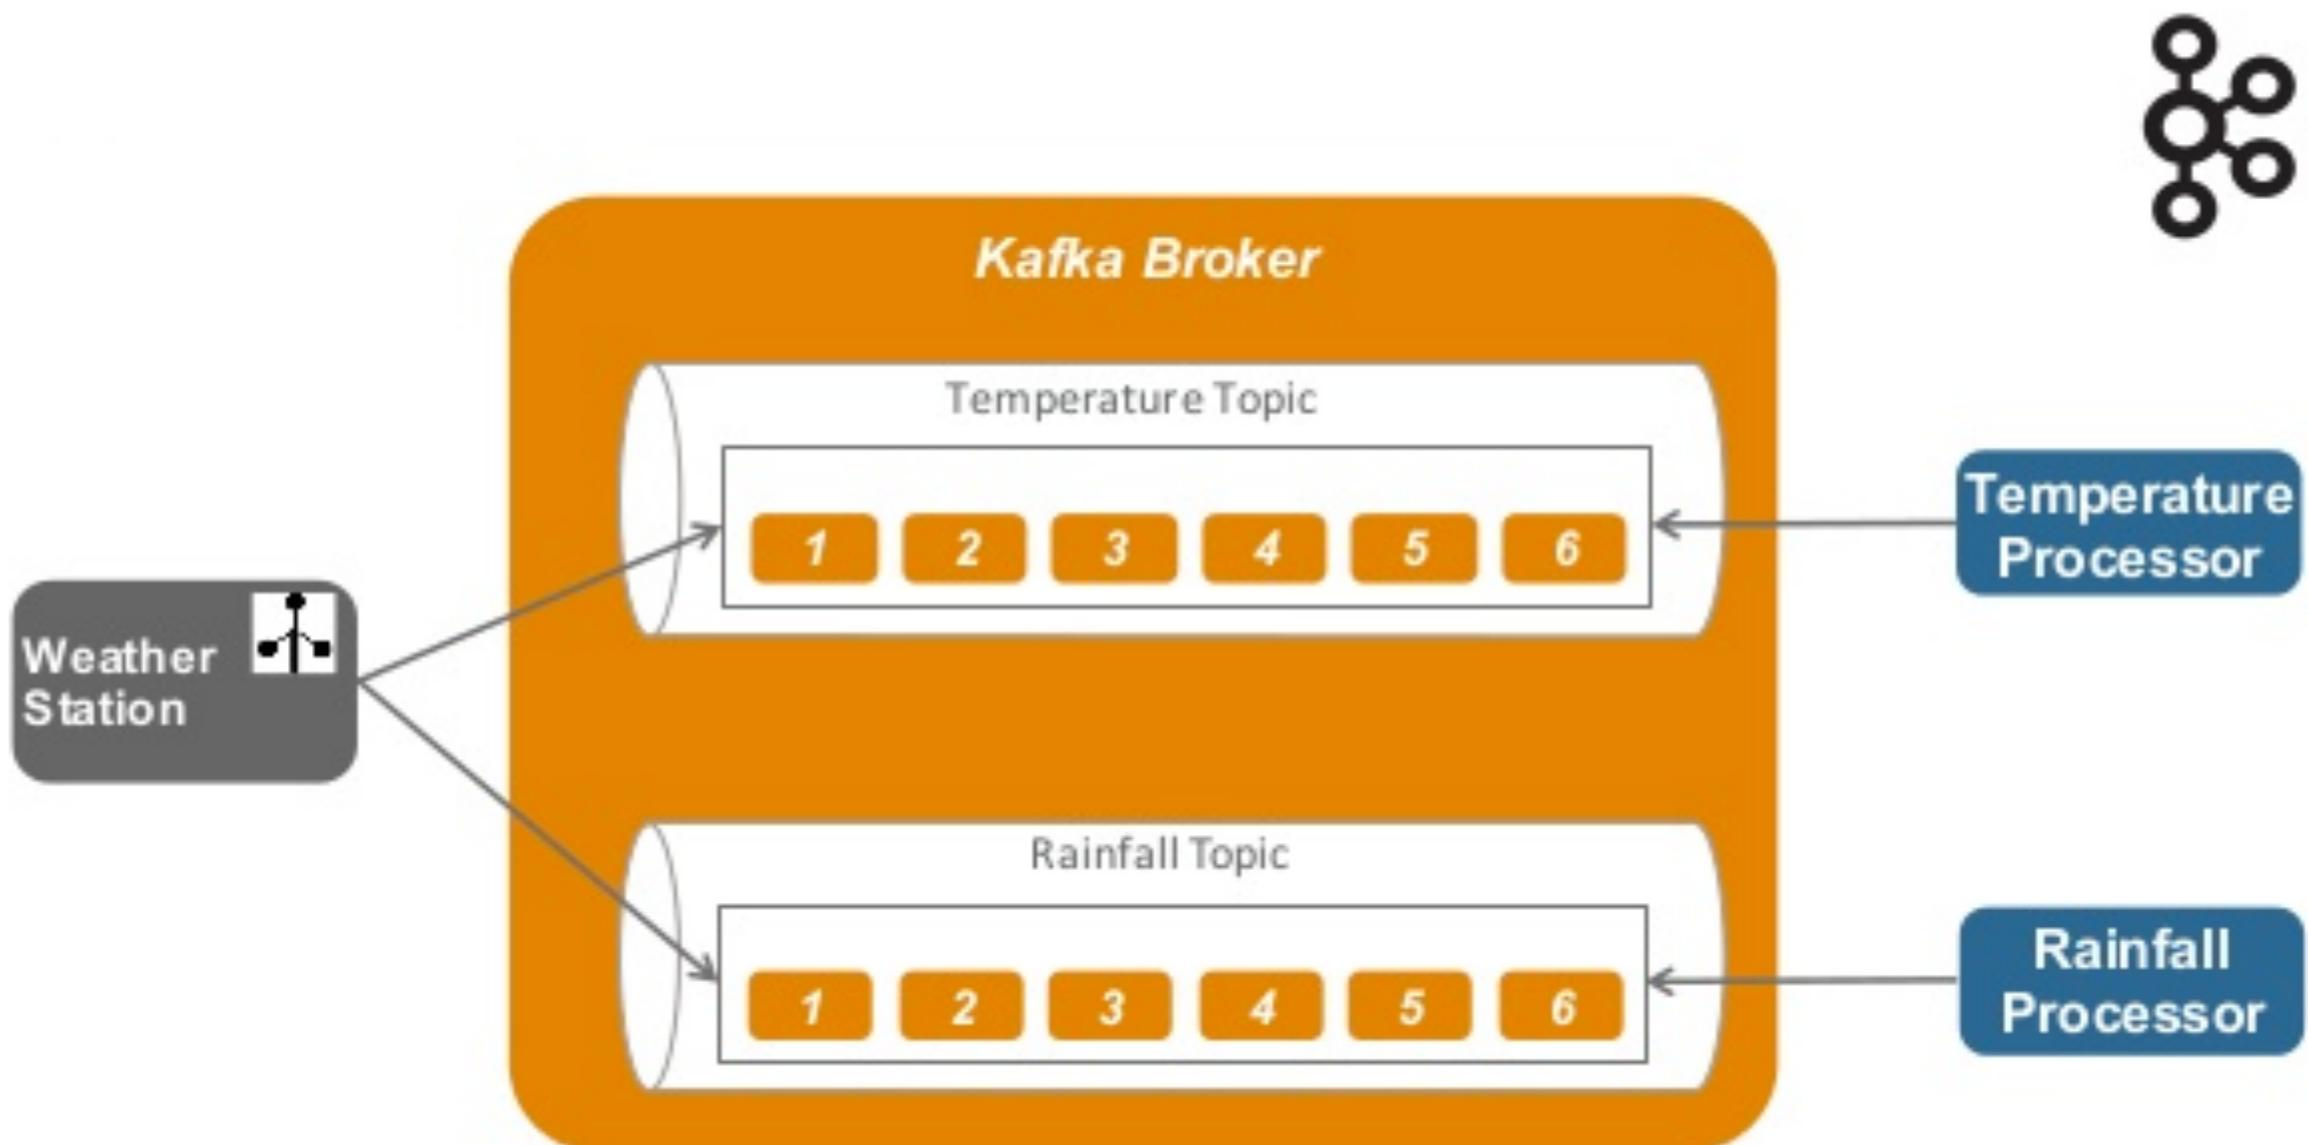
```
(Drama,11)  
(Comedy,14)  
(Action,13)  
(Romance,2)  
(Dance,9)  
(Family,12)  
(Adventure,9)  
(Science,2)  
(Thriller,3)  
(Fantasy,8)  
...
```

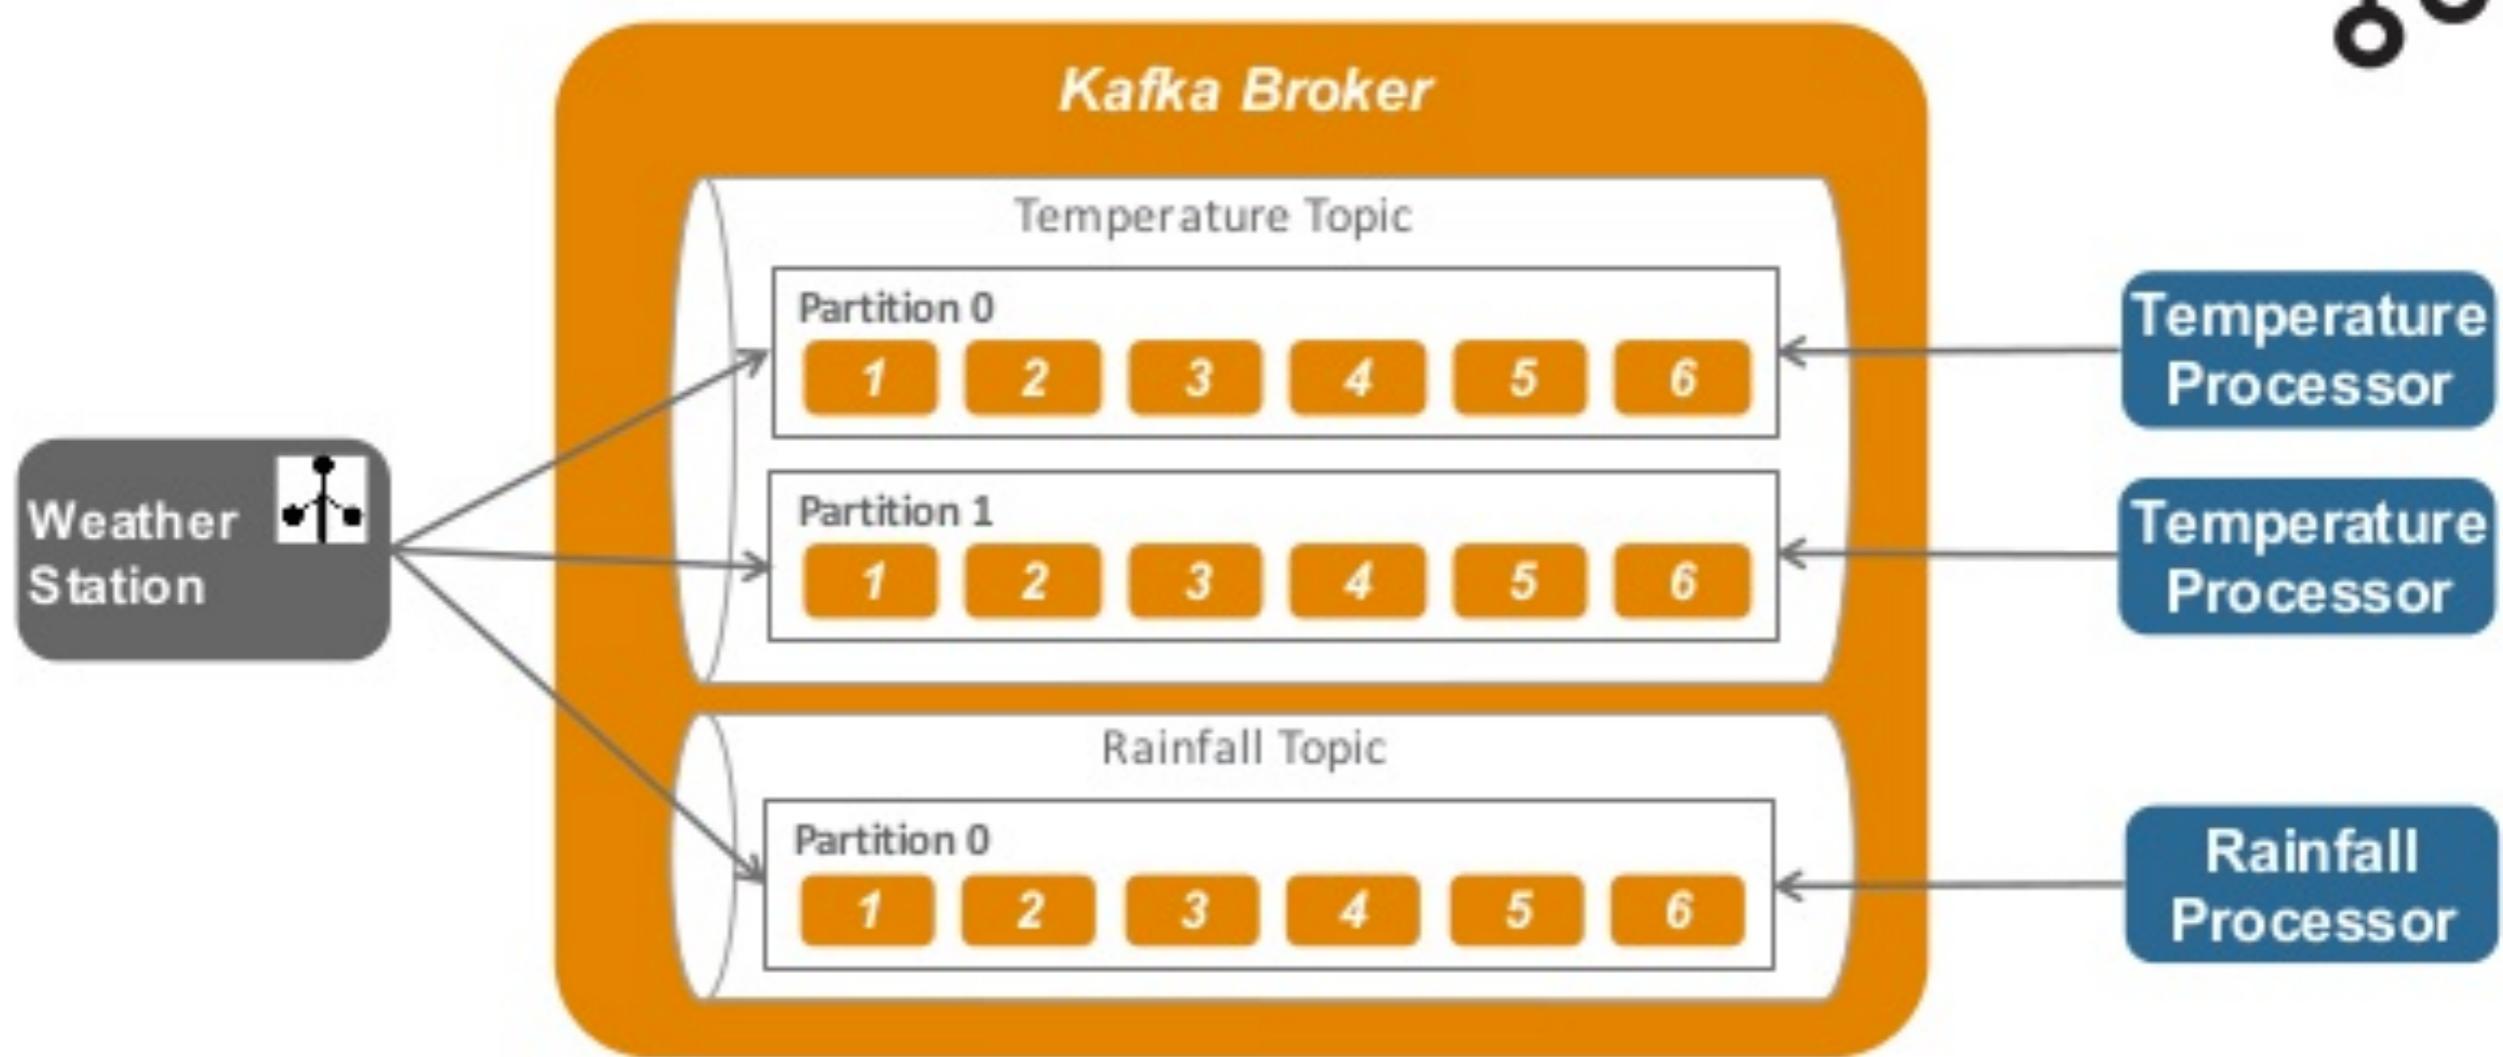


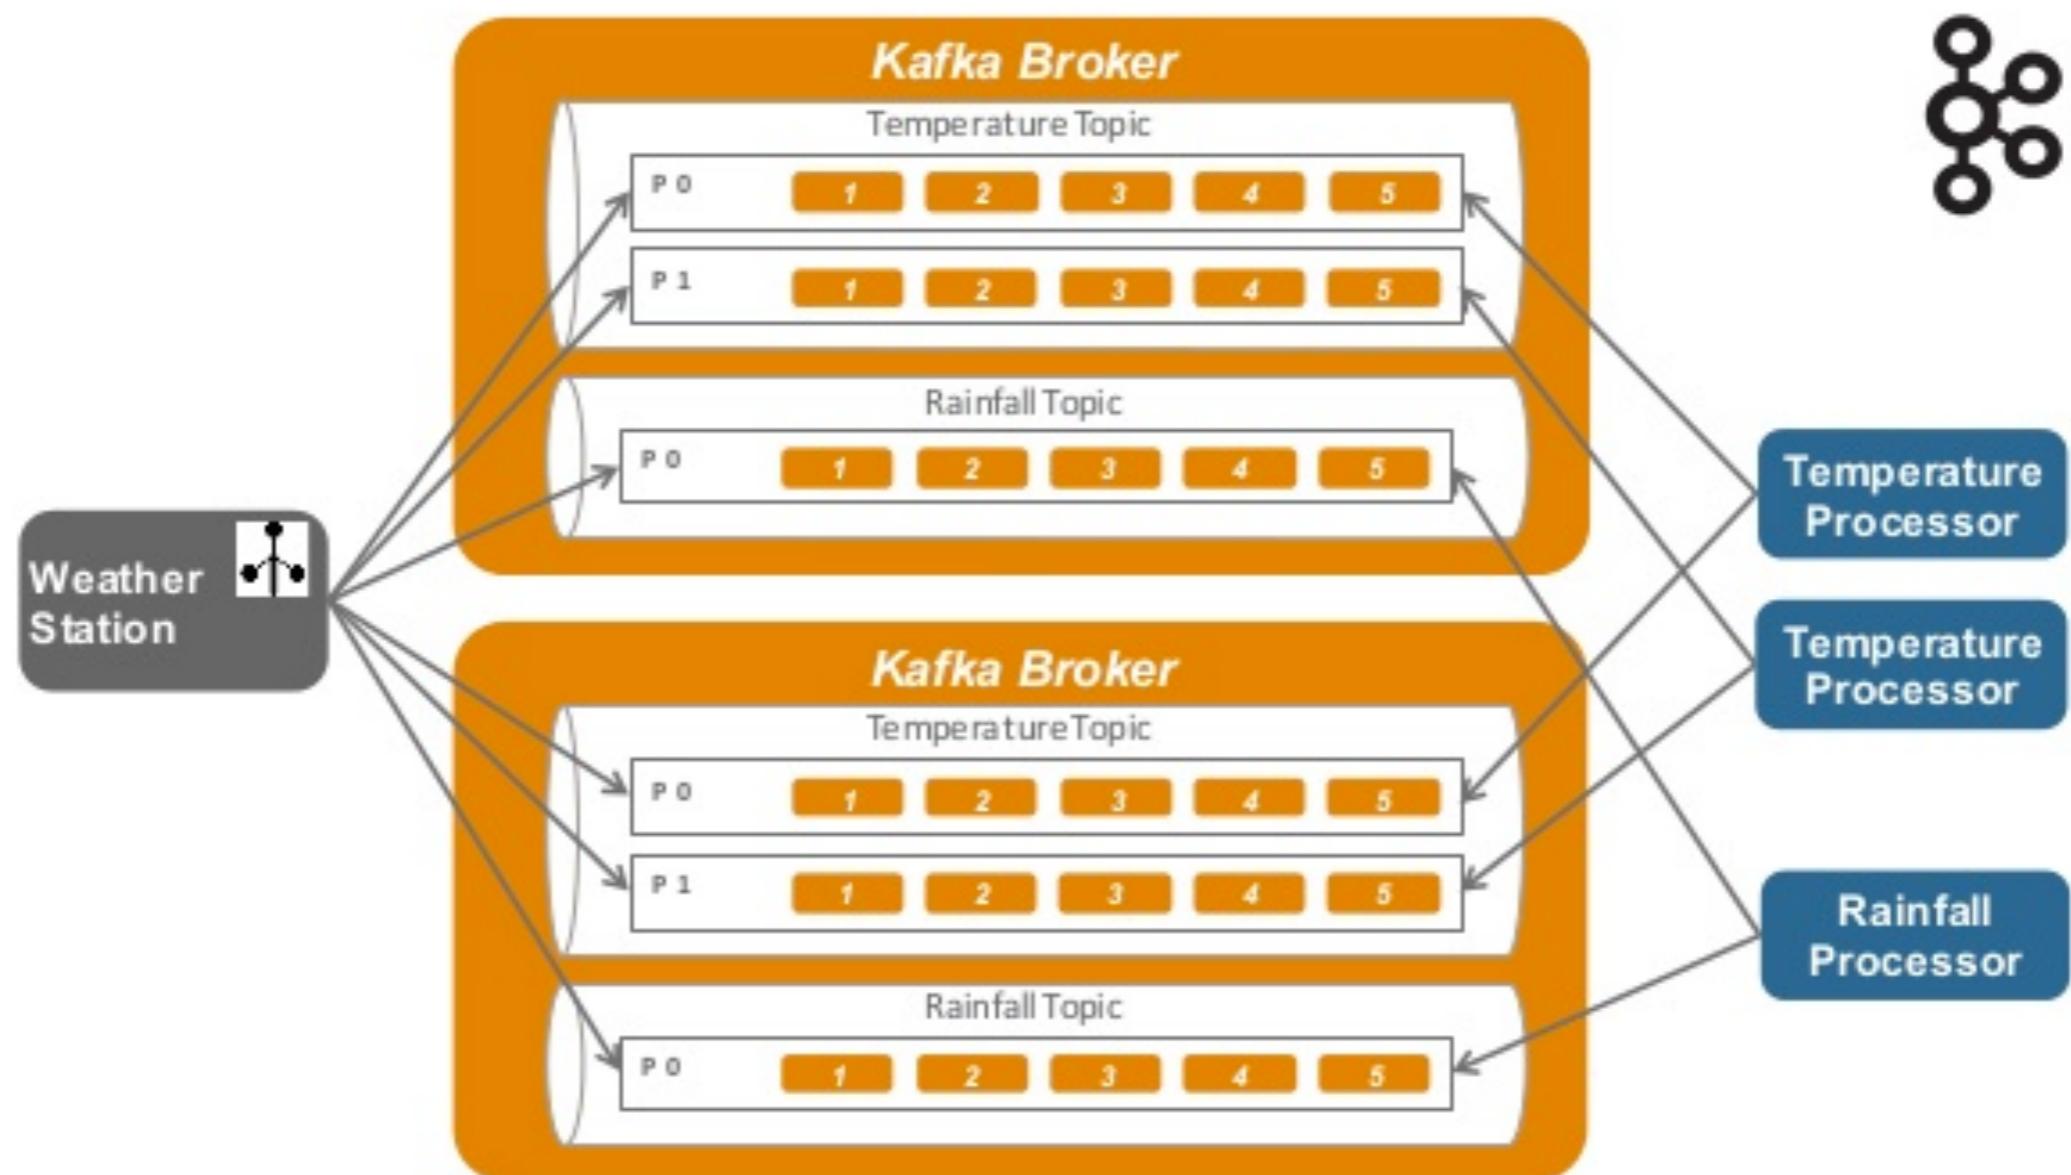
Apache kafka

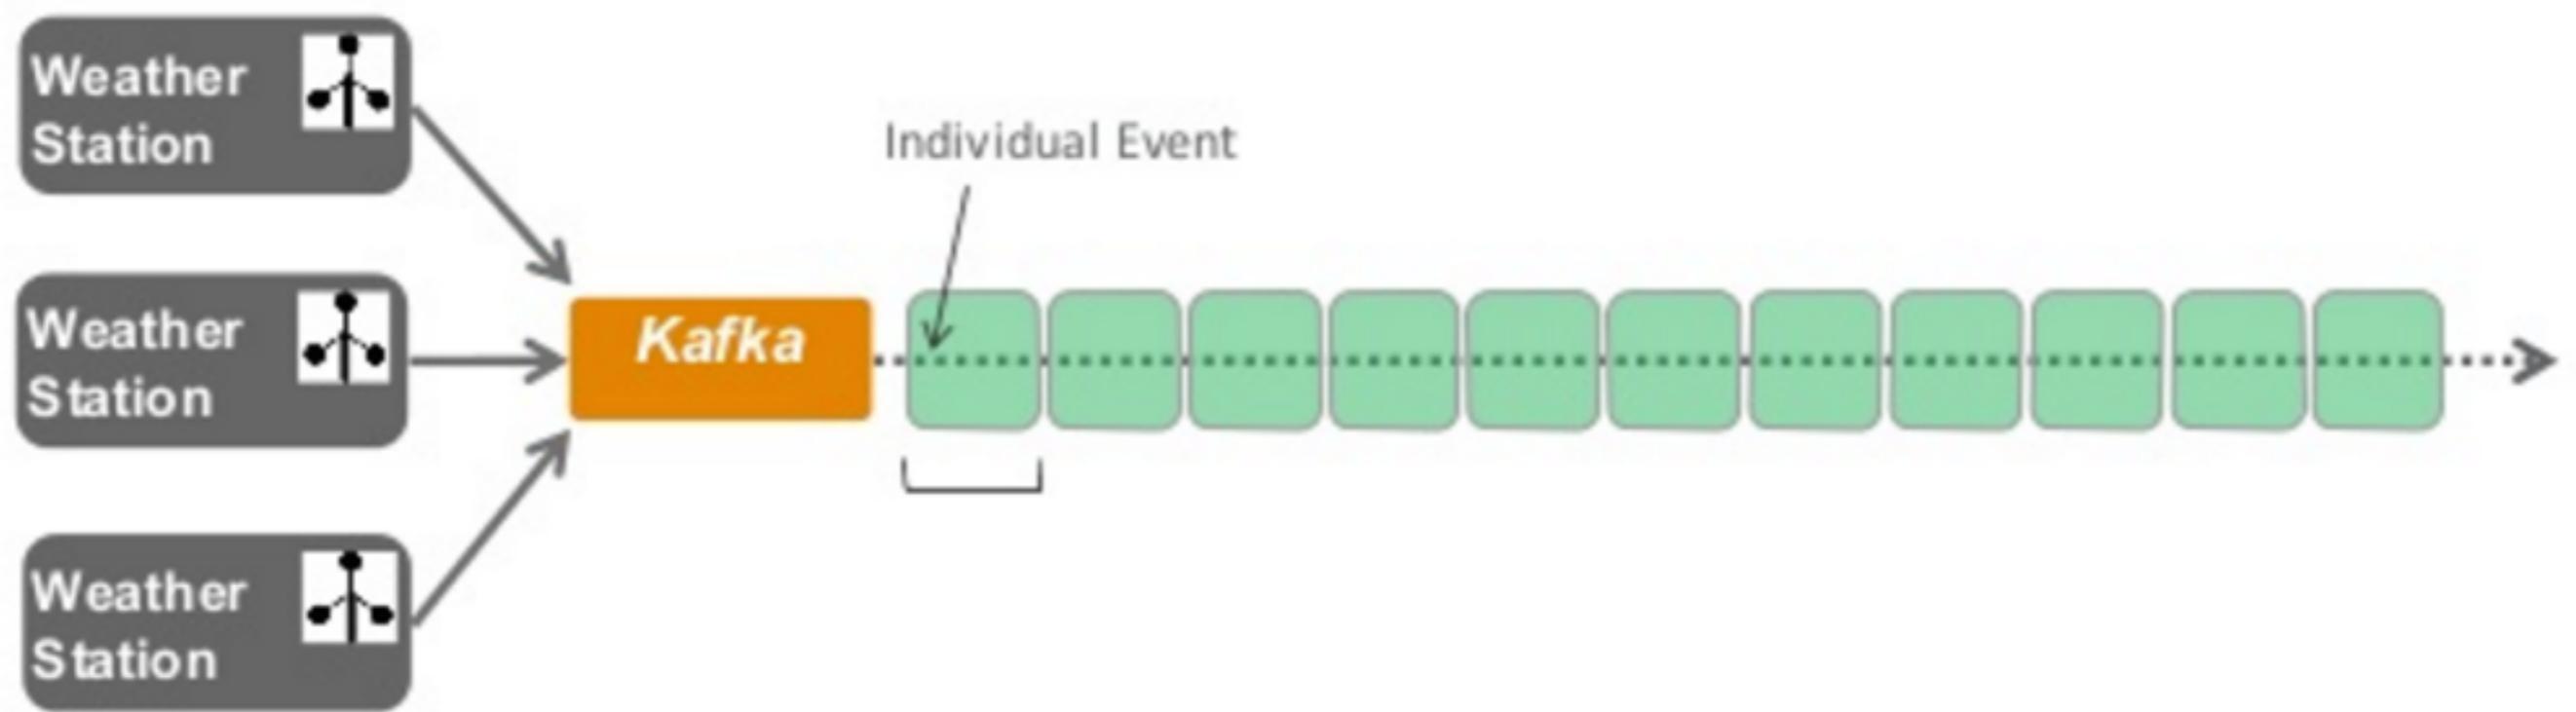
- Distributed publish-subscribe messaging system
- Processing real time activity stream metrics, collections, social media s
- Initially by LinkedIn, now Apache p
- Does not use JMS API and standards
- Kafka maintains feeds of messages in topics





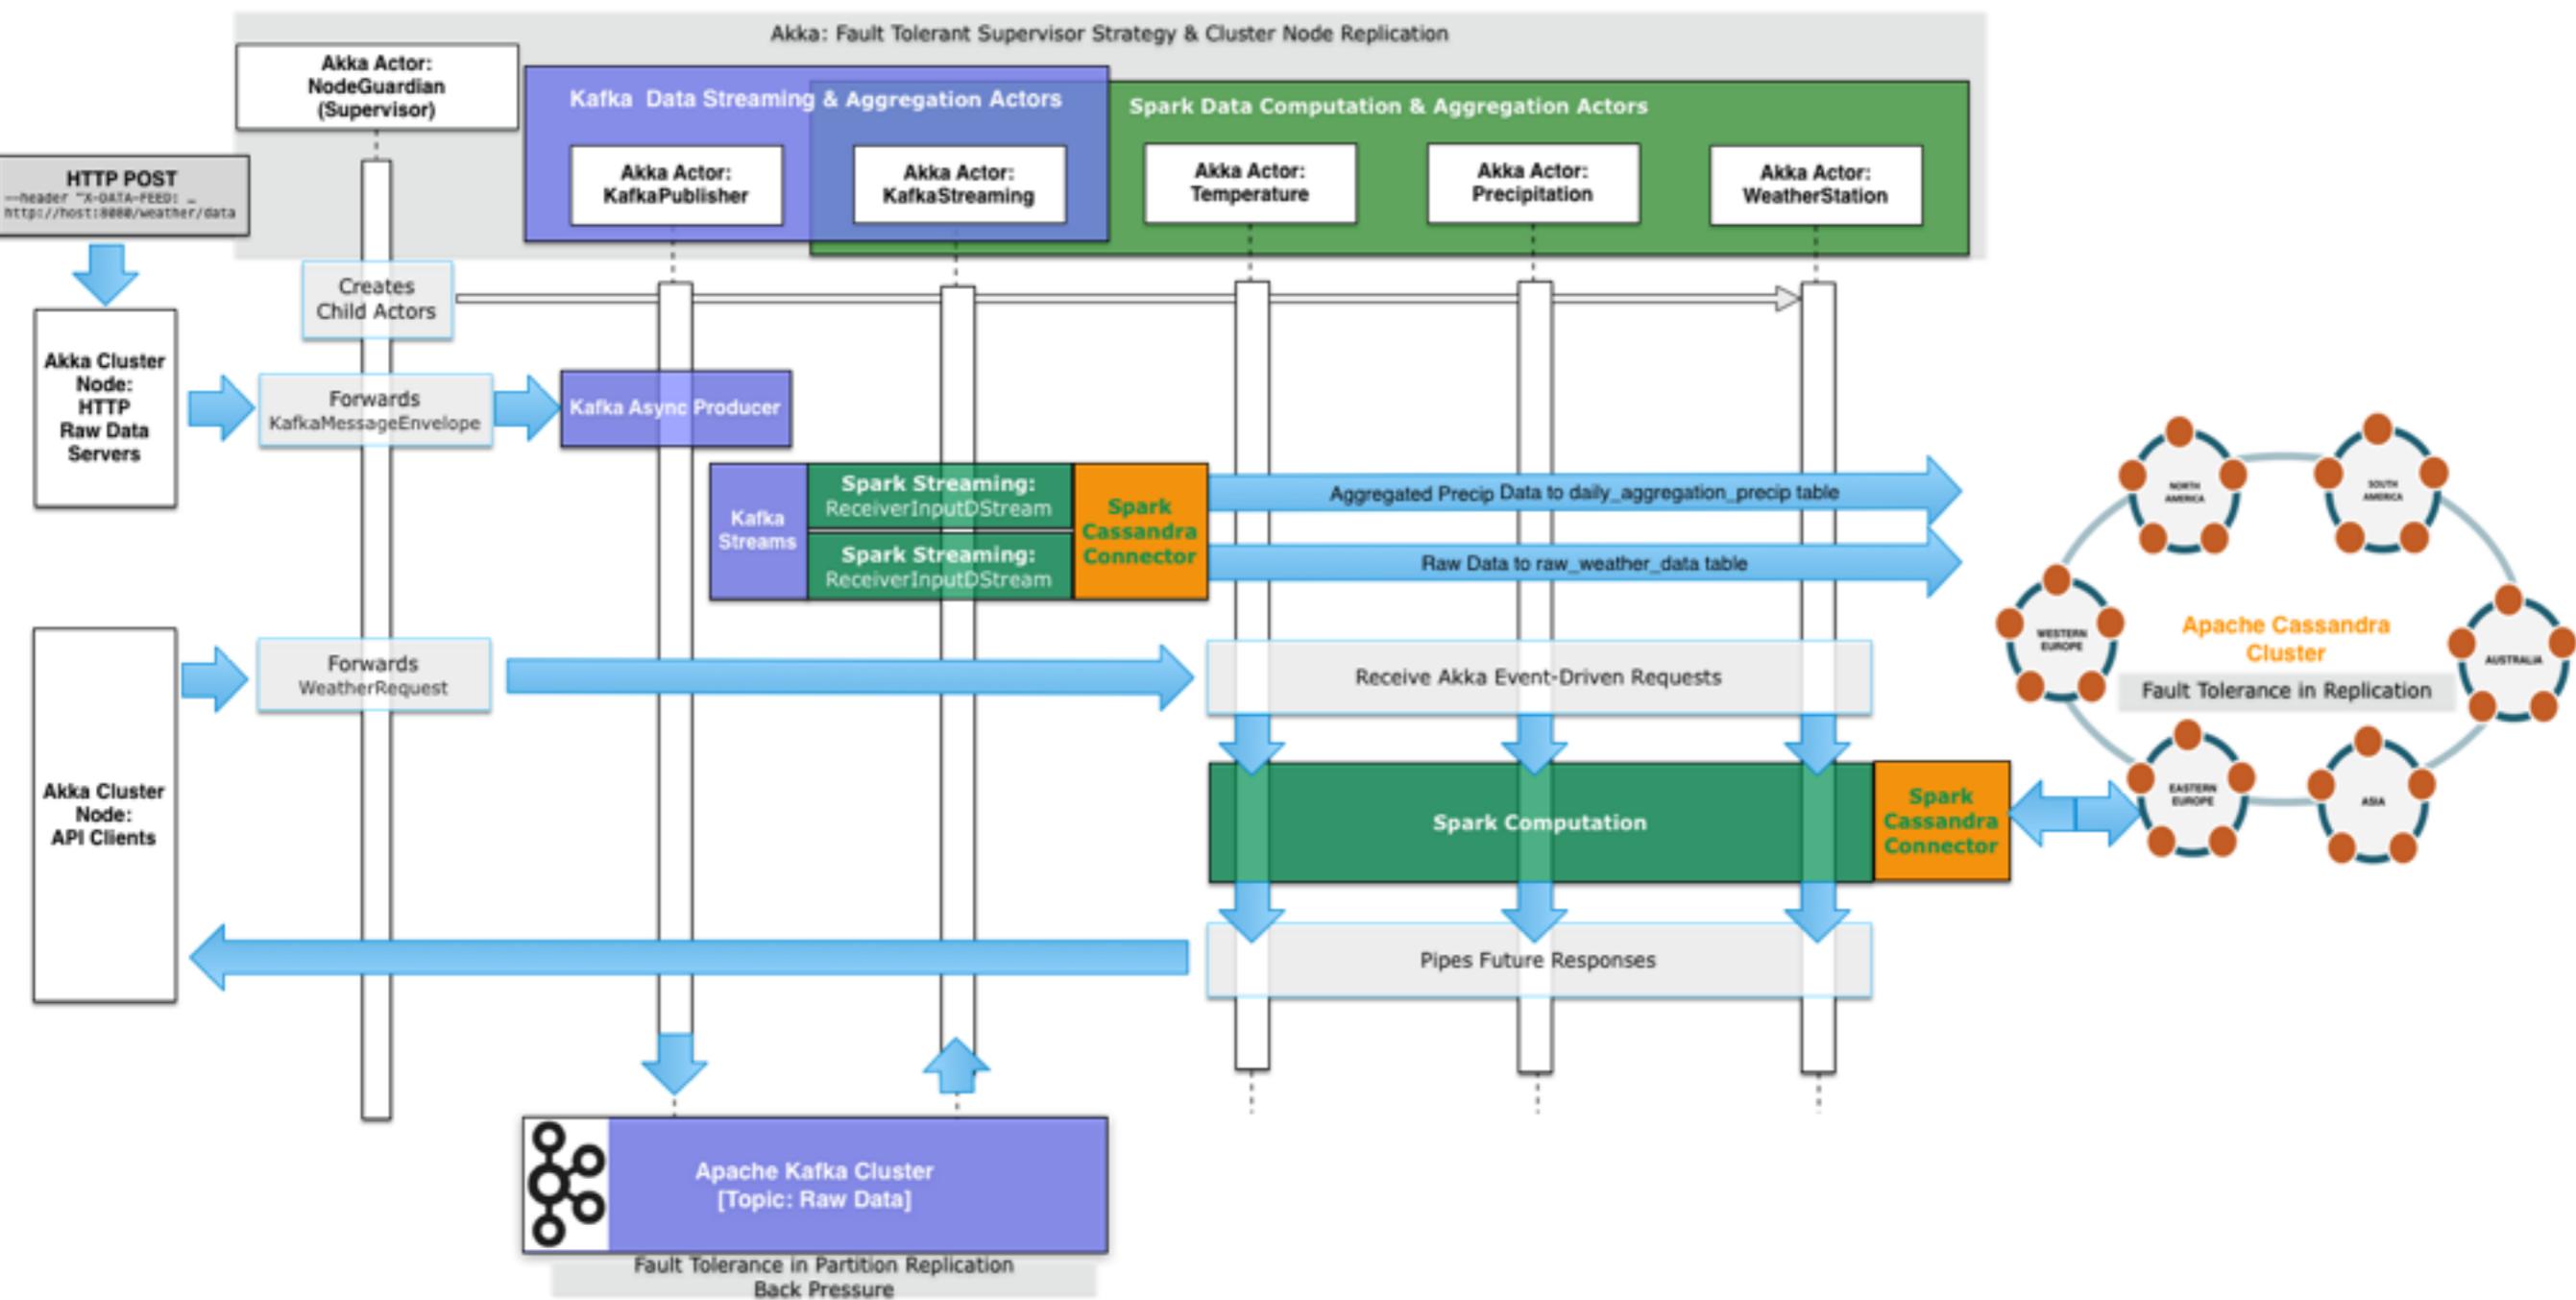








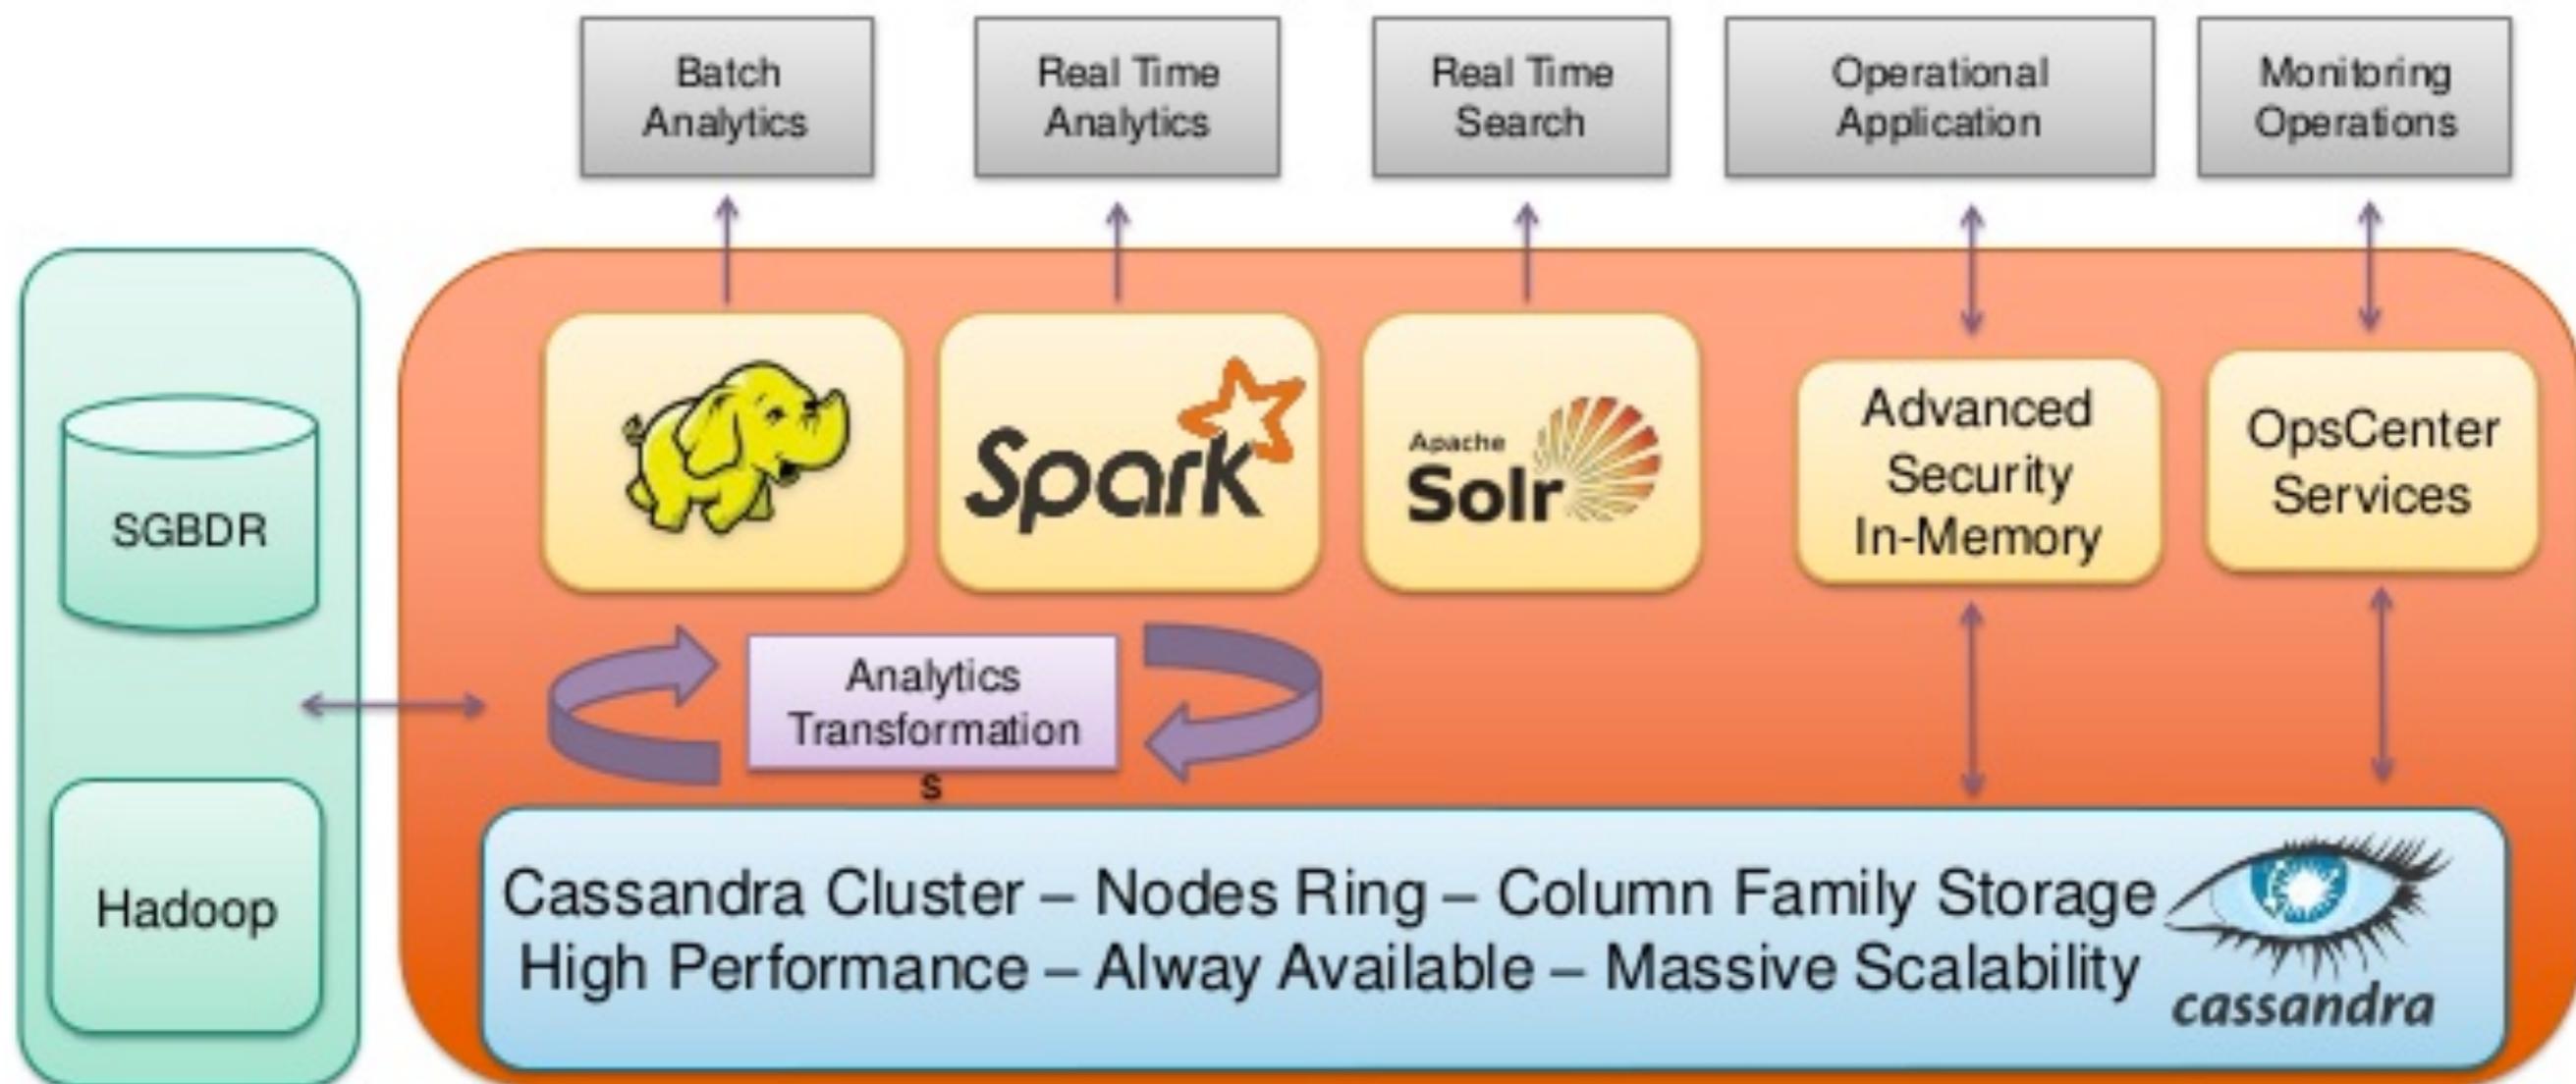
Data Flow Spark Cassandra Kafka Akka



Combining spark and cassandra

- Spark
 - Great for analyzing large amount of data
- Cassandra
 - Great for storing large amount of data

DATASTAX



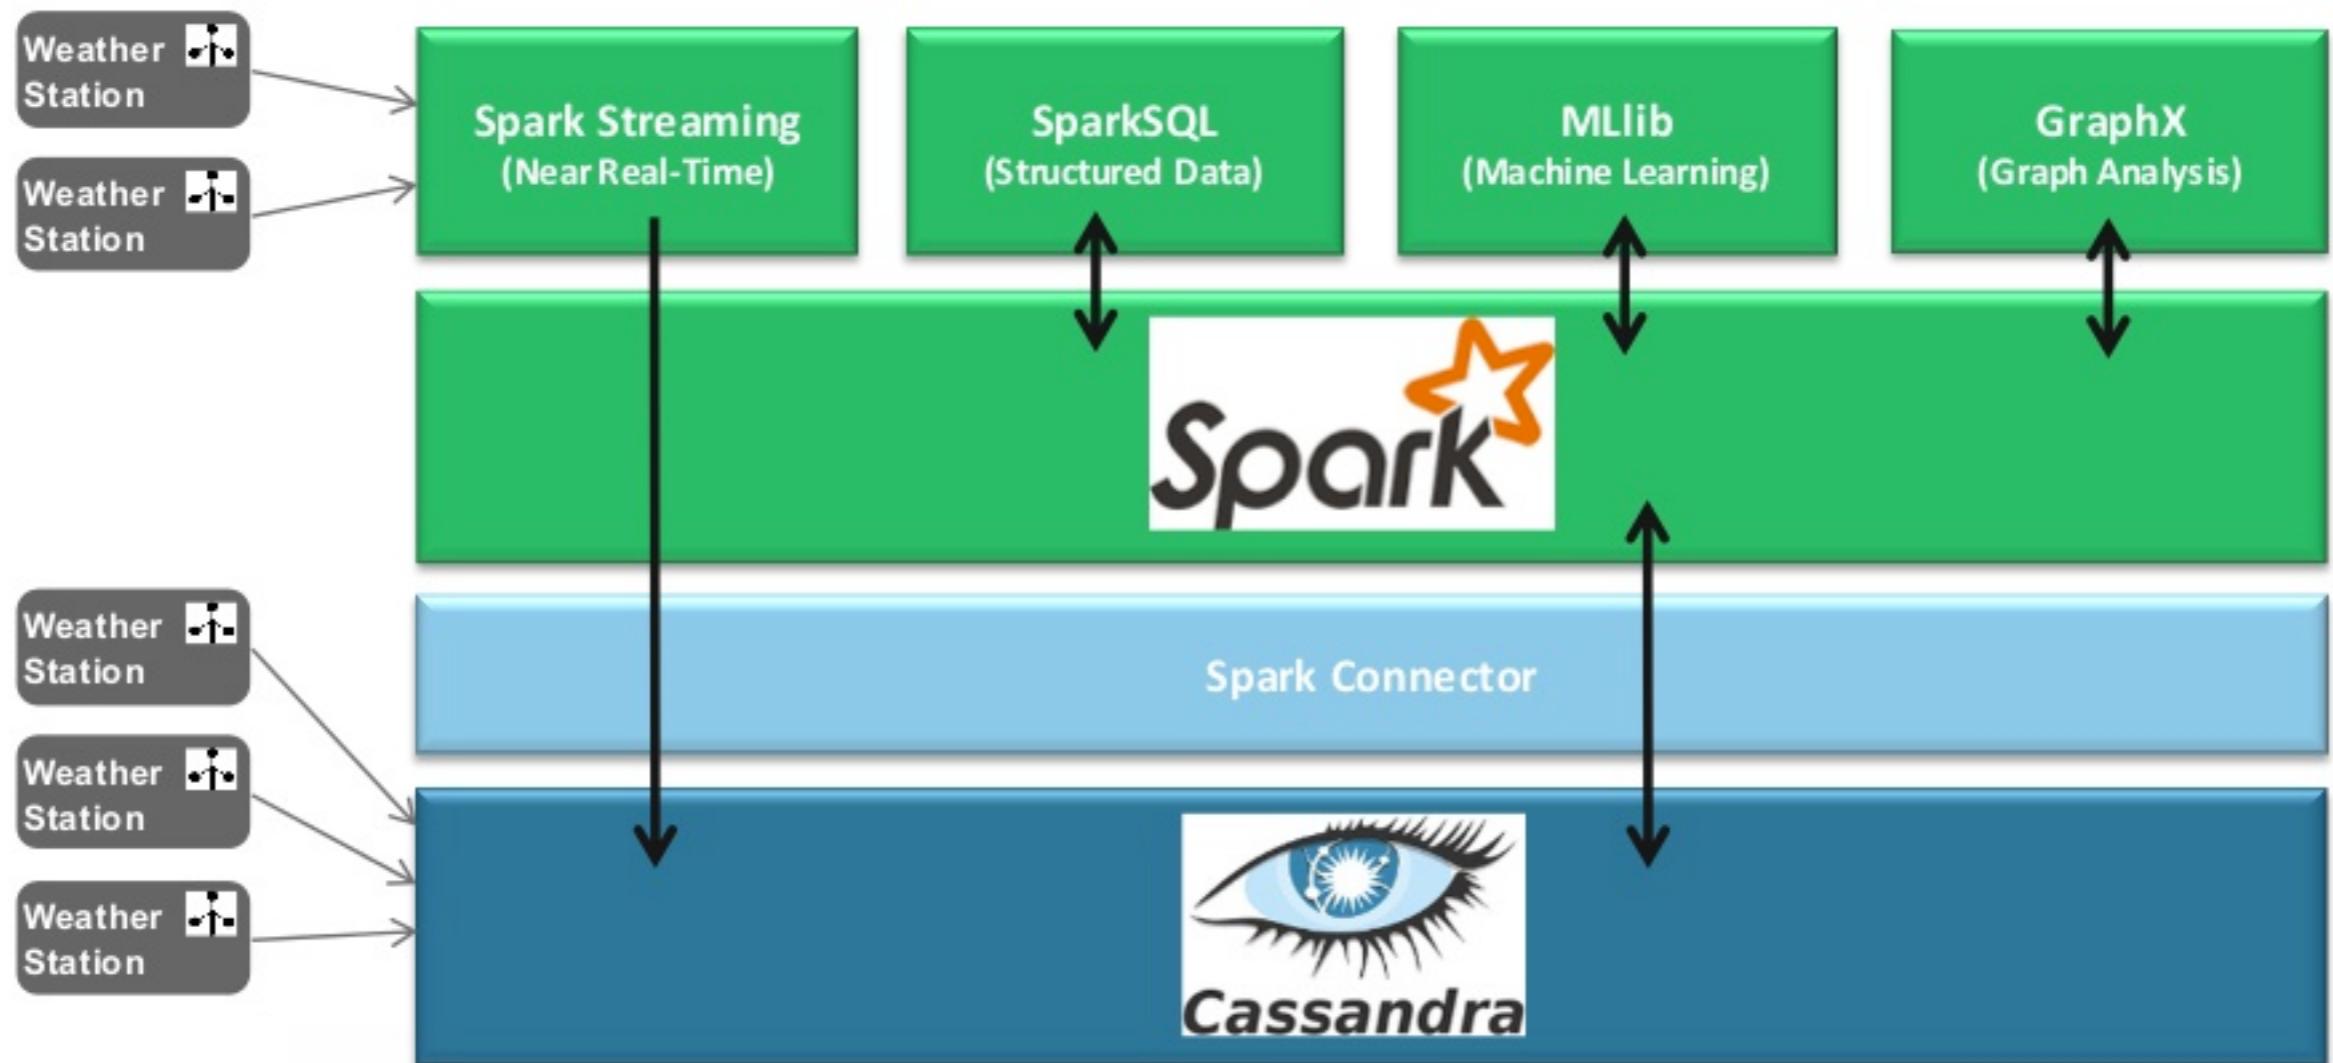
Spark Streaming
(Near Real-Time)

SparkSQL
(Structured Data)

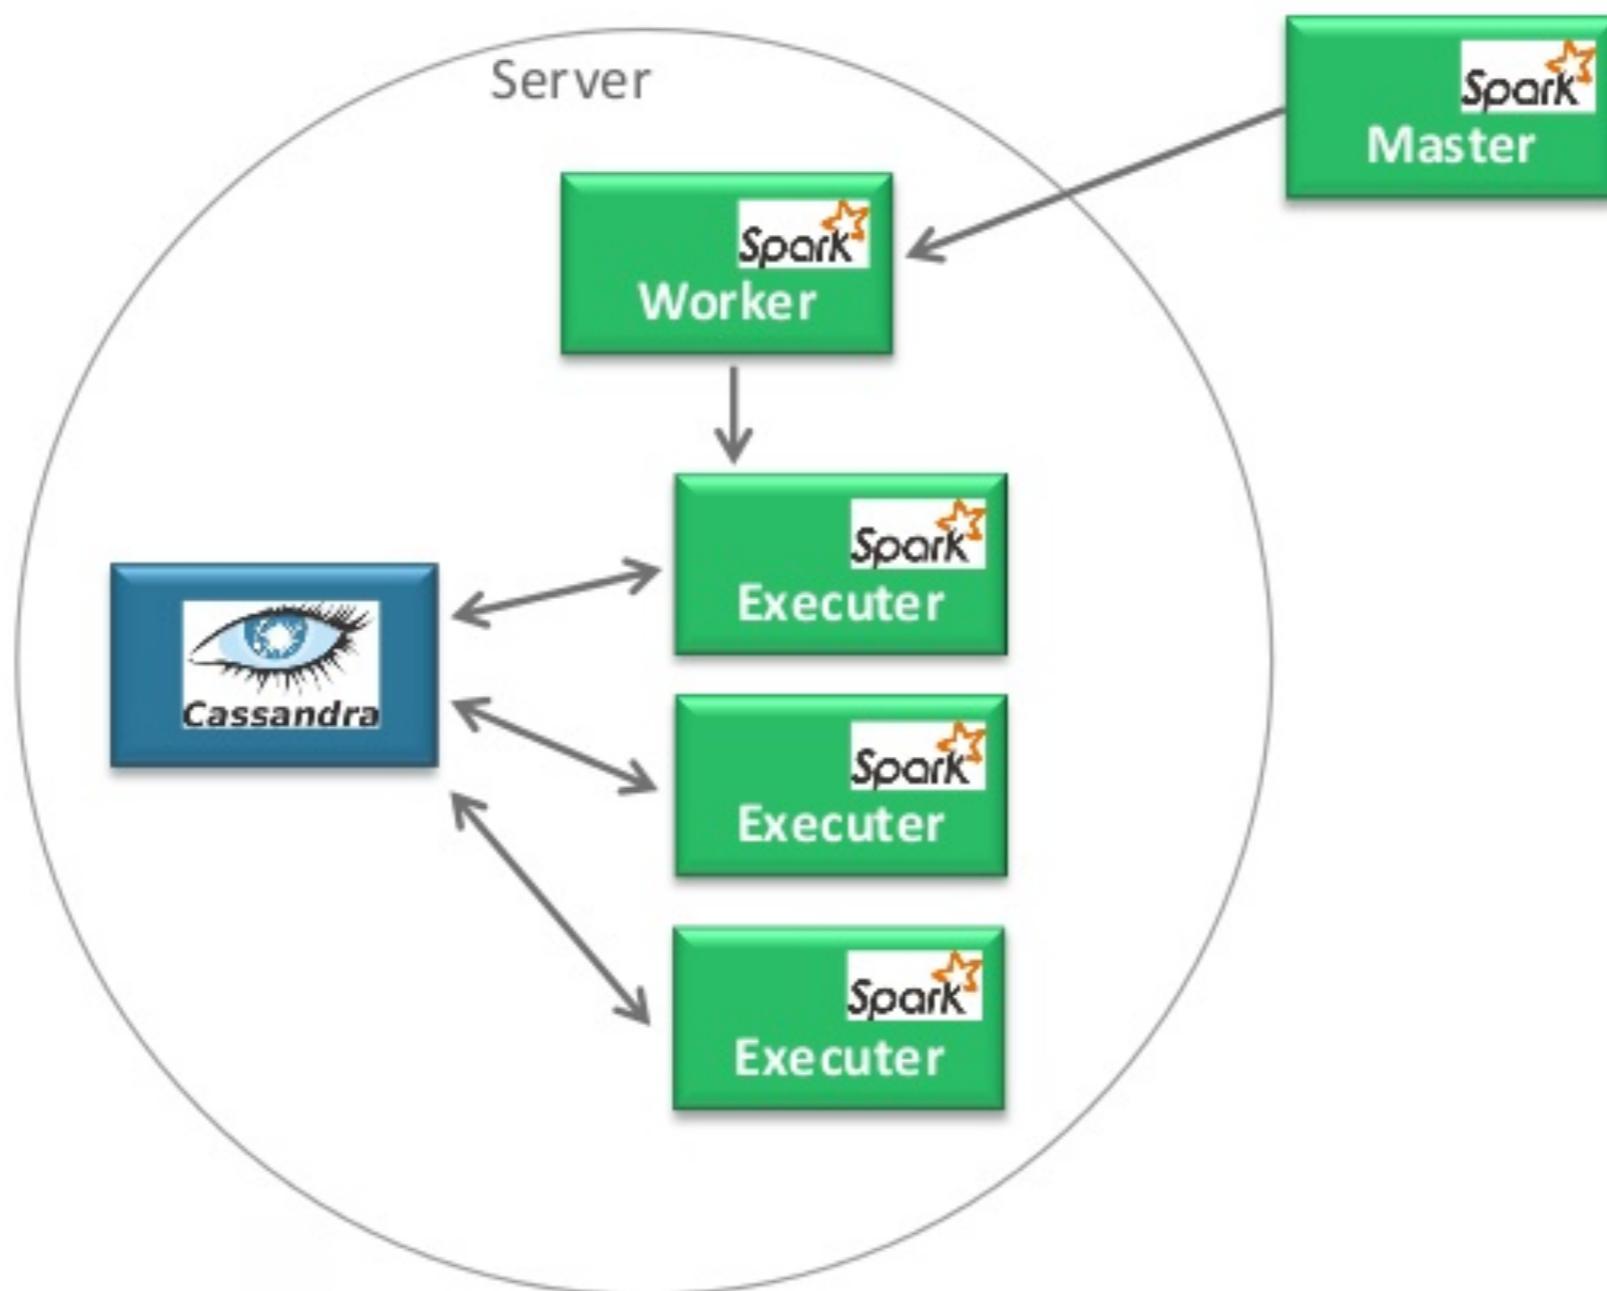
Mlib
(Machine Learning)

GraphX
(Graph Analysis)





Spark and Cassandra architecture

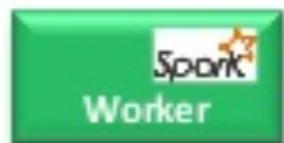




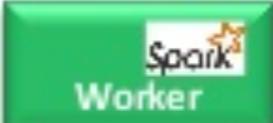
Master

Will only have
to analyze 25%
of data!

0-25



76-100



26-50



51-75



Transactional

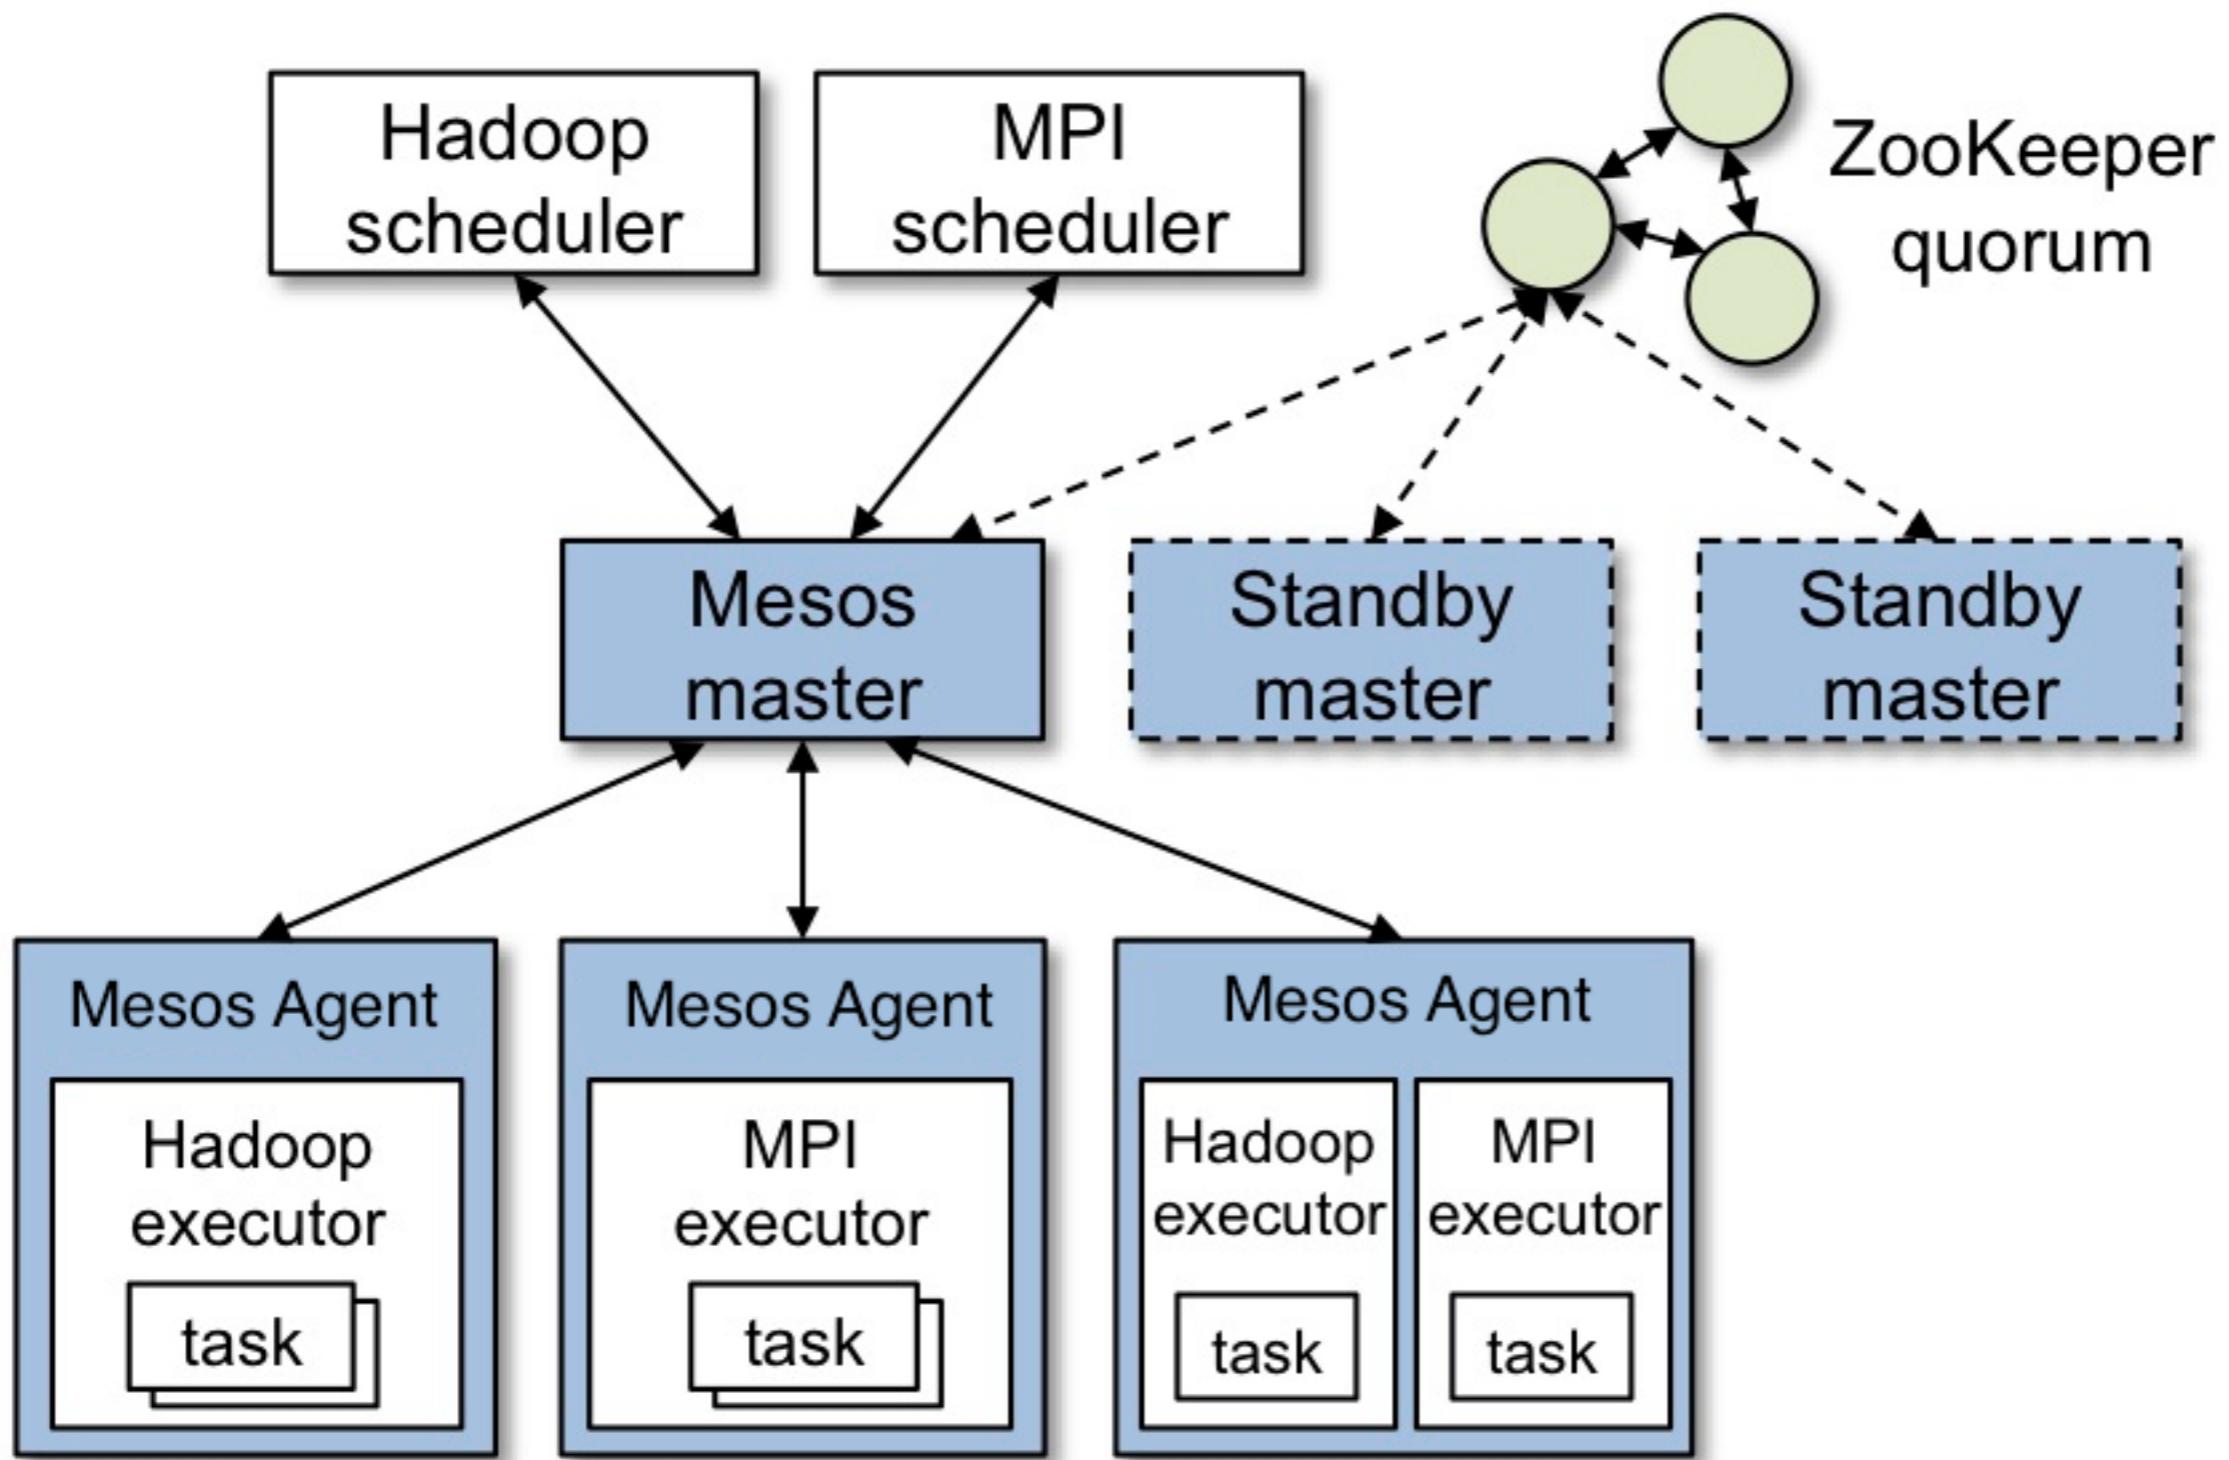


Analytics



cassandra and spark

	Cassandra	Cassandra & Spark
Joins and Unions	No	Yes
Transformations	Limited	Yes
Outside Data Integration	No	Yes
Aggregations	Limited	Yes

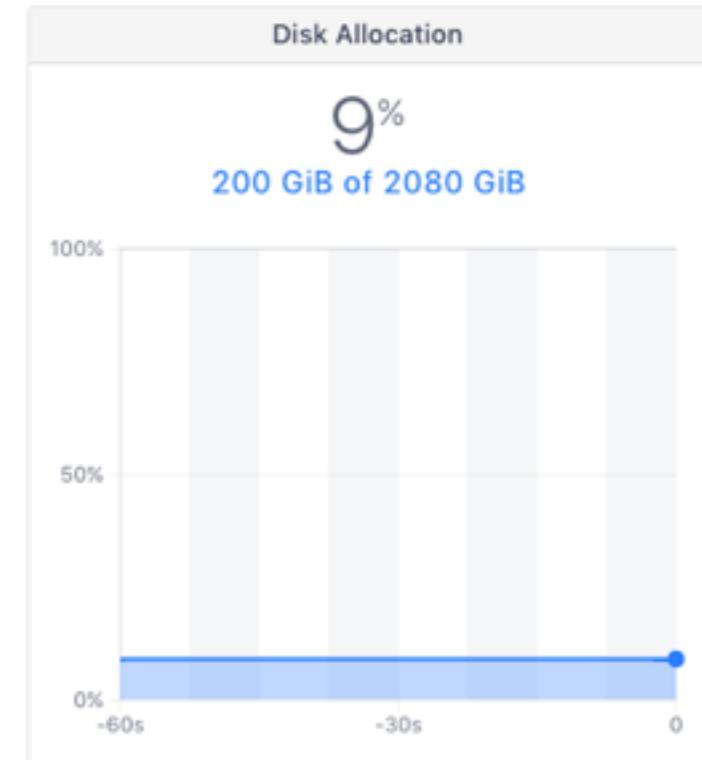
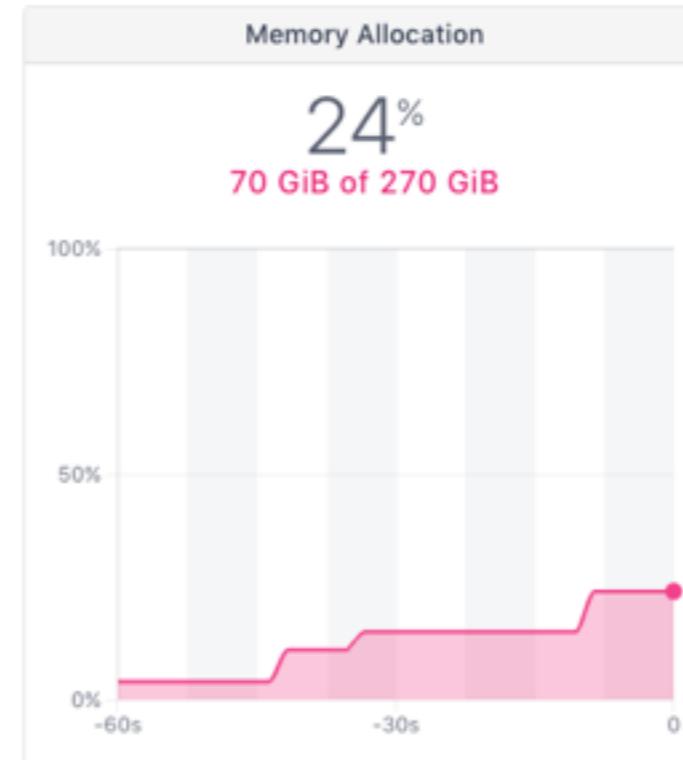
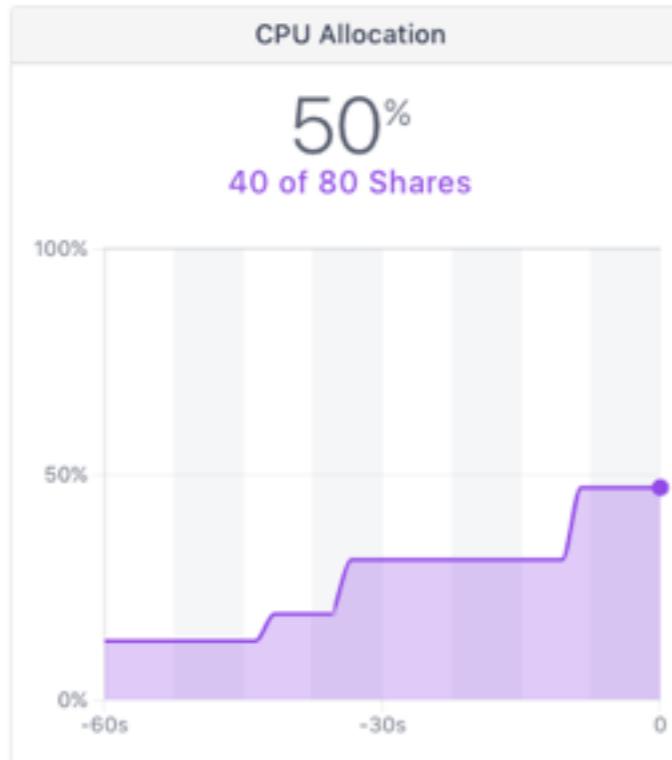


 Dashboard Services Jobs Universe

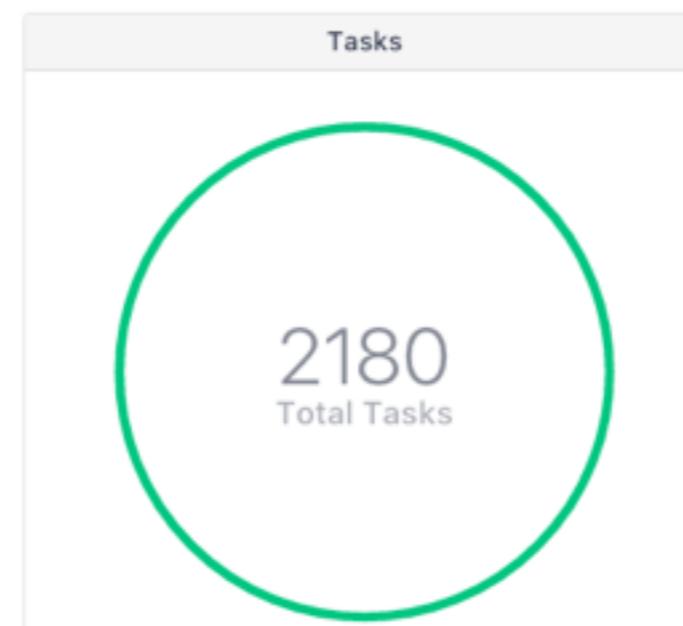
RESOURCES

 Nodes Networking Security

SYSTEM

 Cluster Components Settings Organization

Services Health	
arangodb3	Healthy
cassandra	Healthy
kafka	Healthy
nginx	Healthy
confluent-kafka	Healthy
jenkins	Healthy
linkerd	Healthy



Component Health	
Admin Router Agent	Healthy
Admin Router Master	Healthy
Admin Router Reloader	Healthy
Admin Router Reloader	Healthy
Admin Router Reloader Timer	Healthy

Packages

Dashboard

Services

Jobs

Universe

Packages

Installed

RESOURCES

Nodes

Networking

Security

SYSTEM

Cluster

Components

Settings

Organization

**Apache Spark**

1.0.7-2.1.0

INSTALL PACKAGE**Datastax Ent Max**

1.0.16-5.0.2

INSTALL PACKAGE**Confluent Kafka**

0.9.6-3.1.2

INSTALL PACKAGE**Elastic Stack**

1.0.5-5.2.1

INSTALL PACKAGE**Apache HDFS**

1.0.0-2.6.0

INSTALL PACKAGE**Couchbase**

4.6.0

INSTALL PACKAGE**Apache Flink**

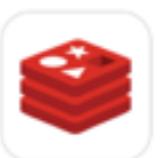
1.2.0-1.0

INSTALL PACKAGE**Alluxio**

1.4.0

INSTALL PACKAGE**Lightbend Reactive**

2.1.0

**Redis**

2.0.7-0.0.1

**Jenkins**

2.0.2-2.22.2

**Basho Riak**

2.0.0

Dashboard

Services ▾

Services

Deployments

Jobs

Universe ▾

RESOURCES

Nodes

Networking ▾

Security ▾

SYSTEM

Cluster

Components

Settings ▾

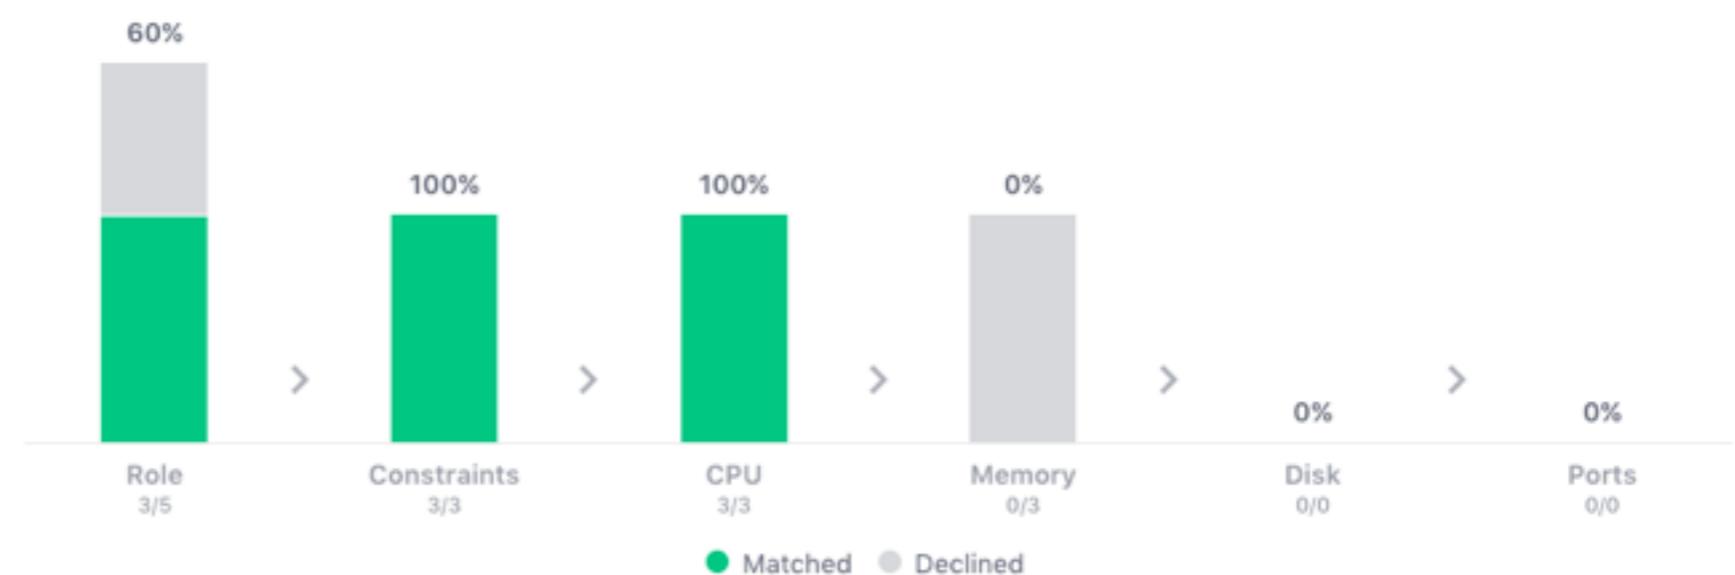
Organization ▾

[Instances](#) [Configuration](#) [Debug](#)

Recent Resource Offers (5)

When you attempt to deploy a service, DC/OS waits for offers to match the resources your service requires. If the offer does not satisfy the requirement, it is declined and DC/OS retries. [Learn more](#).

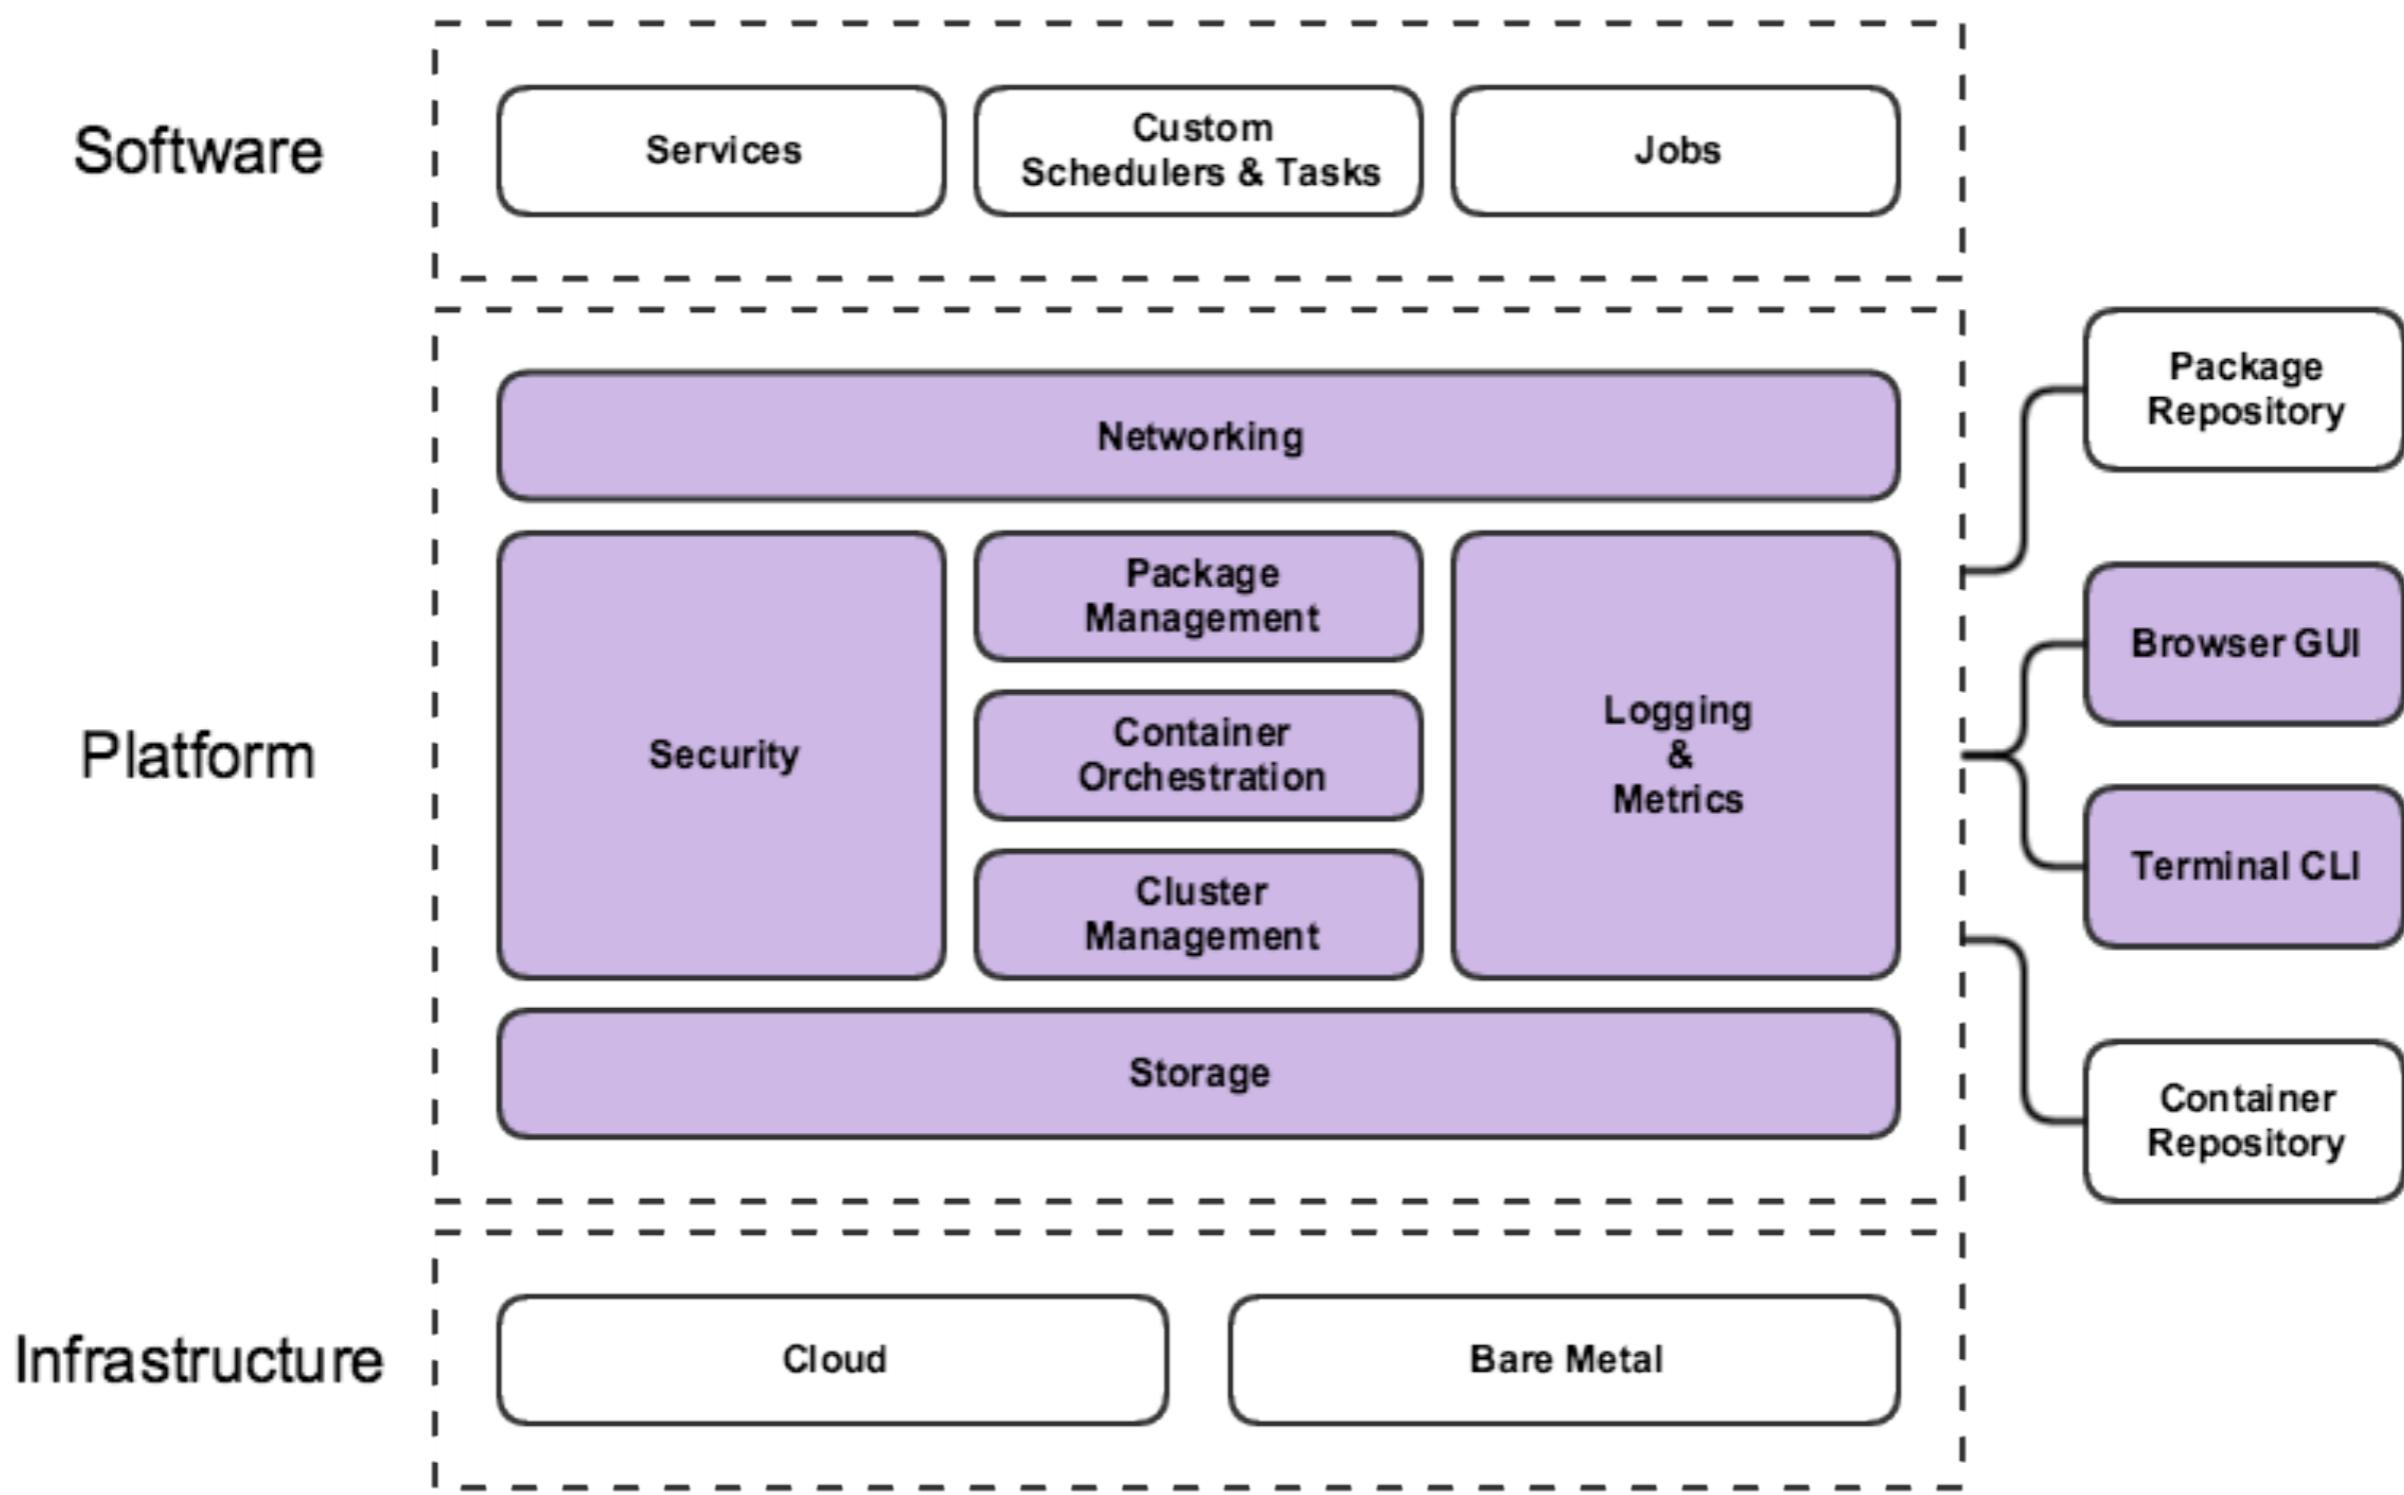
Summary



Details

HOST	ROLE	CONSTRAINT	CPU	MEM	DISK	PORT	RECEIVED
10.0.1.155	✓	✓	✓	✗	✓	✓	2 minutes ago
10.0.3.205	✓	✓	✓	✗	✓	✓	2 minutes ago
10.0.3.241	✓	✓	✓	✗	✓	✓	2 minutes ago
10.0.4.213	✗	✓	✗	✗	✓	✓	2 minutes ago

DC/OS Architecture Layers



SMACK AGENDA

- SMACK stack overview
- Bigdata characteristics
- Cassandra NoSQL database
- Spark Streaming
- Kafka and Akka event sourcing
- Mesos Cluster management
- Mesos Scheduling and execution

MODS

MOBILE & DISRUPTIVE
TECHNOLOGY SUMMIT

October 5-6, 2017

Indian Institute of Science, Bangalore

www.modsummit.com



Register early and get the best discounts

GREAT INDIAN DEVELOPER SUMMITTM



April 23-28, 2018

Indian Institute of Science, Bangalore

www.developersummit.com

